

CS 143A

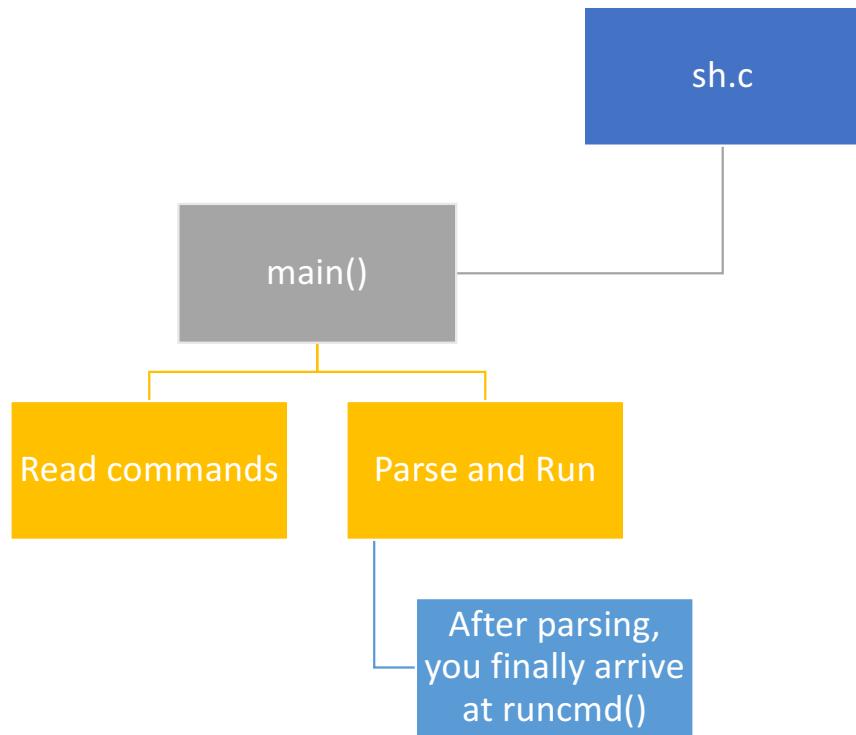
Principles of Operating Systems

Agenda

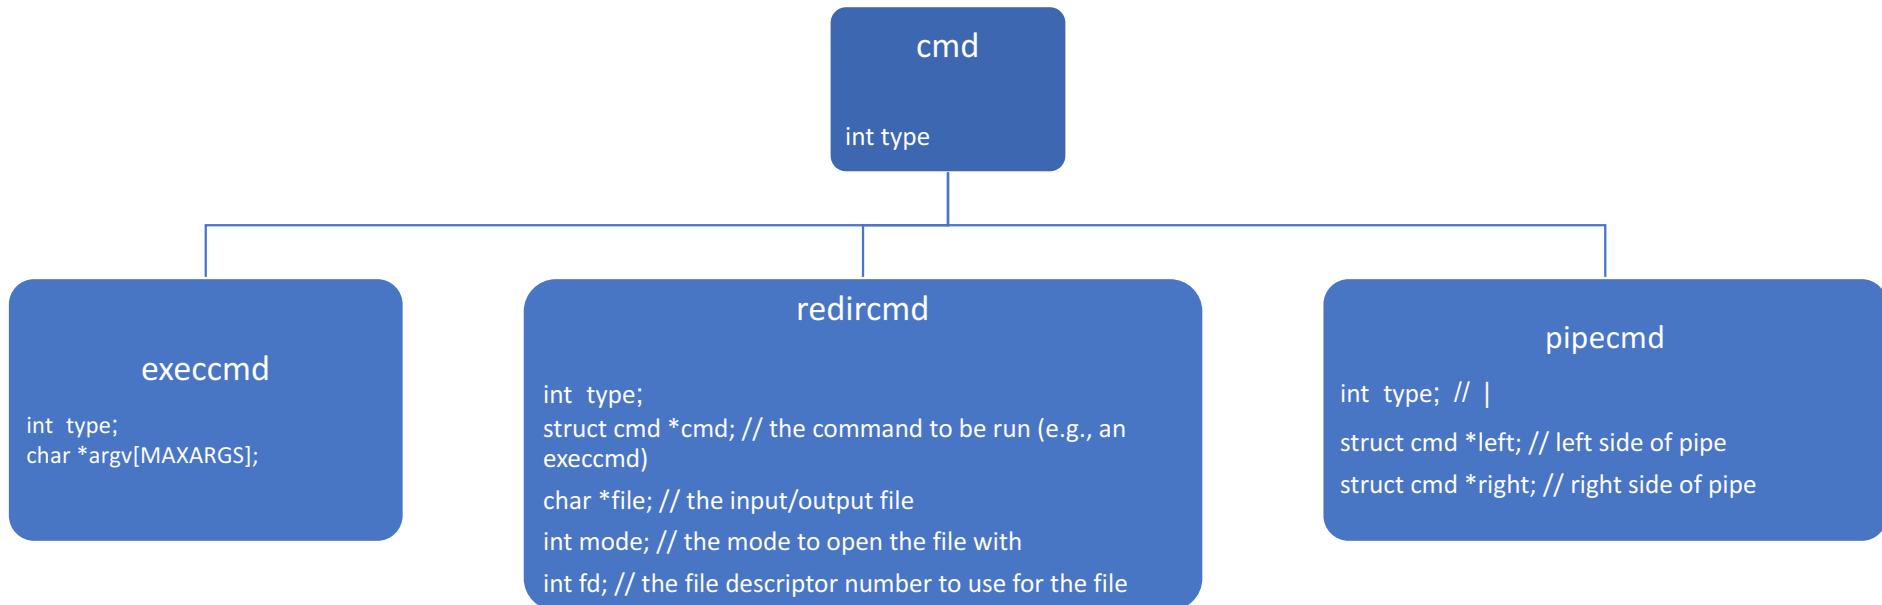
Homework 1 solution

Homework 2 solution

Homework 3 open questions (offline)



Homework 1: Program Structure



Homework 1 : Program Structure

What happens in an **execvp** call:

- Your shell searches in \$PATH
- Typically, if the executable is found, a call to execve will be made.
- Linux prepares a program by allocating memory for it, opening it, scheduling it for execution, initializes memory structures, sets up its arguments and environment from the supplied arguments to the execvp call, finds a handler appropriate for loading the binary, and sets the current task (the execvp caller) as not executing.

Homework 1 Q1: **execcmd**

```
case ' ': ;
    ecmd = (struct execmd*)cmd;
    if(ecmd->argv[0] == 0)
        exit(0);
    // Your code here ...
    if (execvp((const char*)ecmd->argv[0], ecmd->argv) == -1)
        perror("exec: ");
    break;
```

Homework 1 Q1: execmd

Homework 1

Q1: execcmd

Common Errors

```
char path_bin[100] = "/bin/";
char path_usr[100] = "/usr/bin/";

//try argv[0]
execv(ecmd->argv[0], ecmd->argv);
//try /bin/argv[0]
execv(strcat(path_bin, ecmd->argv[0]), ecmd->argv);
//try /usr/bin/argv[0]s
execv(strcat(path_usr, ecmd->argv[0]), ecmd->argv);

fprintf(stderr, " %s eexec error\n", ecmd->argv[0]);
```

Homework 1

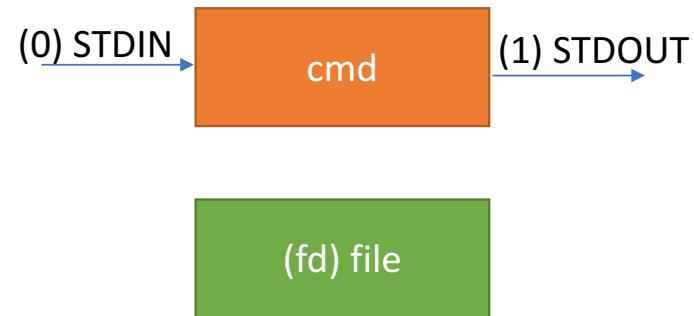
Q1: execcmd

Common Errors

```
case ' ':
    ecmd = (struct execcmd *)cmd;
    if (ecmd->argv[0] == 0)
        exit(0);
    // fprintf(stderr, "exec not implemented\n");
    // Your code here ...
    for (unsigned int i = 0; i < 3; ++i)
    {
        char path[256] = {0};
        strcpy(path, PATH[i]);
        strcat(path, ecmd->argv[0]);
        execv(path, ecmd->argv);
    }
    fprintf(stderr, "%s: Command not found.\n", ecmd->argv[0]);
    break;
```

redircmd

```
int type;  
struct cmd *cmd; // the command to be run (e.g., an  
execcmd)  
char *file; // the input/output file  
int mode; // the mode to open the file with  
int fd; // the file descriptor number to use for the file
```



Default

Homework 1 Q2: **redircmd**

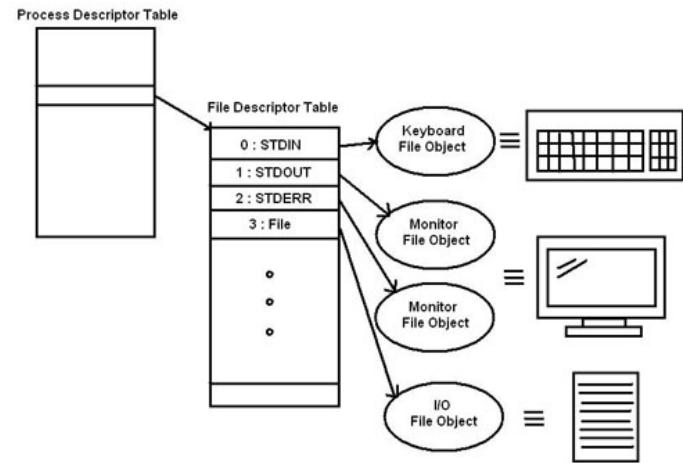
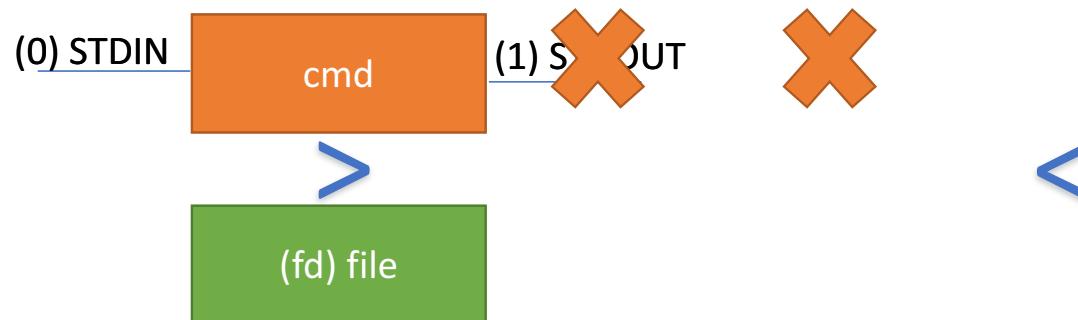


```
struct cmd*
redircmd(struct cmd *subcmd, char *file, int type)
{
    struct redircmd *cmd;

    cmd = malloc(sizeof(*cmd));
    memset(cmd, 0, sizeof(*cmd));
    cmd->type = type;
    cmd->cmd = subcmd;
    cmd->file = file;
    cmd->mode = (type == '<') ? 0_RDONLY : 0_WRONLY|O_CREAT|O_TRUNC;
    cmd->fd = (type == '<') ? 0 : 1;
    return (struct cmd*)cmd;
}
```



Homework 1 Q2: redircmd



```
case REDIR:  
    rcmd = (struct redircmd*)cmd;  
    close(rcmd->fd);  
    if(open(rcmd->file, rcmd->mode) < 0){  
        printf(2, "open %s failed\n", rcmd->file);  
        exit();  
    }  
    runcmd(rcmd->cmd);  
    break;
```

Homework 1 Q2: redircmd

Homework 1 Q2: redircmd with proper permission bits

- (400) S_IRUSR : Read permission bit for the owner of the file.
- (200) S_IWUSR: Write permission bit for the owner of the file.
- (40) S_IRGRP : Read permission bit for the group owner of the file.
- (20) S_IWGRP : Write permission bit for the group owner of the file.
- (4) S_IROTH : Read permission bit for other users. Usually 04.
- (2) S_IWOTH : Write permission bit for other users. Usually 02.

```
case '>':
case '<':
int fd;
rcmd = (struct redircmd *)cmd;

if (rcmd->mode == 0_RDONLY)
{
    fd = open(rcmd->file, rcmd->mode);
}
else
{
    // rw-rw-rw-
    fd = open(rcmd->file, rcmd->mode,
               S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH);
}

if (fd < 0)
{
    perror("open:");
    break;
}
// close either 0 or 1 depending on the redir
close(rcmd->fd);
// dup the fd to the one that is closed
dup2(fd, rcmd->fd);
// run
runcmd(rcmd->cmd);
break;
```

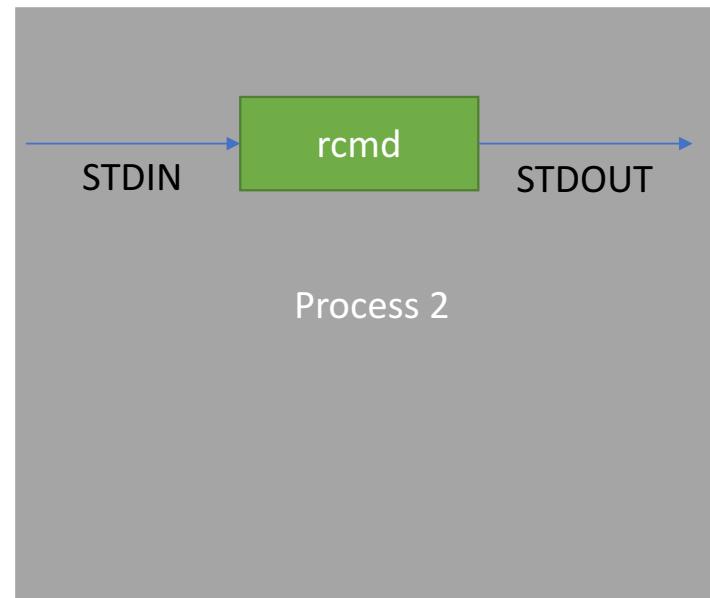
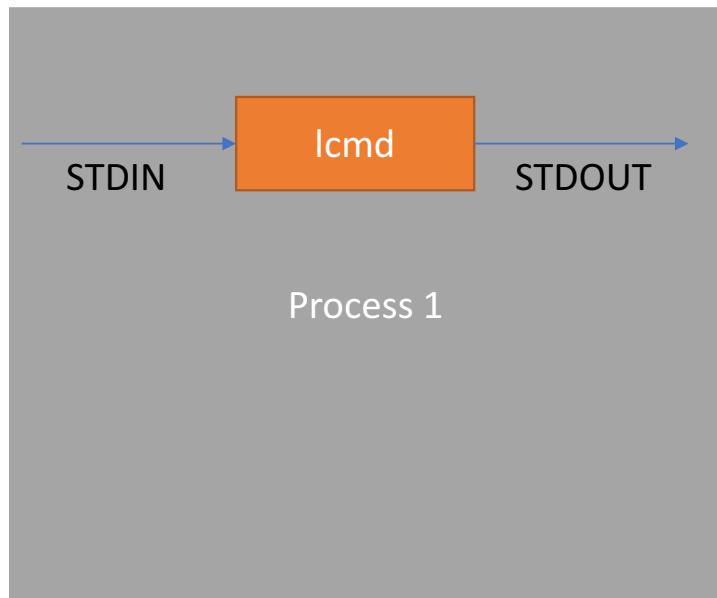
lcmd

rcmd

Homework 1 Q3: pipecmd

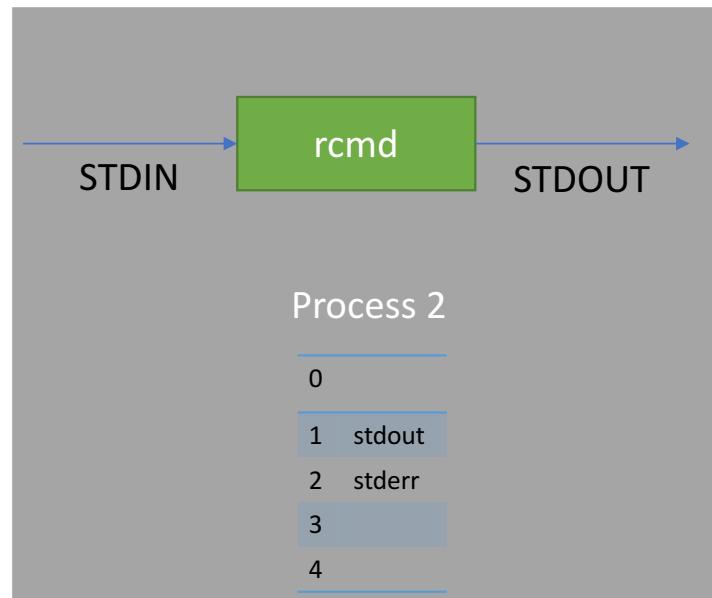
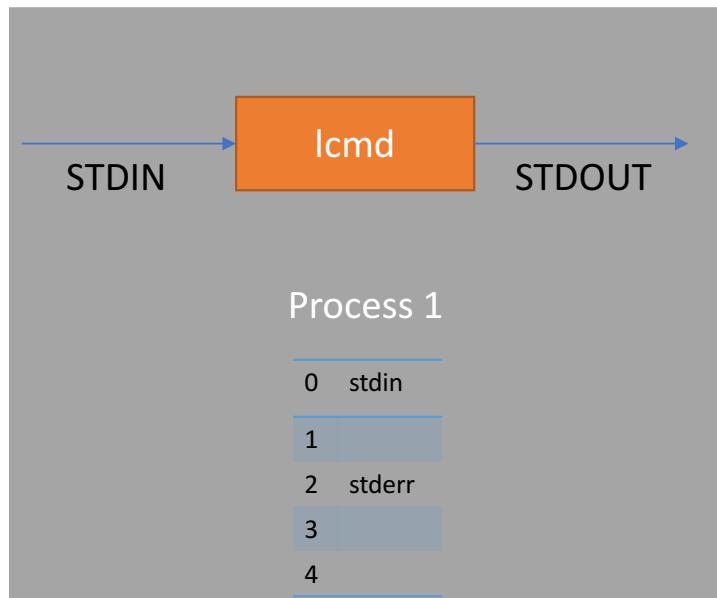


Homework 1 Q3: **pipecmd**



```
case PIPE:  
    if(fork1() == 0){  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        runcmd(pcmd->right);  
    }  
    break;
```

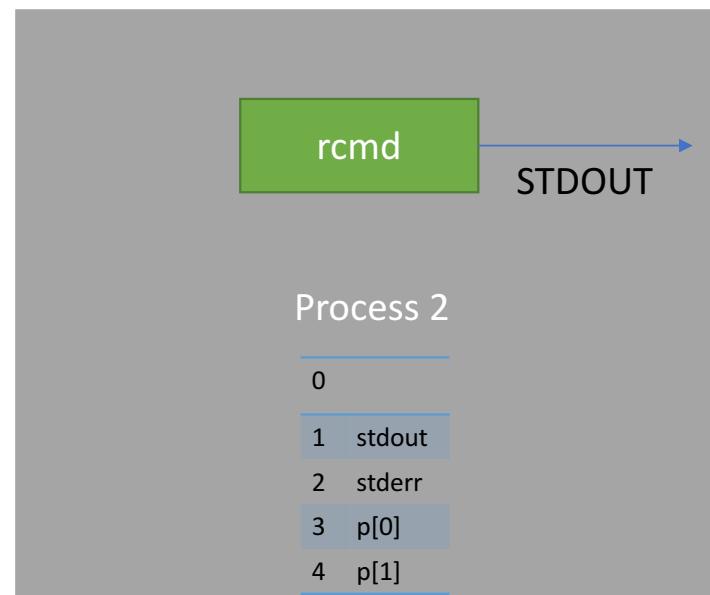
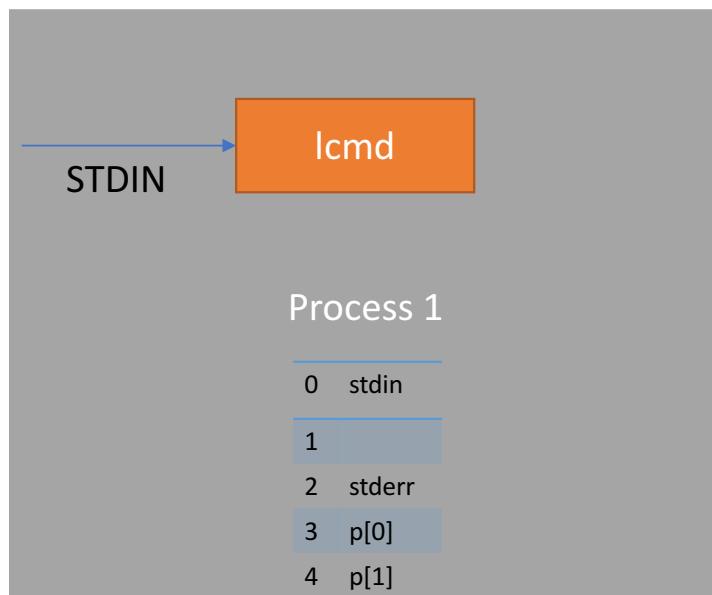
Homework 1 Q3: pipecmd



```
case PIPE:  
  
    if(fork1() == 0){  
        close(1);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        runcmd(pcmd->right);  
    }  
  
    break;
```

Homework 1 Q3: pipecmd

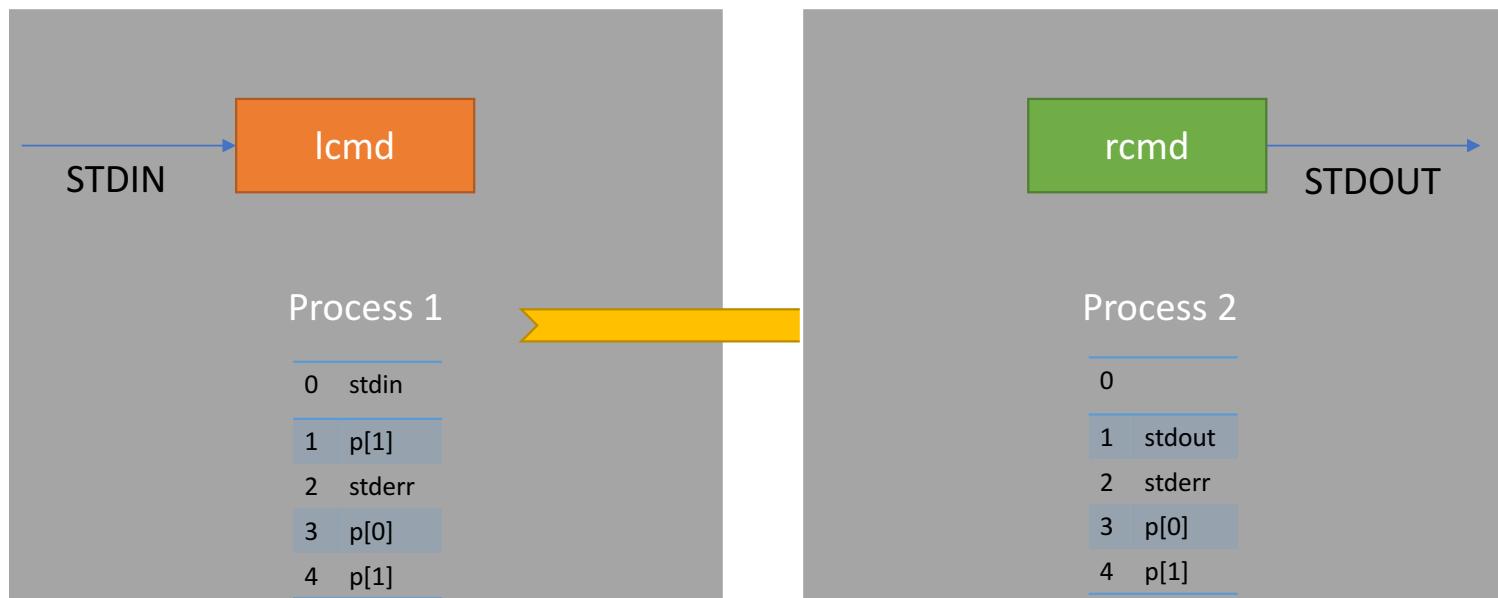
(set for writing) p[1]  p[0] (set for reading)



```
case PIPE:  
    pcmd = (struct pipecmd*)cmd;  
    if(pipe(p) < 0)  
        panic("pipe");  
    if(fork1() == 0){  
        close(1);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        runcmd(pcmd->right);  
    }  
  
    break;
```

Homework 1 Q3: pipecmd

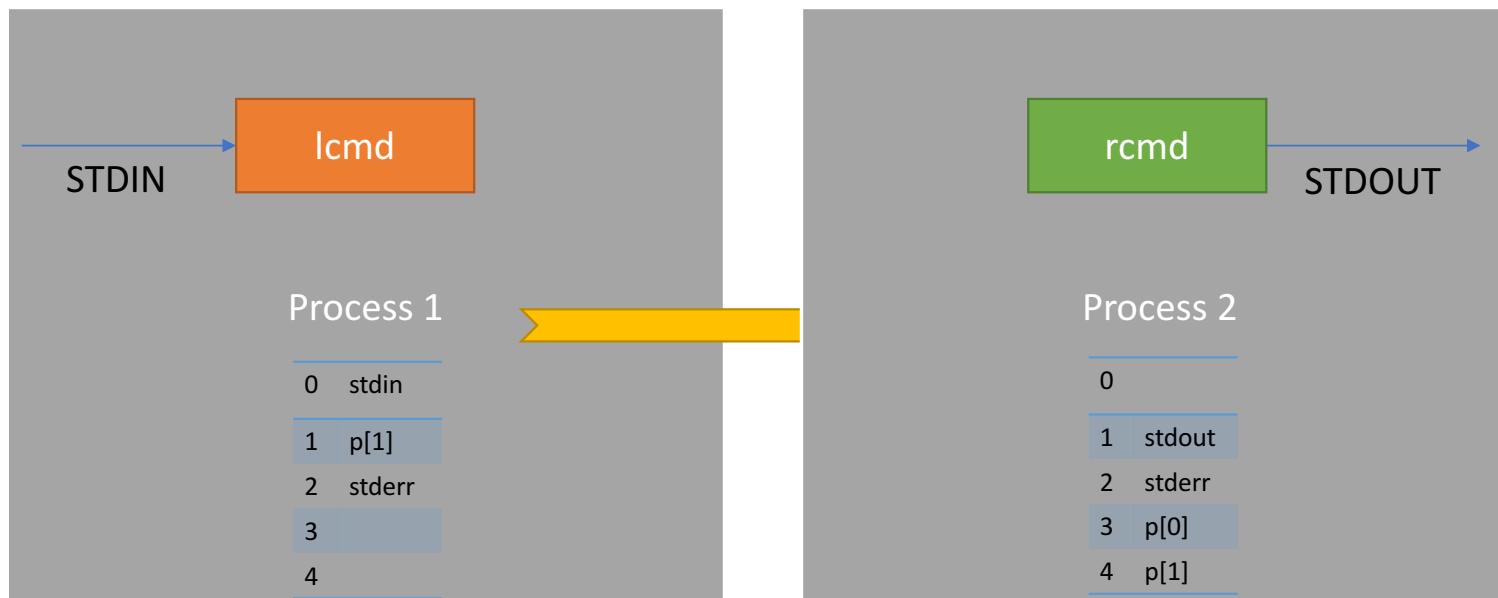
(set for writing) p[1]  p[0] (set for reading)



```
case PIPE:  
    pcmd = (struct pipecmd*)cmd;  
    if(pipe(p) < 0)  
        panic("pipe");  
    if(fork1() == 0){  
        close(1);  
        dup(p[1]);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        runcmd(pcmd->right);  
    }  
  
    break;
```

Homework 1 Q3: pipecmd

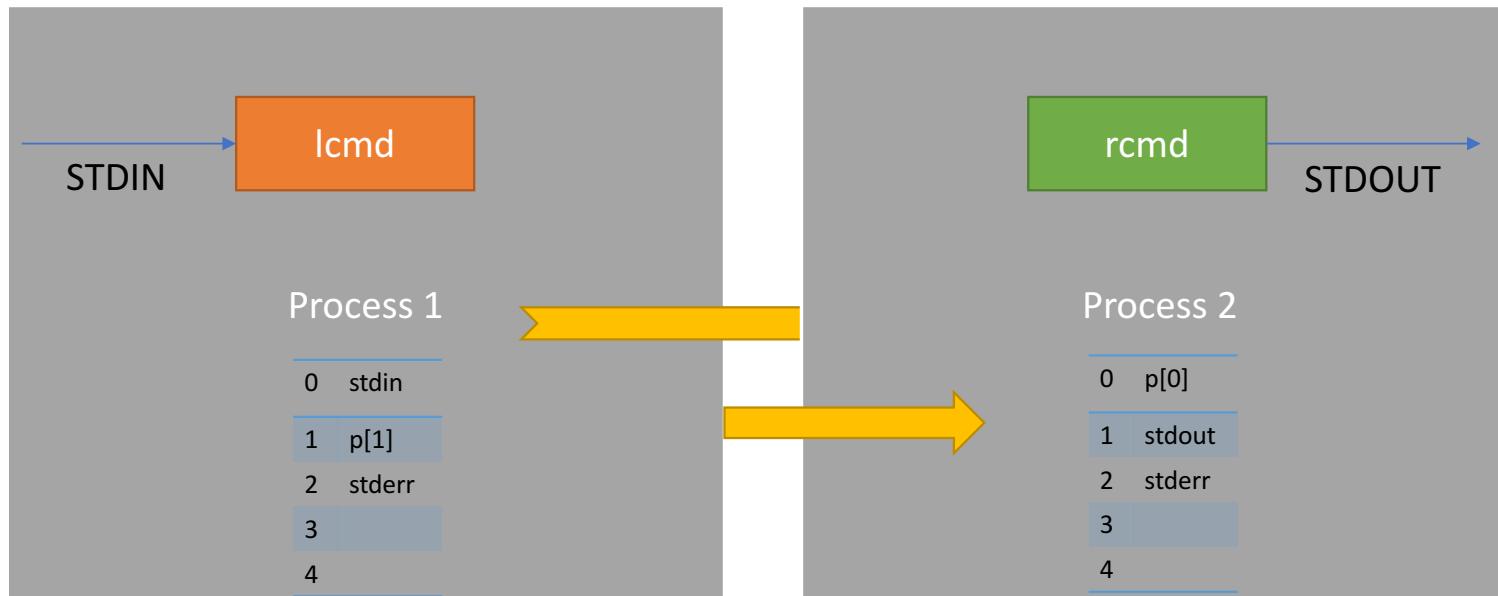
(set for writing) p[1]  p[0] (set for reading)



```
case PIPE:  
    pcmd = (struct pipecmd*)cmd;  
    if(pipe(p) < 0)  
        panic("pipe");  
    if(fork1() == 0){  
        close(1);  
        dup(p[1]);  
        close(p[0]);  
        close(p[1]);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        runcmd(pcmd->right);  
    }  
    break;
```

Homework 1 Q3: pipecmd

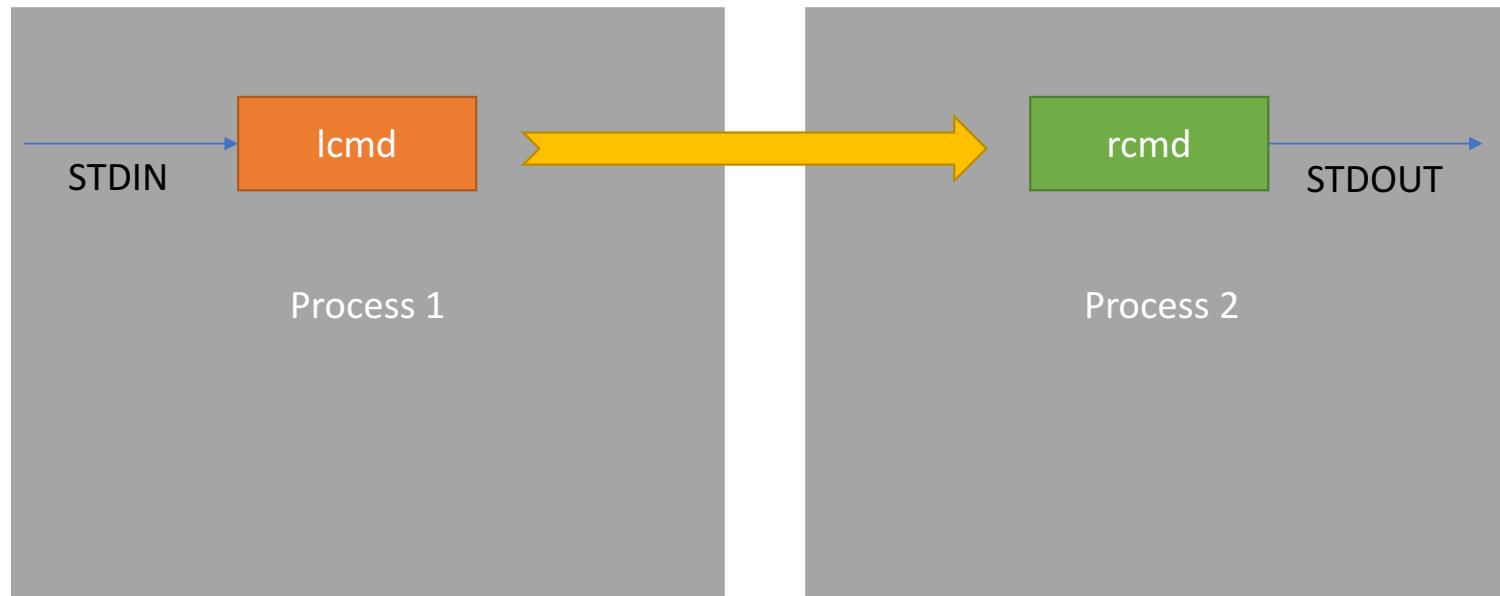
(set for writing) p[1]  p[0] (set for reading)



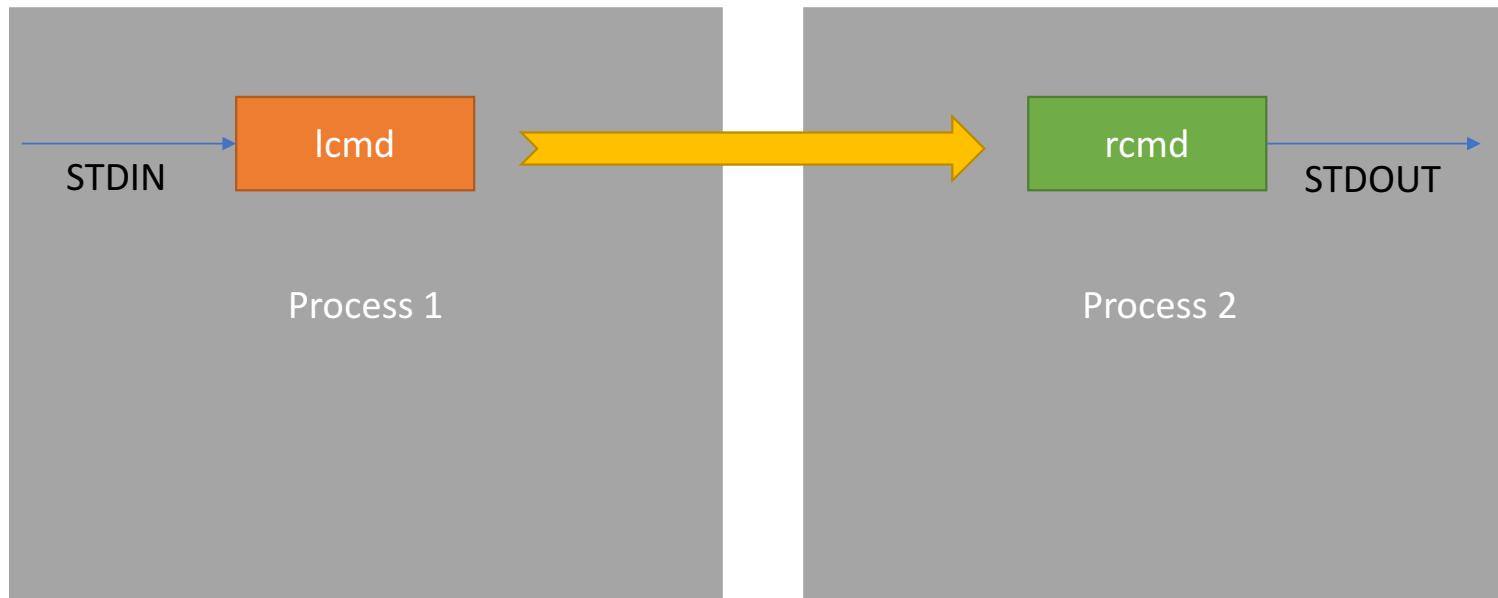
```
case PIPE:  
    pcmd = (struct pipecmd*)cmd;  
    if(pipe(p) < 0)  
        panic("pipe");  
    if(fork1() == 0){  
        close(1);  
        dup(p[1]);  
        close(p[0]);  
        close(p[1]);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        dup(p[0]);  
        close(p[0]);  
        close(p[1]);  
        runcmd(pcmd->right);  
    }  
    break;
```

Homework 1 Q3: pipecmd

(set for writing) p[1]  p[0] (set for reading)



Homework 1 Q3: **pipecmd**



```

case PIPE:
    pcmd = (struct pipecmd*)cmd;
    if(pipe(p) < 0)
        panic("pipe");
    if(fork1() == 0){
        close(1);
        dup(p[1]);
        close(p[0]);
        close(p[1]);
        runcmd(pcmd->left);
    }
    if(fork1() == 0){
        close(0);
        dup(p[0]);
        close(p[0]);
        close(p[1]);
        runcmd(pcmd->right);
    }
    close(p[0]);
    close(p[1]);
    break;
}

```

Homework 1 Q3: pipecmd



```
case PIPE:  
    pcmd = (struct pipecmd*)cmd;  
    if(pipe(p) < 0)  
        panic("pipe");  
    if(fork1() == 0){  
        close(1);  
        dup(p[1]);  
        close(p[0]);  
        close(p[1]);  
        runcmd(pcmd->left);  
    }  
    if(fork1() == 0){  
        close(0);  
        dup(p[0]);  
        close(p[0]);  
        close(p[1]);  
        runcmd(pcmd->right);  
    }  
    close(p[0]);  
    close(p[1]);  
    wait();  
    wait();  
    break;
```

Homework 1 Q3: pipecmd

- Can we do it with one fork ?

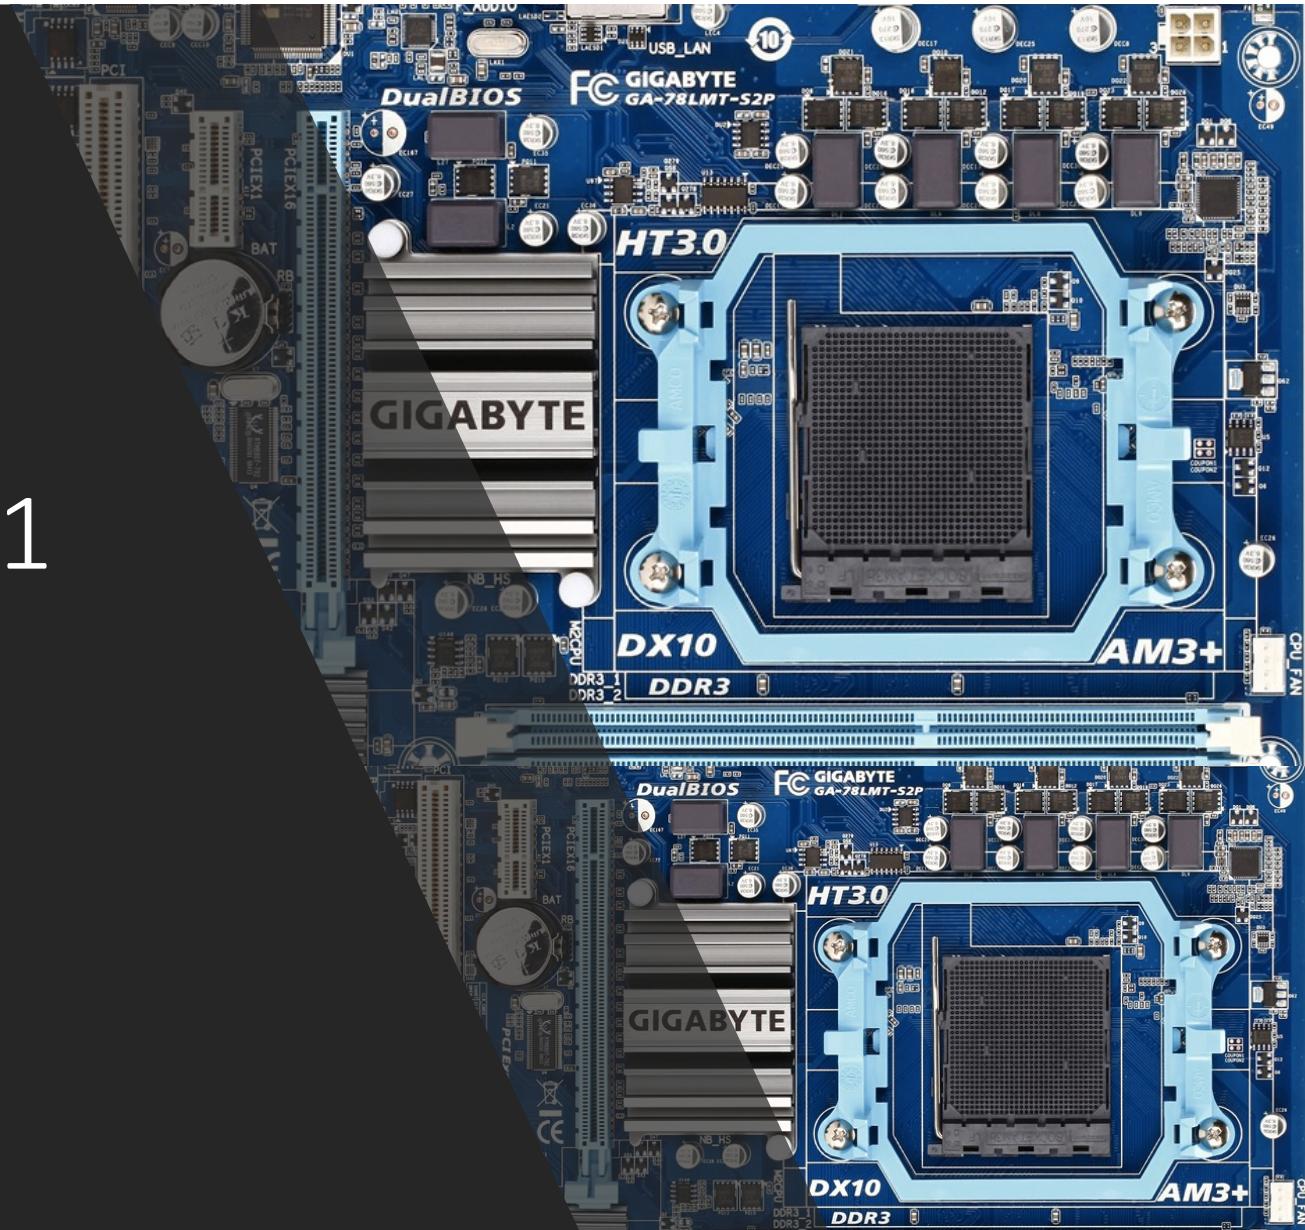
No, that's incorrect, since the left child might fill the pipe before finishing it's input, e.g., if you cat a file that is longer than a pipe. Although it might work for some instances, it won't work for all cases.

So both left and right have to run in parallel

Homework 1 Q3: pipecmd Common Problems

Homework 2: Q1

XV6 - Bootup



Can BIOS load
kernel directly
?

Loading a kernel involves different steps

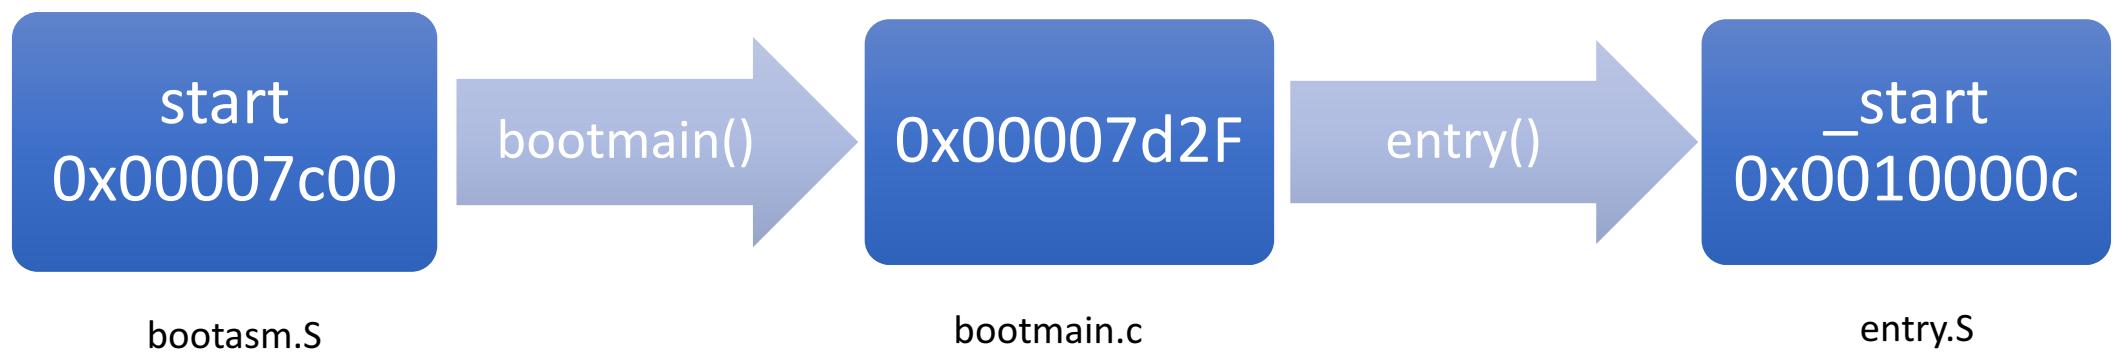
These steps are different for each Operating Systems

The boot loader step also lets you start one of several operating systems off different places on your disk.



0x00007d2F

bootmain.c



Homework 2: Q1

A memory dump table with 10 columns. The first column is labeled "Add. (0x)". A blue arrow labeled "\$eip" points to the value at address 0x0F. The values in the table are:

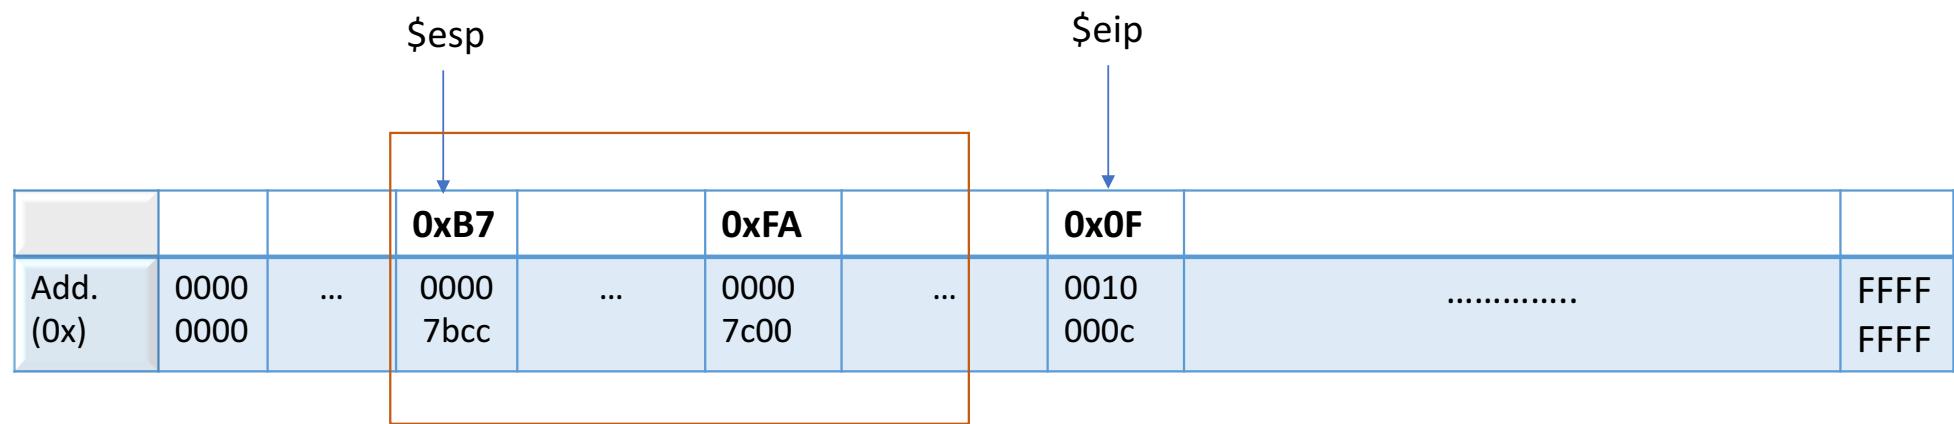
		...	0xFA	...	0x0F	
Add. (0x)	0000 0000	...	0000 7c00	...	0010 000c	FFFF FFFF

Homework 2: Q1

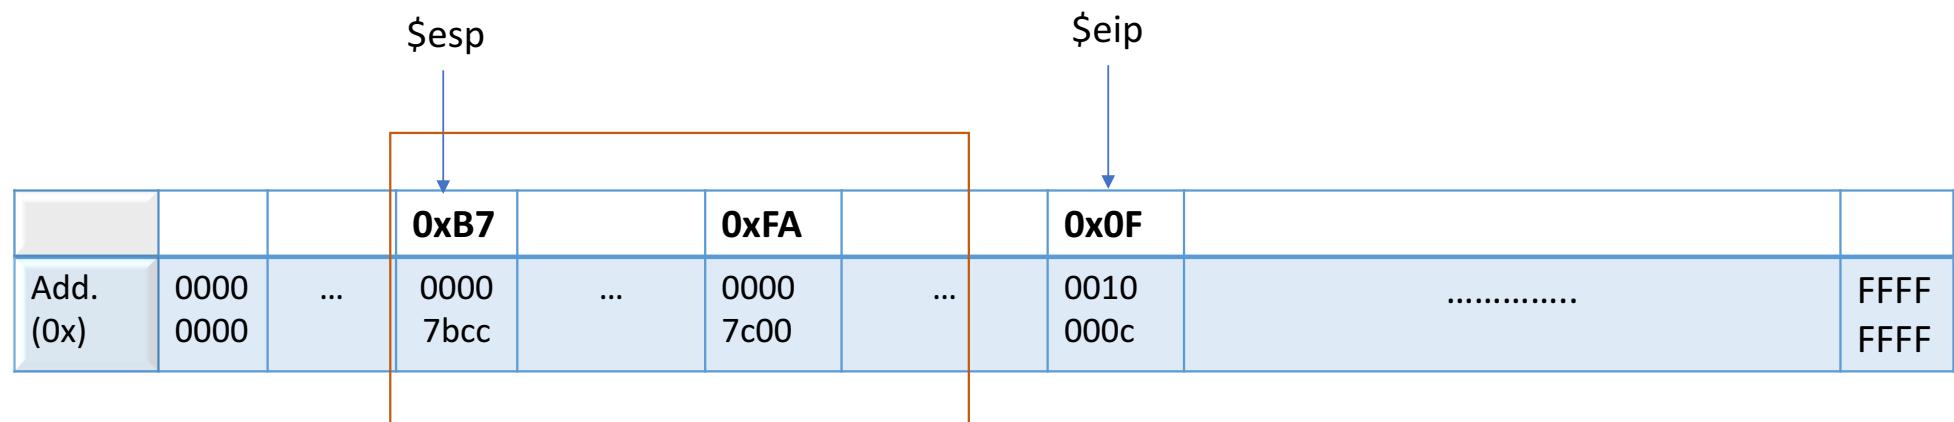
The diagram illustrates a memory dump with two pointers: \$esp and \$eip. The memory is organized into columns representing bytes. The first column is labeled 'Add. (0x)' and contains addresses. The second column contains the byte value 0000. The third column contains three dots (...). The fourth column is labeled 0xB7 and contains the byte value 0000. The fifth column contains three dots (...). The sixth column is labeled 0xFA and contains the byte value 7bcc. The seventh column contains three dots (...). The eighth column is labeled 0x0F and contains the byte value 0000. The ninth column contains three dots (...). The tenth column contains four dots (.....). The eleventh column is labeled FFFF and contains the byte value 000c. The twelfth column is also labeled FFFF.

			0xB7		0xFA		0x0F			
Add. (0x)	0000 0000	...	0000 7bcc	...	0000 7c00	...	0010 000c	FFFF FFFF

Homework 2: Q1



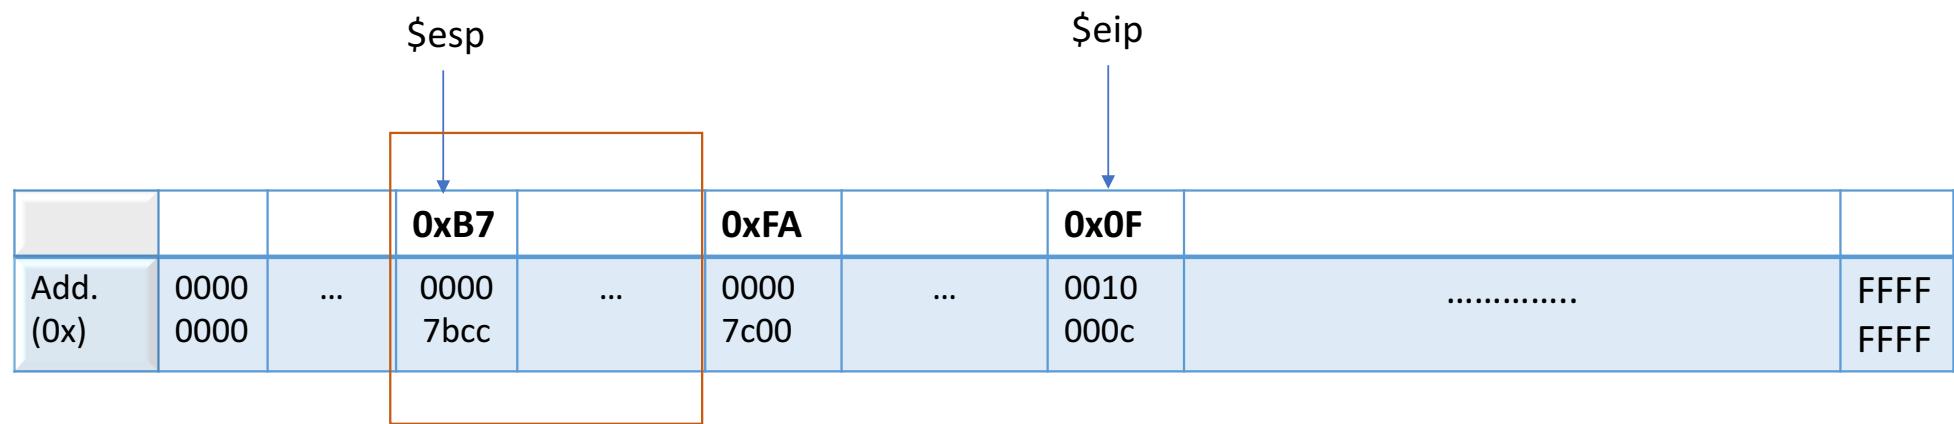
Homework 2: Q1



x/24x

x/24xw

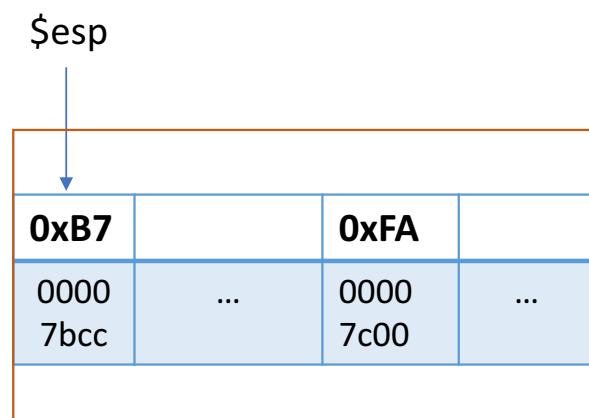
Homework 2: Q1



x/12x

x/12xw

Homework 2: Q1



0x7bcc: 0x00007db7	0x00000000	0x00000000	0x00000000
0x7bdc: 0x00000000	0x00000000	0x00000000	0x00000000
0x7bec: 0x00000000	0x00000000	0x00000000	0x00000000
0x7bfc: 0x00007c4d	0x8ec031fa	0x8ec08ed8	0xa864e4d0
0x7c0c: 0xb0fa7502	0xe464e6d1	0x7502a864	0xe6dfb0fa
0x7c1c: 0x16010f60	0x200f7c78	0xc88366c0	0xc0220f01

Homework 2: Q1

00

00

7D

B7

0xB7		0xFA	
0000 7bcc	...	0000 7c00	...

0x00007bcc	0x00007bcd	0x00007bce	0x00007bcf

0x7bcc: **0x00007db7**
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

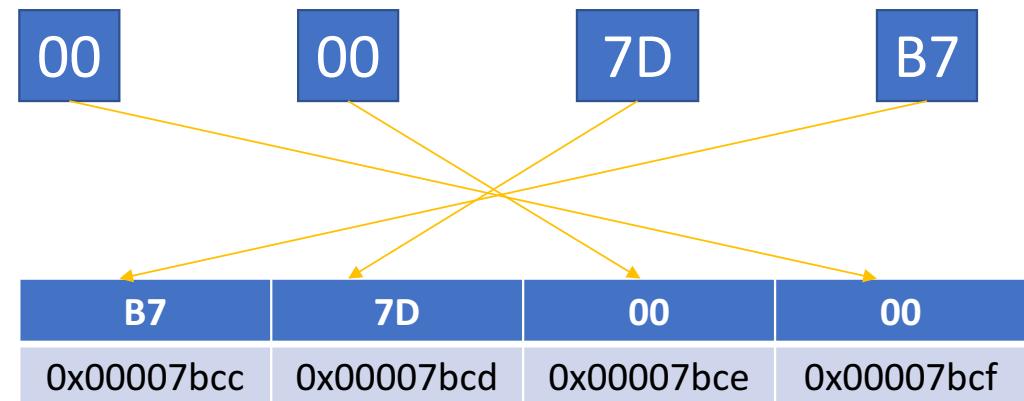
0x00000000
0x00000000
0x00000000
0x8ec031fa
0xe464e6d1
0x200f7c78

0x00000000
0x00000000
0x00000000
0x8ec08ed8
0x7502a864
0xc88366c0

0x00000000
0x00000000
0x00000000
0xa864e4d0
0xe6dfb0fa
0xc0220f01

Homework 2: Q1

0xB7		0xFA	
0000 7bcc	...	0000 7c00	...



0x7bcc: **0x00007db7**
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

0x00000000 0x00000000 0x00000000
0x00000000 0x00000000 0x00000000
0x00000000 0x00000000 0x00000000
0x8ec031fa 0x8ec08ed8 0xa864e4d0
0xe464e6d1 0x7502a864 0xe6dfb0fa
0x200f7c78 0xc88366c0 0xc0220f01

Homework 2: Q1

0xB7		0xFA	
0000 7bcc	...	0000 7c00	...

```
(gdb) x/1xw $esp
0x7bcc: 0x00007db7
(gdb) x/4xb $esp
0x7bcc: 0xb7    0x7d    0x00    0x00
```

0x7bcc: 0x00007db7	0x00000000	0x00000000	0x00000000
0x7bdc: 0x00000000	0x00000000	0x00000000	0x00000000
0x7bec: 0x00000000	0x00000000	0x00000000	0x00000000
0x7bfc: 0x00007c4d	0x8ec031fa	0x8ec08ed8	0xa864e4d0
0x7c0c: 0xb0fa7502	0xe464e6d1	0x7502a864	0xe6dfb0fa
0x7c1c: 0x16010f60	0x200f7c78	0xc88366c0	0xc0220f01

0x7bcc:	0x00007db7	0x00000000	0x00000000	0x00000000
0x7bdc:	0x00000000	0x00000000	0x00000000	0x00000000
0x7bec:	0x00000000	0000	0x00000000	0x00000000
0x7bfc:	0x00007c4d	7c00	0x8ec031fa	0x8ec08ed8
0x7c0c:	0xb0fa7502		0xe464e6d1	0xa864e4d0
0x7c1c:	0x16010f60		0x200f7c78	0xe6dfb0fa
				0xc0220f01

Instructions opcodes

```
.globl start
start:
    cli                                # BIOS enabled interrupts; disable
    7c00: fa

    # Zero data segment registers DS, ES, and SS.
    xorw    %ax,%ax                      # Set %ax to zero
    7c01: 31 c0                         xor    %eax,%eax
    movw    %ax,%ds                      # -> Data Segment
    7c03: 8e d8                         mov    %eax,%ds
    movw    %ax,%es                      # -> Extra Segment
    7c05: 8e c0                         mov    %eax,%es
    movw    %ax,%ss                      # -> Stack Segment
    7c07: 8e d0                         mov    %eax,%ss
```

0x7bcc:	0x00007db7	0x00000000	0x00000000	0x00000000
0x7bdc:	0x00000000	0x00000000	0x00000000	0x00000000
0x7bec:	0x00000000	0000 0x00000000	0x00000000	0x00000000
0x7bfc:	0x00007c4d	7c00 0x8ec031fa	0x8ec08ed8	0xa864e4d0
0x7c0c:	0xb0fa7502	0xe464e6d1	0x7502a864	0xe6dfb0fa
0x7c1c:	0x16010f60	0x200f7c78	0xc88366c0	0xc0220f01



0x7bcc: 0x00007db7
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

0x00000000
0x00000000
0x00000000
0x8ec031fa
0xe464e6d1
0x200f7c78

0x00000000
0x00000000
0x00000000
0x8ec08ed8
0x7502a864
0xc88366c0

0x00000000
0x00000000
0x00000000
0xa864e4d0
0xe6dfb0fa
0xc0220f01

start
0x00007c00

bootmain()

0x00007d2F

entry()

_start
0x0010000c

bootasm.S

bootmain.c

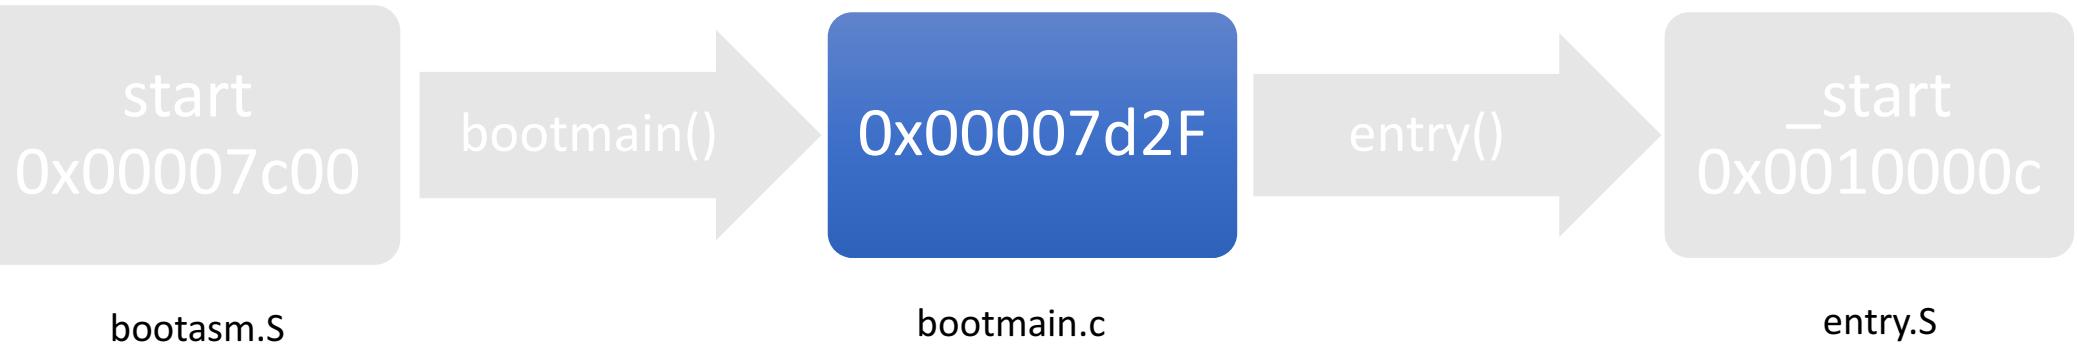
entry.S

Address of the next instruction to execute after return from call

0x7bcc: 0x00007db7
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

```
call    bootmain
7c48: e8 e2 00 00 00          call    7d2f <bootmain>

# If bootmain returns (it shouldn't), trigger a Bochs
# breakpoint if running under Bochs, then loop.
movw   $0x8a00, %ax           # 0x8a00 -> port 0x8a00
7c4d: 66 b8 00 8a            mov    $0x8a00,%ax
movl   0xe464e6d1             0x7502a864             0xe6a1b07a
                           0x200f7c78             0xc88366c0
                                         0xc0220f01
```



1. Push the value of EBP (0) onto the stack, and then copy the value of ESP into EBP

0x7bcc:	0x00007db7	0x00000000	0x00000000	0x00000000
0x7bdc:	0x00000000	0x00000000	0x00000000	0x00000000
0x7bec:	0x00000000	0x00000000	0x00000000	0x00000000
0x7bfc:	0x00007c4d	0x8ec031fa	0x8ec08ed8	0xa864e4d0
0x7c0c:	0xb0fa7502	0xe464e6d1	0x7502a864	0xe6dfb0fa
0x7c1c:	0x16010f60	0x200f7c78	0xc88366c0	0xc0220f01

start
0x00007c00

bootmain()

0x00007d2F

entry()

_start
0x0010000c

bootasm.S

bootmain.c

entry.S

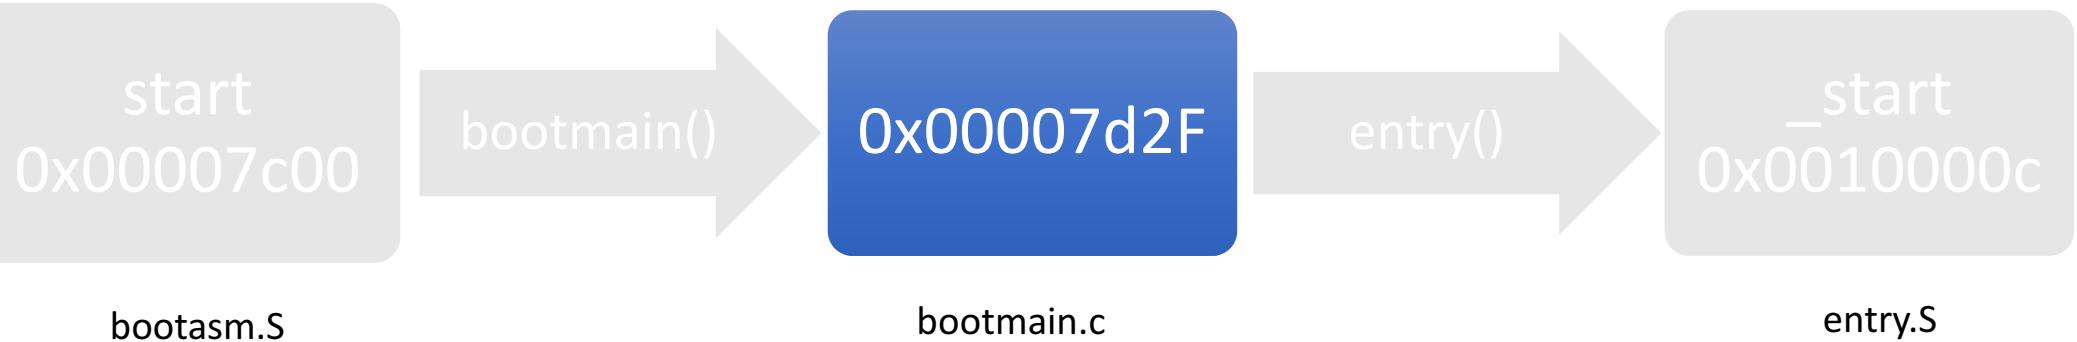
1. Push the value of EBP (0) onto the stack, and then copy the value of ESP into EBP
2. Saves three registers: %edi, %esi, and %ebx

0x7bcc:	0x00007db7
0x7bdc:	0x00000000
0x7bec:	0x00000000
0x7bfc:	0x00007c4d
0x7c0c:	0xb0fa7502
0x7c1c:	0x16010f60

0x00000000
0x00000000
0x00000000
0x8ec031fa
0xe464e6d1
0x200f7c78

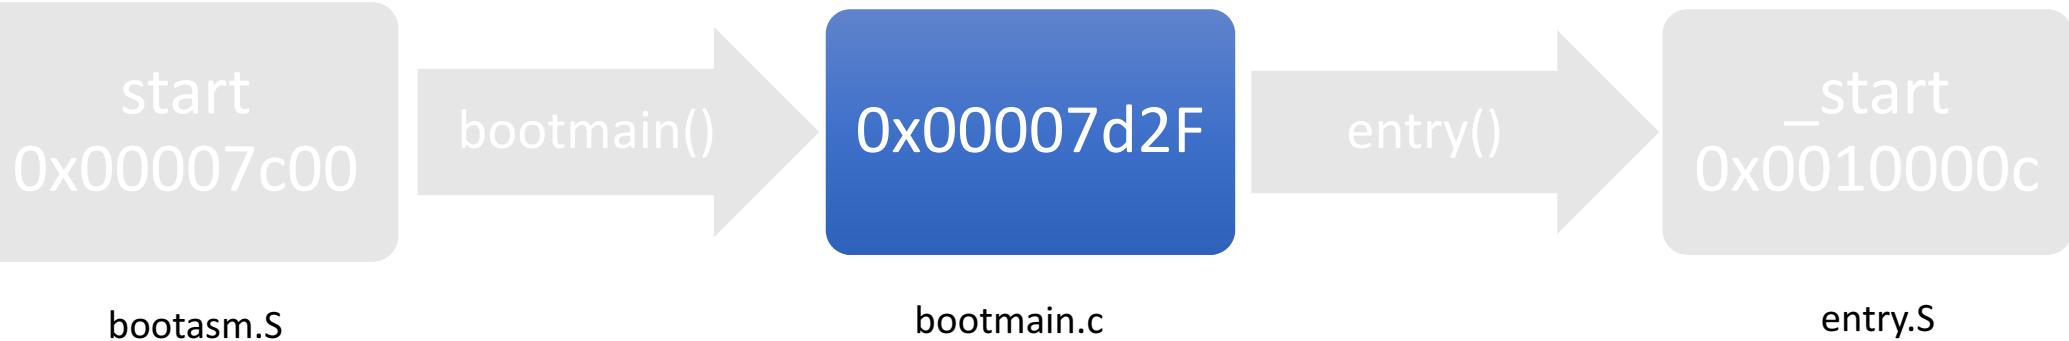
0x00000000
0x00000000
0x00000000
0x8ec08ed8
0x7502a864
0xc88366c0

0x00000000
0x00000000
0x00000000
0xa864e4d0
0xe6dfb0fa
0xc0220f01



sub \$0x1c,%esp

<code>0x7bcc:</code>	<code>0x00007db7</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bdc:</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bec:</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bfc:</code>	<code>0x00007c4d</code>	<code>0x8ec031fa</code>	<code>0x8ec08ed8</code>	<code>0xa864e4d0</code>
<code>0x7c0c:</code>	<code>0xb0fa7502</code>	<code>0xe464e6d1</code>	<code>0x7502a864</code>	<code>0xe6dfb0fa</code>
<code>0x7c1c:</code>	<code>0x16010f60</code>	<code>0x200f7c78</code>	<code>0xc88366c0</code>	<code>0xc0220f01</code>



- Why 28bytes and not 20bytes ?
- It's an optimization. The compiler allocates the space for both local variables and for the arguments to the `readseg((uchar*)elf, 4096, 0)` function call that is about to be called

<code>0x7bcc:</code>	<code>0x00007db7</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bdc:</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bec:</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>	<code>0x00000000</code>
<code>0x7bfc:</code>	<code>0x00007c4d</code>	<code>0x8ec031fa</code>	<code>0x8ec08ed8</code>	<code>0xa864e4d0</code>
<code>0x7c0c:</code>	<code>0xb0fa7502</code>	<code>0xe464e6d1</code>	<code>0x7502a864</code>	<code>0xe6dfb0fa</code>
<code>0x7c1c:</code>	<code>0x16010f60</code>	<code>0x200f7c78</code>	<code>0xc88366c0</code>	<code>0xc0220f01</code>

start
0x00007c00

bootmain()

0x00007d2F

entry()

_start
0x0010000c

bootasm.S

bootmain.c

entry.S

0x7bcc: 0x00007db7
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

```
// Call the entry point from the ELF header.  
// Does not return!  
entry = (void(*)(void))(elf->entry);  
entry();  
| 7db1: ff 15 18 00 01 00      call   *0x10018  
}|  
| 7db7: 83 c4 1c              add    $0x1c,%esp  
| . . .
```

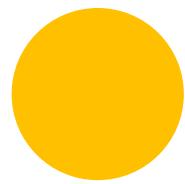
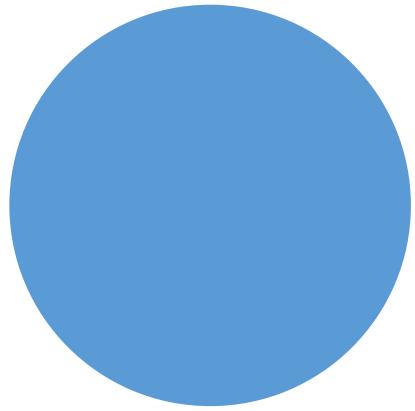


0x7bcc: 0x00007db7
0x7bdc: 0x00000000
0x7bec: 0x00000000
0x7bfc: 0x00007c4d
0x7c0c: 0xb0fa7502
0x7c1c: 0x16010f60

0x00000000
0x00000000
0x00000000
0x8ec031fa
0xe464e6d1
0x200f7c78

0x00000000
0x00000000
0x00000000
0x8ec08ed8
0x7502a864
0xc88366c0

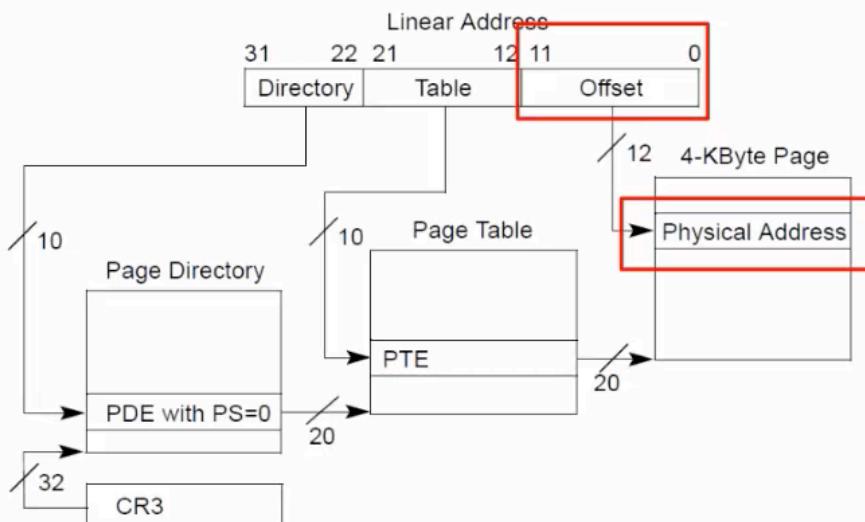
0x00000000
0x00000000
0x00000000
0xa864e4d0
0xe6dfb0fa
0xc0220f01



Homework 2 |
Question 2 |

Lecture 5&6 : Address Translation

Page translation



- Virtual address being divided into three parts and the widths of each part
- Page directory indexing (with flags)
- Page table indexing
- Final physical address composition

References

- <http://web.cs.ucla.edu/classes/fall08/cs111/scribe/4/index.html>
- <https://linux.die.net/man/3/open>
- https://www.gnu.org/software/libc/manual/html_node/Permission-Bits.html