

COMPSCI 143A: Principles of Operating Systems

Lecture 1: Introduction

Anton Burtsev
September, 2018

Class details

- Undergraduate
 - 240 students
- Instructor: Anton Burtsev
- Meeting time: 2:00pm-3:20pm (Tue/Thu)
 - Discussions: 5:00pm-5:50pm (Mon)
 - Regular discussion sections
 - Feel free to stop by my office with questions (DBH 3066)
- 4 Tas
- Web page
 - <http://www.ics.uci.edu/~aburtsev/143A>

More details

- 4-5 homeworks
 - Implement a shell
 - Explain what's on the stack
 - Implement a system call
 - Change file system layout
- Midterm
- Final
- Grades are curved
 - Homework: 60%, midterm exam: 15%, final exam: 25% of your grade.
 - You can submit late homework 3 days after the deadline for 60% of your grade

This course

- Inspired by
 - MIT 6.828: Operating System Engineering
<https://pdos.csail.mit.edu/6.828/2016/>
 - Adapted for undergraduate students
- We will use xv6
 - Relatively simple OS kernel (only 9K lines of code)
 - Reasonably complete UNIX kernel
 - <https://pdos.csail.mit.edu/6.828/2016/xv6.html>
- xv6 comes with a book
 - <https://pdos.csail.mit.edu/6.828/2016/xv6/book-rev9.pdf>
- And source code printout
 - <https://pdos.csail.mit.edu/6.828/2016/xv6/xv6-rev9.pdf>

Another Book

“Operating Systems: Three Easy Pieces”
(OSTEP) Remzi H. Arpaci-Dusseau and Andrea
C. Arpaci-Dusseau

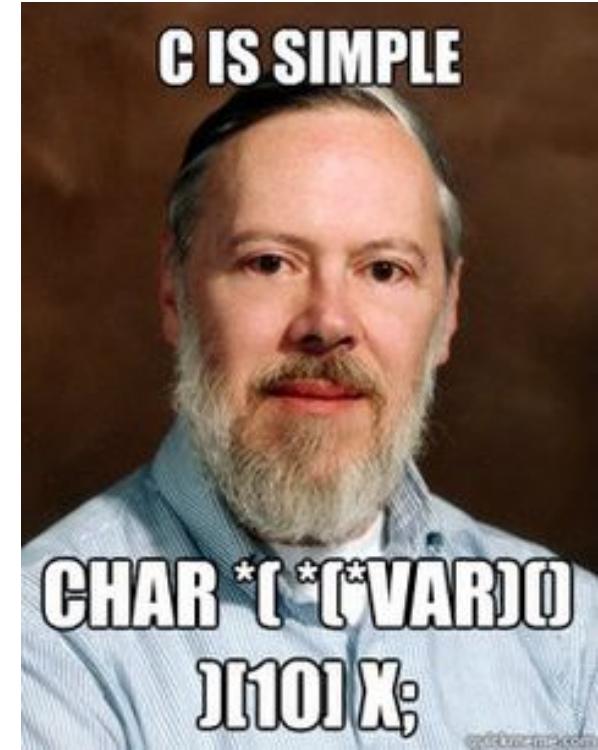
- Free online version
<http://pages.cs.wisc.edu/~remzi/OSTEP/>

Course organization

- Lectures
 - High level concepts and abstractions
- Reading
 - Xv6 book + source code
 - Bits of OSTEP book
- Homeworks
 - Coding real parts of the xv6 kernel
- Design riddles
 - Understanding design tradeoffs, explaining parts of xv6

Prerequisites

- Solid C coding skills
 - Xv6 is written in C
 - You need to read, code and debug
 - All homeworks are in C
 - Many questions will require explaining xv6 code
- Be able to work and code in Linux/UNIX
- Some assembly skills



How to succeed?

- Read the source

What is an operating system?

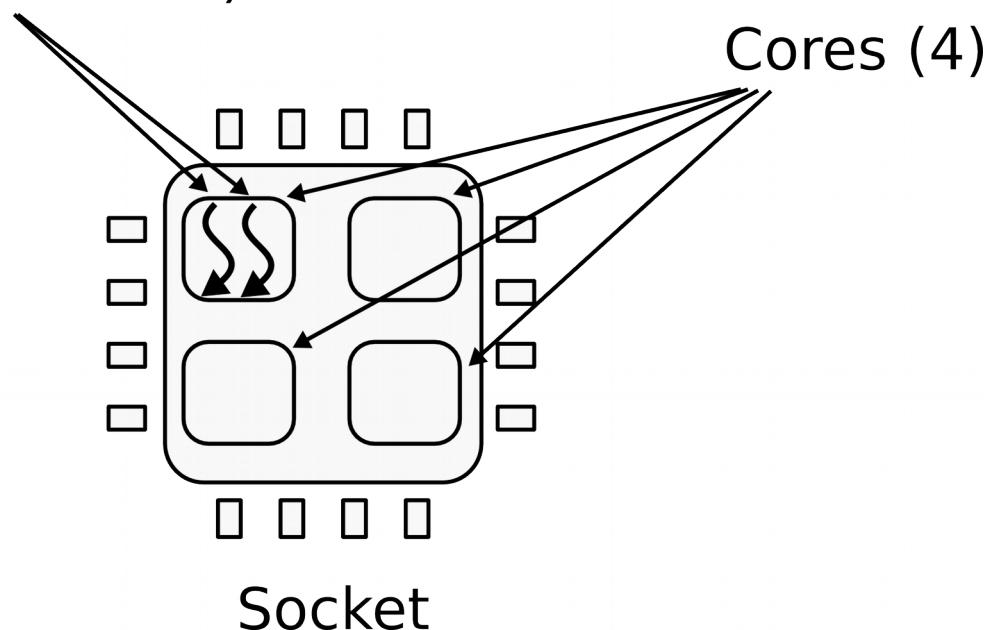
PC Hardware

CPU

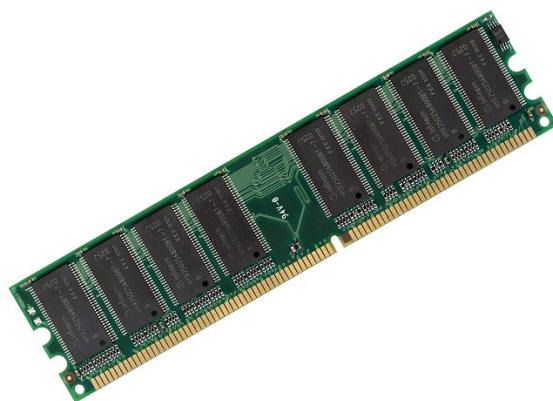
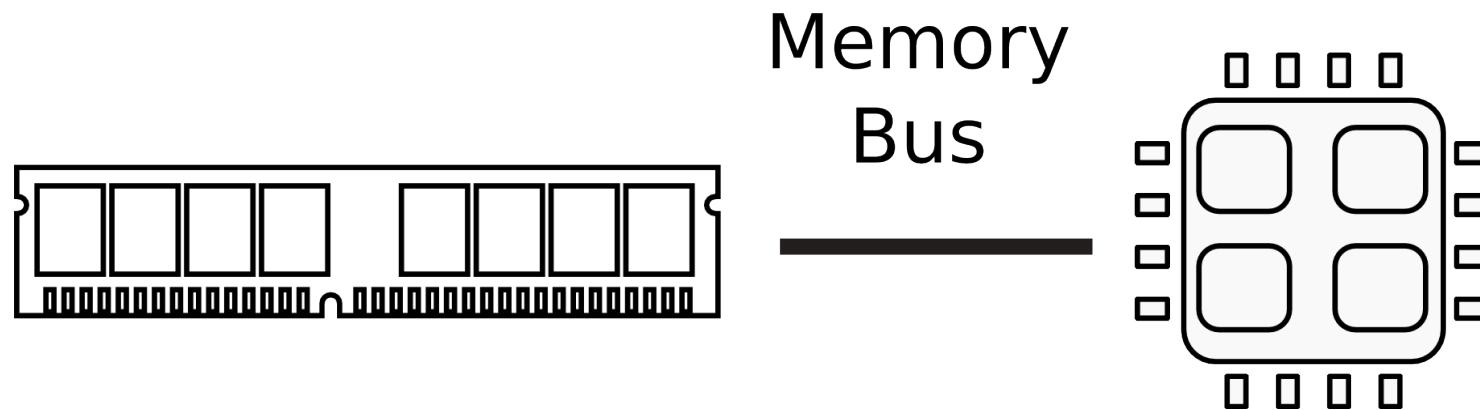
- 1 CPU socket
 - 4 cores
 - 2 logical (HT) threads each



Hyper-Threading
(logical threads)



Memory



Memory abstraction

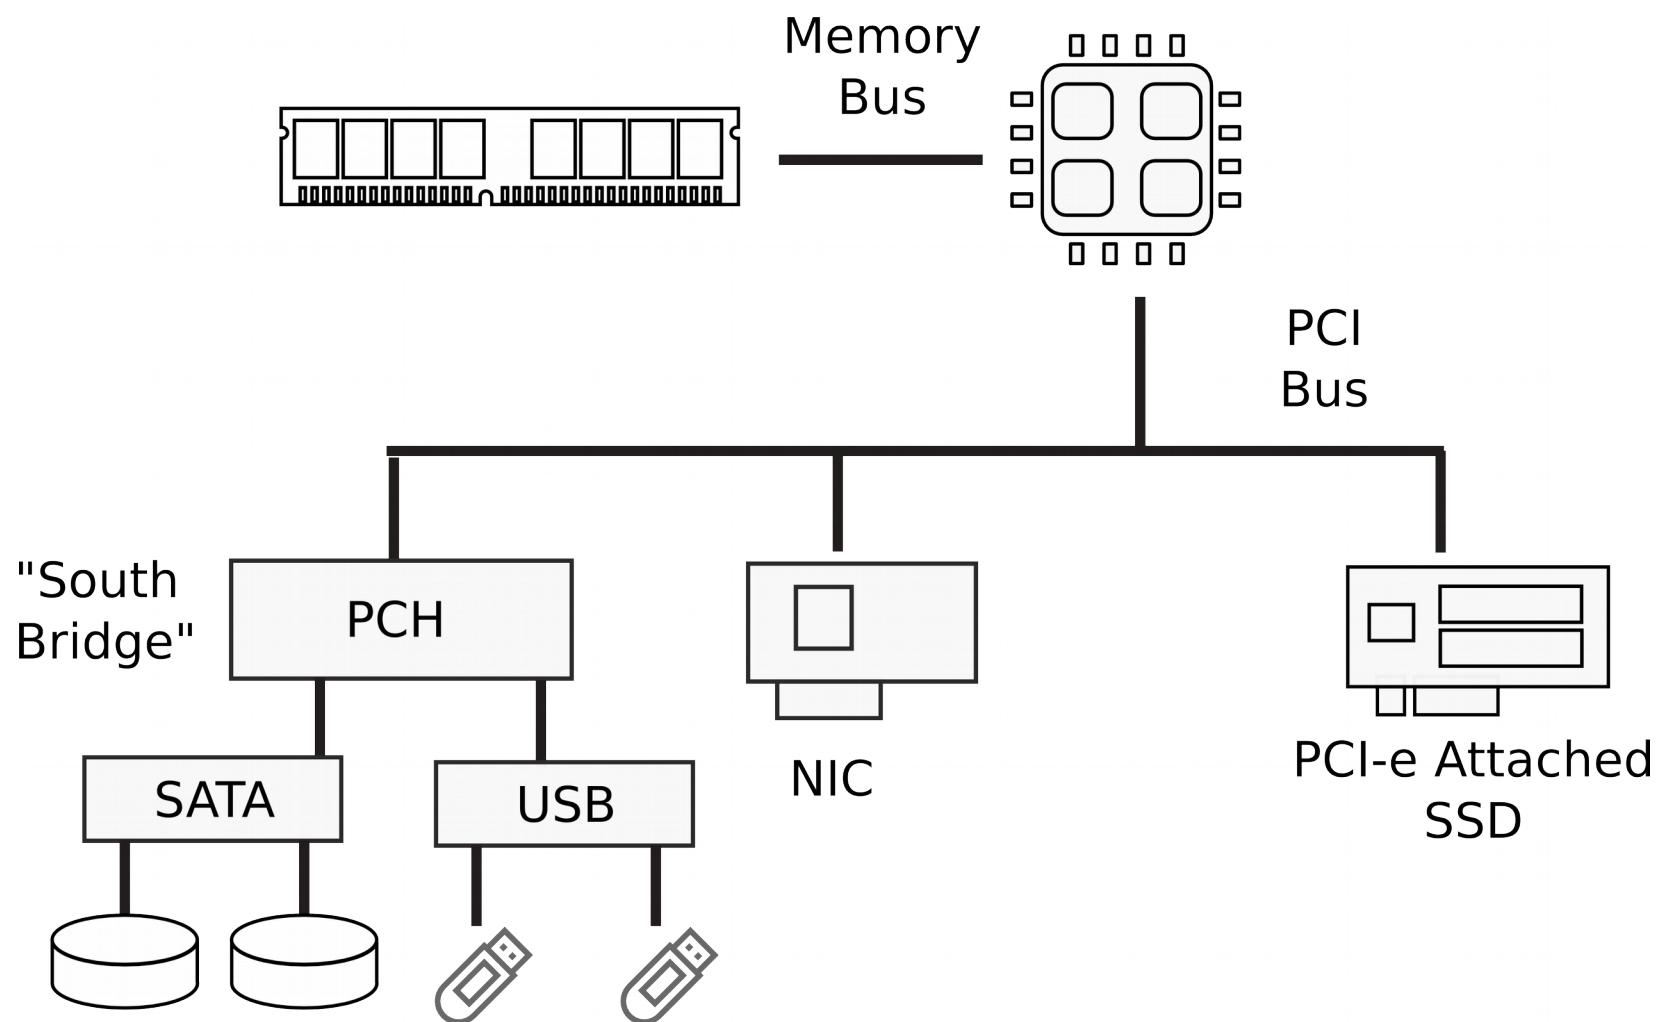
$\text{WRITE}(addr, value) \rightarrow \emptyset$

Store *value* in the storage cell identified by *addr*.

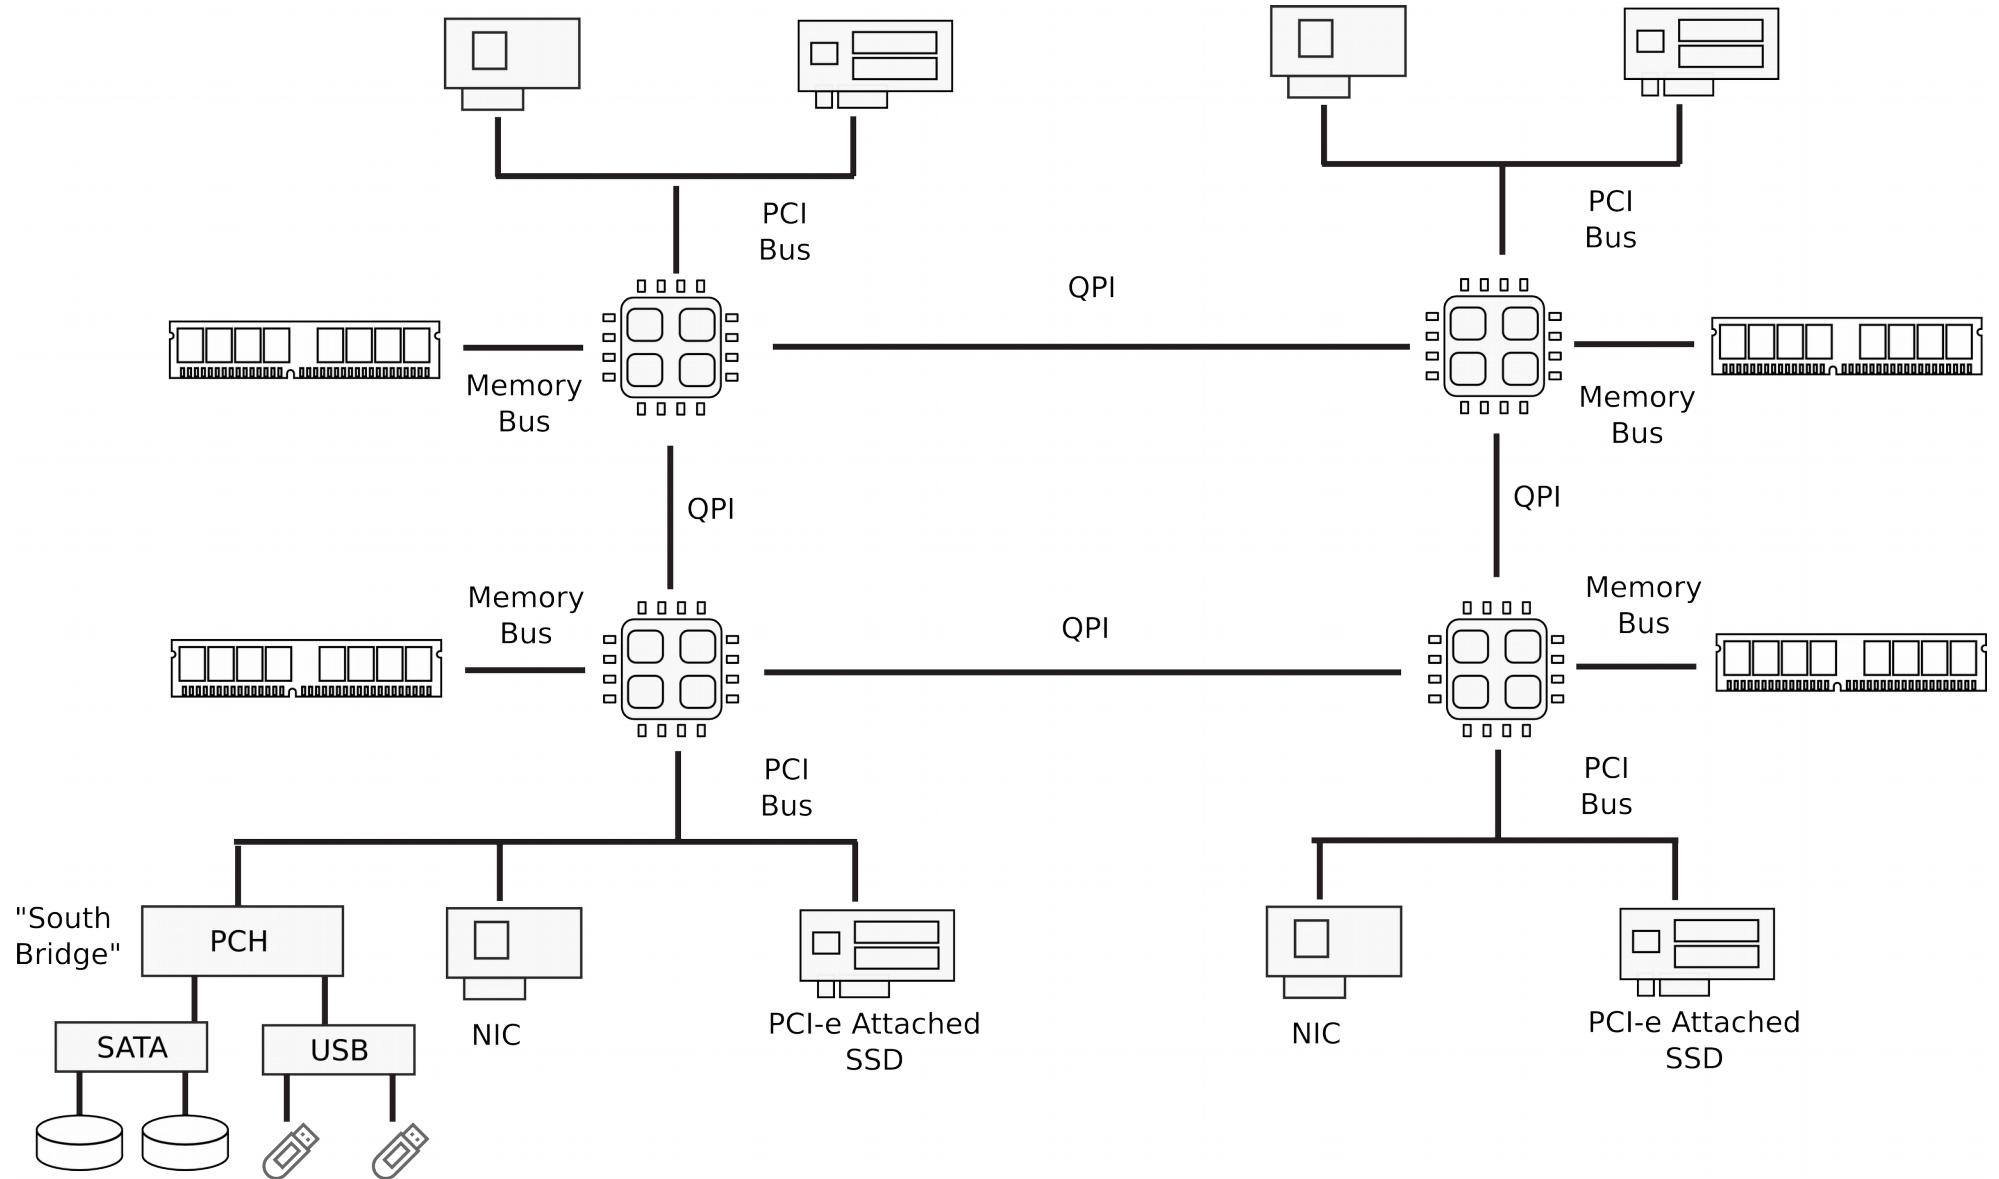
$\text{READ}(addr) \rightarrow value$

Return the *value* argument to the most recent WRITE call referencing *addr*.

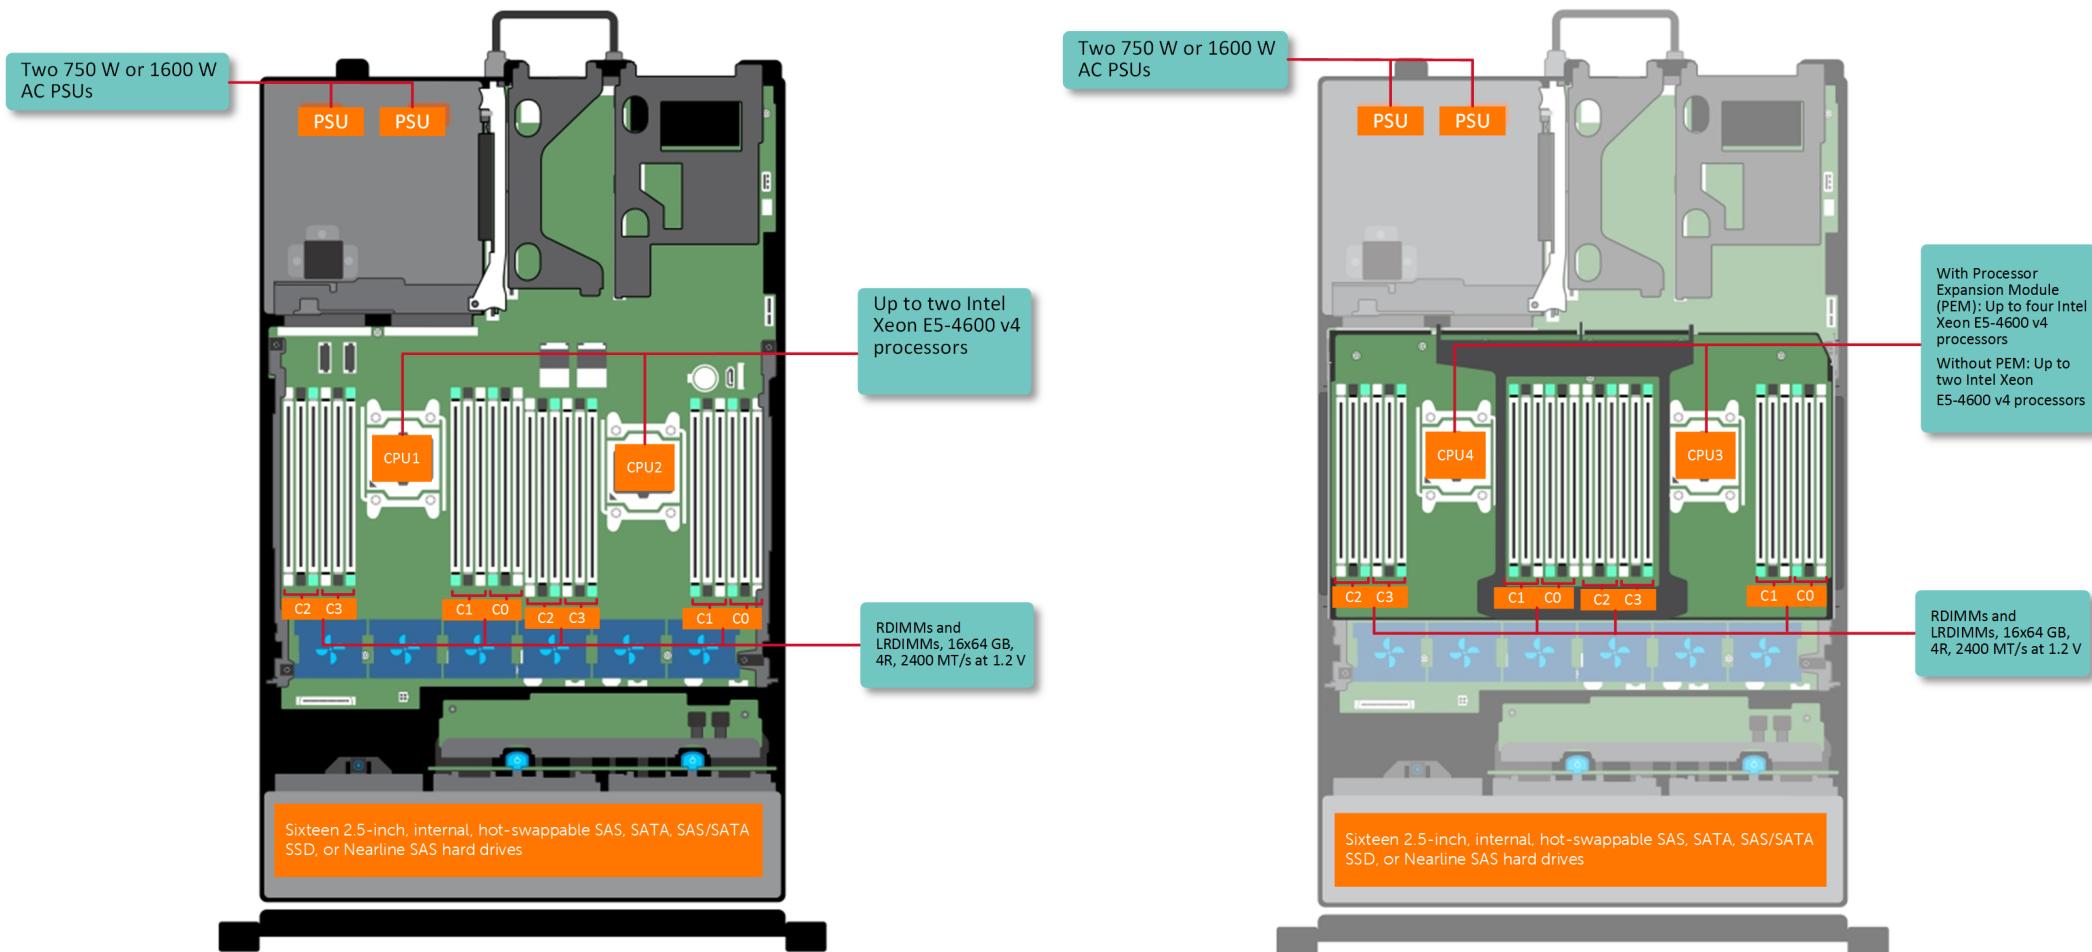
I/O Devices



Multi-socket machines



Dell R830 4-socket server



Dell Poweredge R830 System Server with 2 sockets on the main floor and 2 sockets on the expansion



http://www.dell.com/support/manuals/us/en/19/poweredge-r830/r830_om/supported-configurations-for-the-poweredge-r830-system?guid=guid-01303b2b-f884-4435-b4e2-57bec2ce225a&lang=en-us

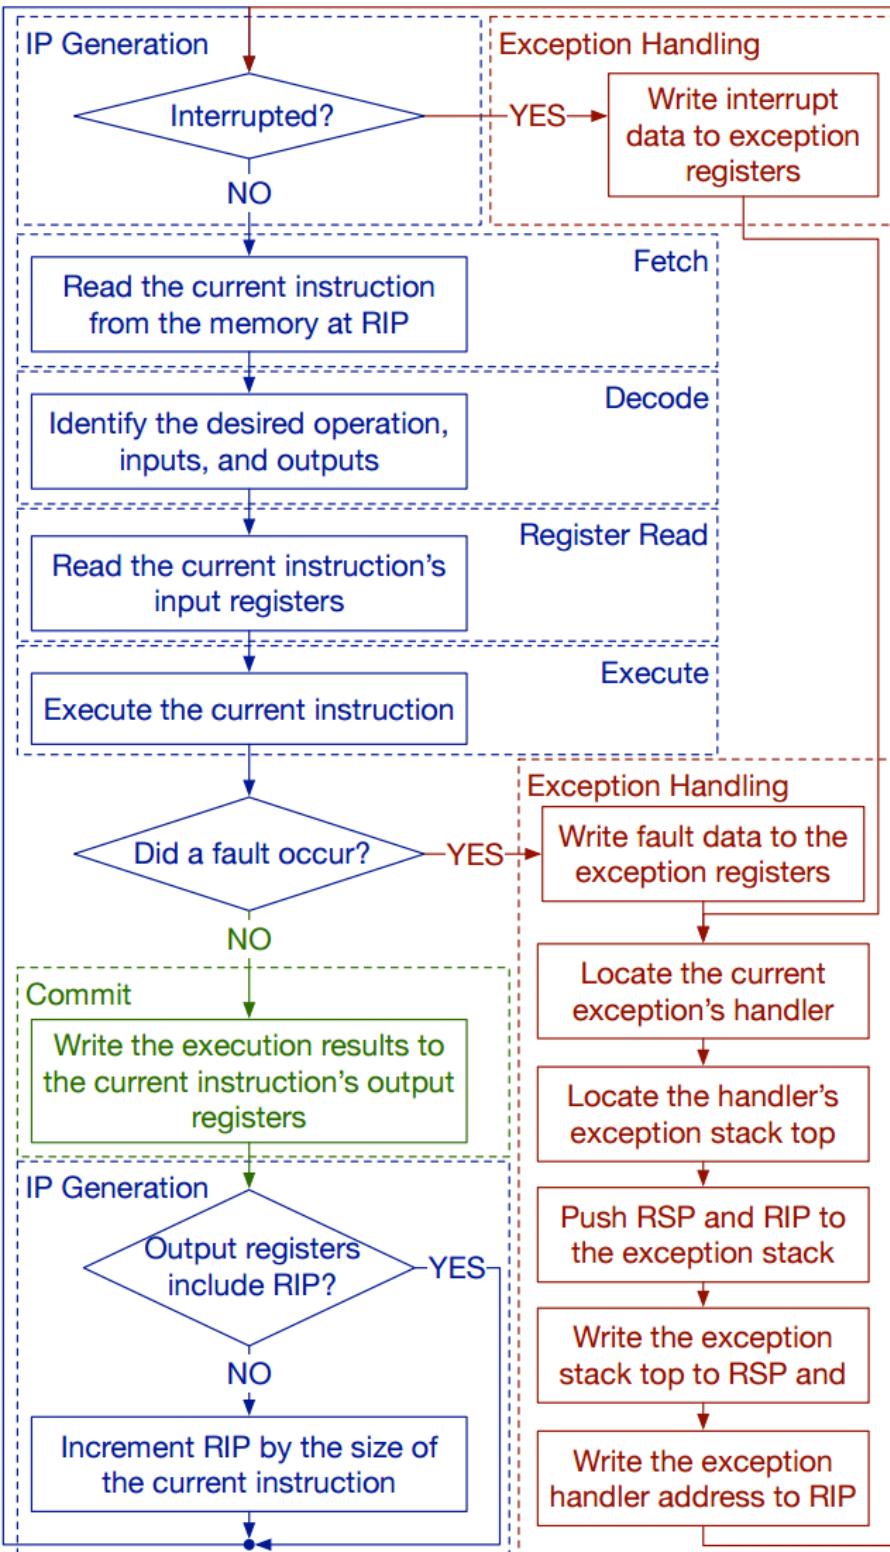
What does CPU do internally?

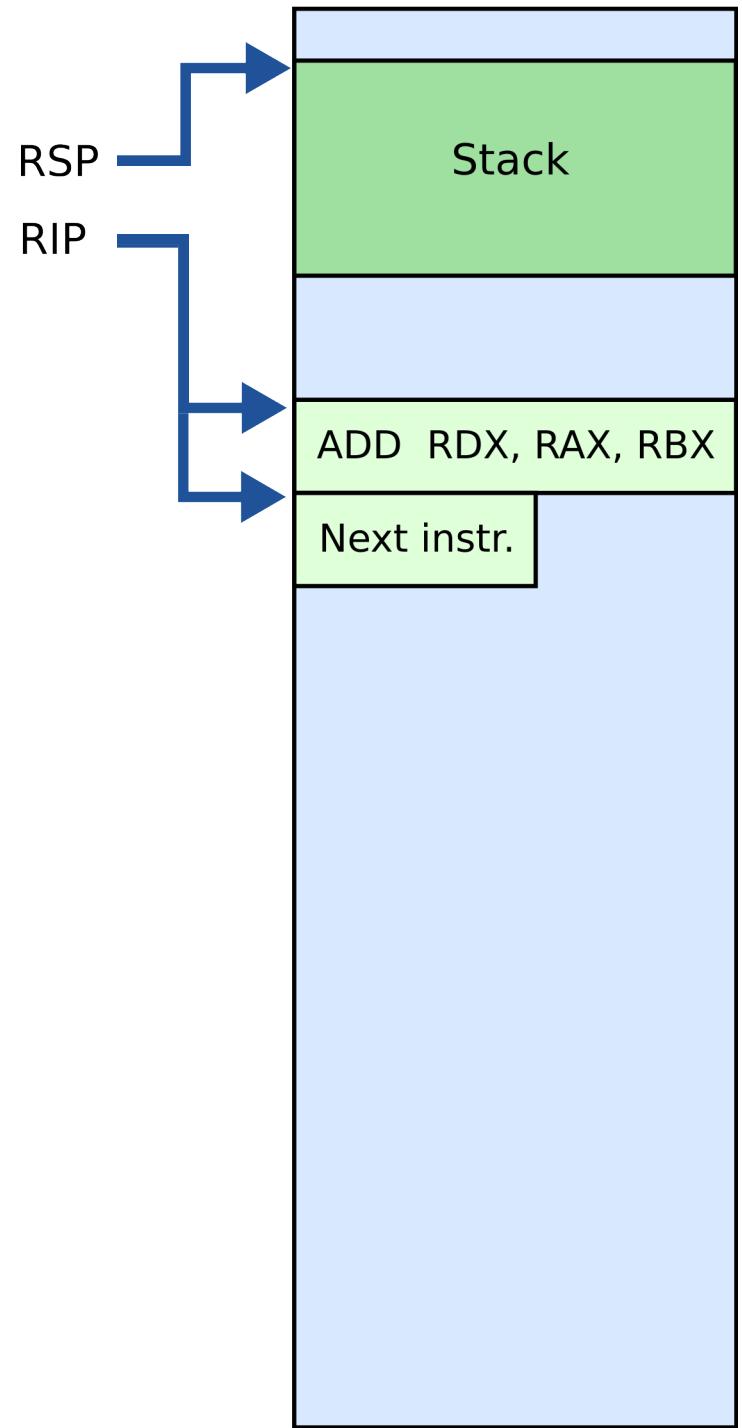
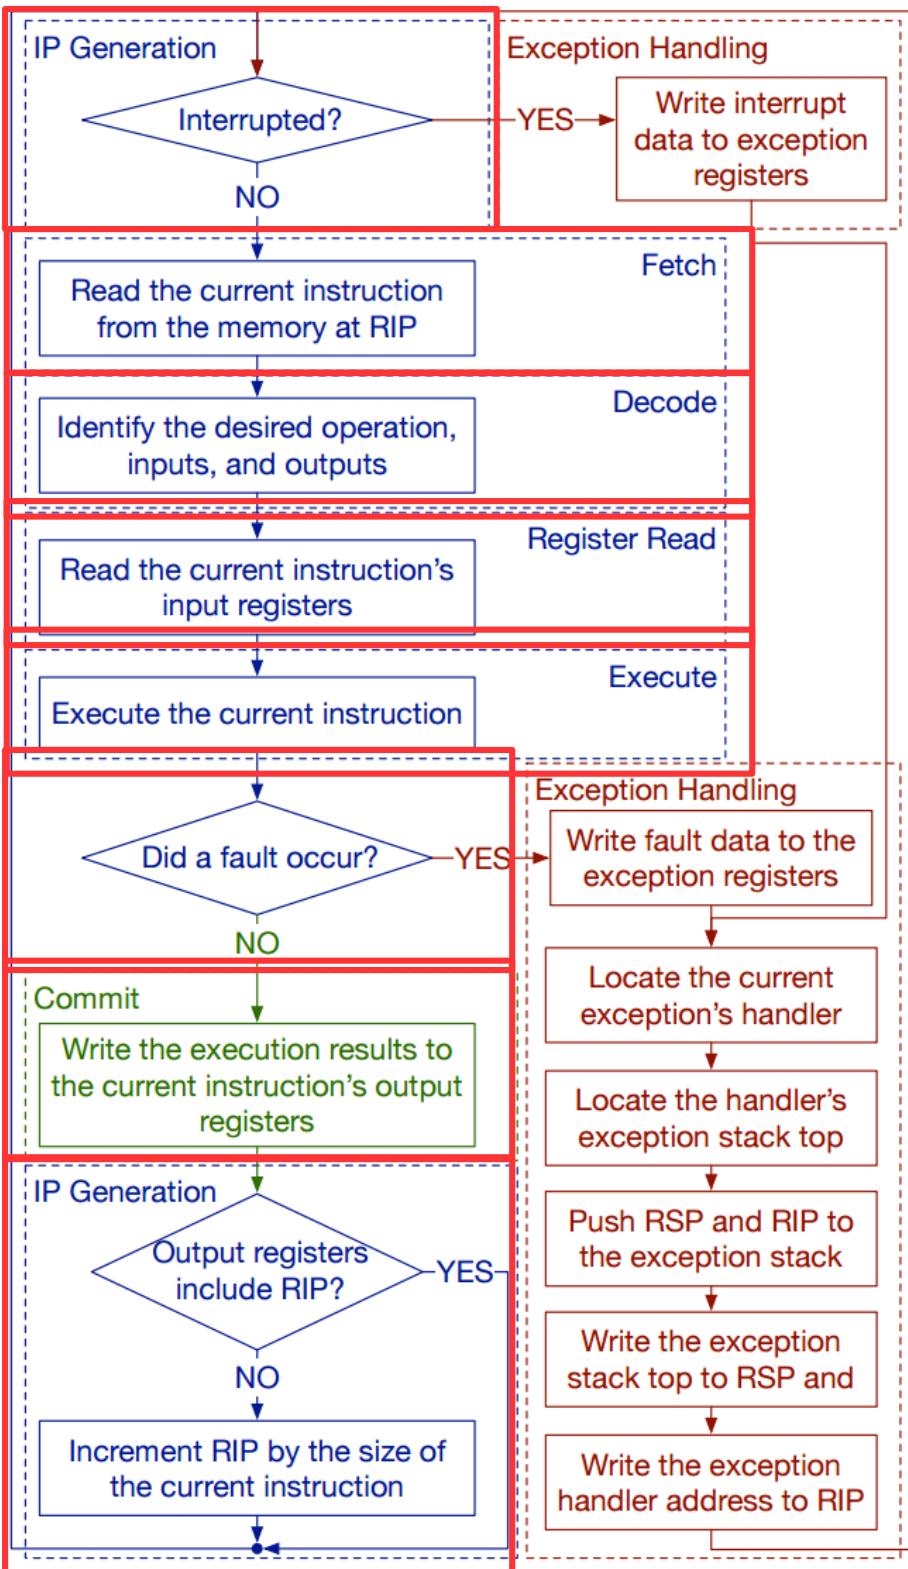
CPU execution loop

- CPU repeatedly reads instructions from memory
- Executes them
- Example

ADD EDX, EAX, EBX

// EDX = EAX + EBX

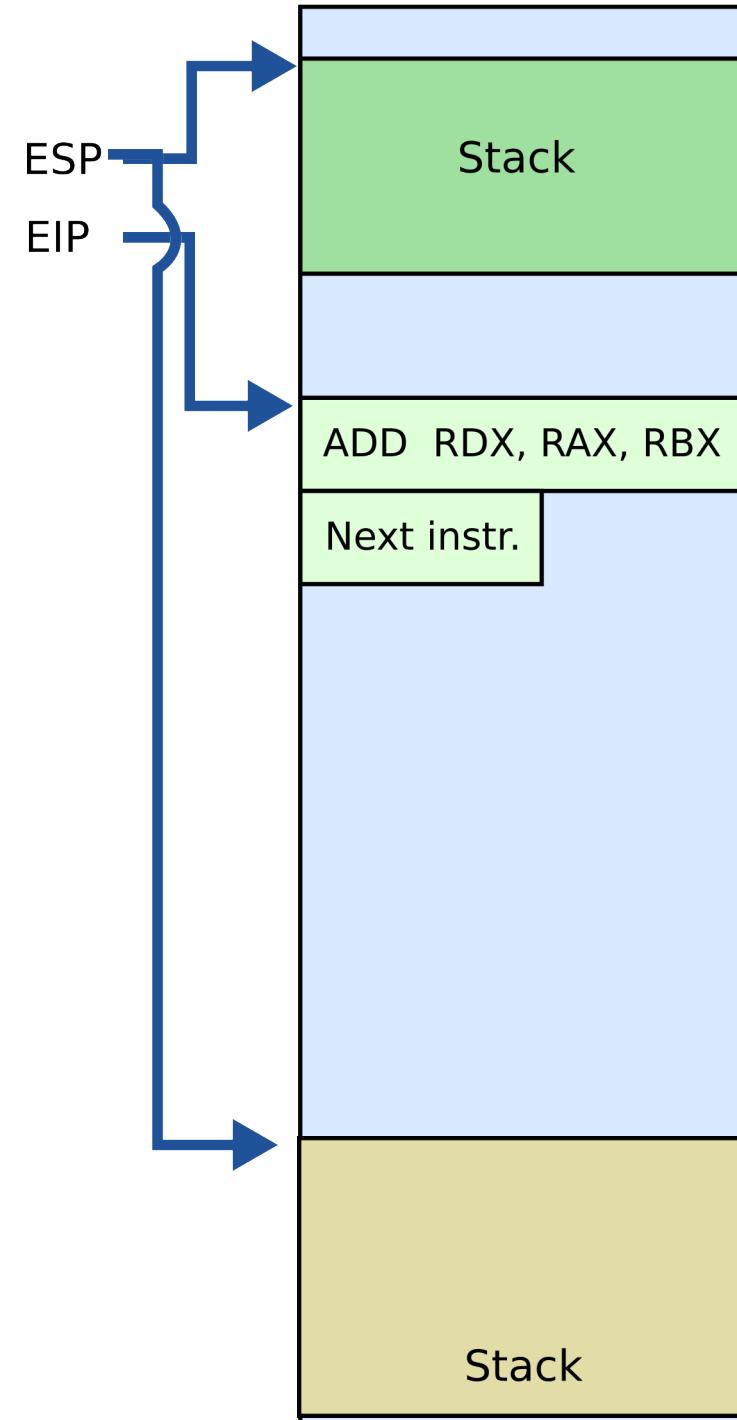




What is stack?

Stack

- It's just a region of memory
 - Pointed by a special register ESP
- You can change ESP
 - Get a new stack



Why do we need stack?

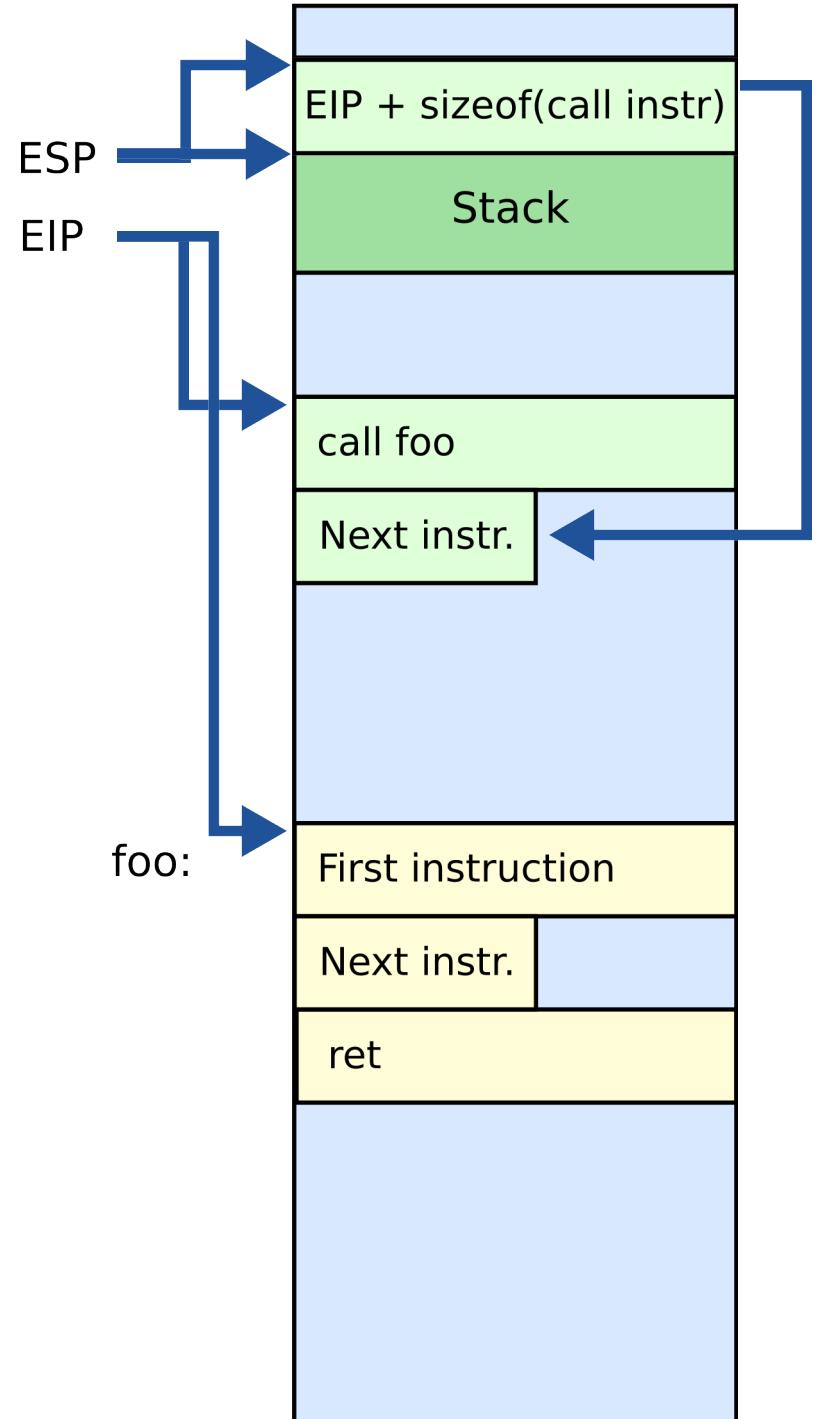
Calling functions

```
// some code...  
foo();  
// more code..
```

- Stack contains information for how to return from a subroutine
 - i.e., foo()

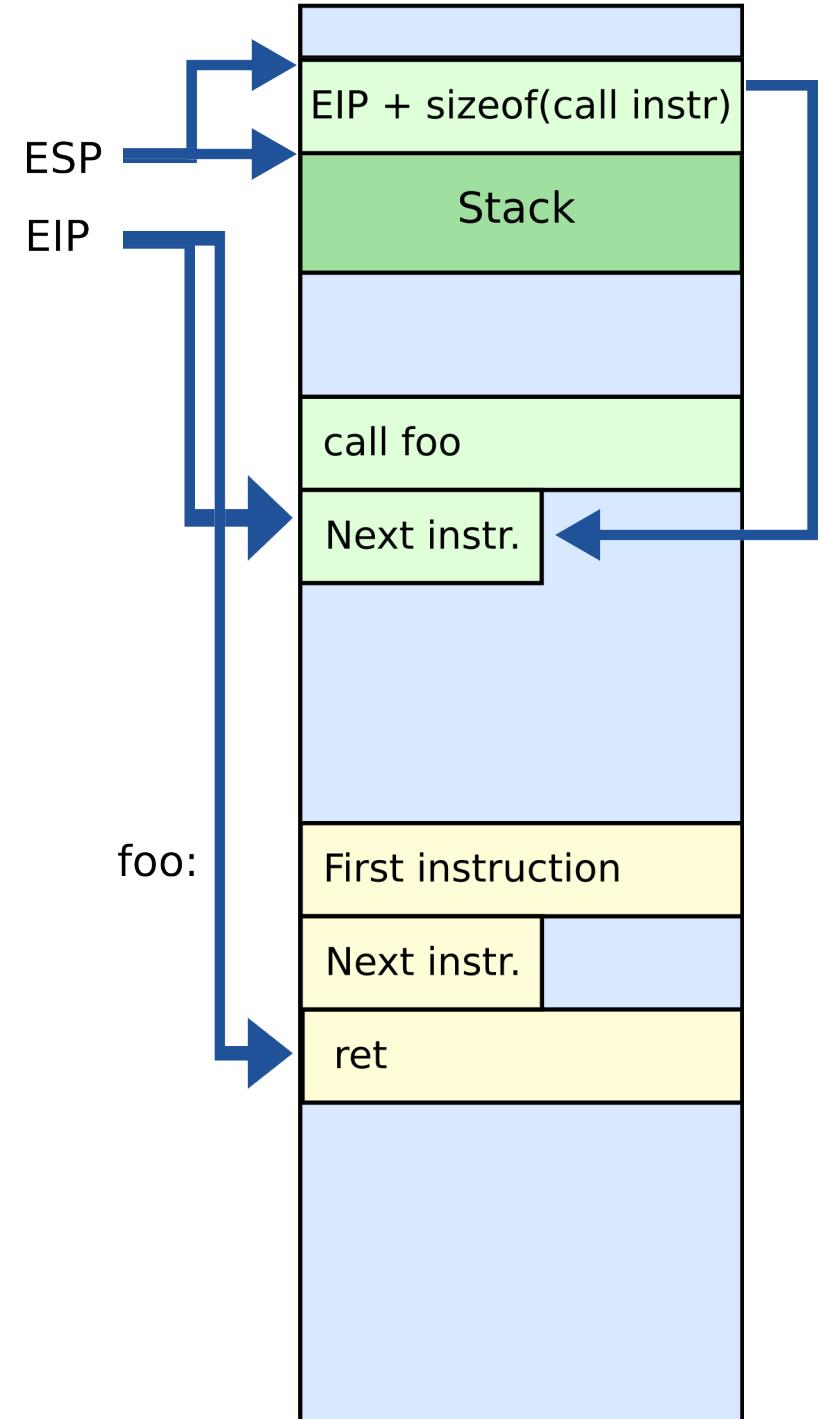
Stack

- Main purpose:
 - Store the return address for the current procedure
 - Caller pushes return address on the stack
 - Callee pops it and jumps



Stack

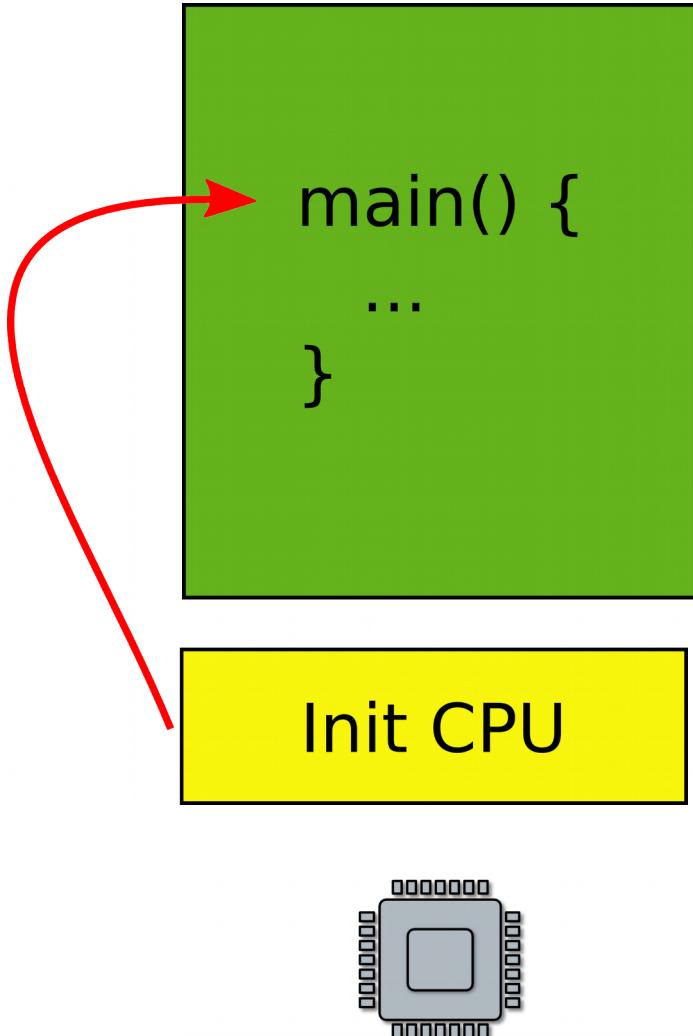
- Main purpose:
 - Store the return address for the current procedure
 - Caller pushes return address on the stack
 - Callee pops it and jumps



Simple observation

- Hardware executes instructions one by one

Goal: Run your code on a piece of hardware



- Read CPU manual
- A tiny boot layer
 - Initialize CPU
 - Jump to the entry point of your program
 - main()
- **This can be the beginning of your OS!**

How do you learn a new programming language?

Hello world

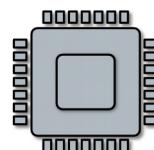
```
printf("Hello world\n");
```

Print out a string

- On the screen or serial line

```
printf() {  
    ...  
    if (vga) {  
        asm("mov <magic number 1>, char");  
    } else if (serial) {  
        asm("out <magic number 2>, char");  
    }  
    ...  
}
```

OS



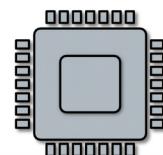
A more general interface

- First device driver

```
printf() {  
    ...  
    putchar(char);  
    ...  
}
```



Console Driver



Device drivers

- Abstract hardware
 - Provide high-level interface
 - Hide minor differences
 - Implement some optimizations
 - Batch requests
- Examples
 - Console, disk, network interface
 - ...virtually any piece of hardware you know

OS is like a library that provides a collection of useful functions

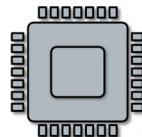
Goal: Want to run two programs

```
main() {  
    ...  
    yield()  
}
```

```
main() {  
    ...  
    yield()  
}
```

- What does it mean?
 - Only one CPU
 - Run one, then run another one

Save/restore



Very much like car sharing

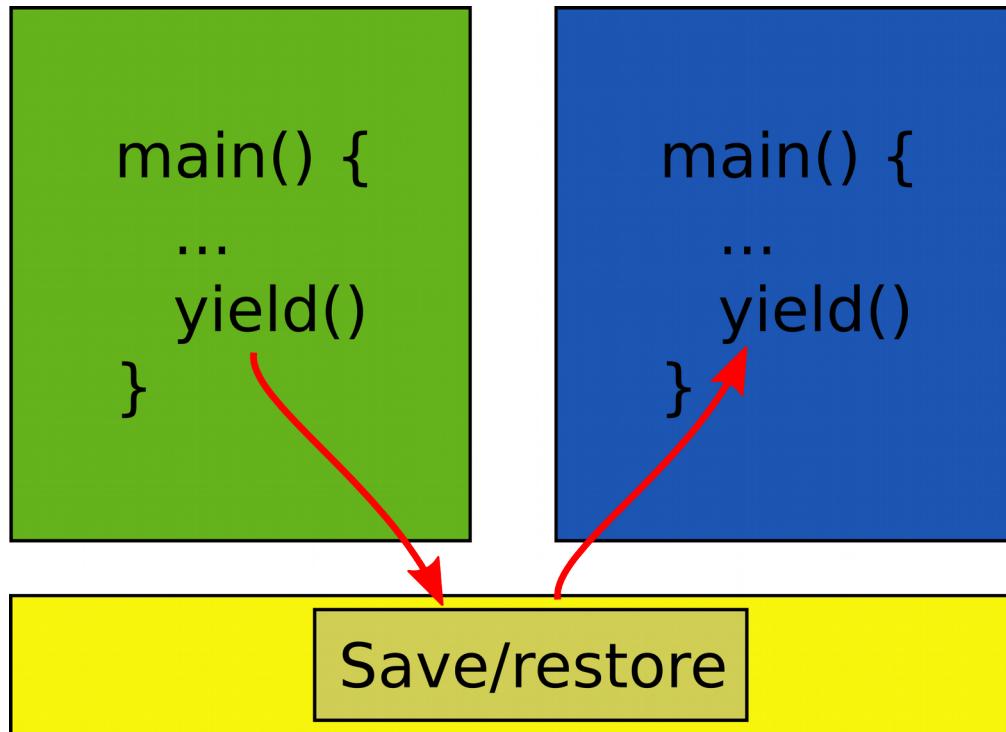


Car rental

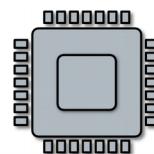
Time sharing

- Programs use CPU in turns
 - One program runs
 - Then OS takes control
 - Launches another program
 - Then another program runs
 - OS takes control again
 - ...

Goal: Want to run two programs



- Exit into the kernel periodically
- Context switch
 - Save state of one program
 - Restore state of another program



What is this state?

State of the program

- Roughly it's
 - Registers
 - Memory
- Plus some state (data structures) in the kernel associated with the program
 - Information about files opened by the program, i.e. file descriptors
 - Information about network flows
 - Information about address space, loaded libraries, communication channels to other programs, etc.

Saving and restoring state

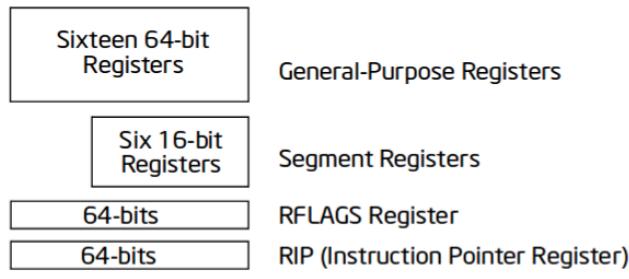
- Note that you do not really have to save/restore in-kernel state on the context switch
 - It's in the kernel already, i.e., in some part of the memory where kernel keeps its data structures
 - You only have to switch from using one to using another
 - i.e., instead of using the file descriptor table (can be as simple as array) for program X start using at file descriptor table for program Y

Saving and restoring state

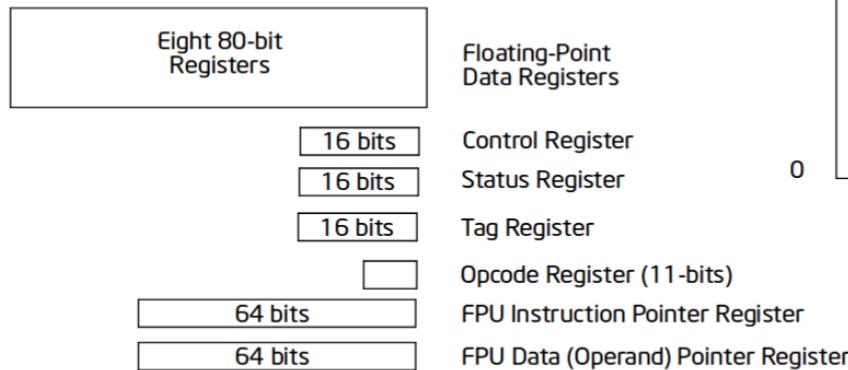
- All you have to save are internal structures of the CPU, i.e.
 - Registers
 - Note CPU has more registers than just
 - General registers, i.e., EAX, EBX, ...
 - 8 general registers in x86 32bit mode
 - 16 general registers in x86 64bit mode

Intel x86 64bit Execution Environment

Basic Program Execution Registers



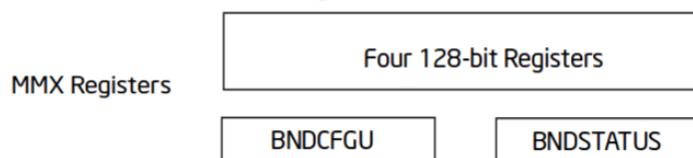
FPU Registers



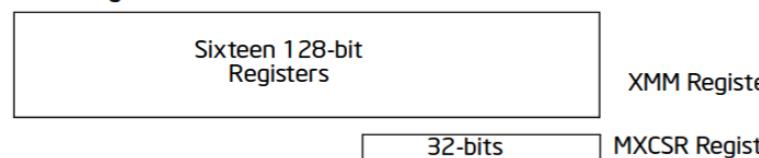
MMX Registers



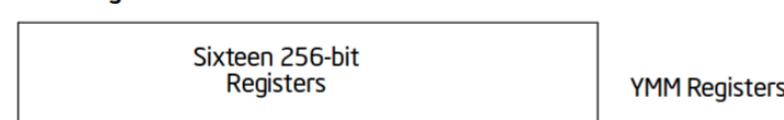
Bounds Registers



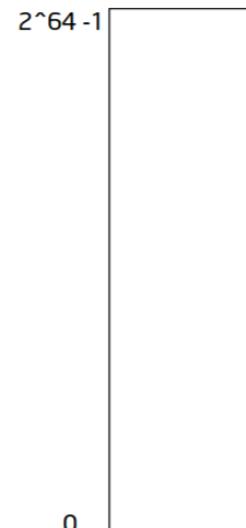
XMM Registers



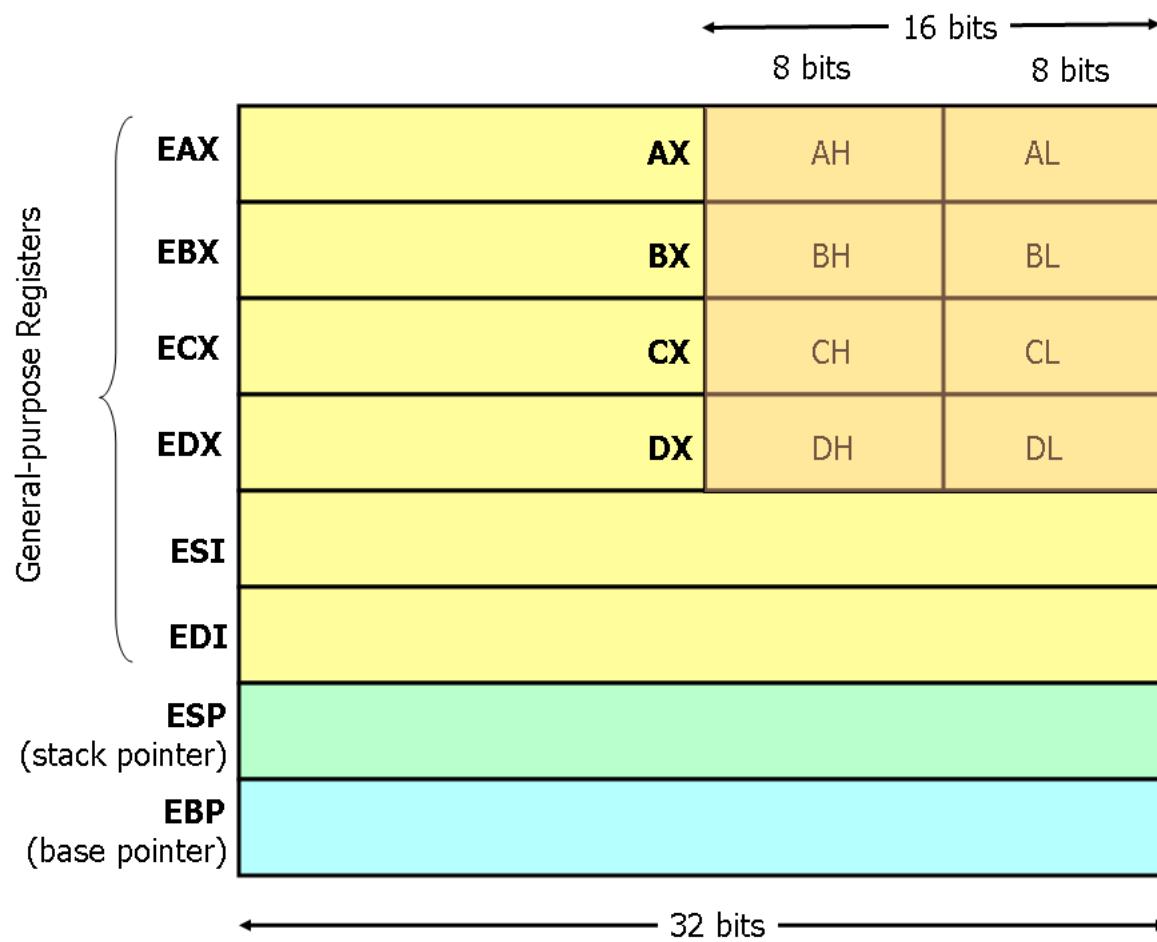
YMM Registers



Address Space



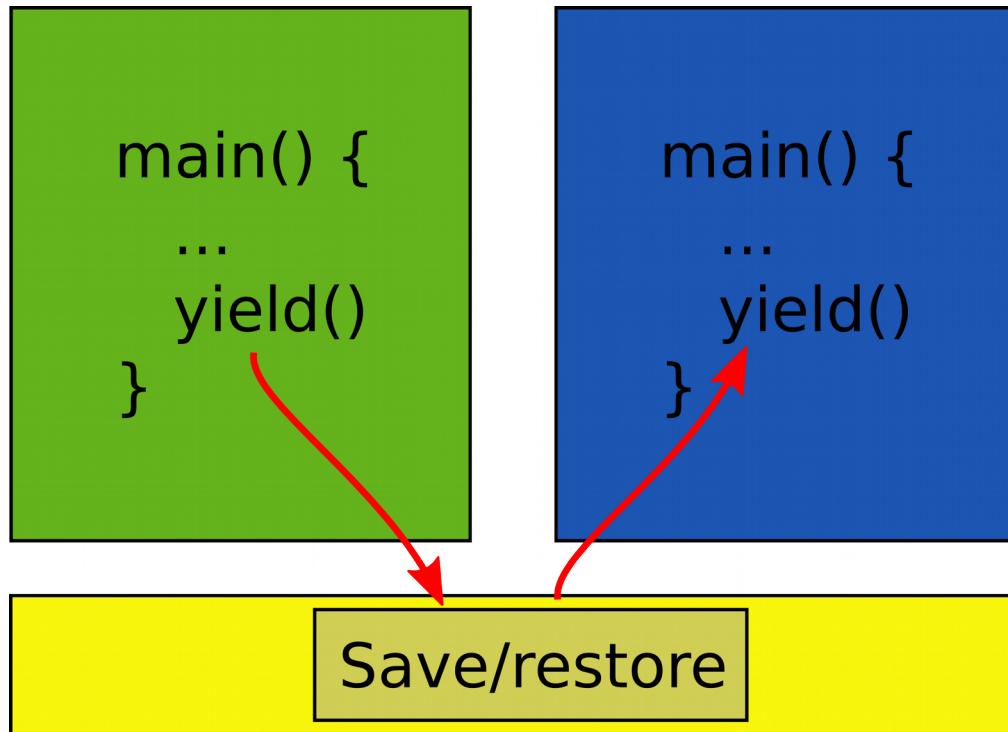
General registers



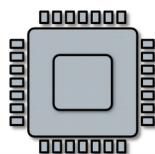
More registers...

- This is a bit misleading...
- CPU also has registers that describe state of
 - Segments
 - Page tables
 - Interrupt tables
 - Etc.
- If they don't change you don't have to save/restore them

But anyway... if you want to run two programs

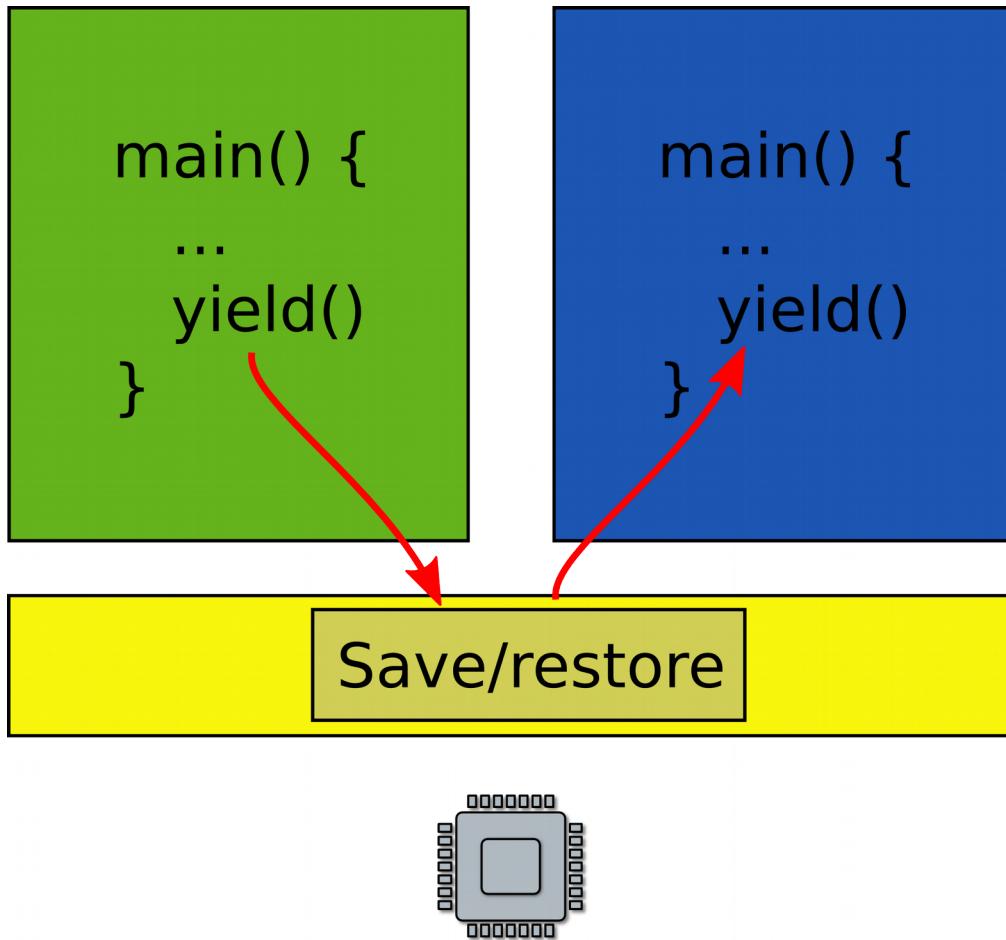


- Exit into the kernel periodically
- Context switch
 - Save state of one program
 - Restore state of another program



What about memory?

- Two programs, one memory?



Time-share memory

- Well you can copy in and out the state of the program into a region of memory where it can run
 - Similar to time-sharing the CPU

Time-share memory

- Well you can copy in and out the state of the program into a region of memory where it can run
 - Similar to time-sharing the CPU
- What do you think is wrong with this approach?

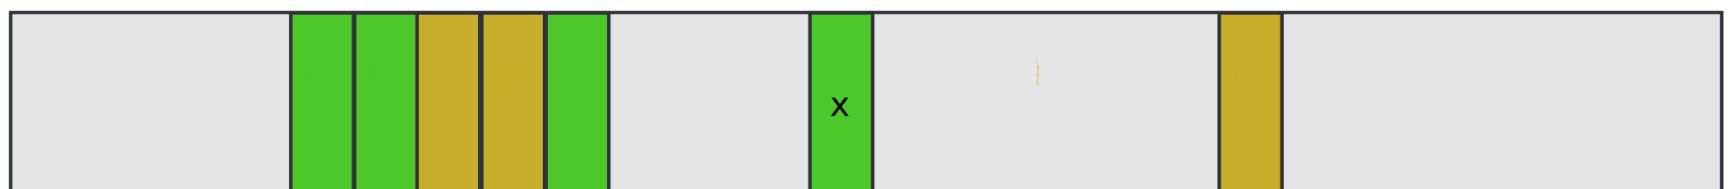
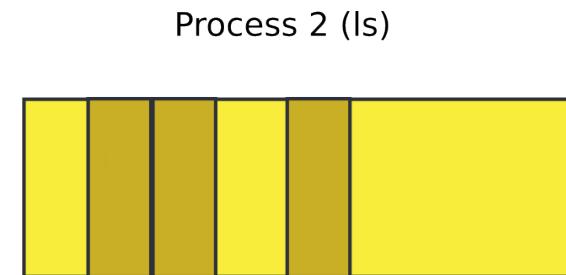
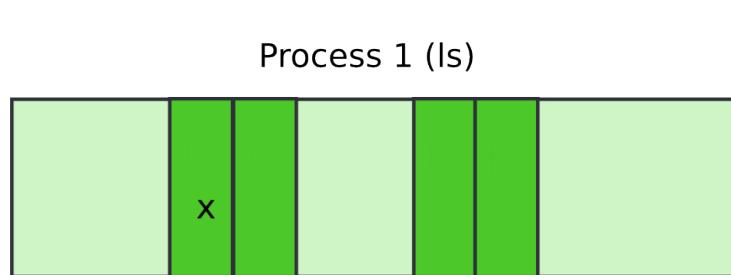
Time-share memory

- Well you can copy in and out the state of the program into a region of memory where it can run
 - Similar to time-sharing the CPU
- What do you think is wrong with this approach?
 - Unlike registers the state of the program in memory can be large
 - Takes time to copy it in and out

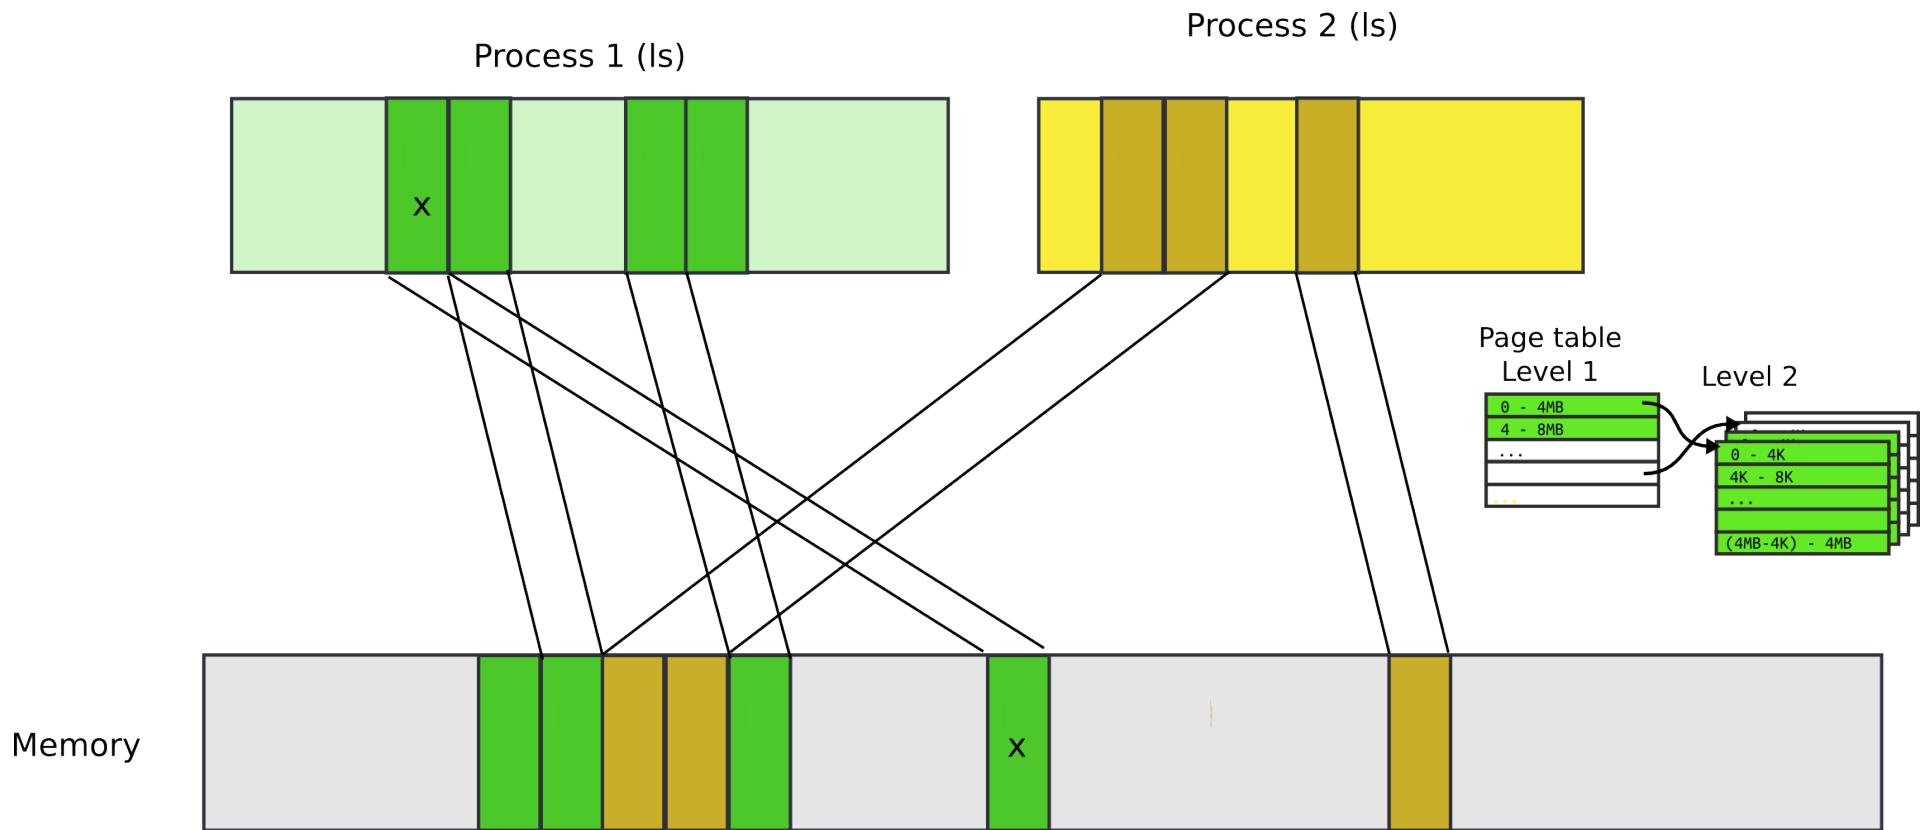
Space sharing: virtual address spaces

- Illusion of a private memory for each application
 - Keep a description of an address space
 - In one of the registers
- OS maintains description of address spaces
 - Switches between them

Address spaces and paging



Address spaces and paging

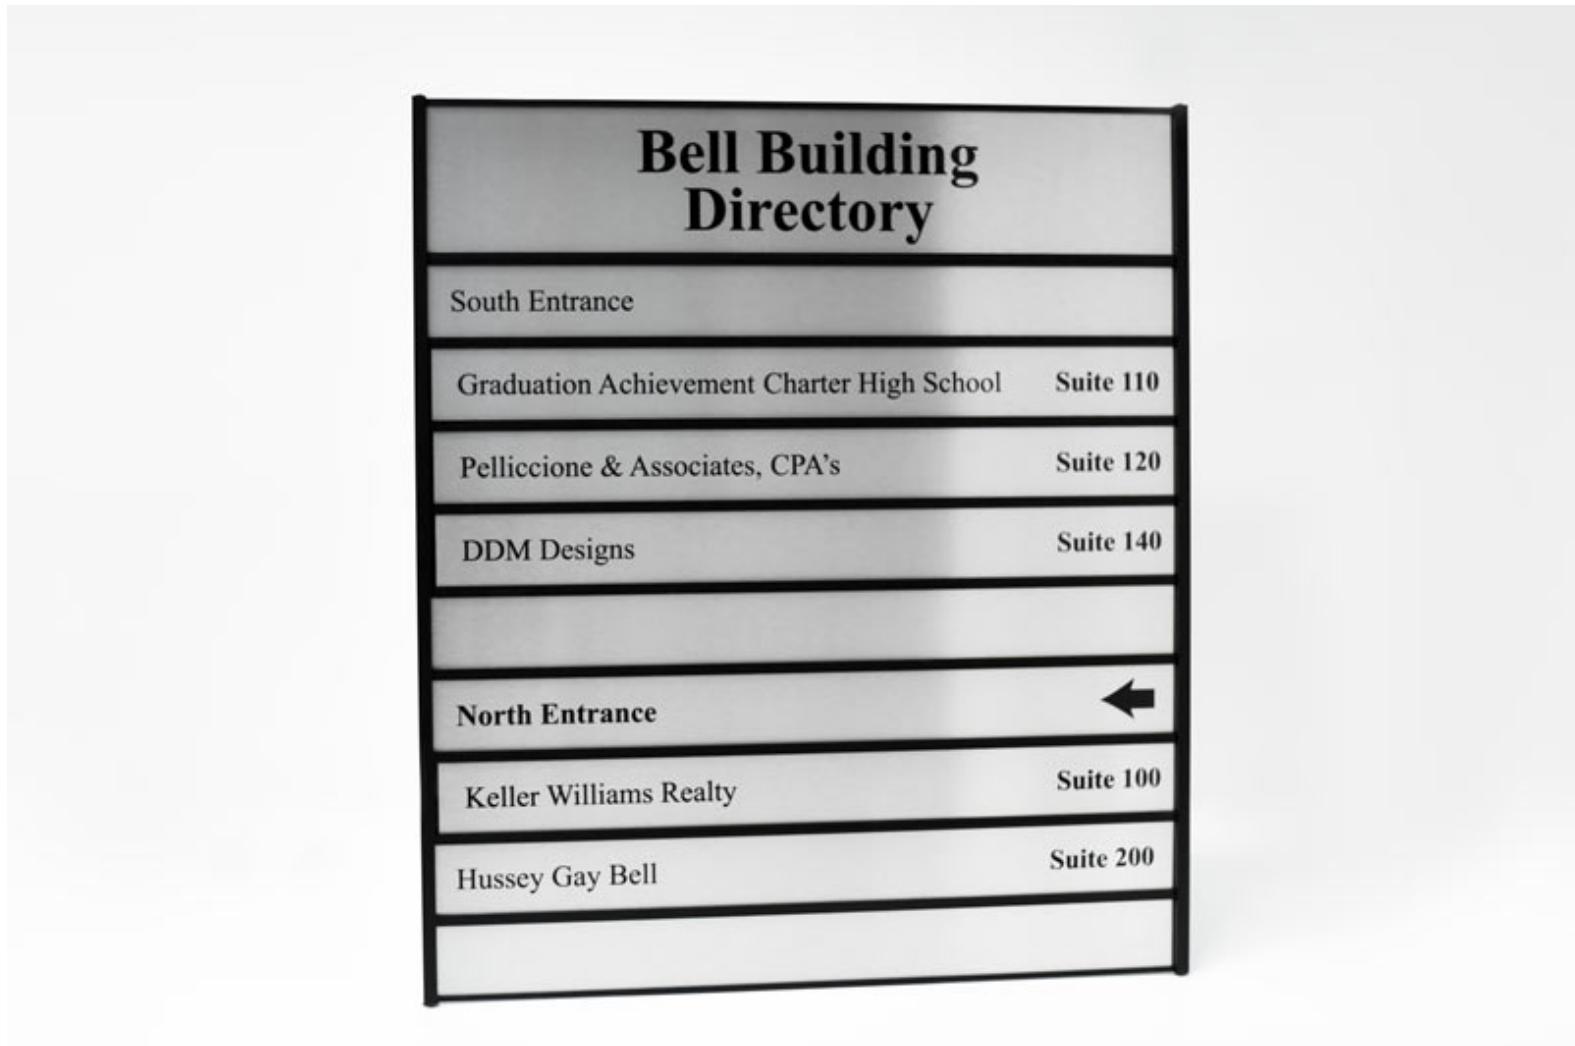


Paging idea

- Break up memory into 4096-byte chunks called pages
 - Modern hardware supports 2MB, 4MB, and 1GB pages
- Independently control mapping for each page of linear address space

Notice the main difference: time-sharing vs space sharing

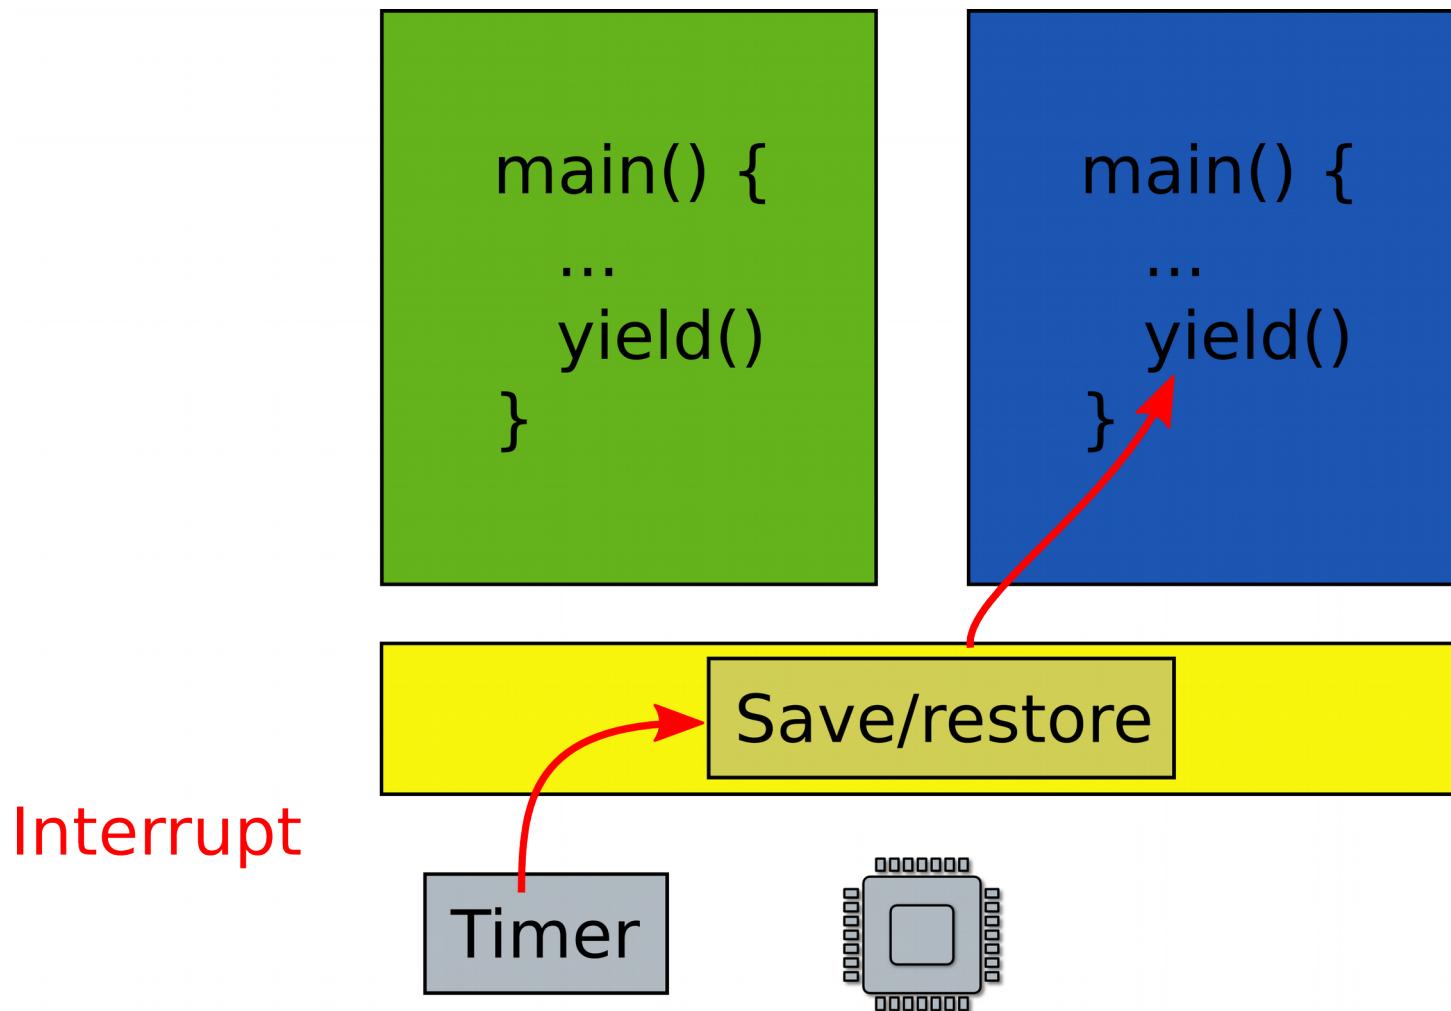
Space sharing is like renting a some rooms in an office building



Staying in control

Staying in control

- What if one program fails to release the CPU?
- It will run forever. Need a way to preempt it. How?

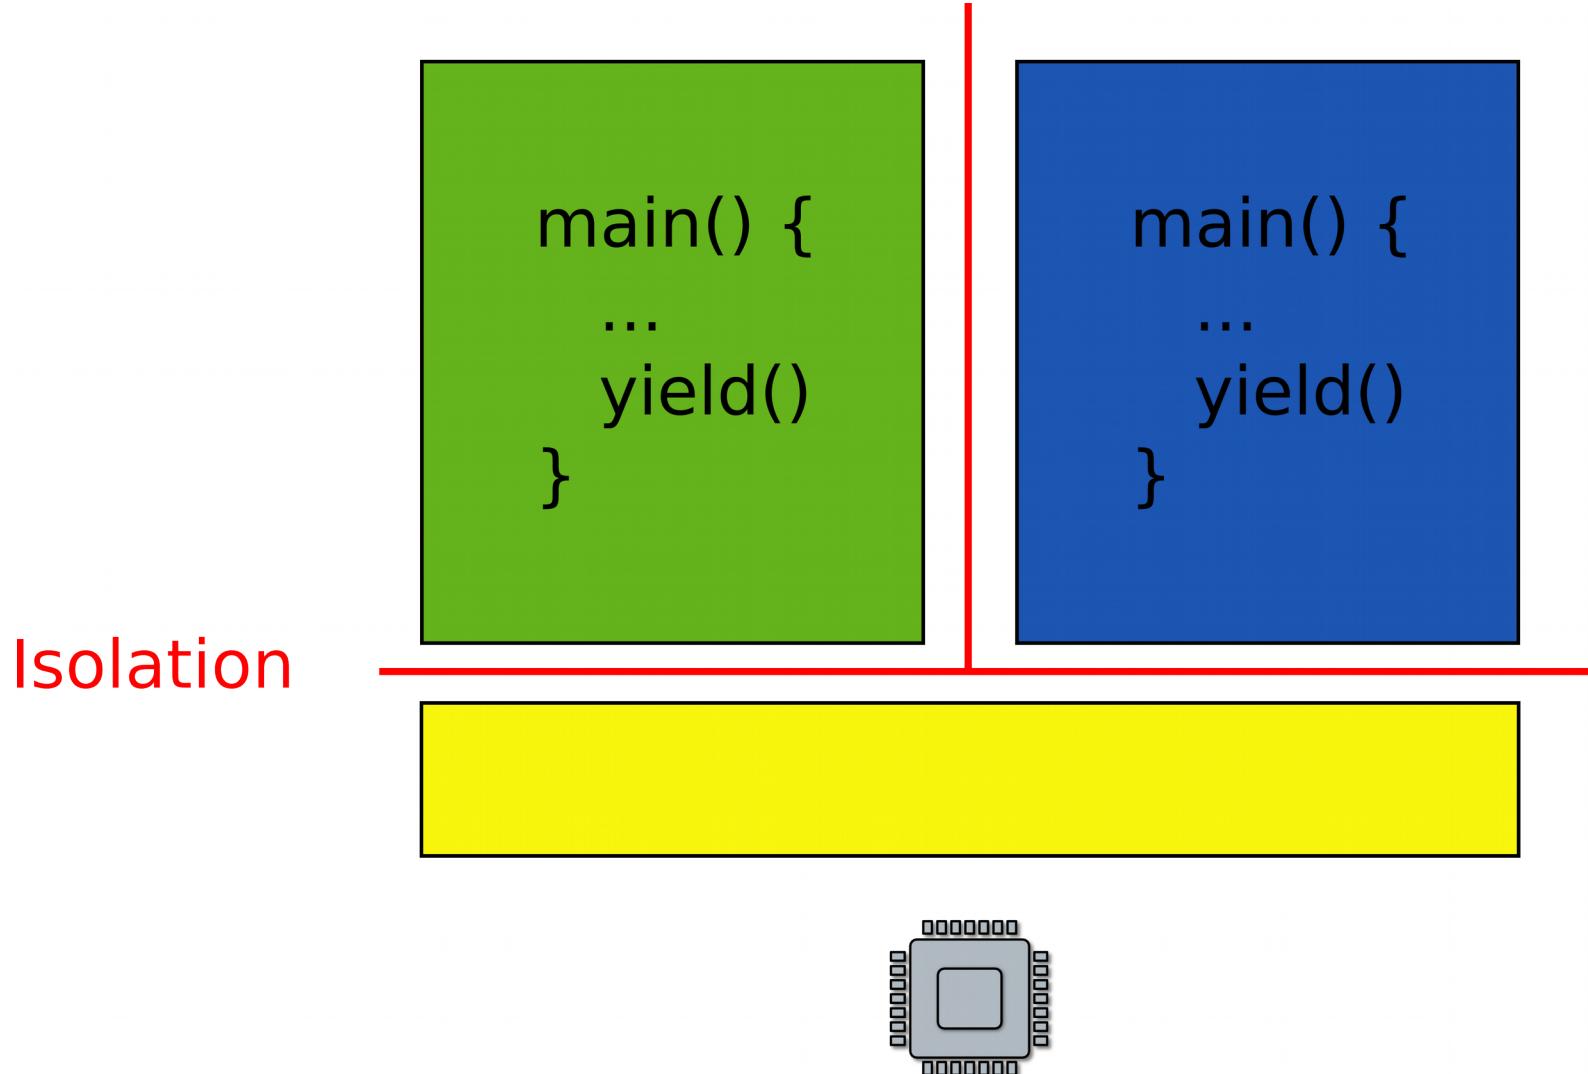


Scheduling

- Pick which application to run next
 - And for how long
- Illusion of a private CPU for each task
 - Frequent context switching

Isolation

- What if one faulty program corrupts the kernel?
- Or other programs?



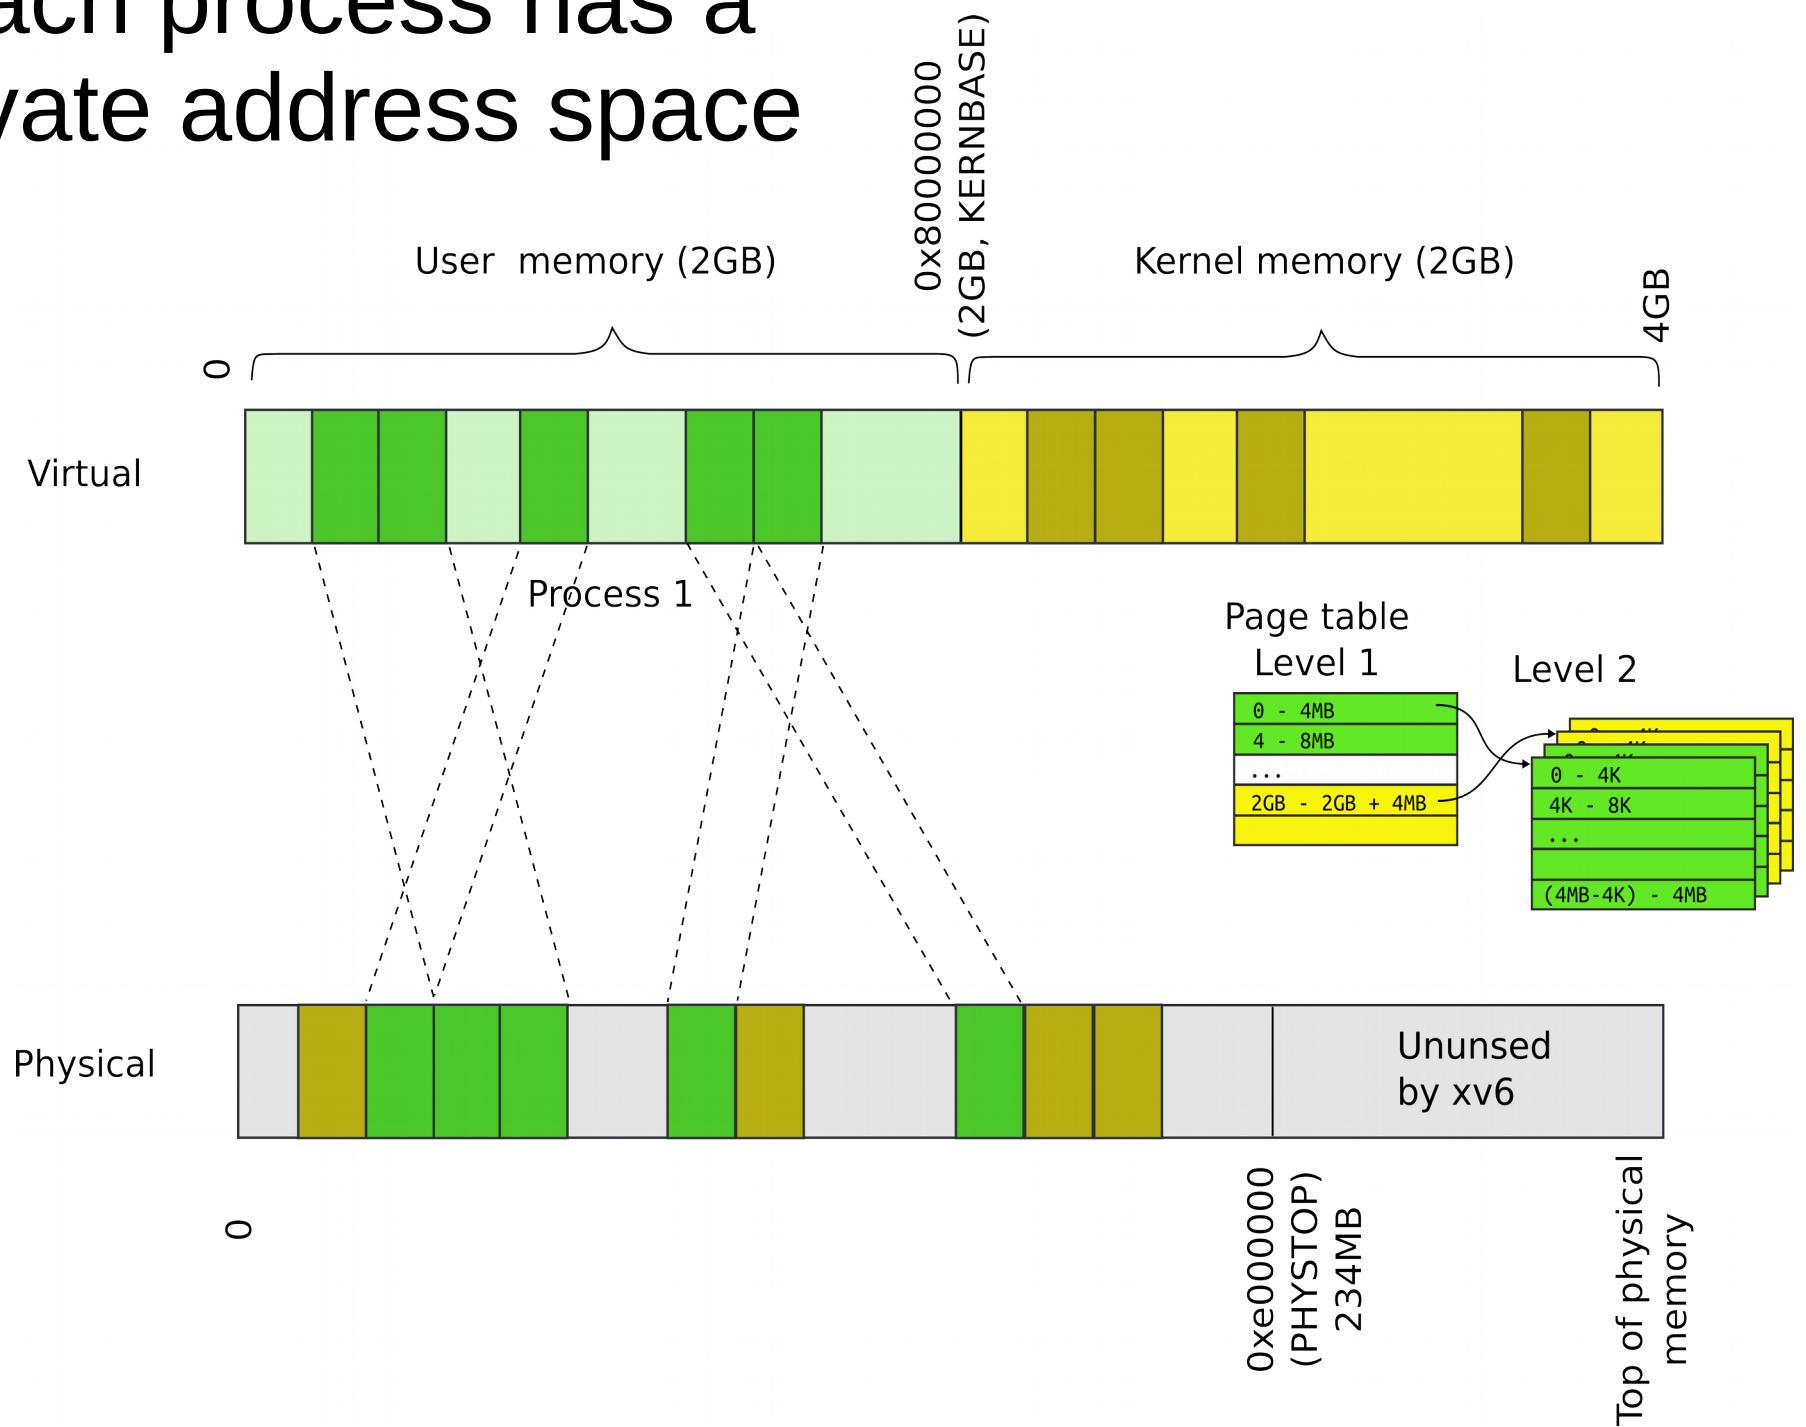
No isolation: open space office



Isolated rooms

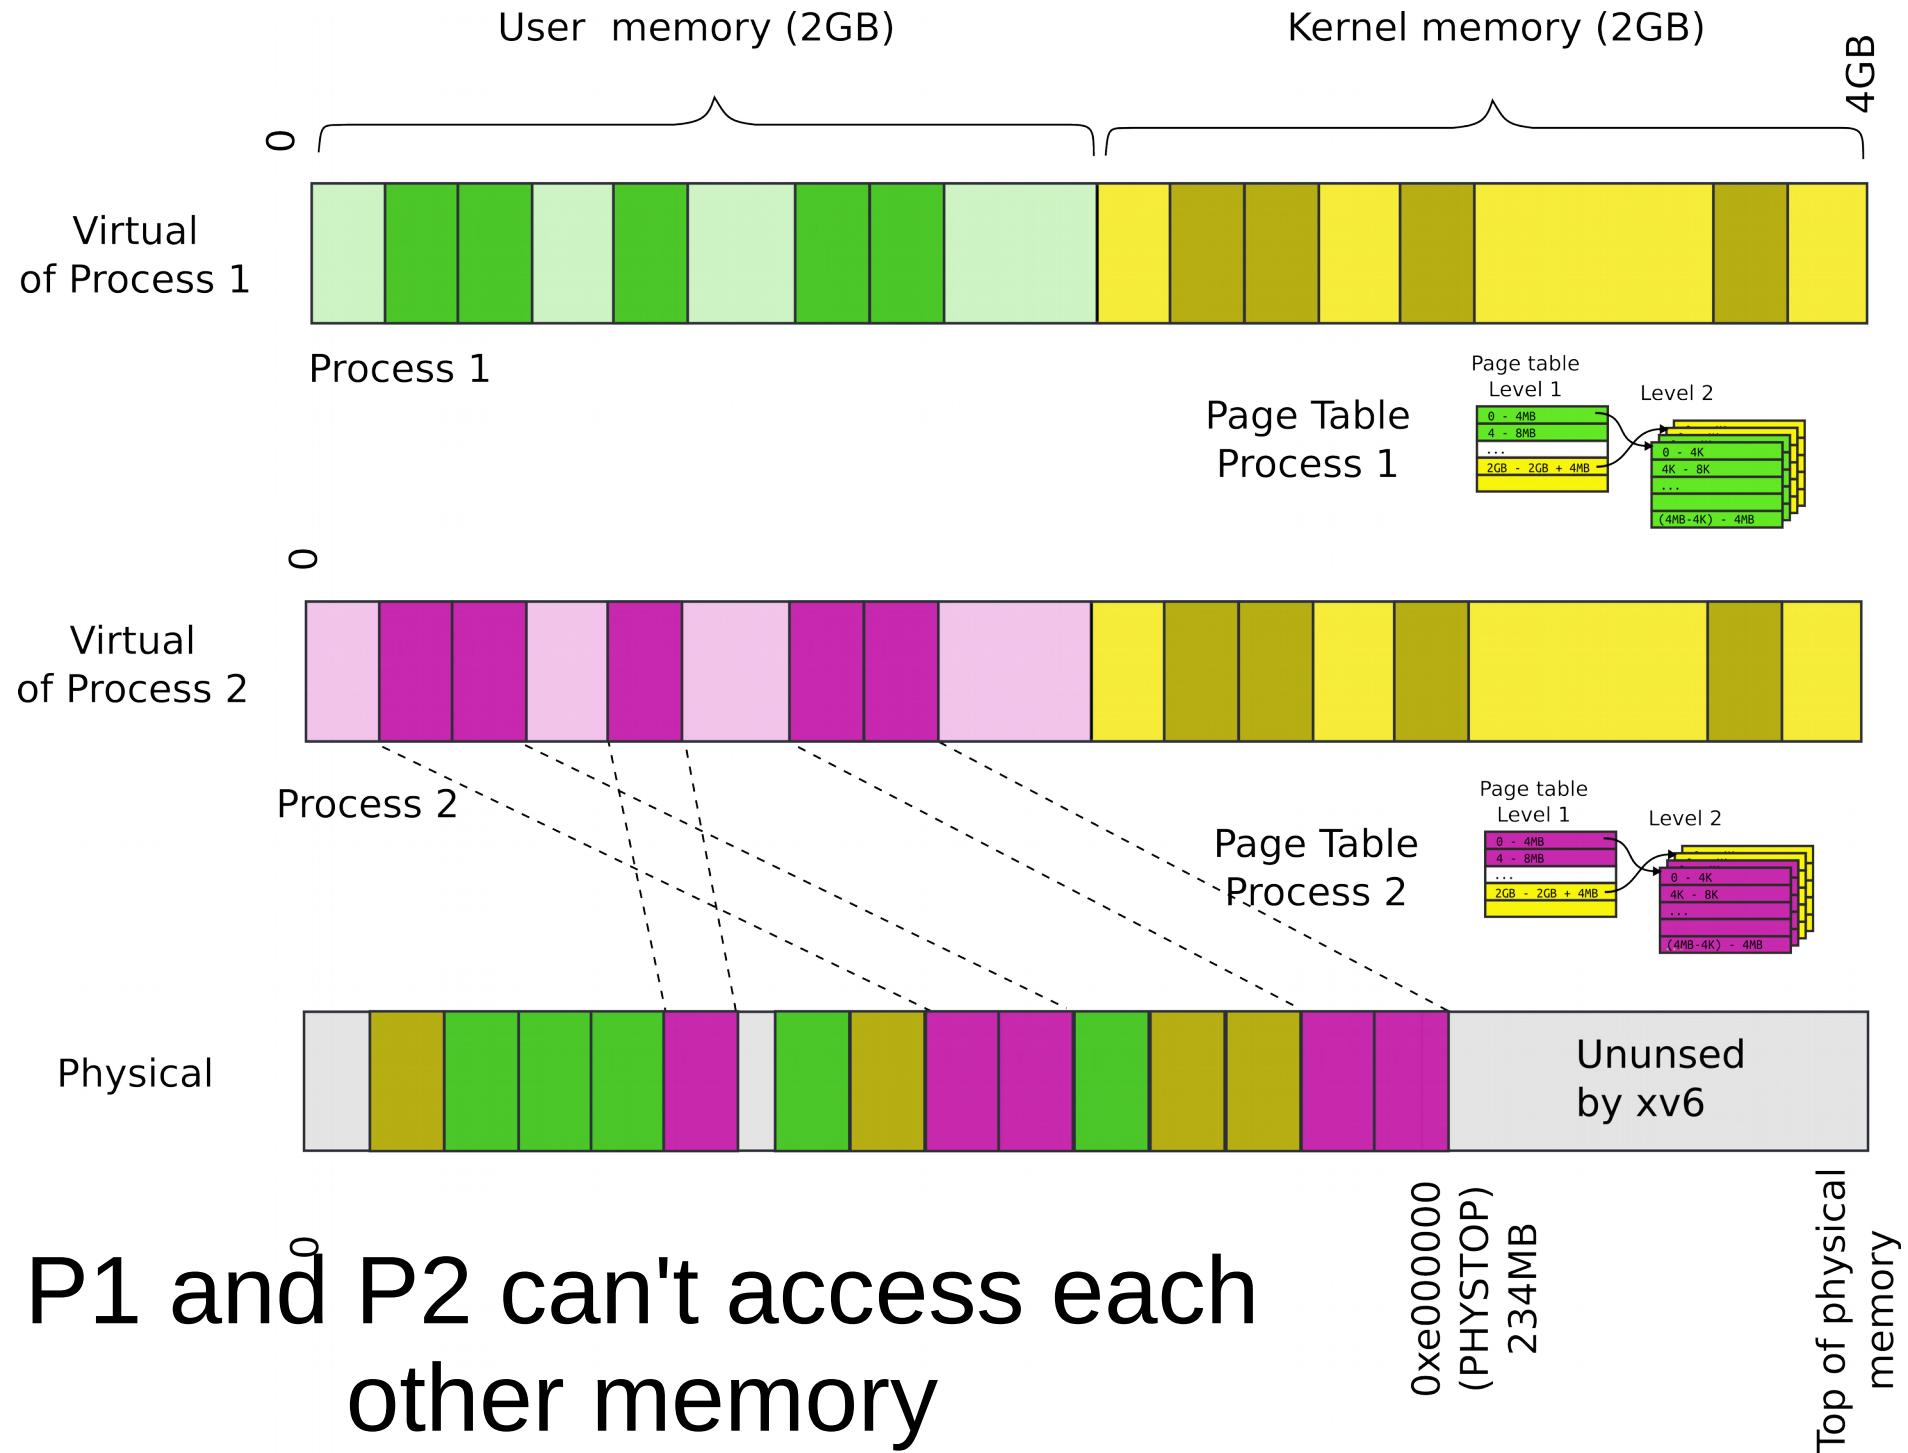


Each process has a private address space

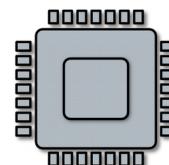
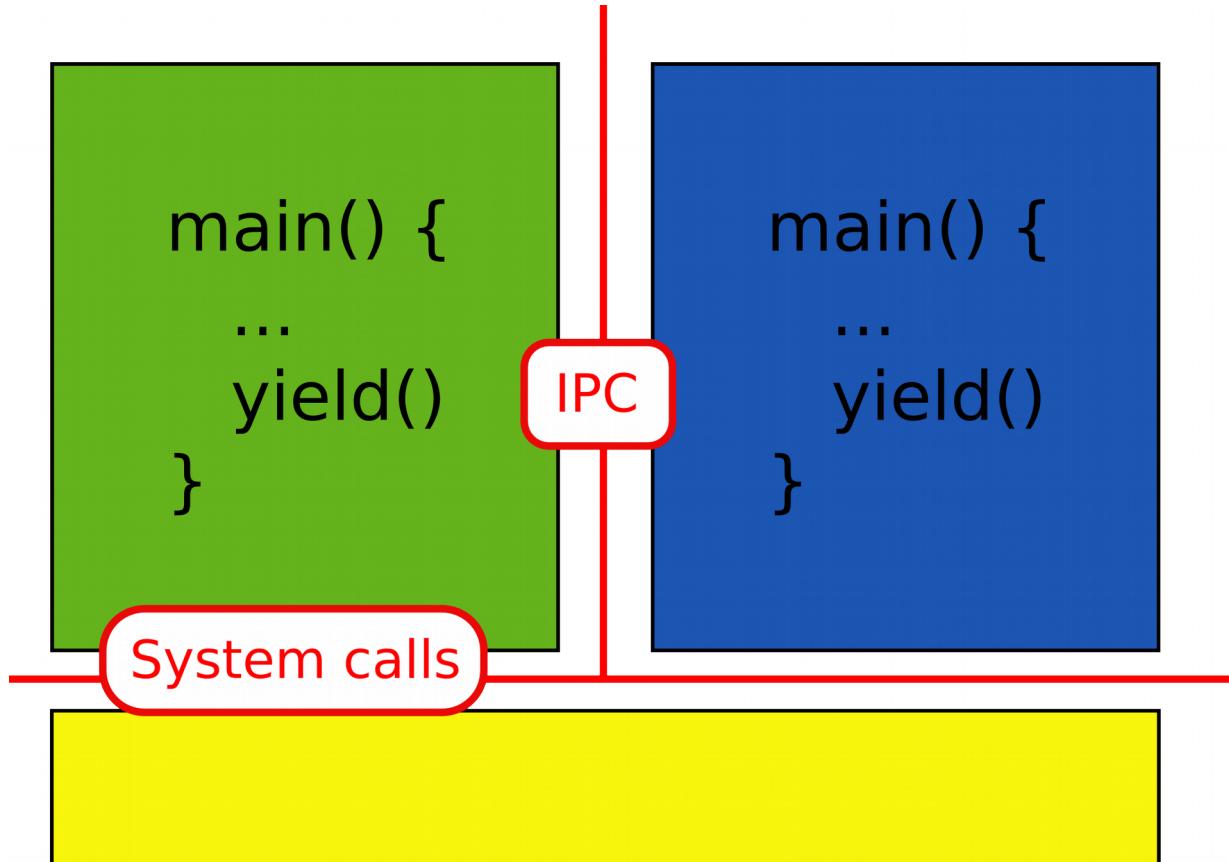


Each process maps the kernel

- It's not strictly required
 - But convenient for system calls
 - No need to change the page table when process enters the kernel with a system call
 - **Things are much faster!**



- What about communication?
- Can we invoke a function in a kernel?



Files and network

- What if you want to save some data to a file?

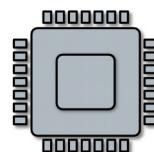
- What if you want to save some data?
- Permanent storage
 - E.g., disks
- But disks are just arrays of blocks
 - `wrtie(block_number, block_data)`
- Files
 - High level abstraction for saving data
 - `fd = open("contacts.txt");`
 - `fpritnf(fd, "Name:%s\n", name);`

Remember our console driver

- Print a string on the screen or serial line

```
printf() {  
    ...  
    if (vga) {  
        asm("mov <magic number 1>, char");  
    } else if (serial) {  
        asm("out <magic number 2>, char");  
    }  
    ...  
}
```

OS



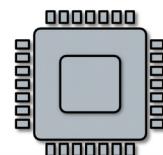
A more general interface

- First device driver

```
printf() {  
    ...  
    putchar(char);  
    ...  
}
```



Console Driver



- File system and block device provide similar abstractions
- Permanent storage
 - E.g., disks
- But disks are just arrays of blocks
 - `wrtie(block_number, block_data)`
- Files
 - High level abstraction for saving data
 - `fd = open("contacts.txt");`
 - `fpritnf(fd, "Name:%s\n", name);`

File system and block layer

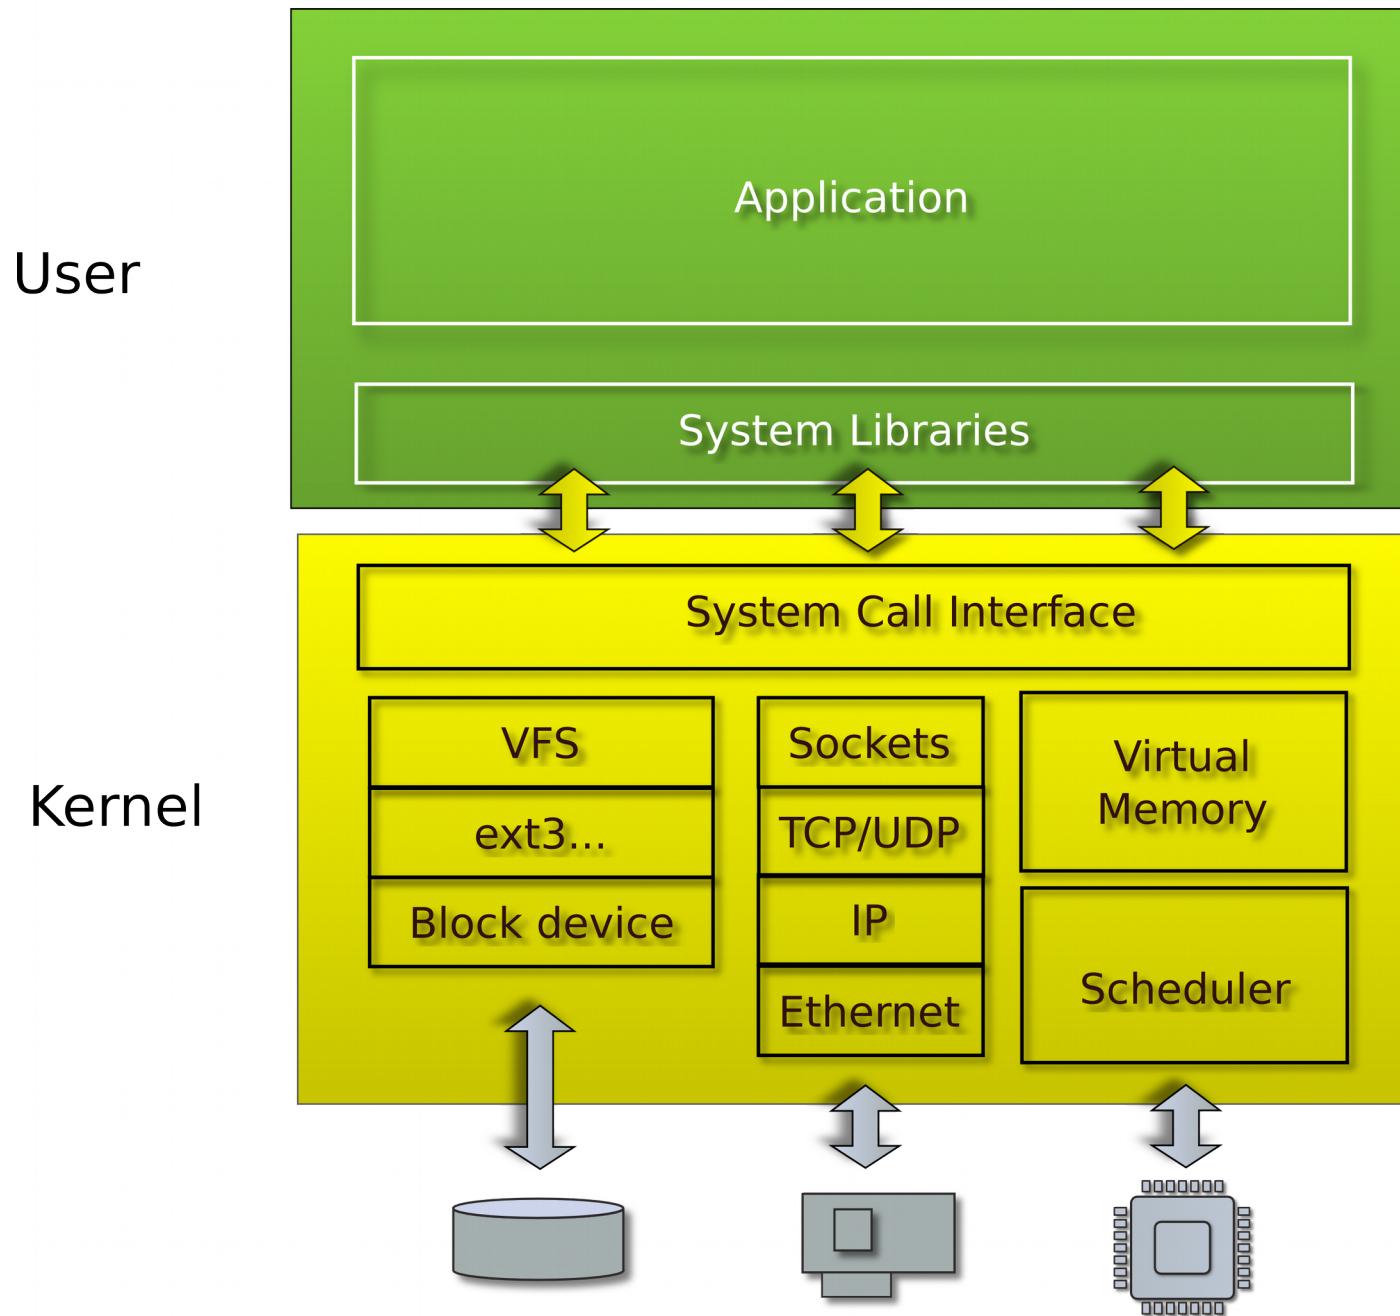
System calls	File descriptors
Pathnames	Recursive lookup
Directories	Directory inodes
Files	Inodes and block allocator
Transactions	Logging
Blocks	Buffer cache

- Reliable storage on top of raw disc blocks
- Disks are just arrays of blocks
`wrtie(block_number, block_data)`
- Human readable names (files)
 - High level abstraction for saving data
`fd = open("contacts.txt");`
`fpritnf(fd, "Name:%s\n", name);`

What if you want to send data over the network?

- Similar idea
 - Send/receive Ethernet packets (Level 2)
 - Two low level
- Sockets
 - High level abstraction for sending data

- Linux/Windows/Mac



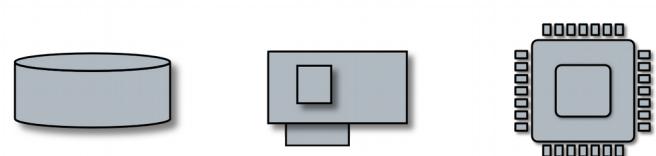
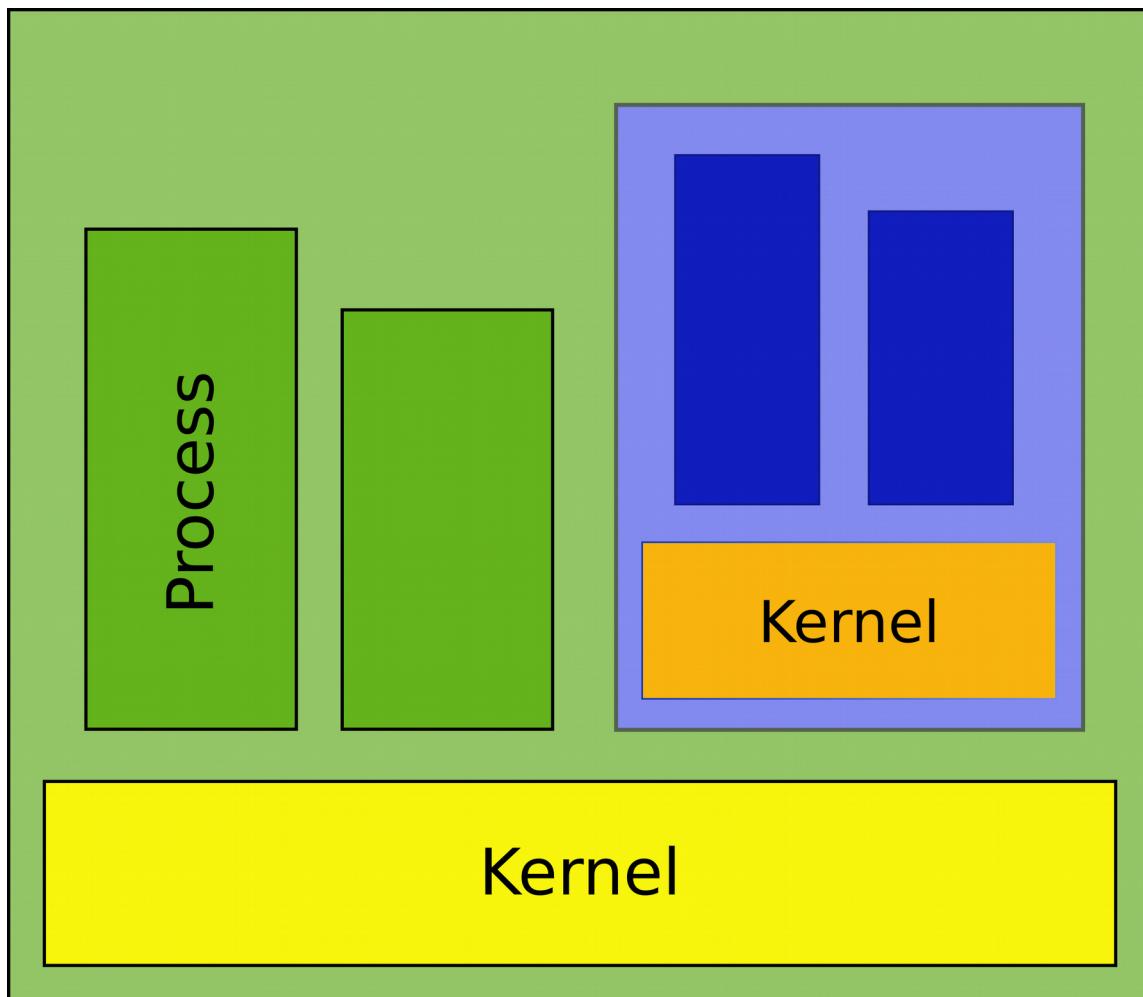
Recap

- Run multiple programs
 - Each has illusion of a private memory and CPU
 - Context switching
 - Isolation and protection
 - Management of resources
 - Scheduling (management of CPU)
 - Memory management (management of physical memory)
- High-level abstractions for I/O
 - File systems
 - Multiple files, concurrent I/O requests
 - Consistency, caching
 - Network protocols
 - Multiple virtual network connections

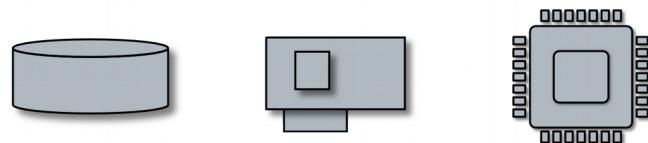
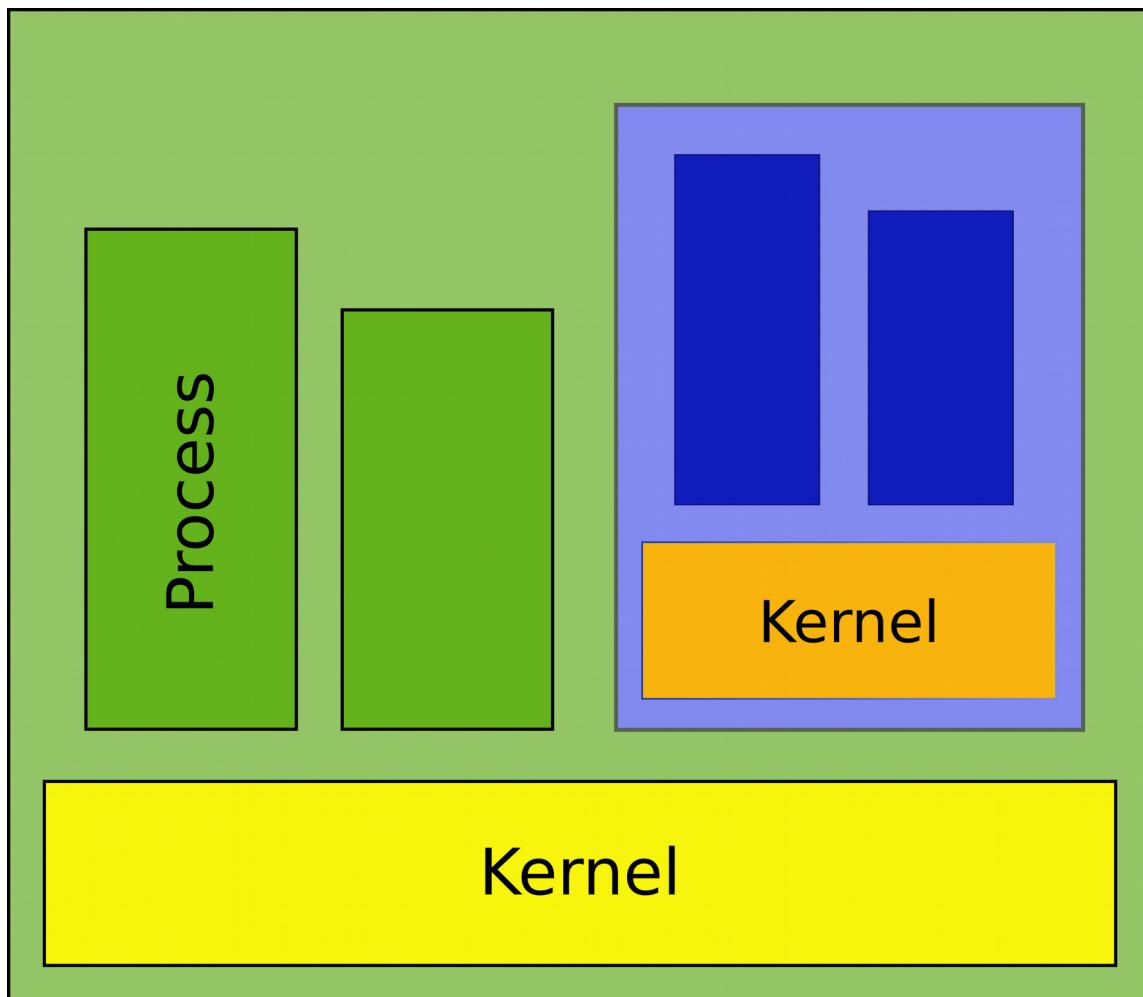
Questions?

Virtualization

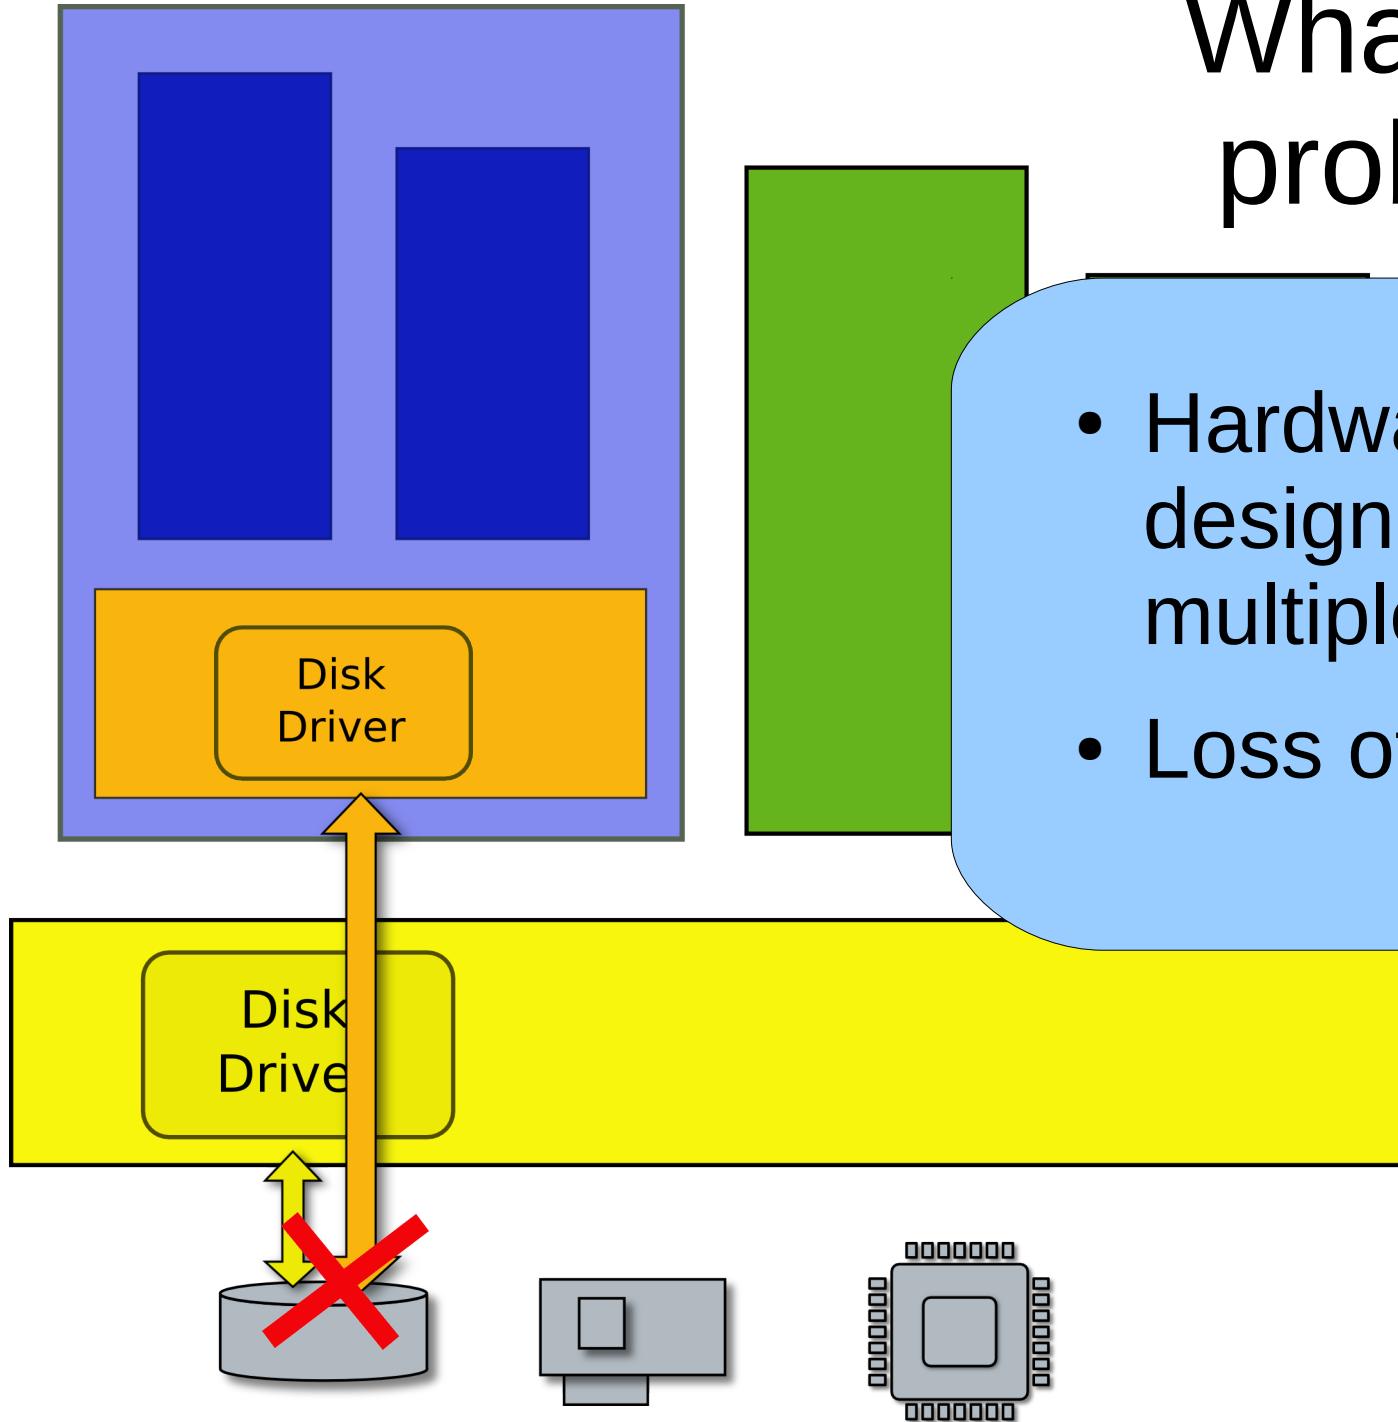
- Want to run a Windows application on Linux?



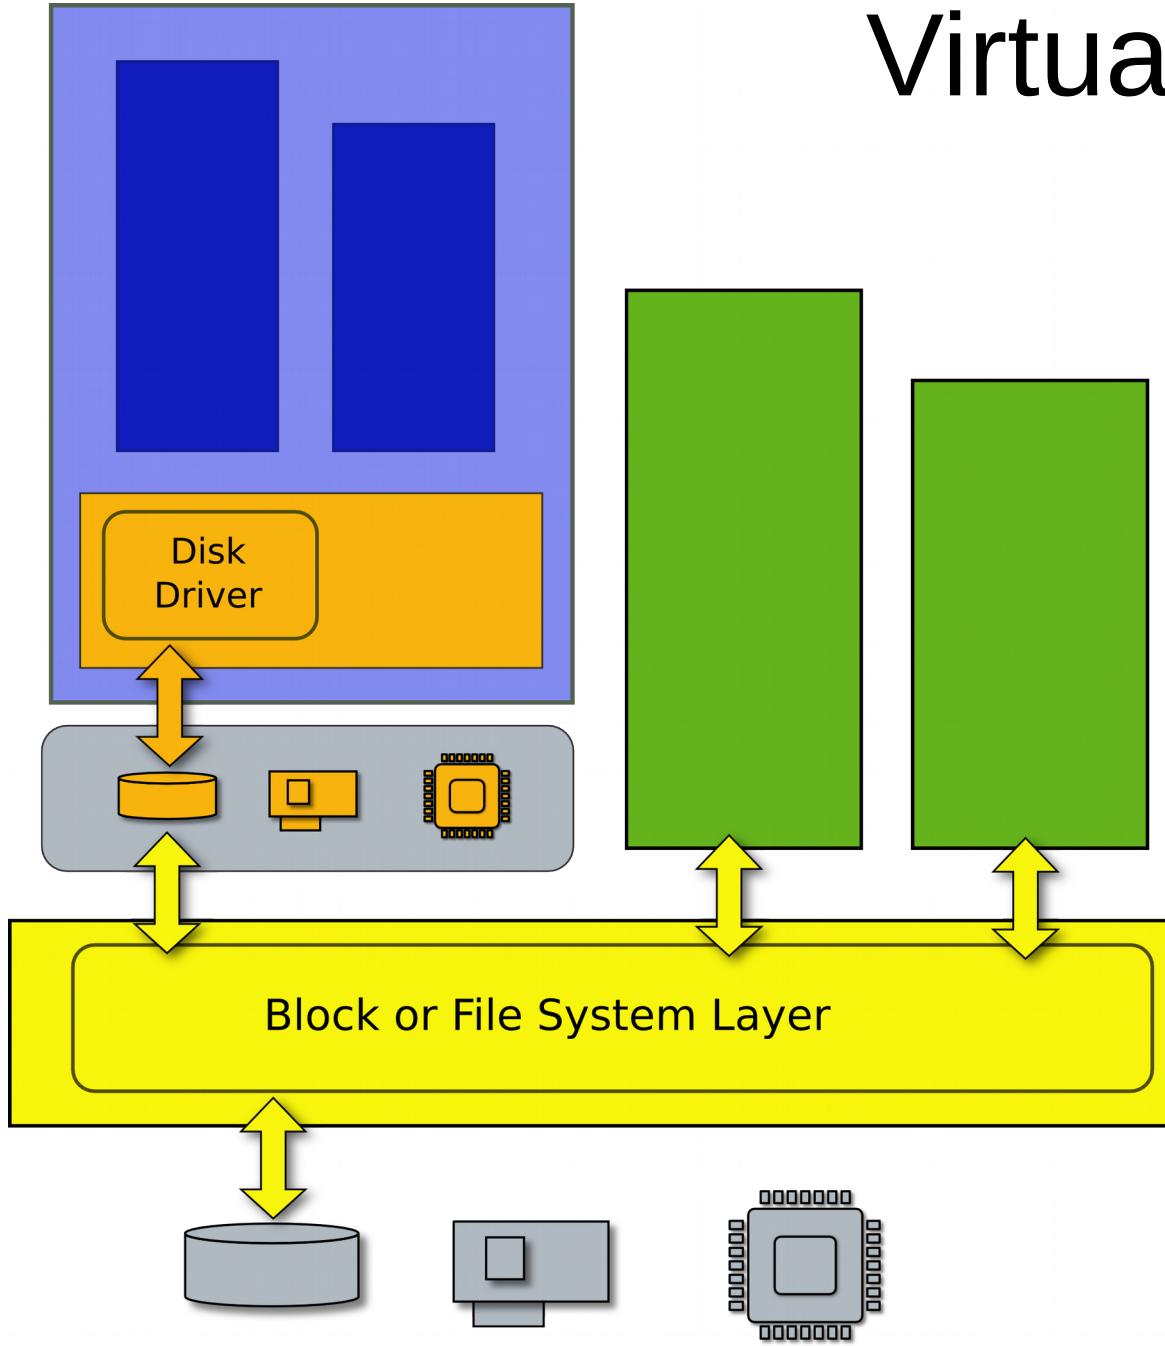
- Want to run a Windows application on Linux?



What is the problem?



Virtual machine



Efficient duplicate
of a real machine

- Compatibility
- Performance
- Isolation

Trap and emulate

