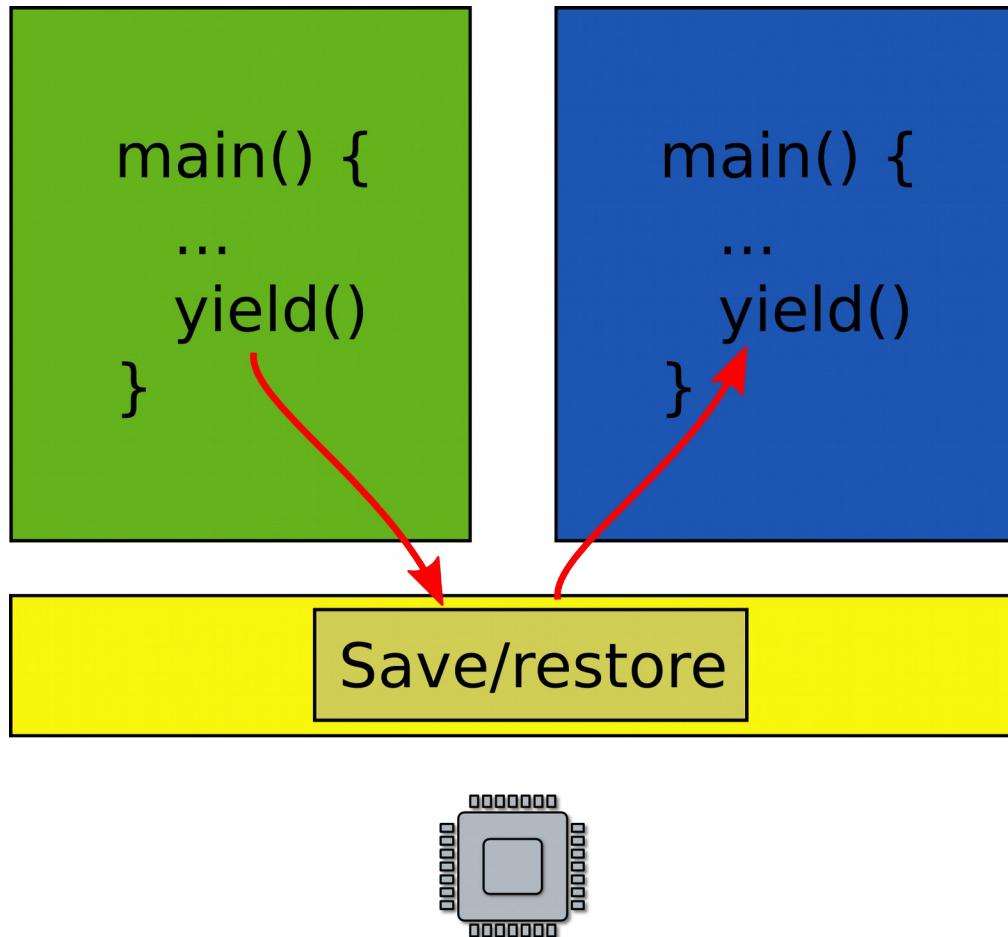


# cs5460/6460: Operating Systems

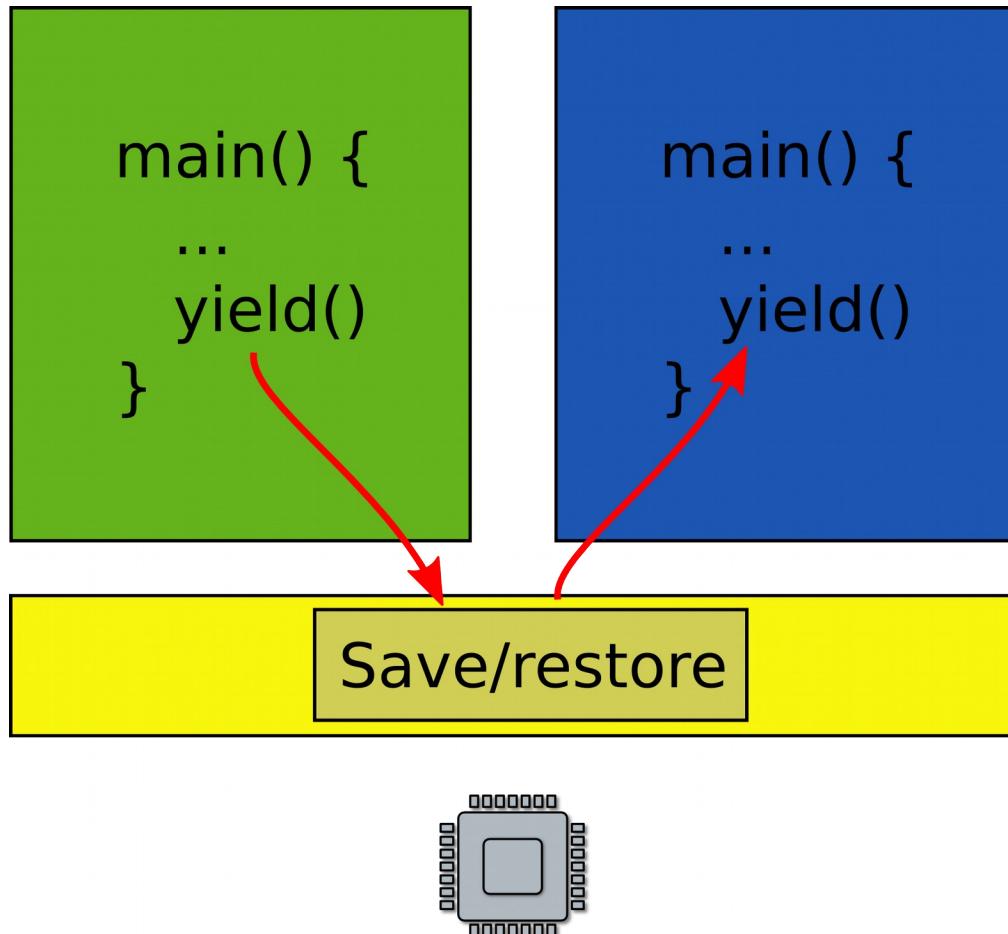
## Address translation (Segmentation and Paging)

Anton Burtsev  
February, 2023

# Two programs one memory



# Two programs one memory

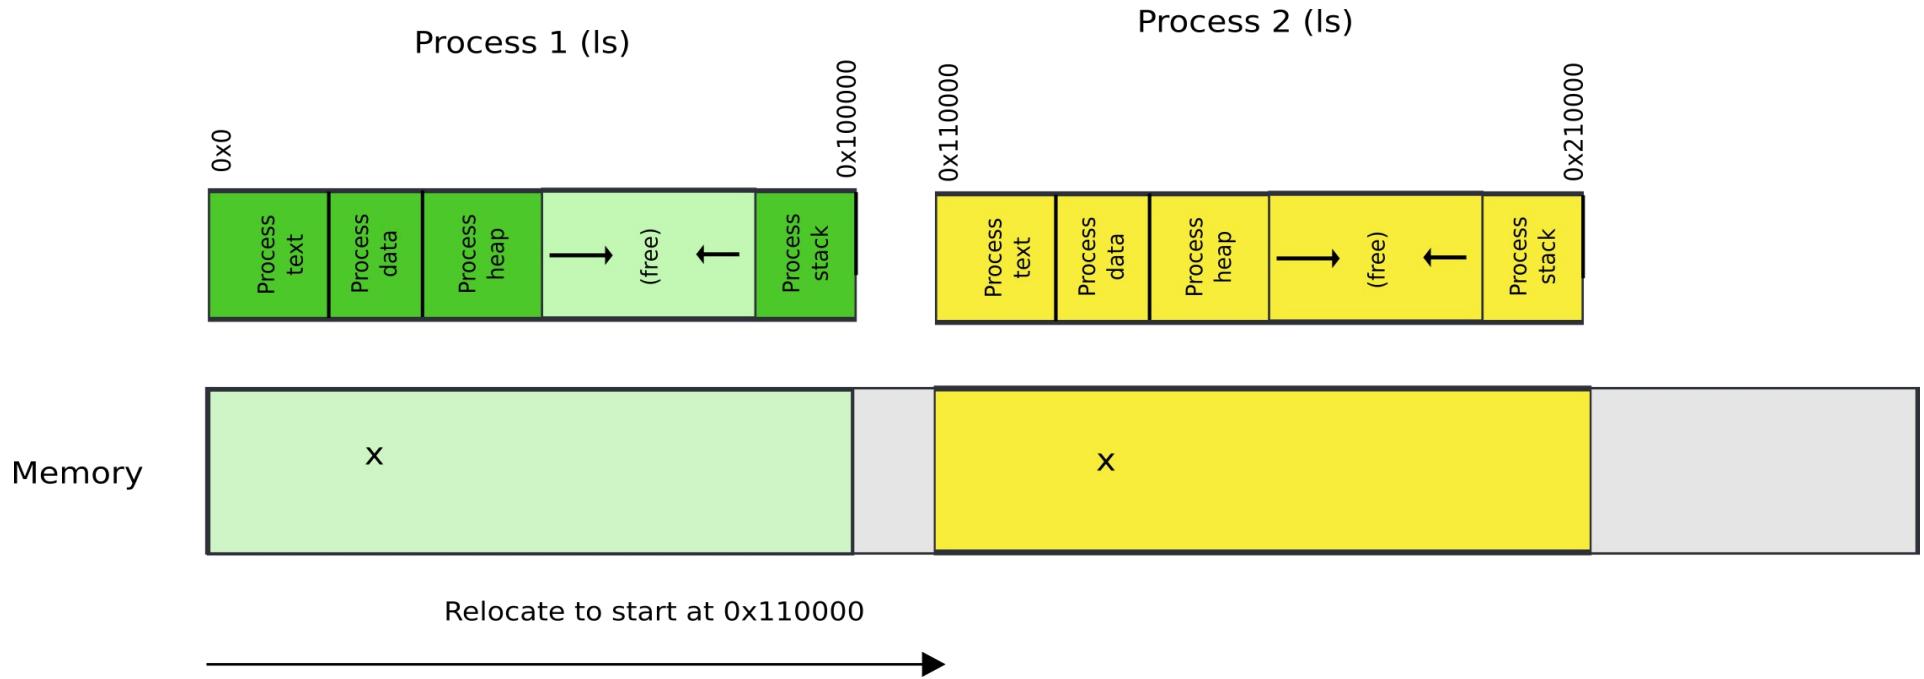


- How can we do this?

# Relocation

- One way to achieve this is to relocate program at different addresses
  - Remember relocation (from linking and loading)

# Relocate binaries to work at different addresses



- One way to achieve this is to relocate program at different addresses
  - Remember relocation (from linking and loading)
  - **This works! But not ideal.**
- What is the problem?

- One way to achieve this is to relocate program at different addresses
  - Remember relocation (from linking and loading)
  - **This works! But not ideal.**
- What is the problem?
  - Isolation can be enforced in software
  - Software Fault Isolation (SFI)
    - Google NaCl (Chrome Sandbox)
    - WASM (Web Assembly, another sandbox standard)

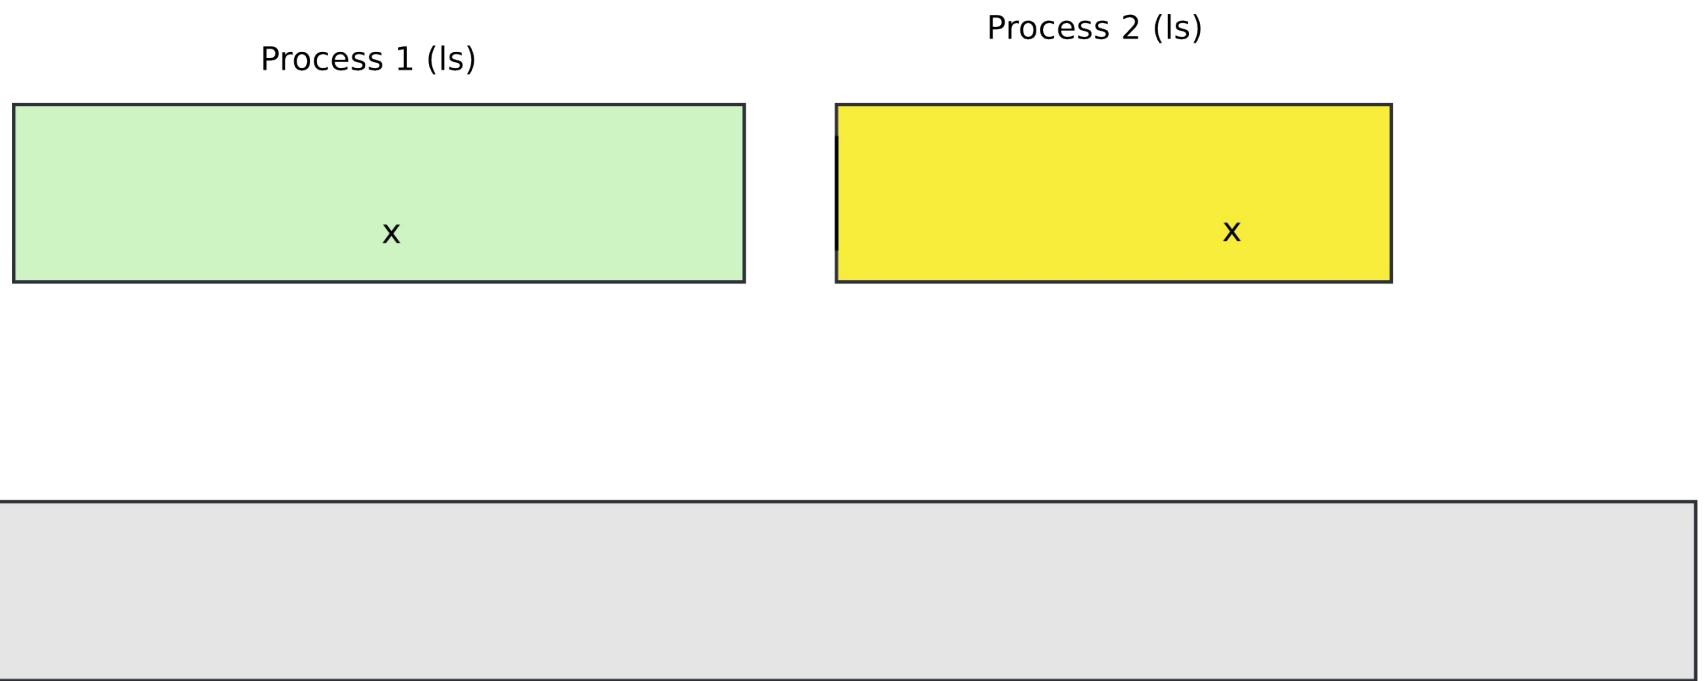
- Another way is to ask for hardware support

# Segmentation

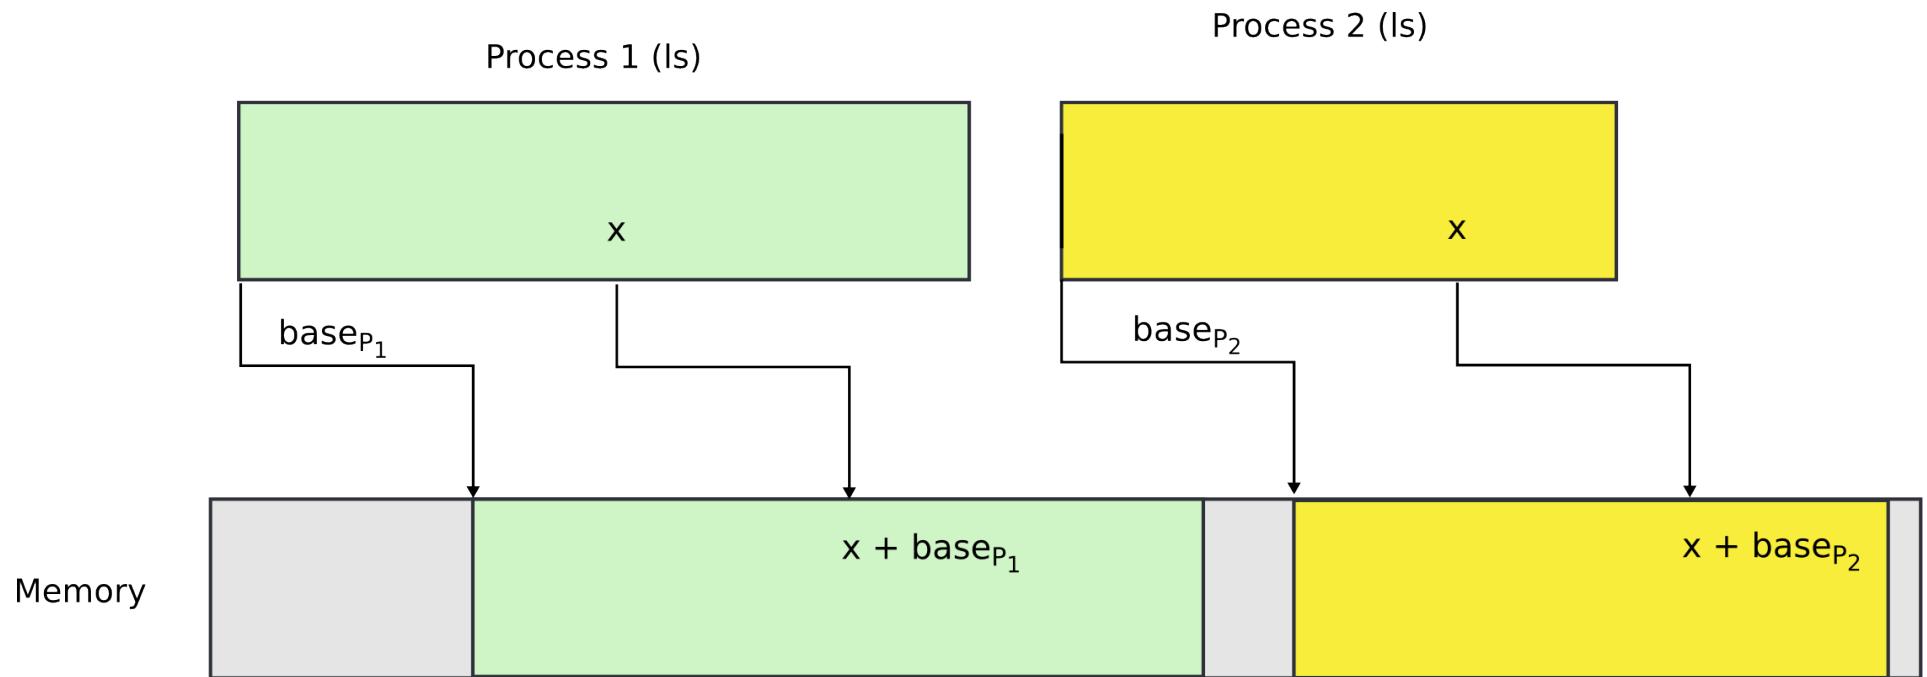
# What are we aiming for?

- Illusion of a private address space
  - Identical copy of an address space in multiple programs
  - Simplifies software architecture
    - One program is not restricted by the memory layout of the others

# Two processes, one memory?



# Two processes, one memory?



- We want hardware to add base value to every address used in the program

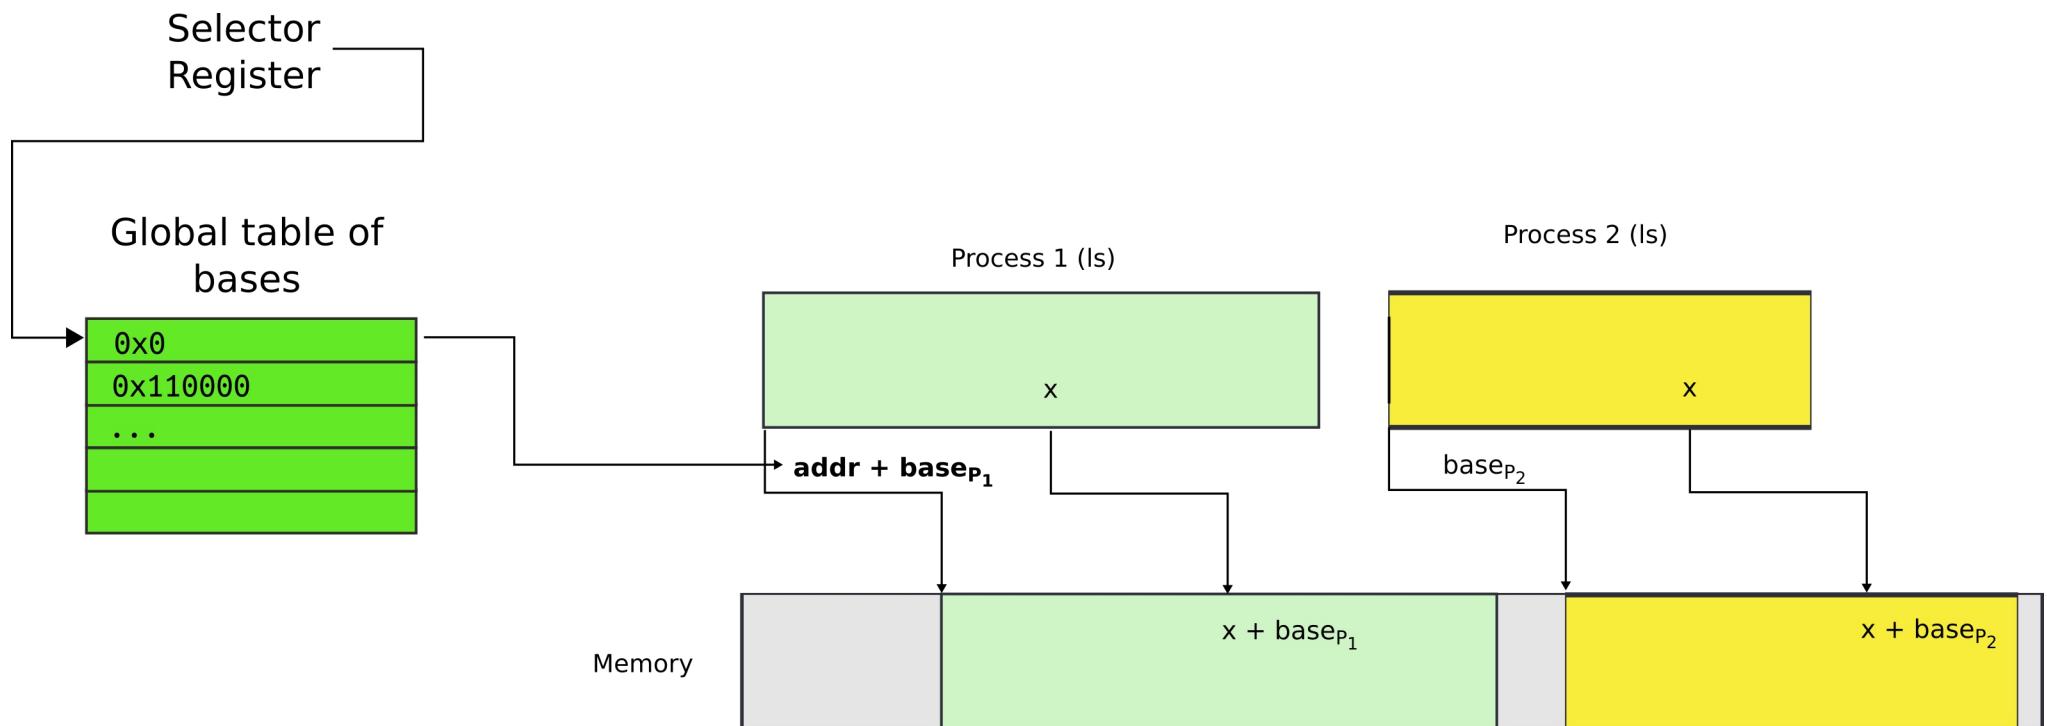
# Seems easy

- One problem
  - Where does this base address come from?

# Seems easy

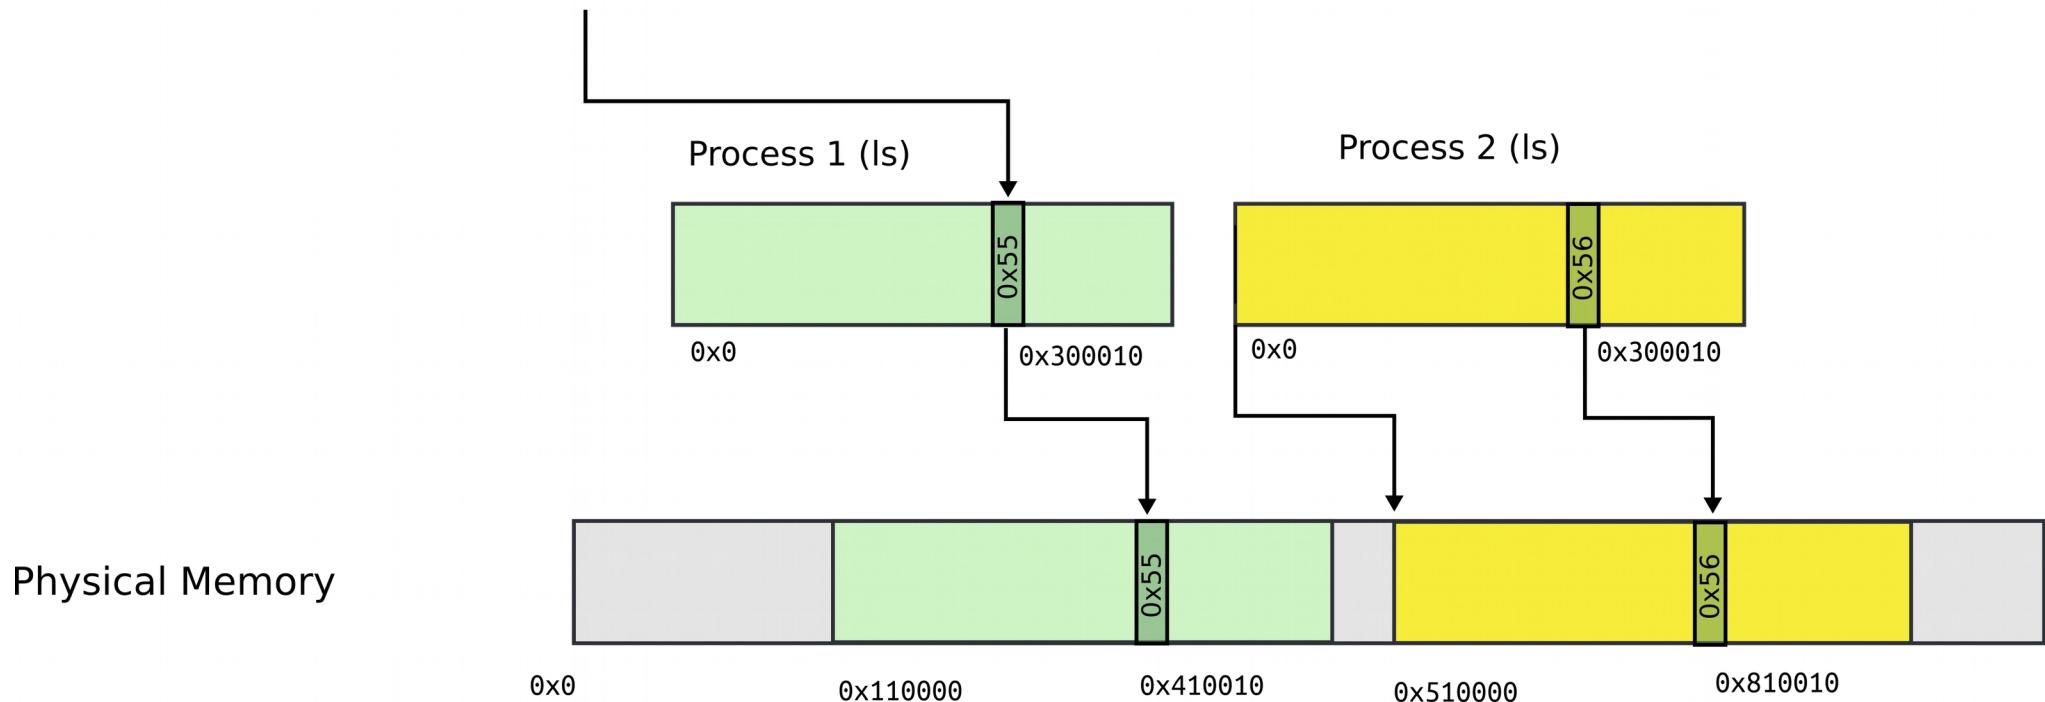
- One problem
  - Where does this base address come from?
  - Hardware can maintain a table of base addresses
    - One base for each process
  - Dedicate a special register to keep an index into that table

- One problem
  - Where does this base address come from?
  - Hardware can maintain a table of base addresses
    - One base for each process
  - Dedicate a special register to keep an index into that table



# Segmentation: example

```
mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0x0  
EBX = 0x300010
```

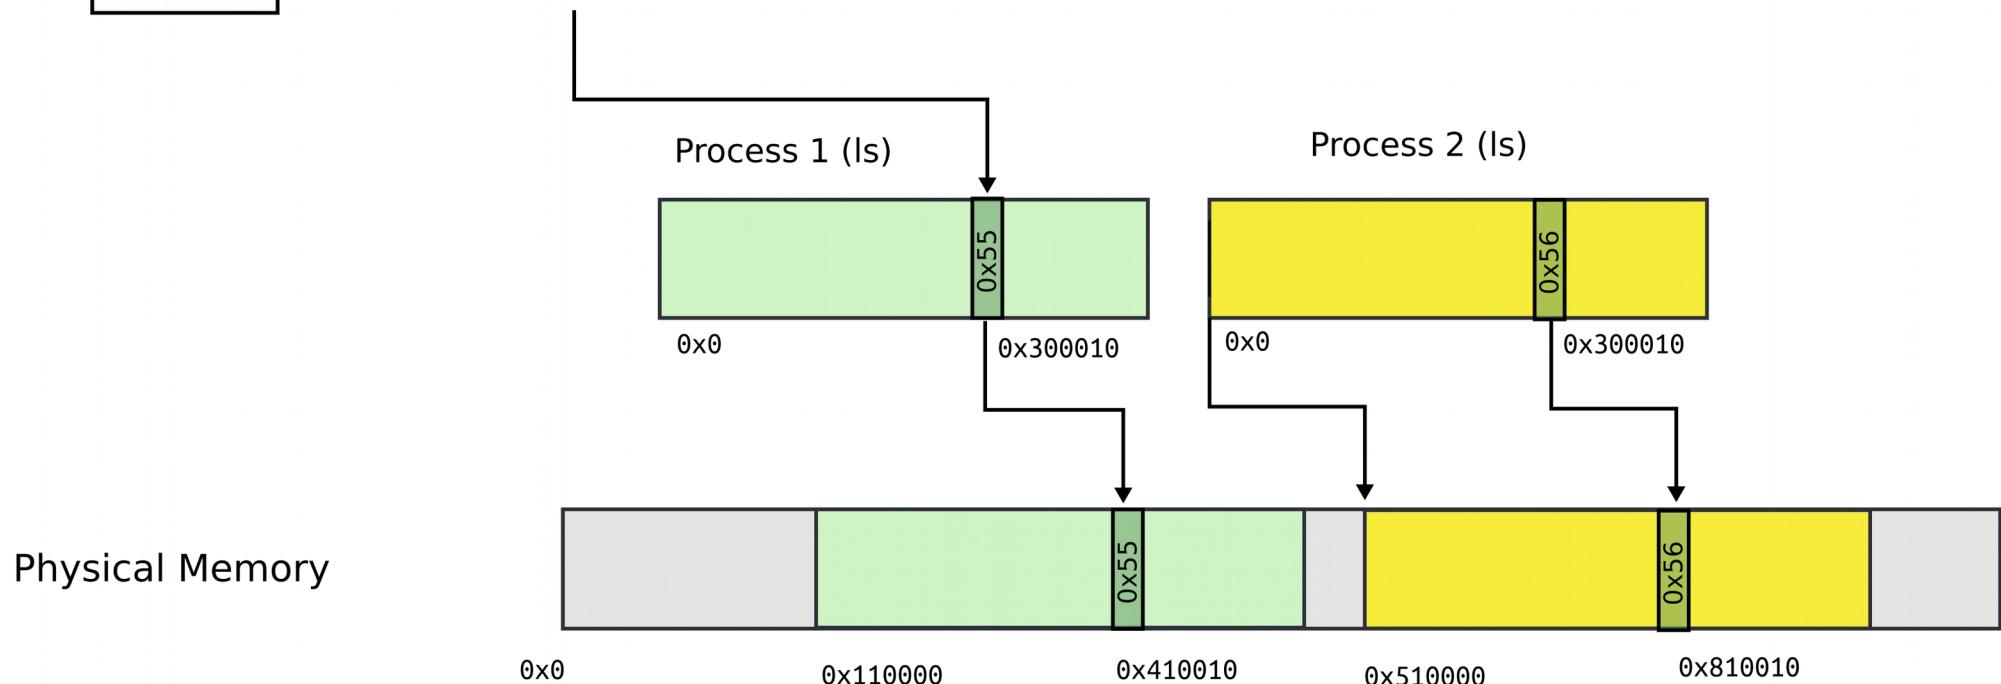


# Segmentation: address consists of two parts

**Segment register  
(CS, SS, DS, ES, FS, GS)**

0x1

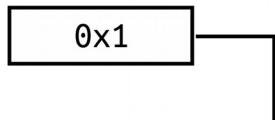
mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0x0  
**EBX = 0x300010, DS = 0x1**



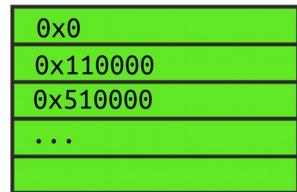
- Segment register contains segment selector
- General registers contain offsets
- Intel calls this address: “logical address”

# Segmentation: Global Descriptor Table

Segment register  
(CS, SS, DS, ES, FS, GS)

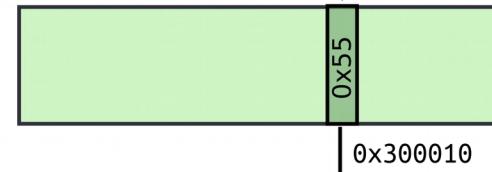


Global Descriptor Table  
(table of segment sizes and bases)



mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0x0  
EBX = 0x300010, DS = 0x1

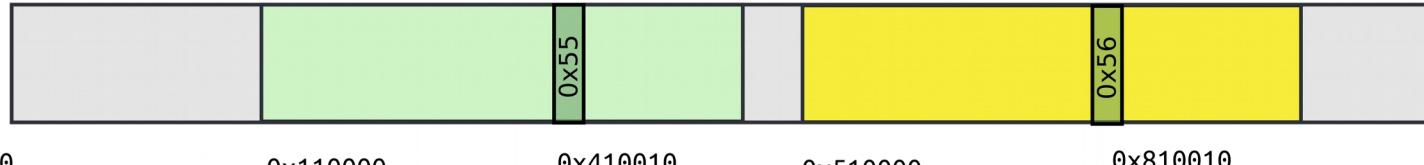
Process 1 (ls)



Process 2 (ls)

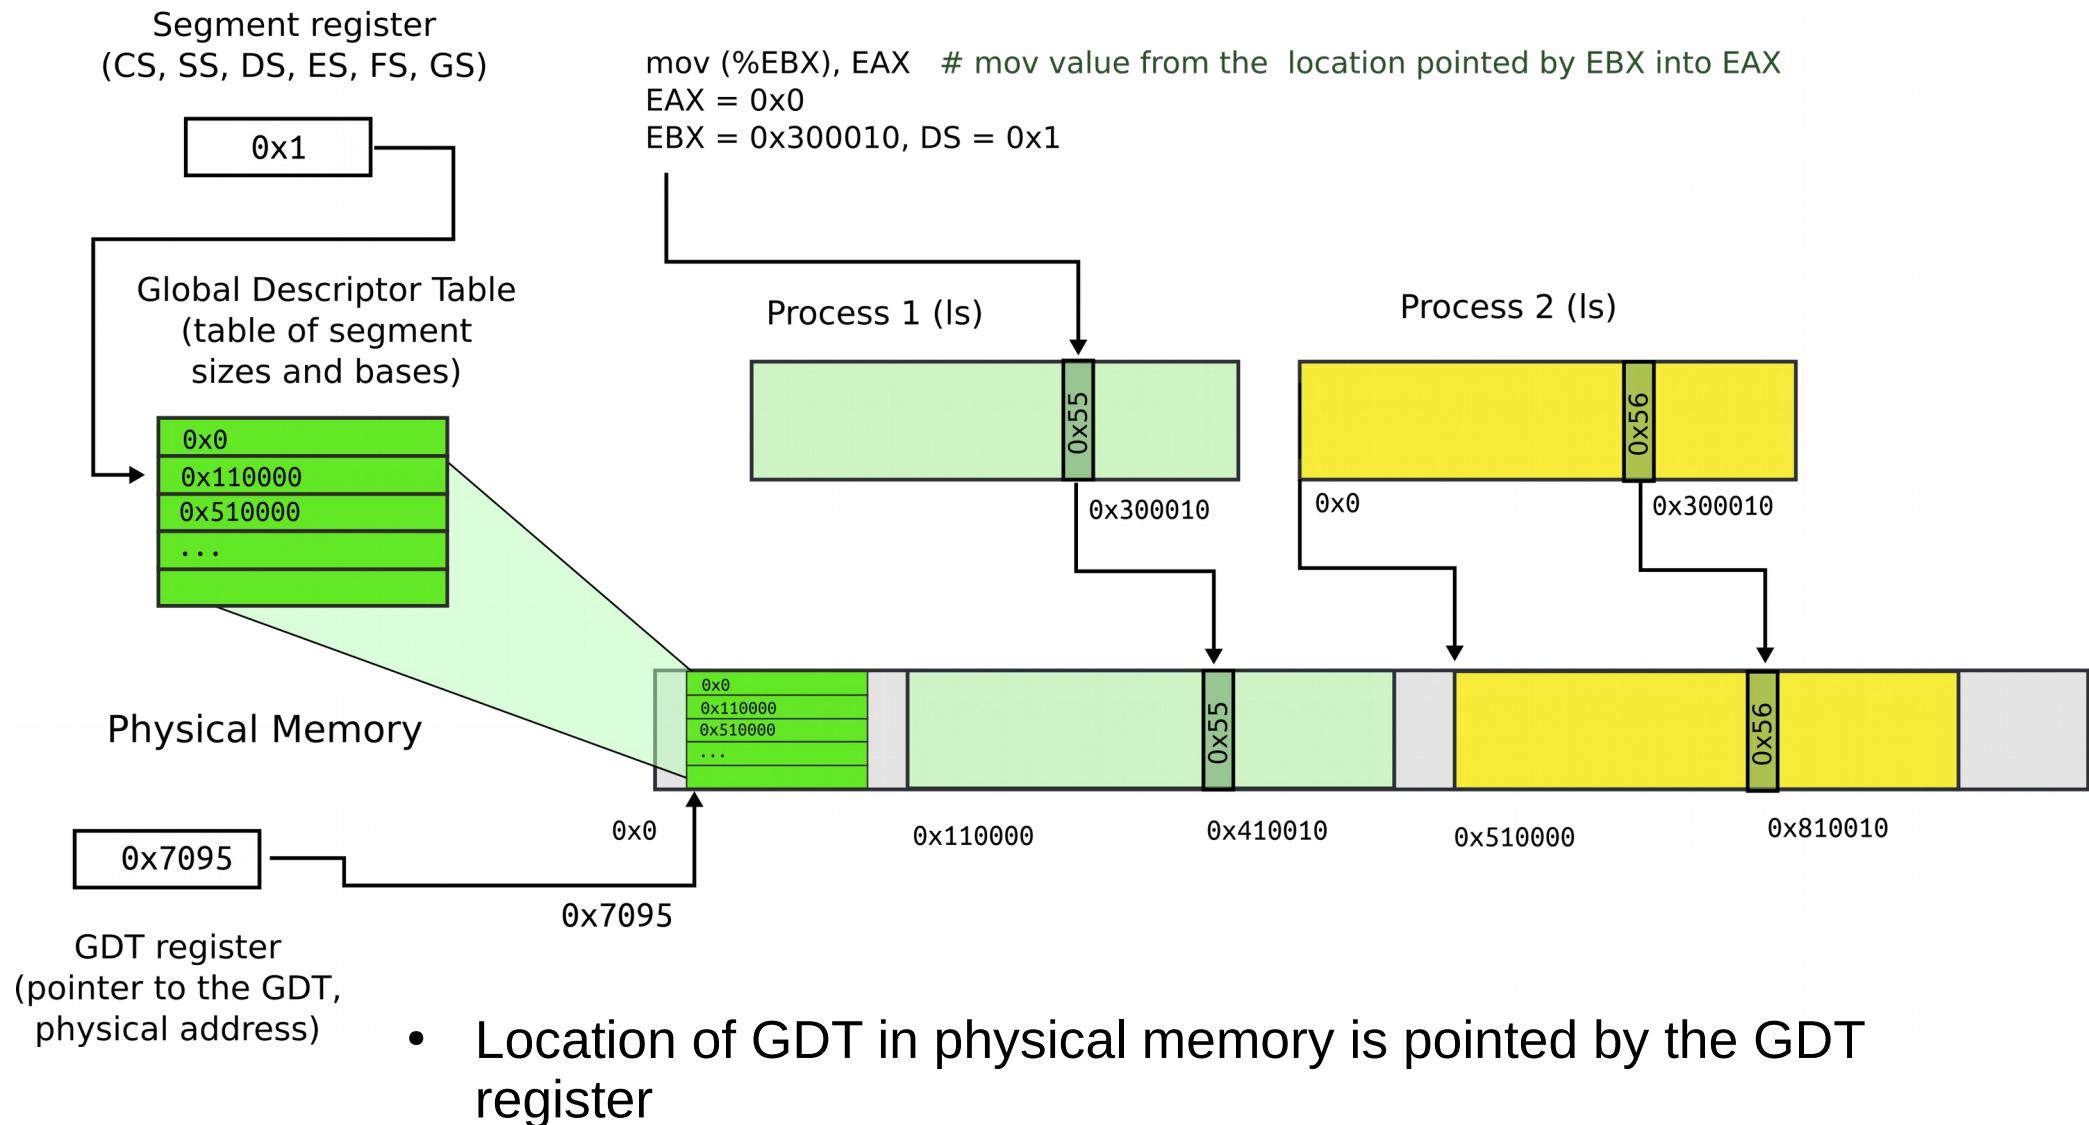


Physical Memory

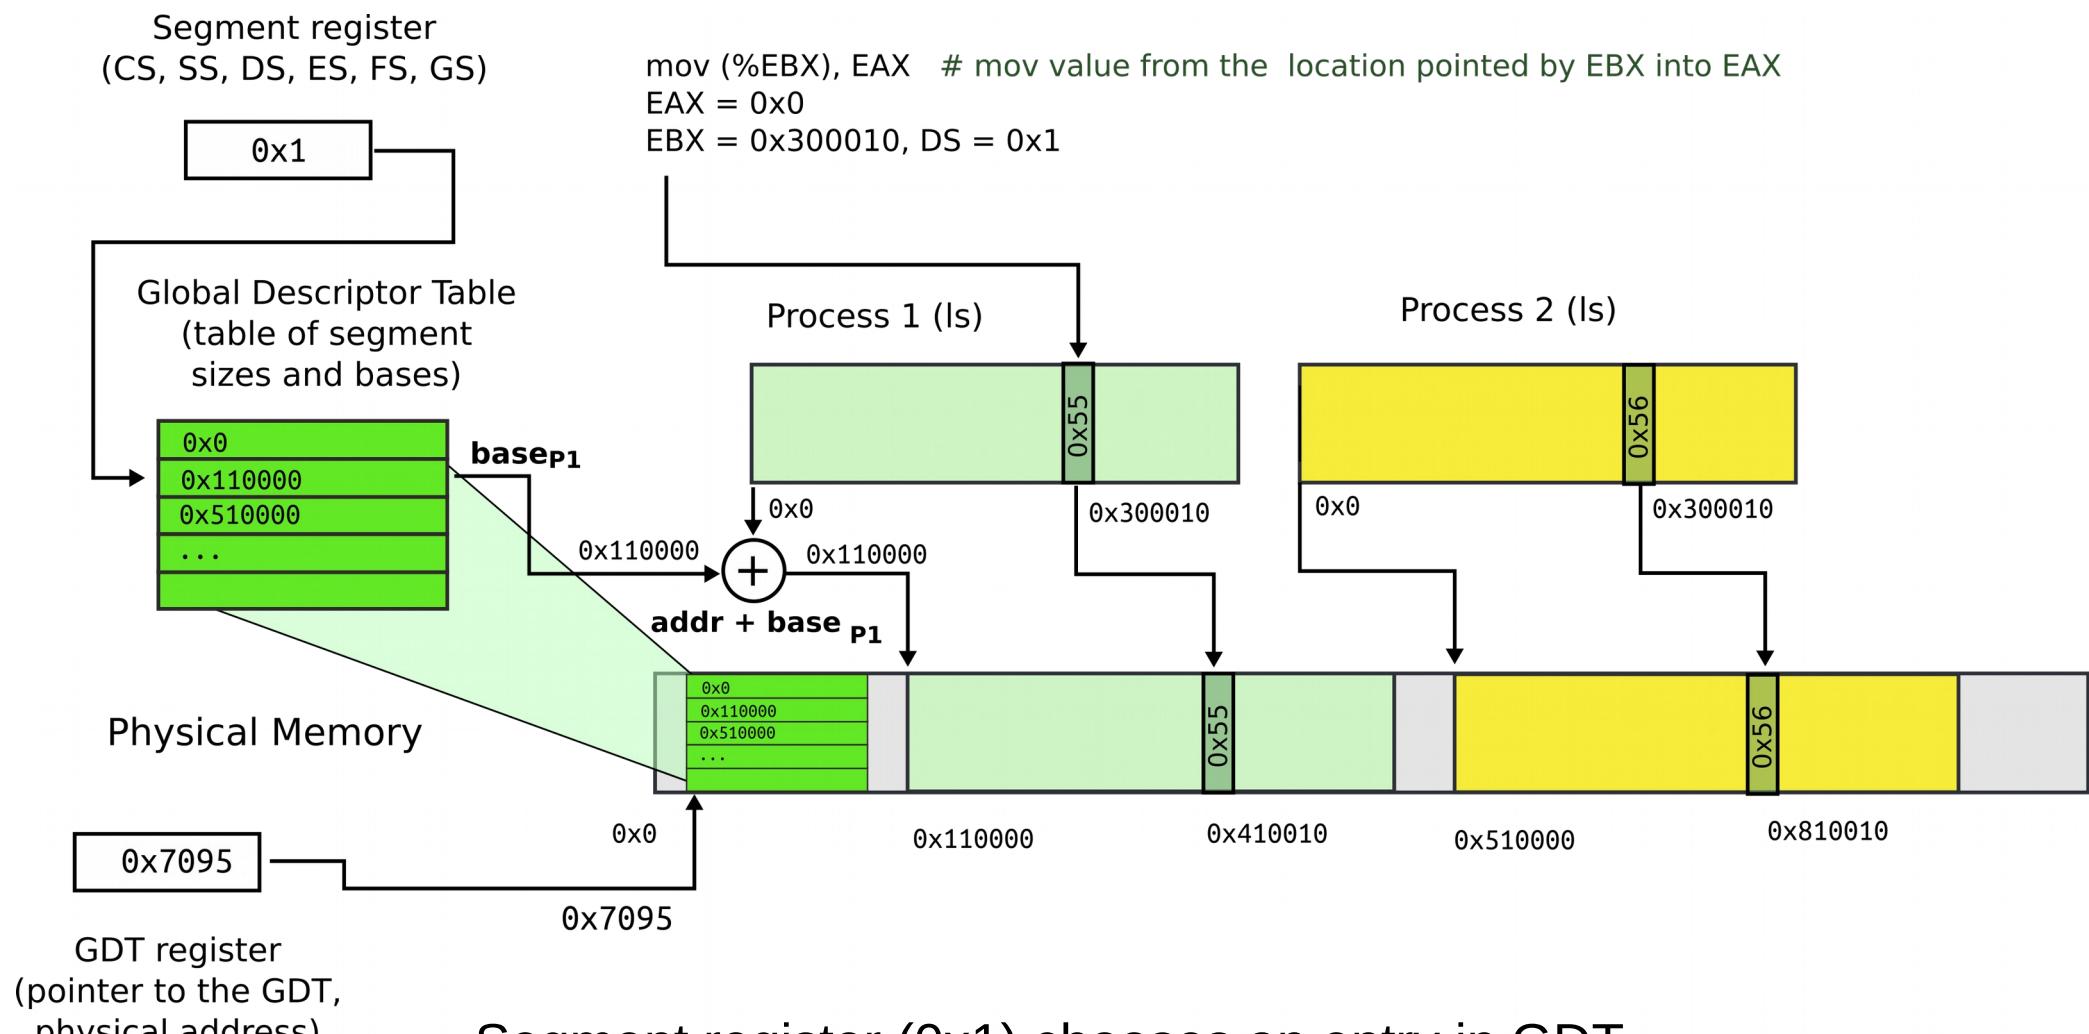


- GDT is an array of segment descriptors
  - Each descriptor contains base and limit for the segment
  - Plus access control flags

# Segmentation: Global Descriptor Table

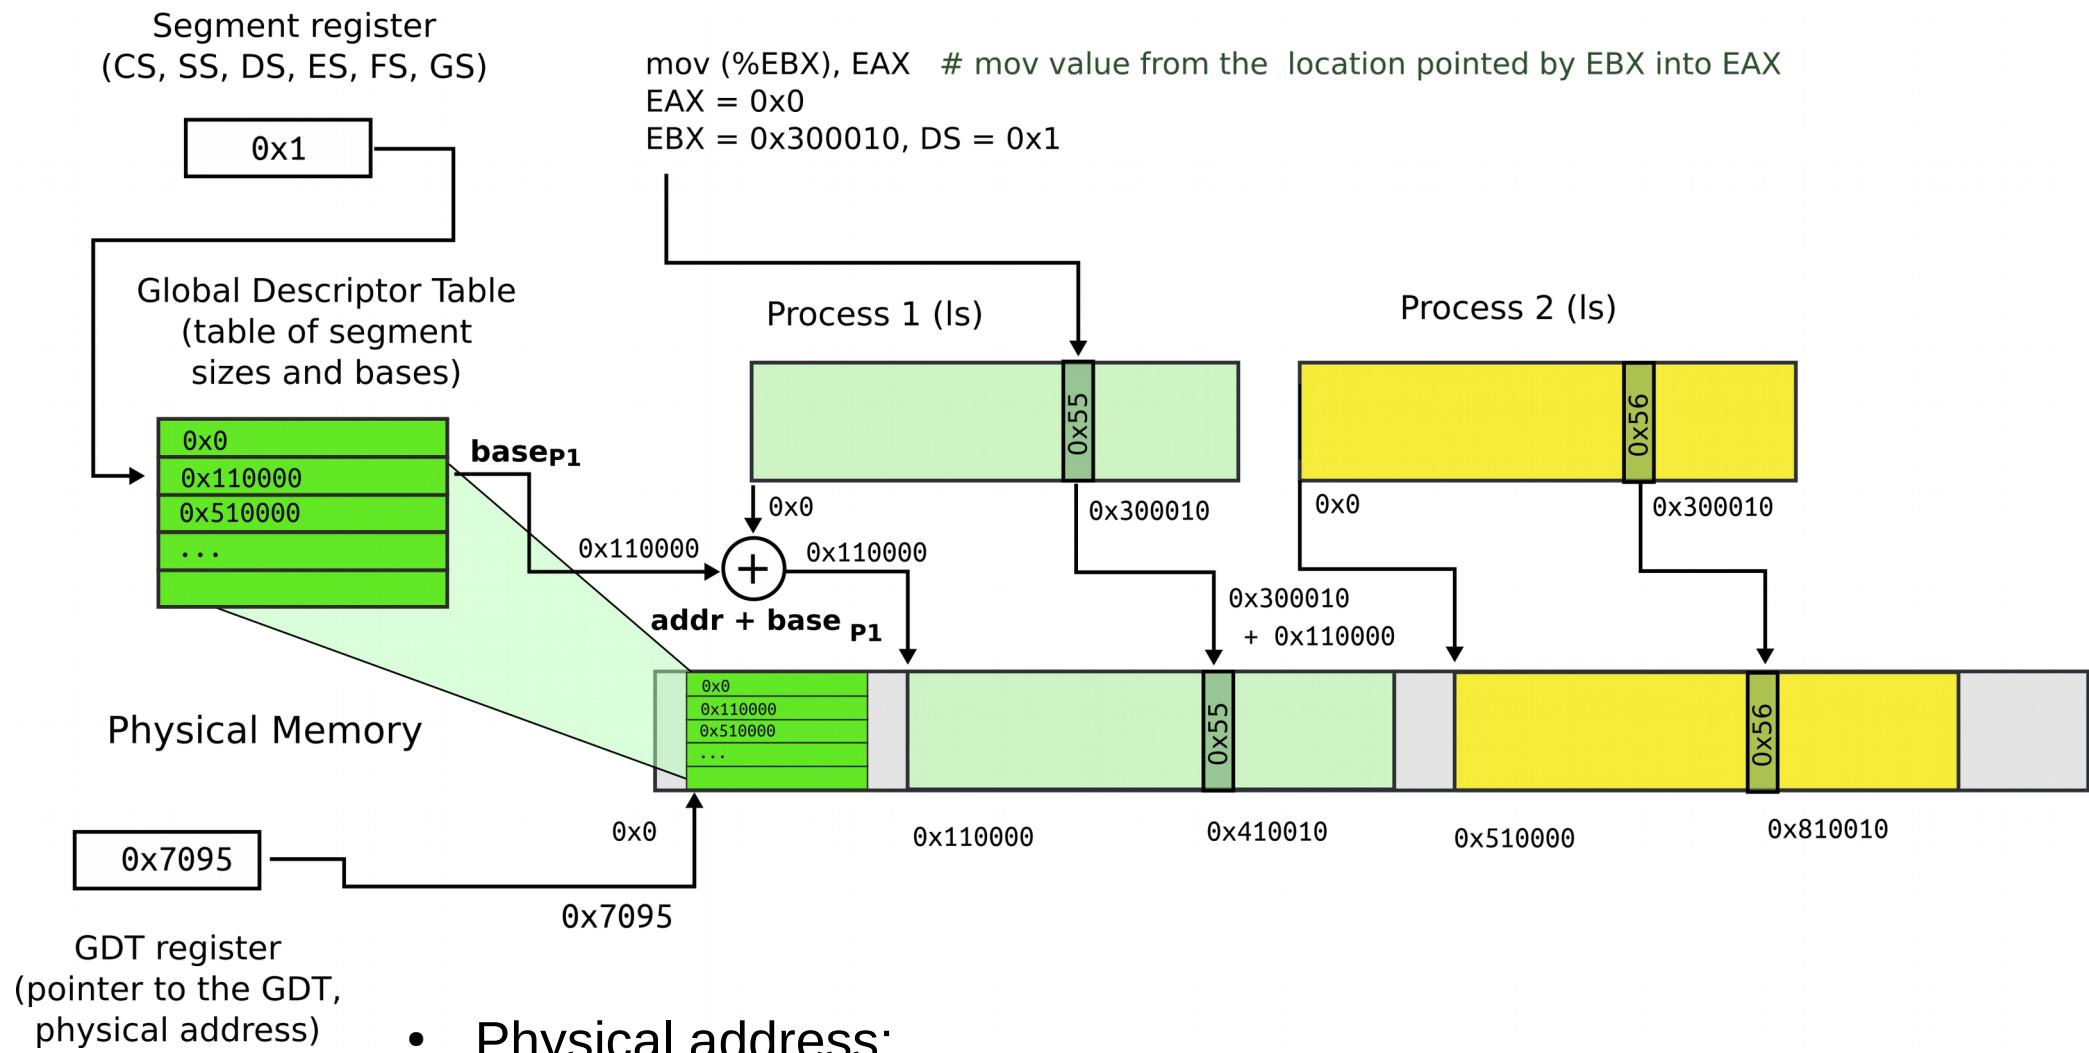


# Segmentation: base + offset



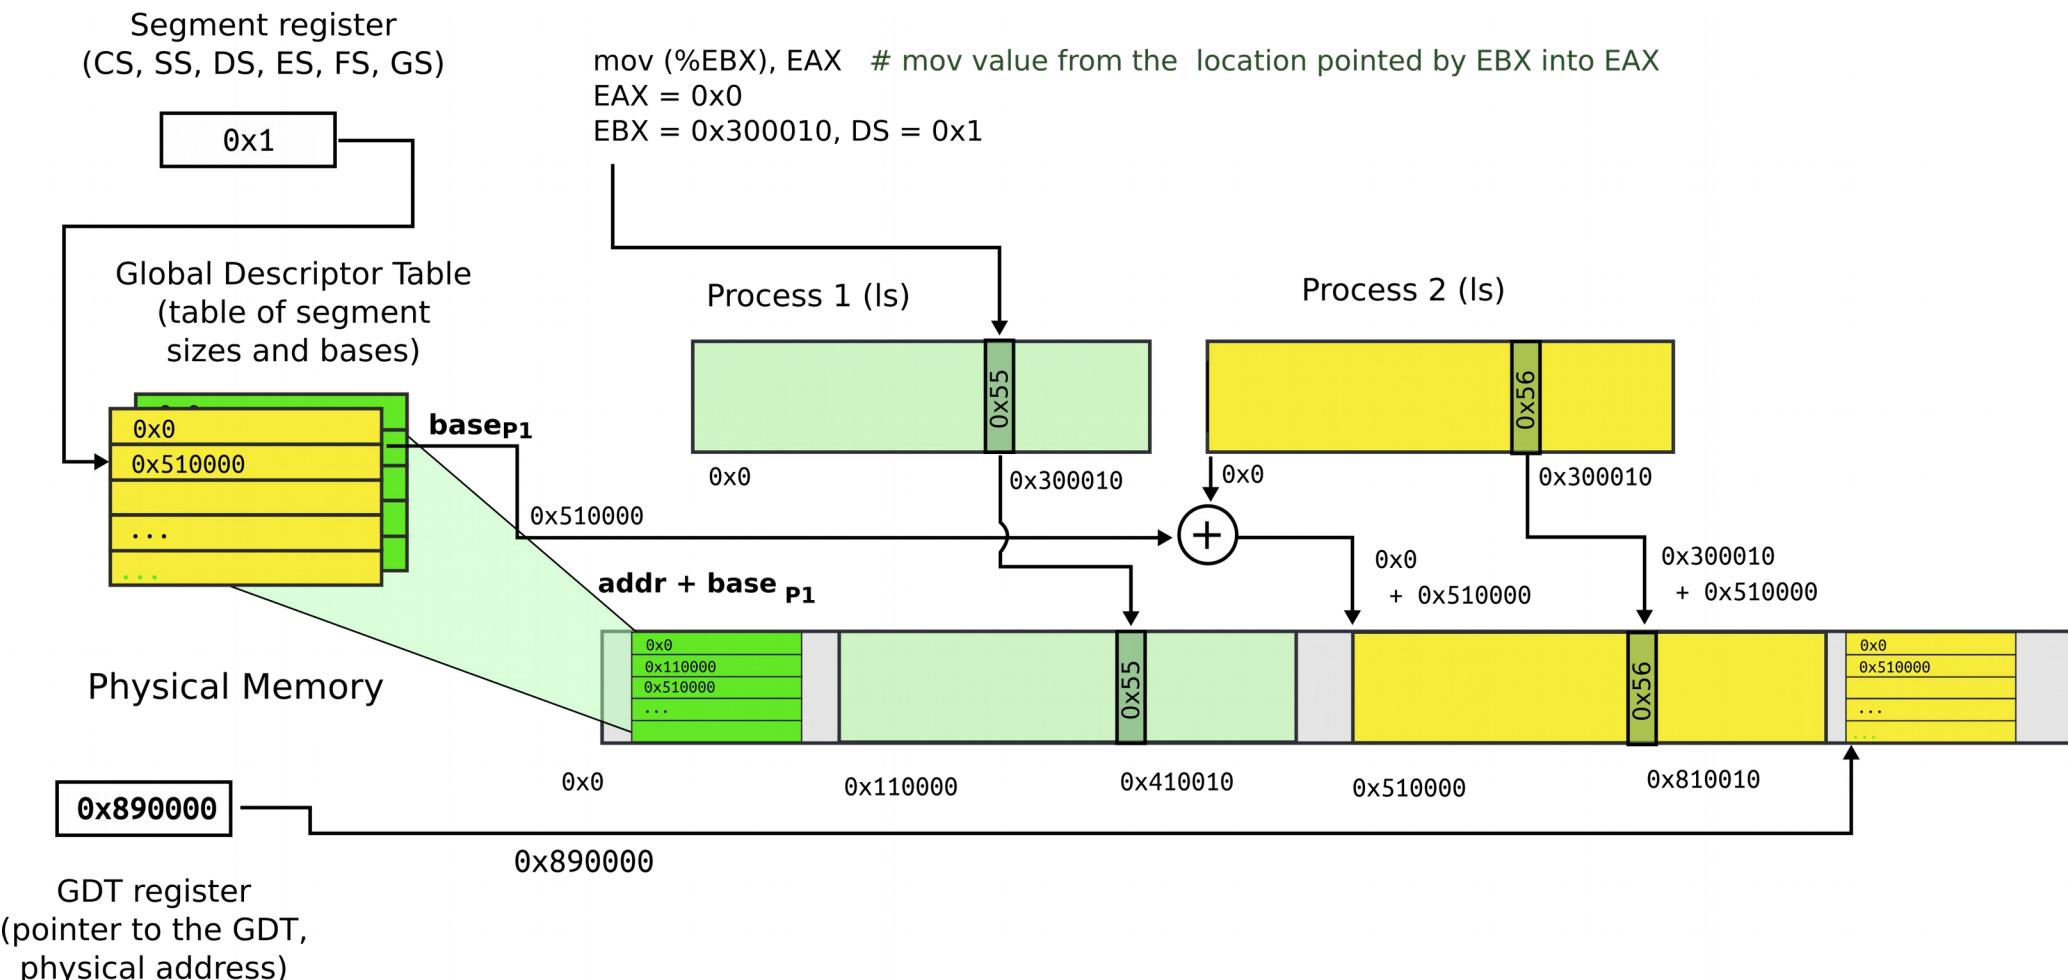
- Segment register (0x1) chooses an entry in GDT
  - This entry contains base of the segment (0x110000) and limit (size) of the segment (not shown)

# Segmentation: base + offset



- Physical address:
    - $0x410010 = 0x300010$  (offset) +  $0x110000$  (base)
    - Intel calls this address “linear”

# Segmentation: process 2

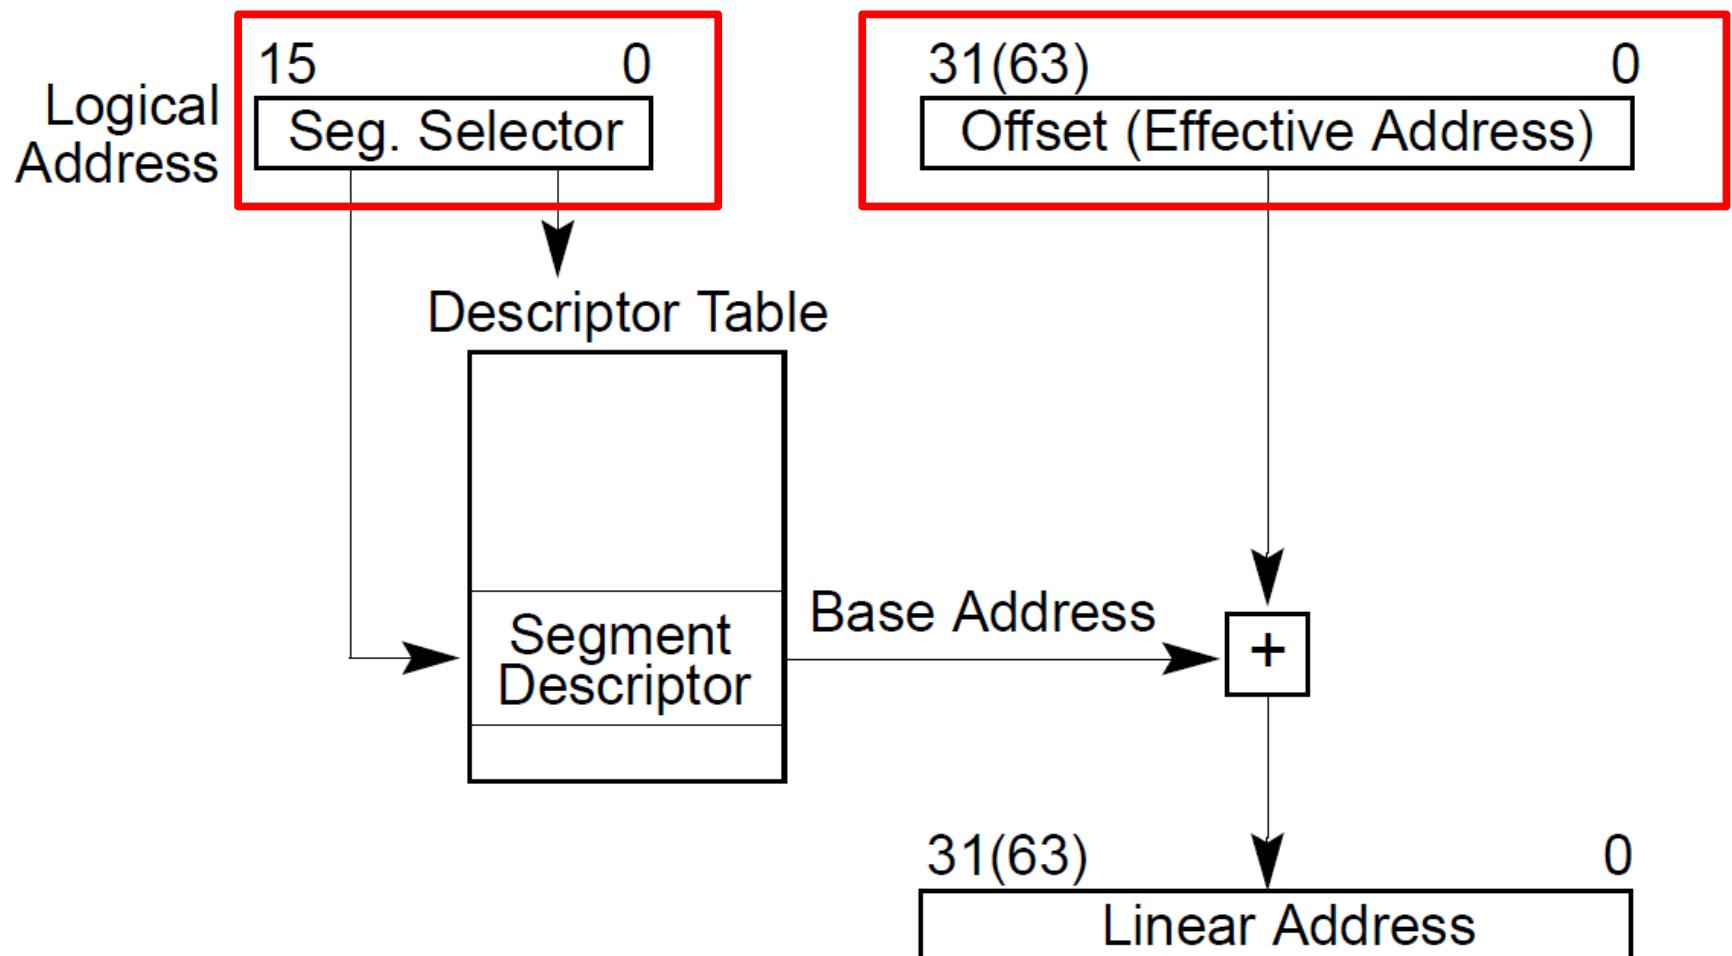


- Each process has a private GDT
    - OS will switch between GDTs

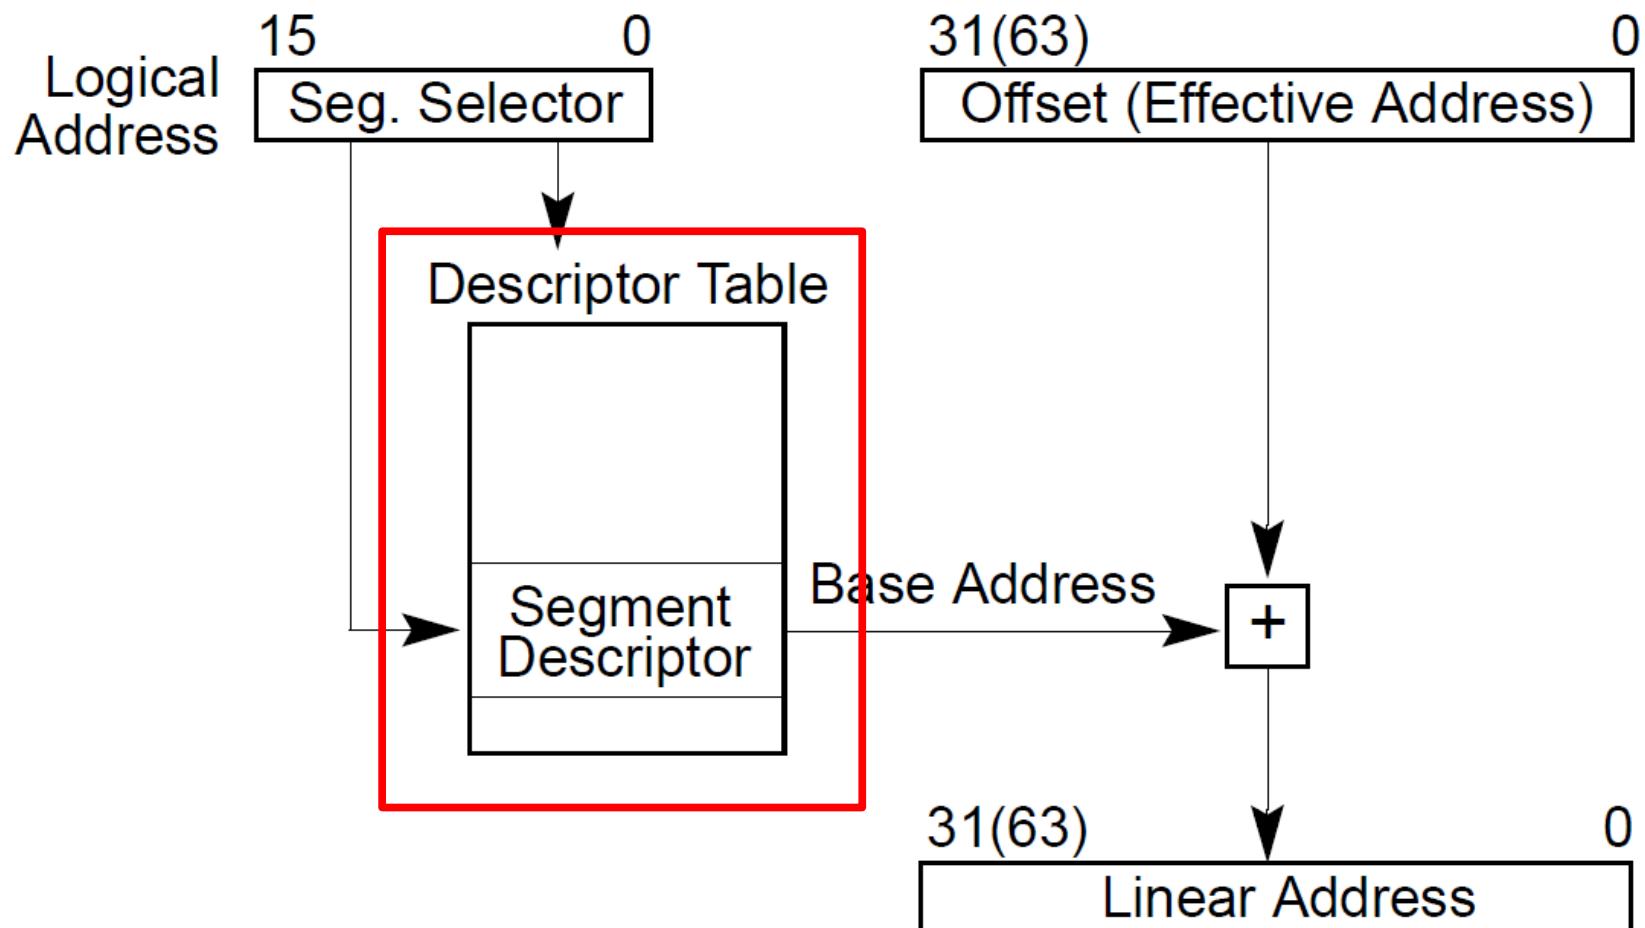
New addressing mode:  
“logical addresses”

# All addresses are logical address

- They consist of two parts
  - Segment selector (16 bit) + offset (32 bit)



- Segment selector (16 bit)
  - Is simply an index into an array (Descriptor Table)
  - That holds segment descriptors
    - Base and limit (size) for each segment



# Elements of the descriptor table are segment descriptors

- Base address

- 0 – 4 GB

- Limit (size)

- 0 – 4 GB

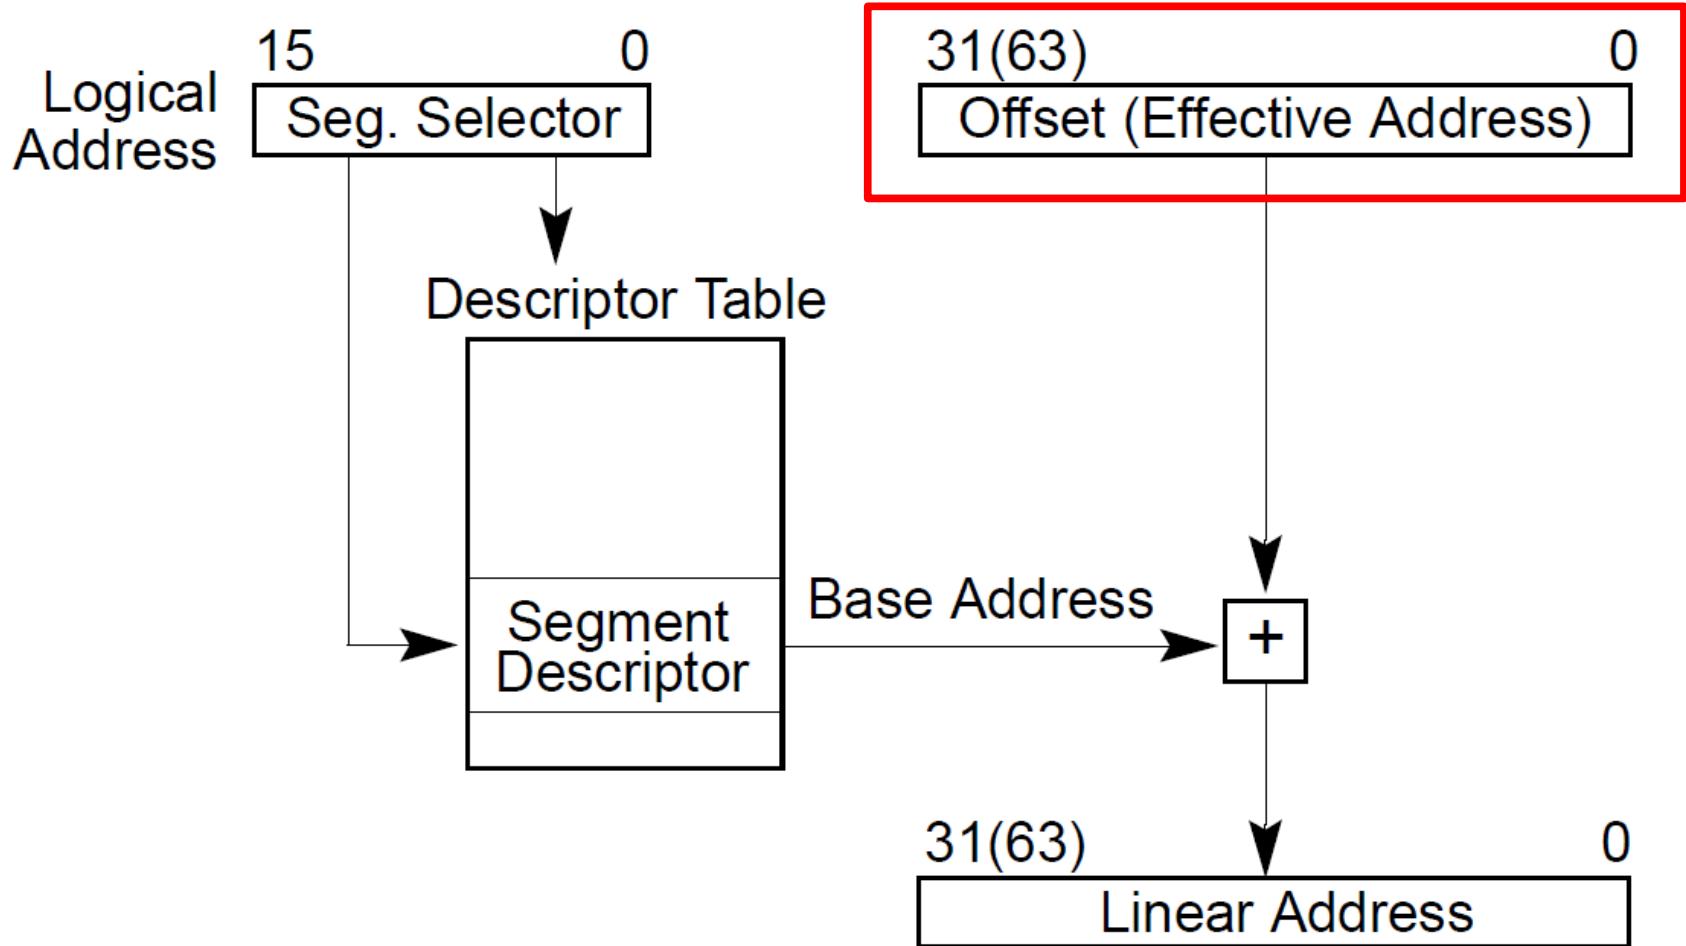
- Access rights

- Executable, readable, writable

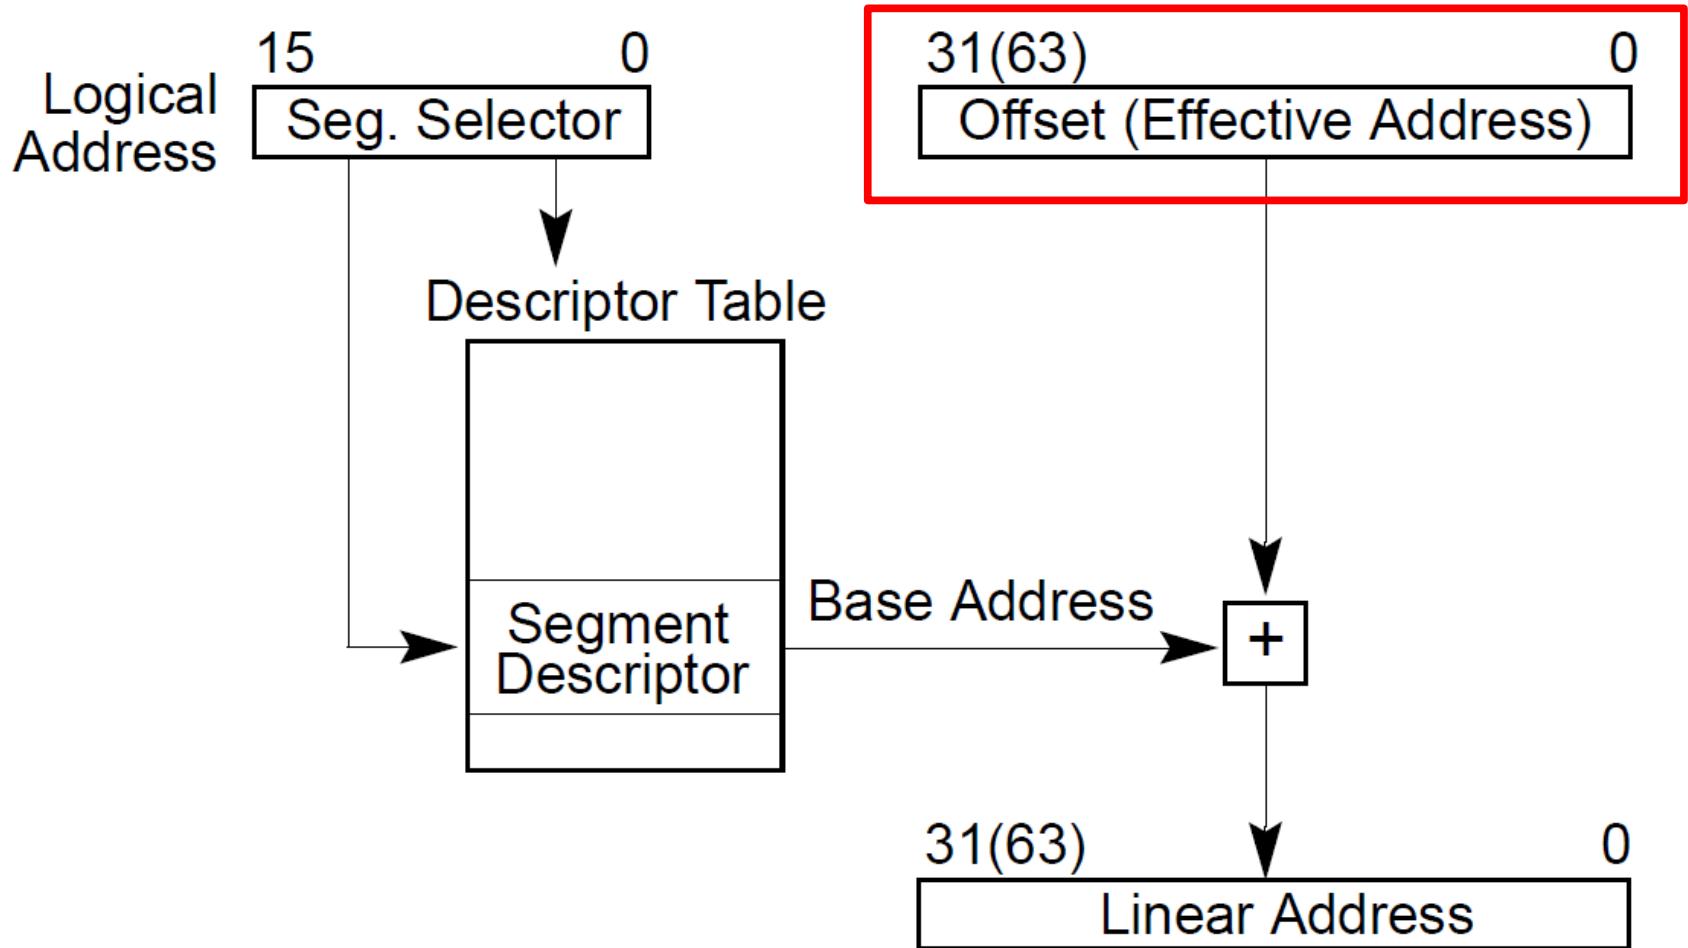
- Privilege level (0 - 3)

Access	Limit
Base Address	

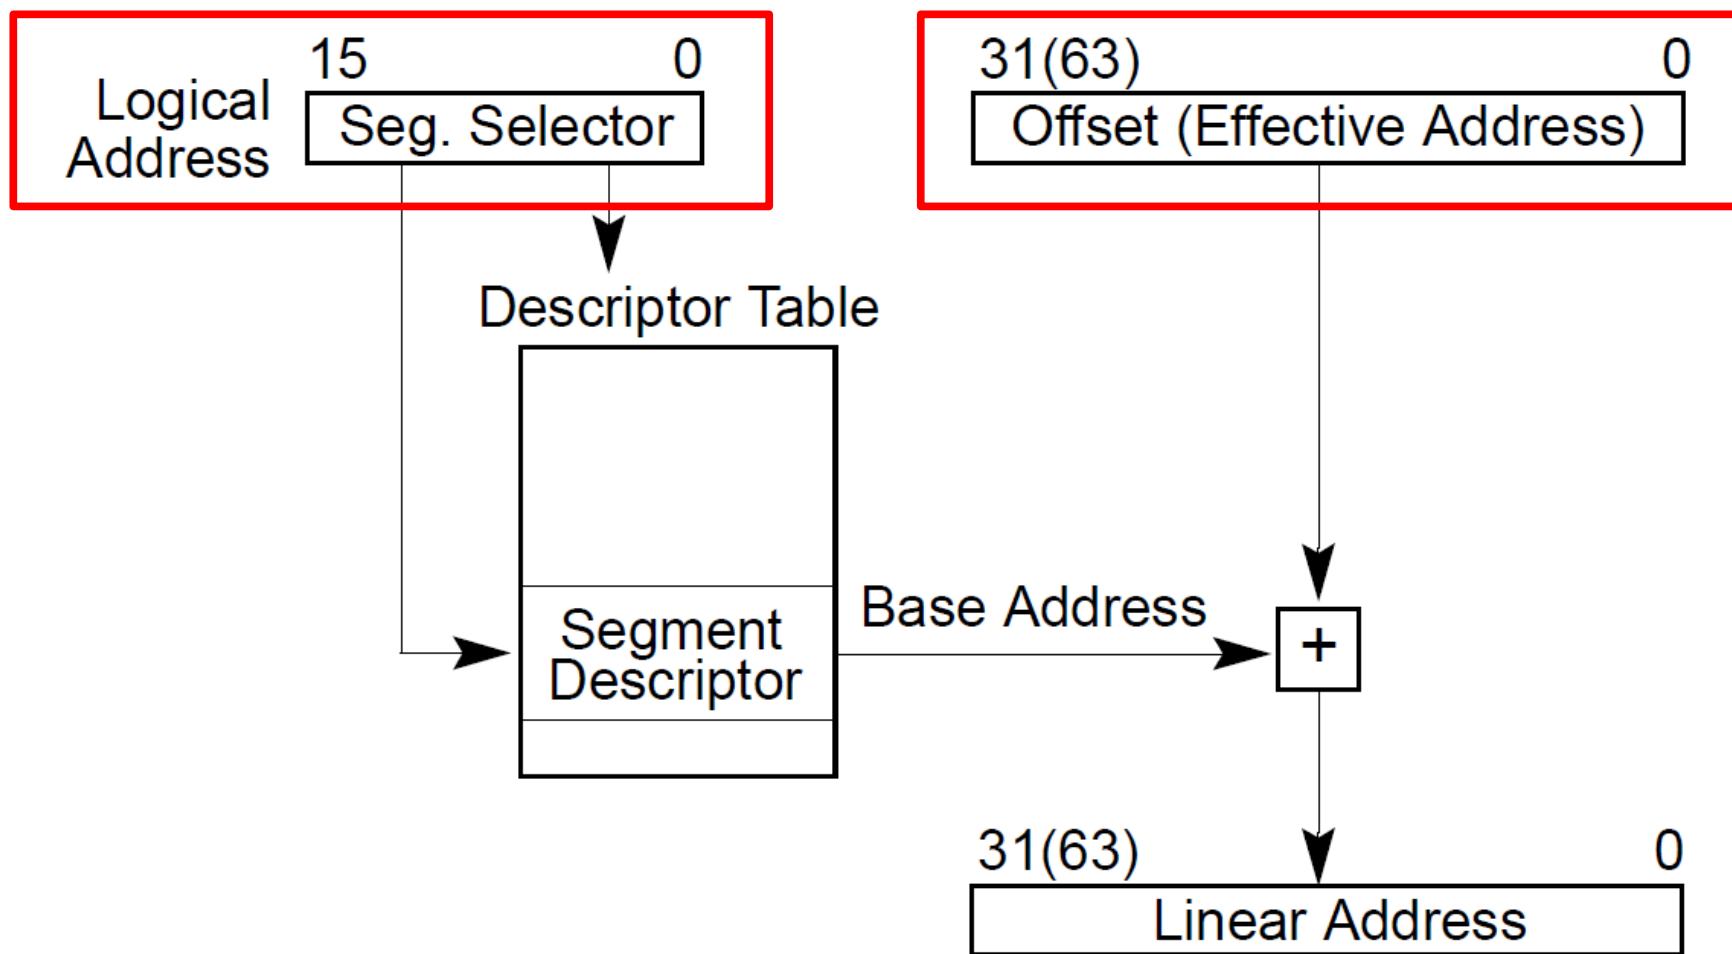
- Offsets into segments (x in our example) or “Effective addresses” are in registers



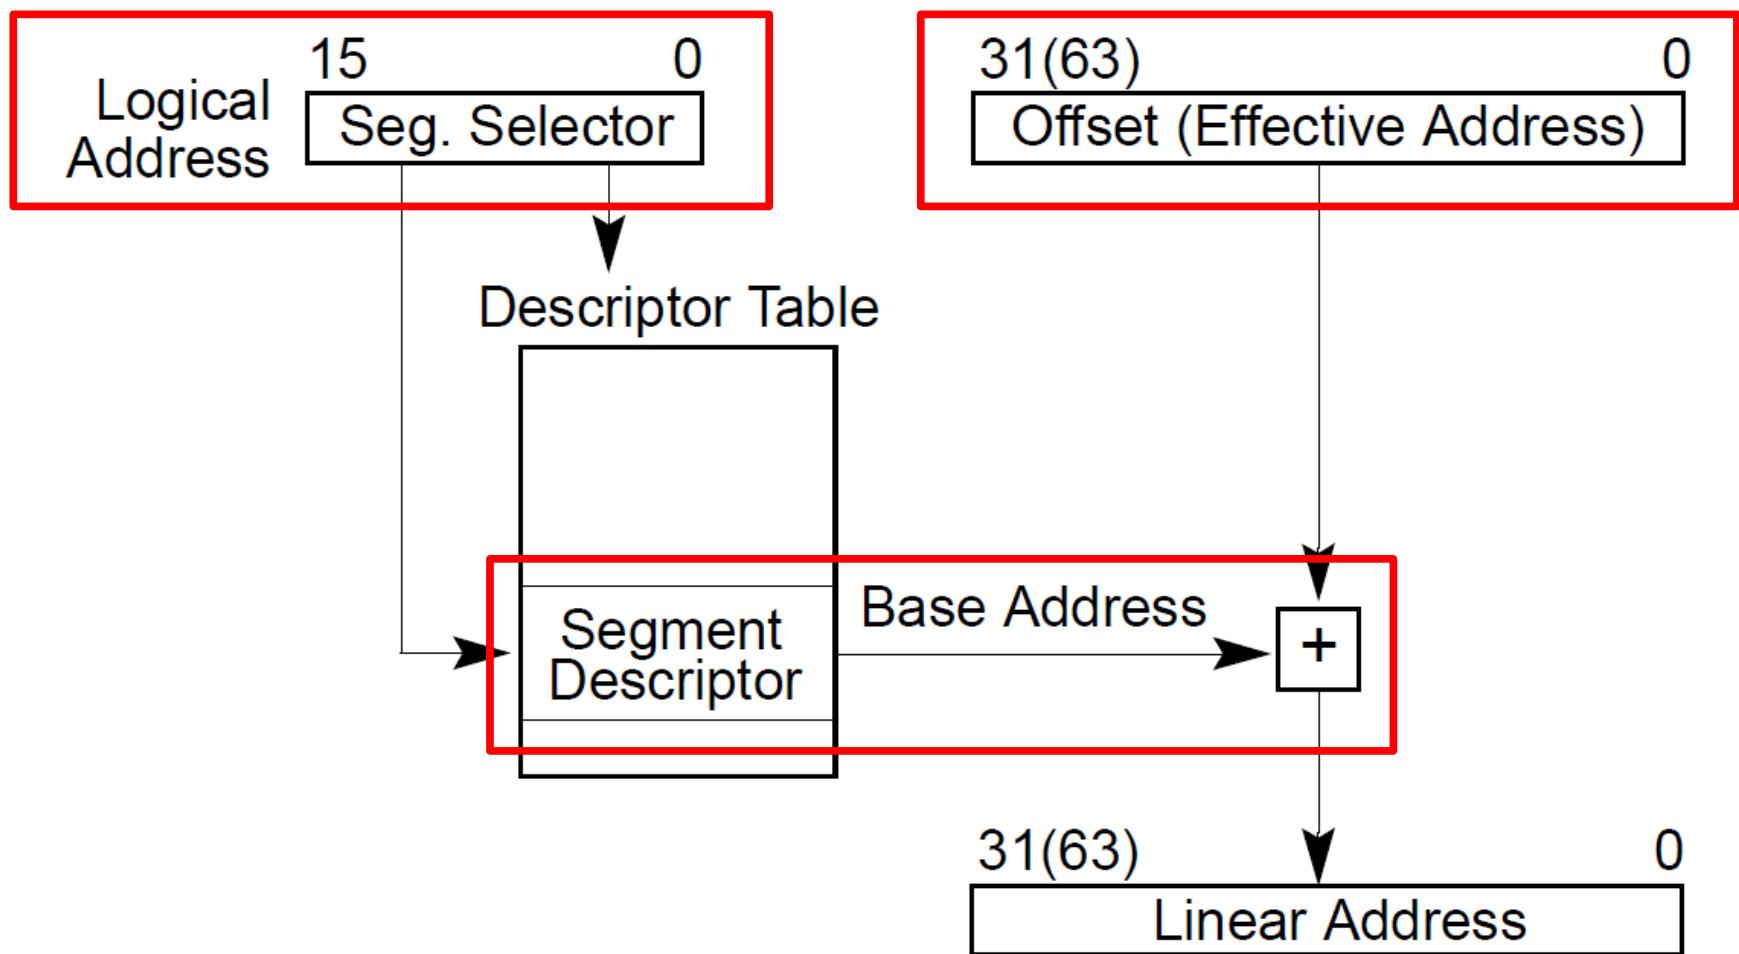
- Logical addresses are translated into physical
  - $\text{Effective address} + \text{DescriptorTable}[selector].Base$



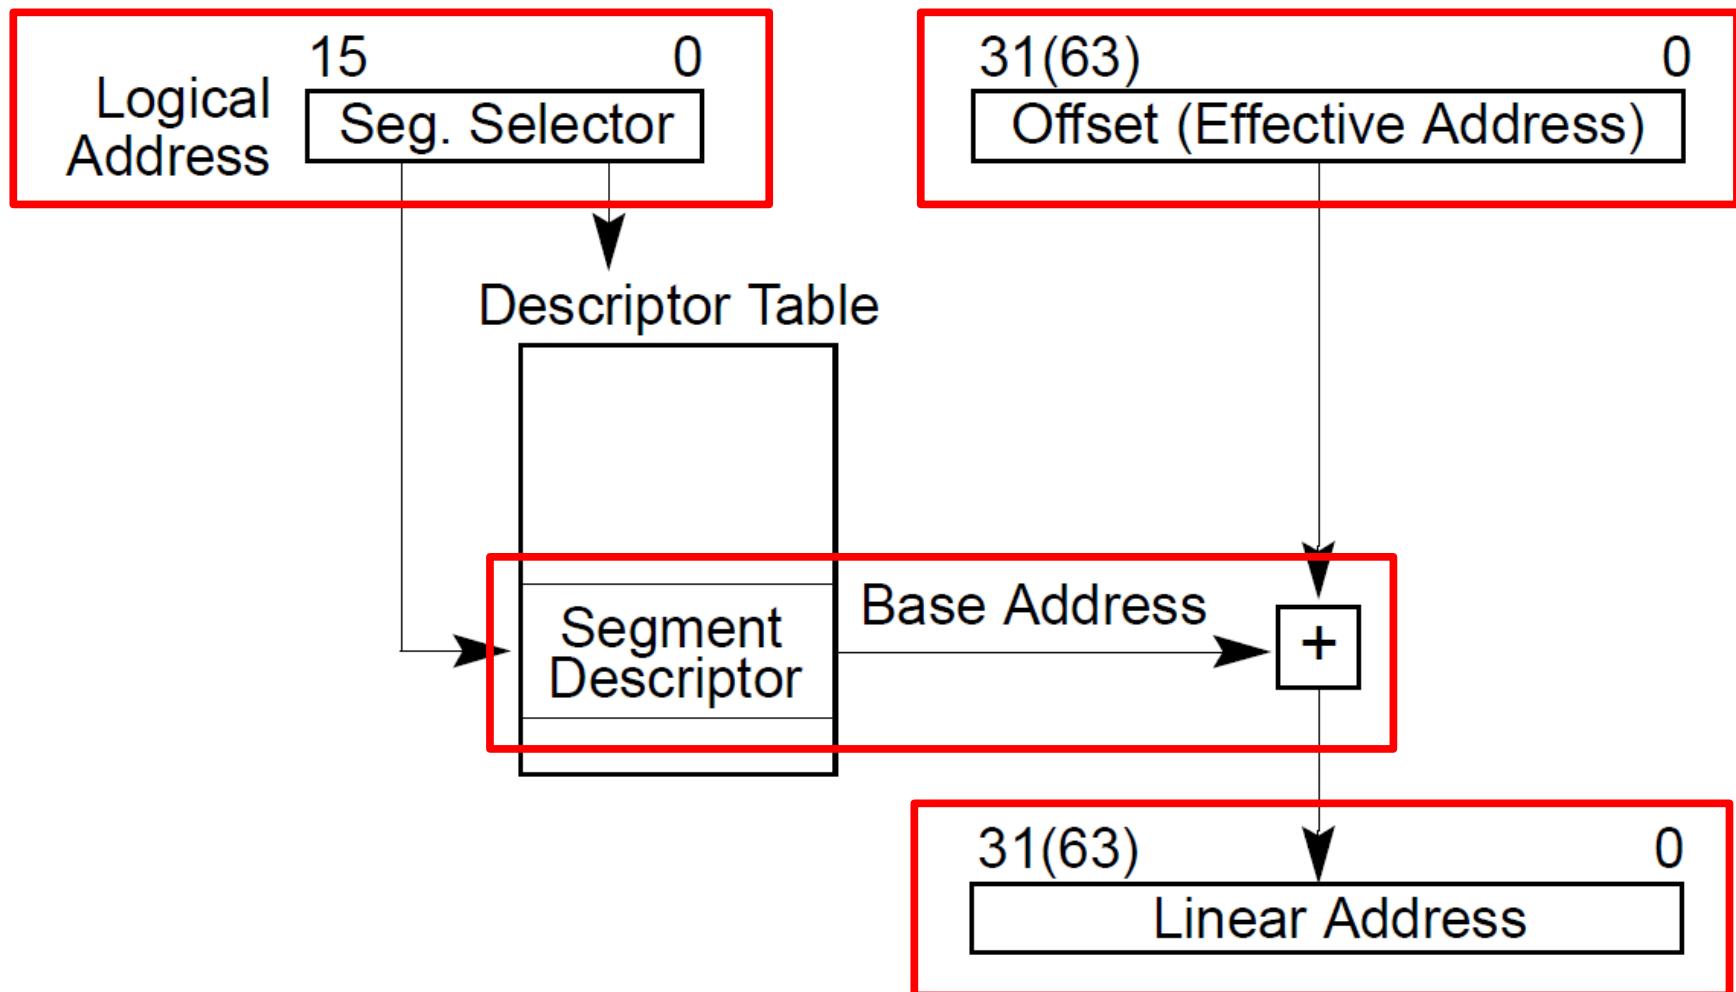
- Logical addresses are translated into physical
  - Effective address + DescriptorTable[selector].Base



- Logical addresses are translated into physical
  - Effective address + DescriptorTable[selector].Base

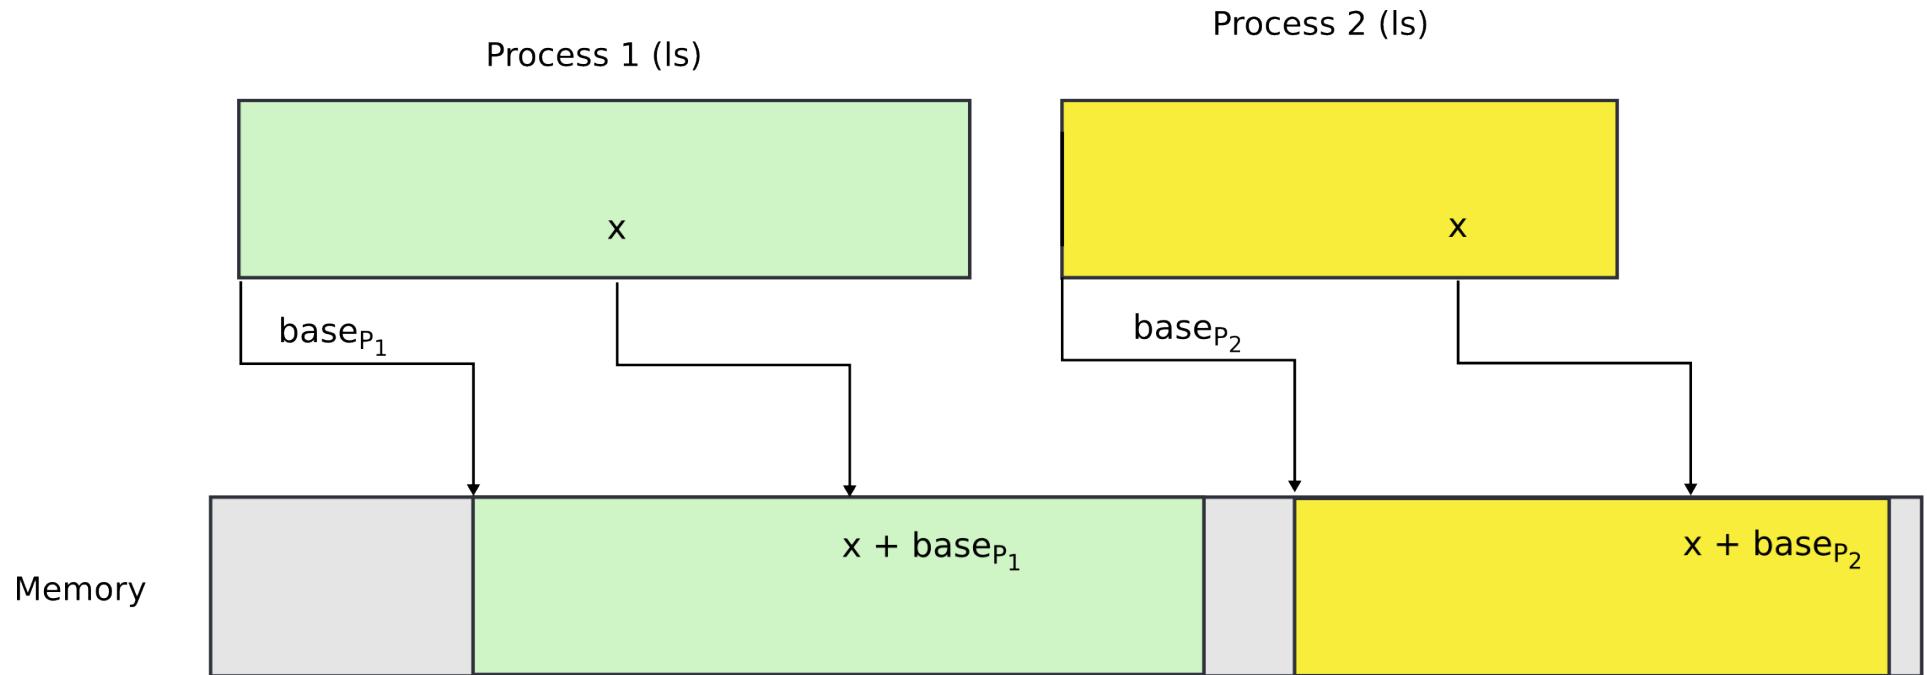


- Logical addresses are translated into physical
  - Effective address + DescriptorTable[selector].Base



- *Physical address =*

$$\text{Effective address} + \text{DescriptorTable}[selector].Base$$
- Effective addresses (or offsets) are in registers
- **Selector is in a special register**



# Segment registers

- Hold 16 bit segment selectors
  - Indexes into GDT
- Segments are associated with one of three types of storage
  - Code
  - Data
  - Stack

# Segmented programming (not real)

```
static int x = 1;           ds:x = 1; // data
int y; // stack           ss:y;      // stack
if (x) {                  if (ds:x) {
    y = 1;                ss:y = 1;
    printf ("Boo");       cs:printf(ds:"Boo");
} else                      } else
    y = 0;                ss:y = 0;
```

# Programming model

- Segments for: code, data, stack, “extra”
  - A program can have up to 6 total segments
  - Segments identified by registers: cs, ds, ss, es, fs, gs
- Prefix all memory accesses with desired segment:
  - `mov eax, ds:0x80` (load offset 0x80 from data into eax)
  - `jmp cs:0xab8` (jump execution to code offset 0xab8)
  - `mov ss:0x40, ecx` (move ecx to stack offset 0x40)

# Programming model, cont.

- This is cumbersome,
- Instead the idea is: infer code, data and stack segments from the instruction type:
  - Control-flow instructions use code segment (jump, call)
  - Stack management (push/pop) uses stack
  - Most loads/stores use data segment
- Extra segments (es, fs, gs) must be used explicitly

# Code segment

- Code
  - CS register
  - EIP is an offset inside the segment stored in CS
- Can only be changed with
  - procedure calls,
  - interrupt handling, or
  - task switching

# Data segment

- Data
  - DS, ES, FS, GS
  - 4 possible data segments can be used at the same time

# Stack segment

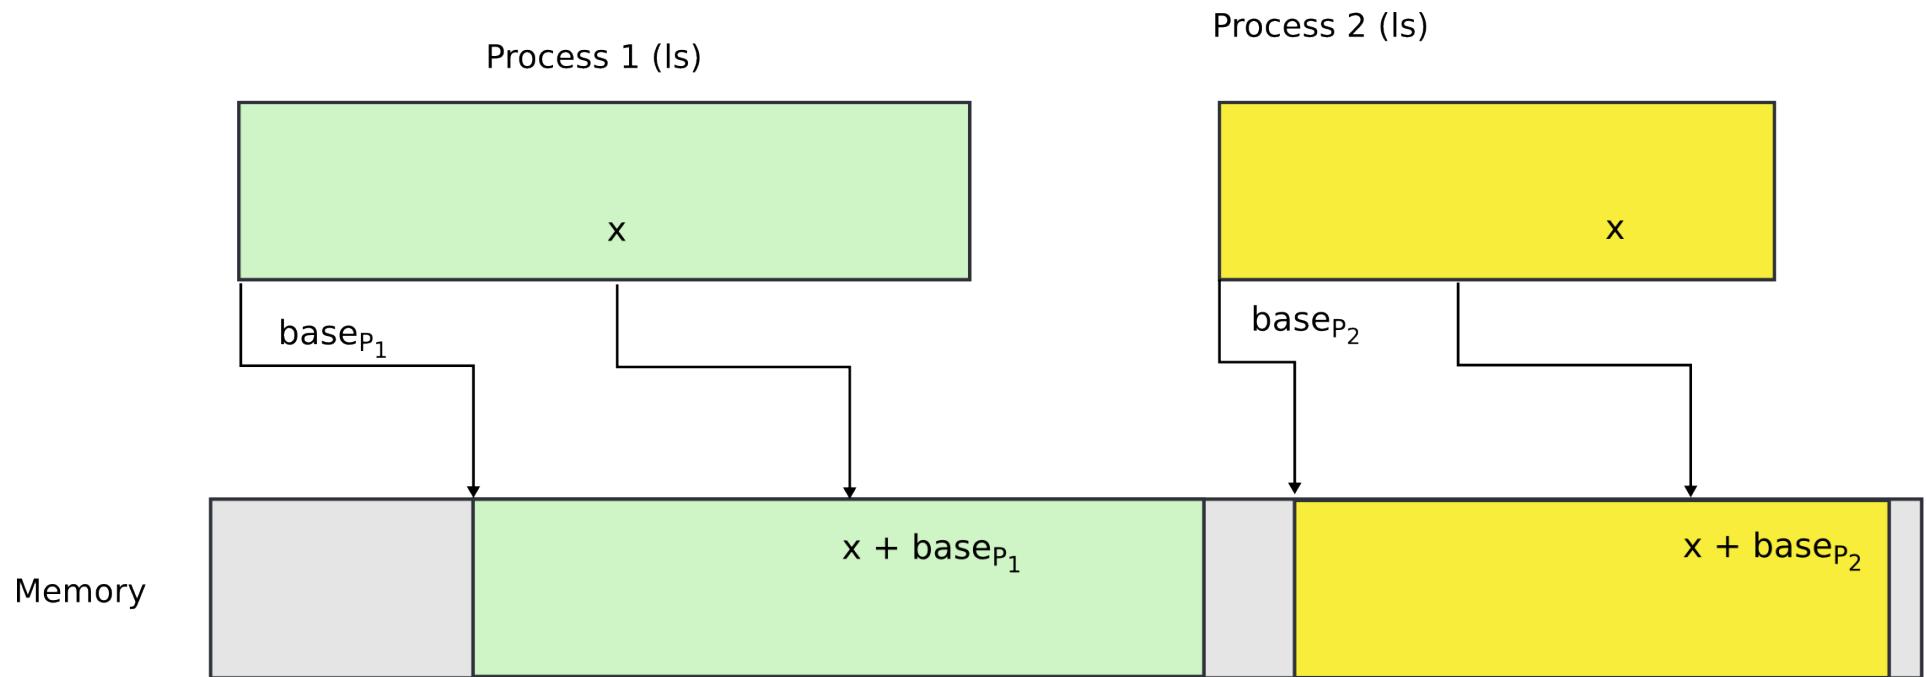
- Stack
  - SS
- Can be loaded explicitly
  - OS can set up multiple stacks
  - Of course, only one is accessible at a time

# Segmentation: what did we achieve

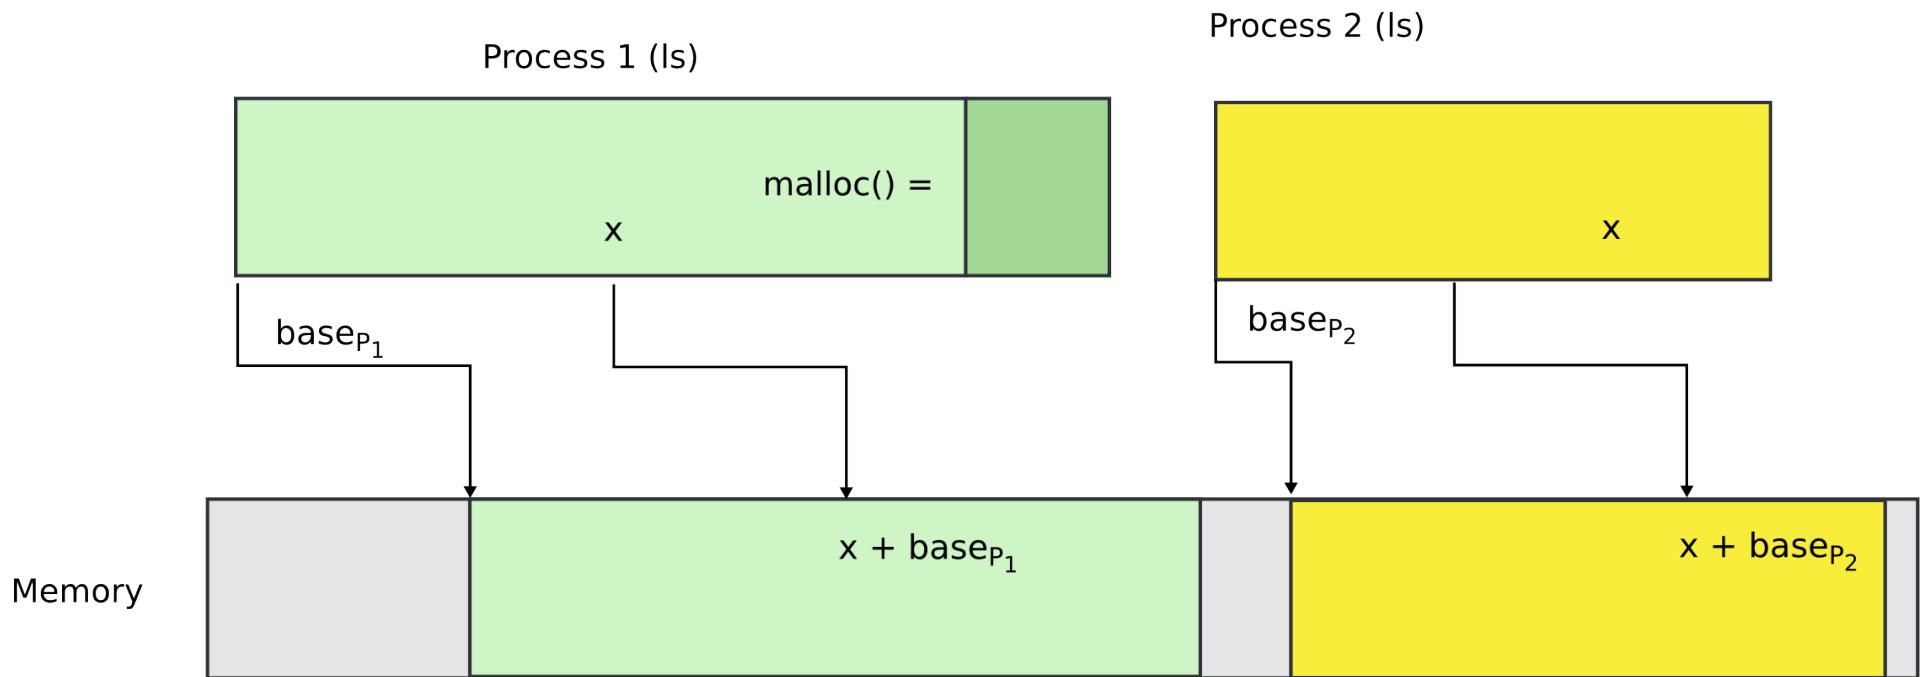
- Illusion of a private address space
  - Identical copy of an address space in multiple programs
    - We can implement `fork()`
- Isolation
  - Processes cannot access memory outside of their segments

Segmentation works for isolation, i.e., it does provide programs with illusion of private memory

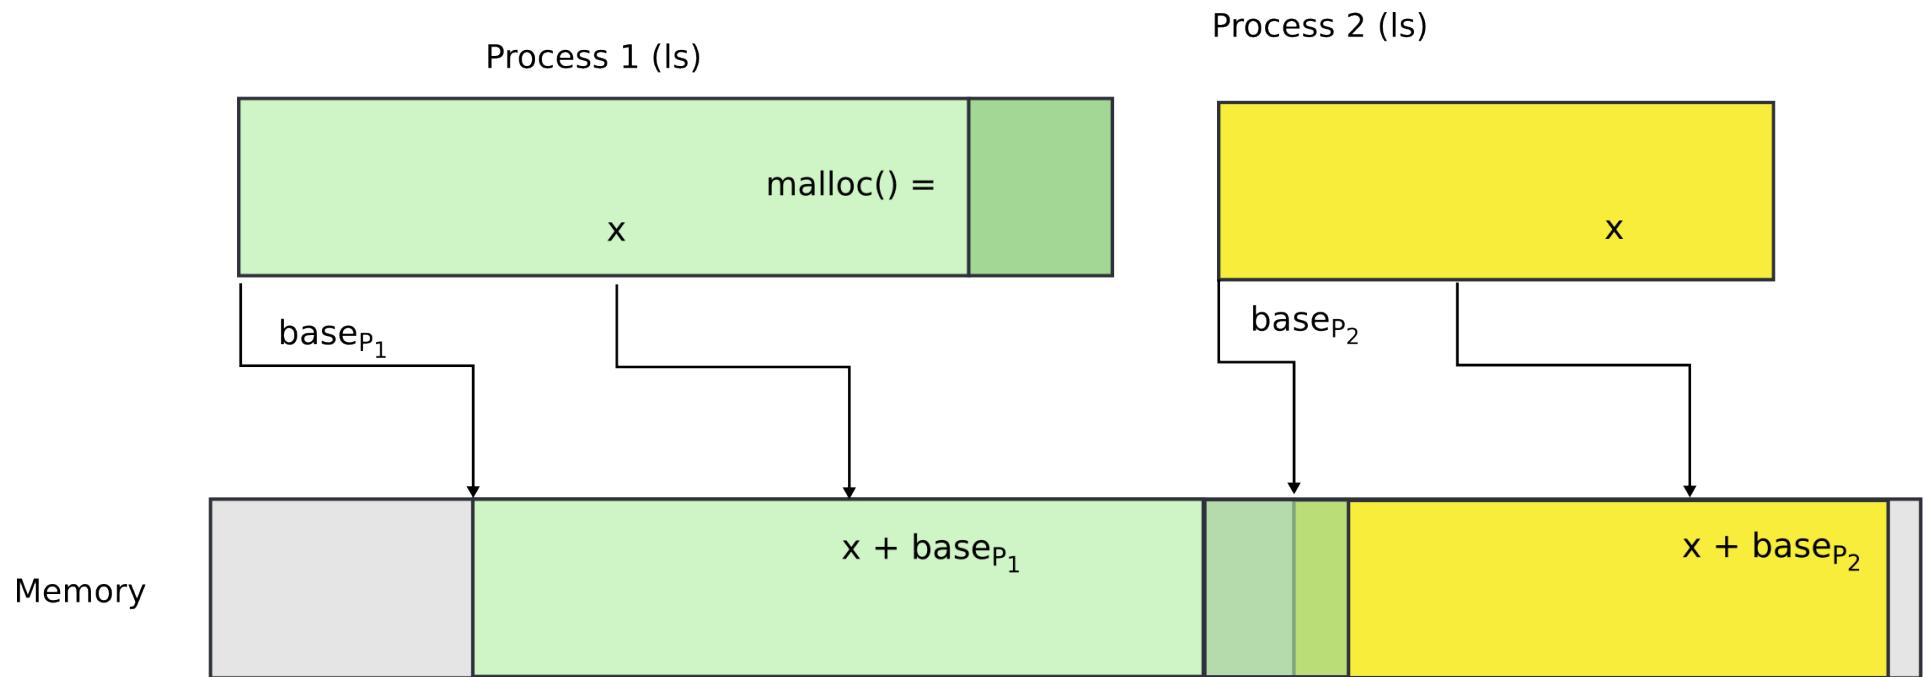
# Segmentation is ok... but



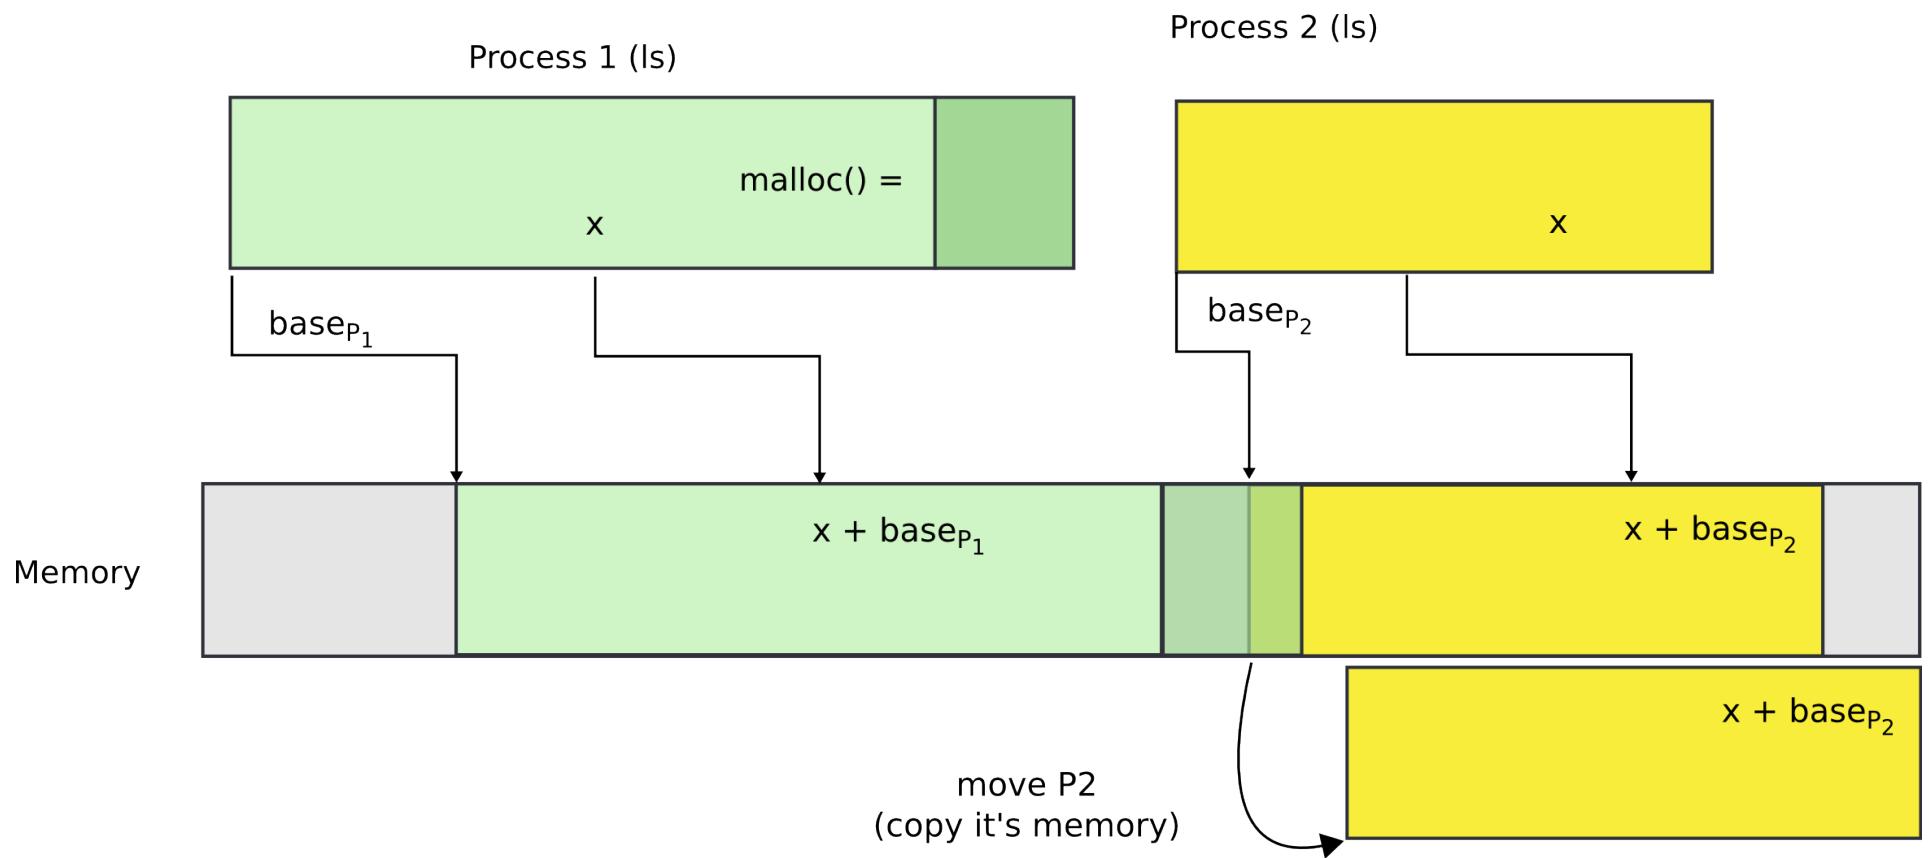
# What if process needs more memory?



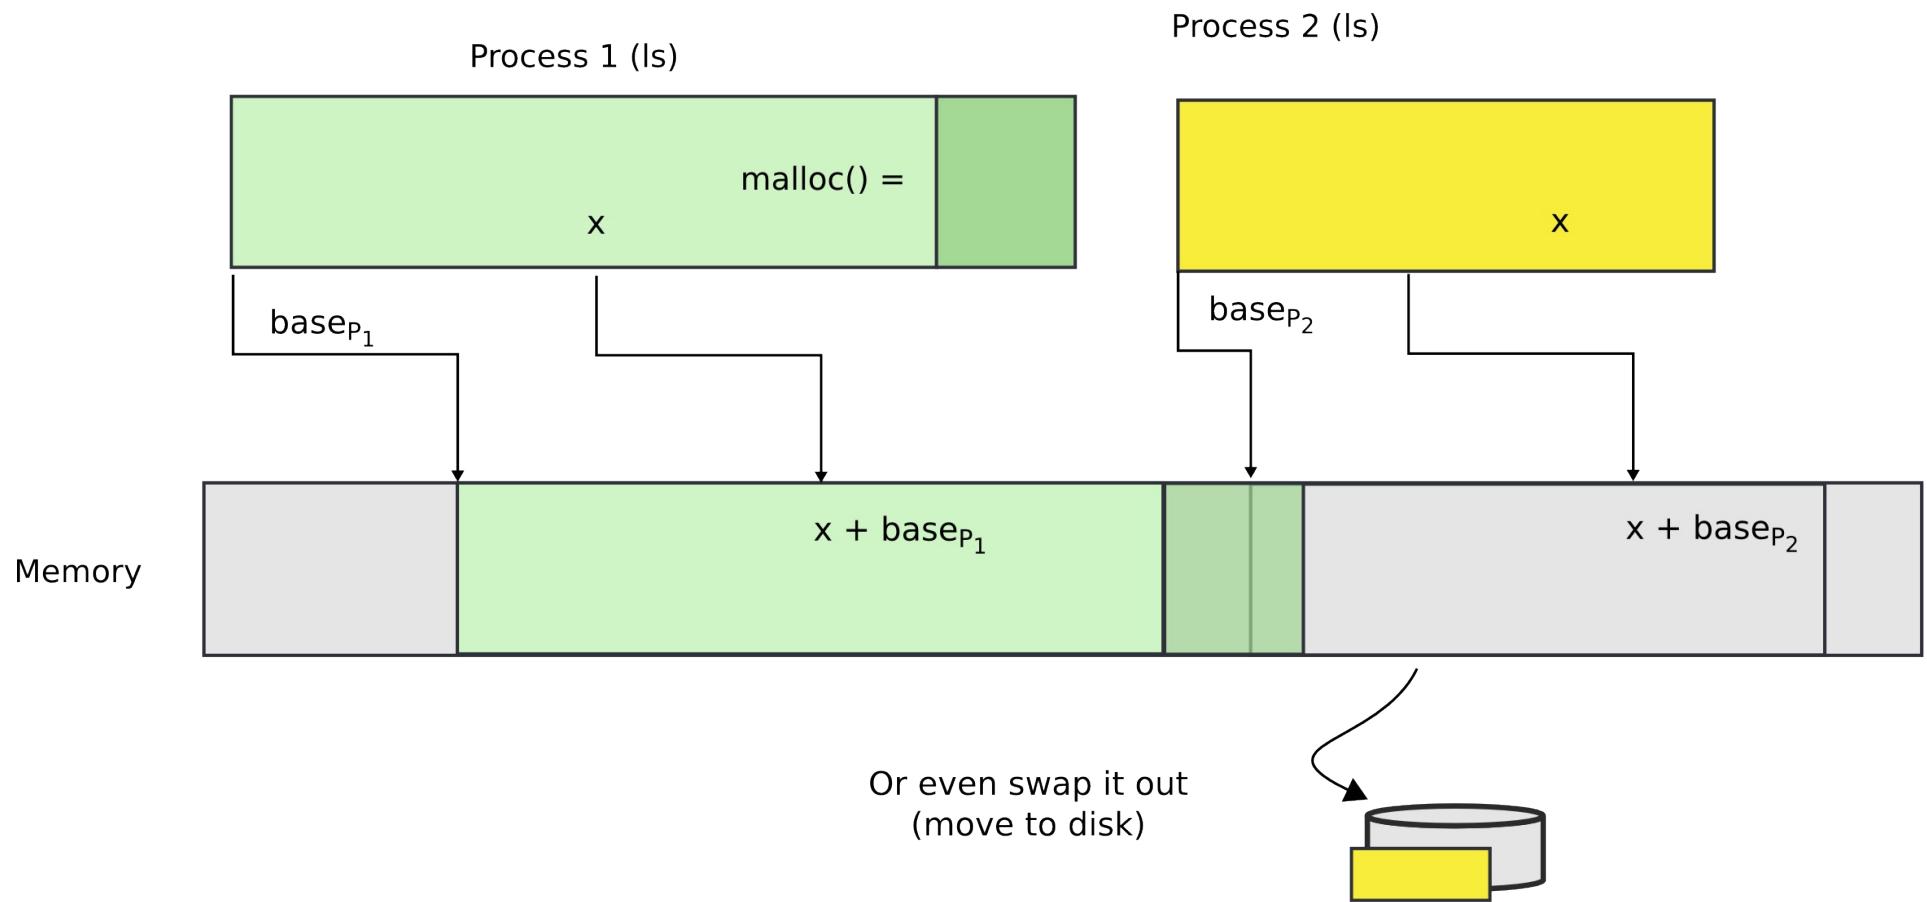
# What if process needs more memory?



# You can move P2 in memory



# Or even swap it out to disk

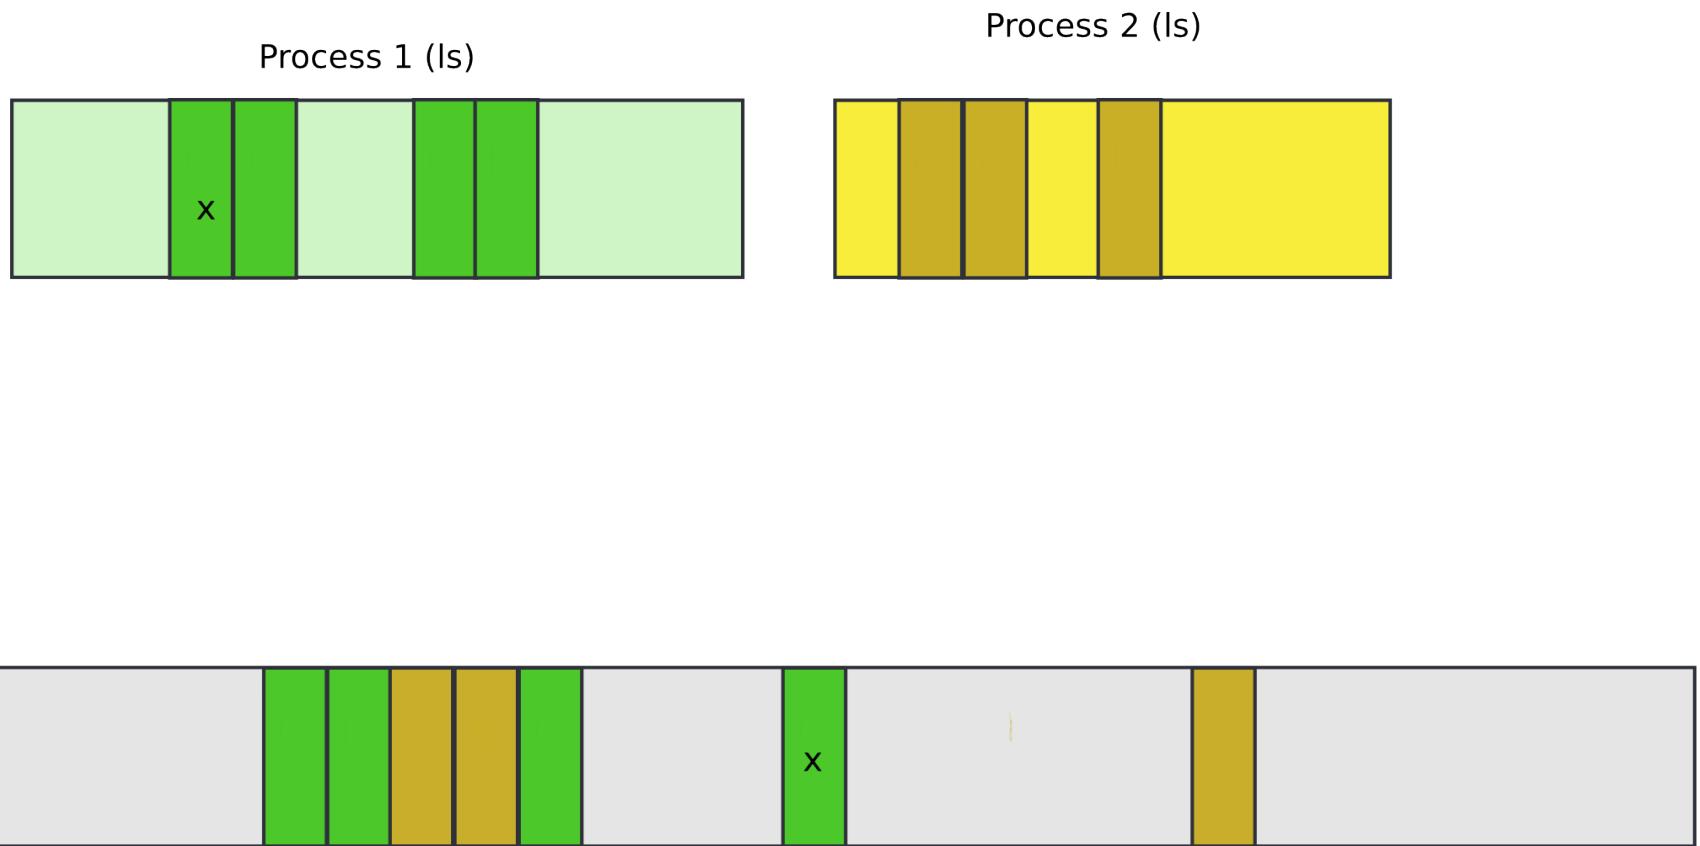


# Problems with segments

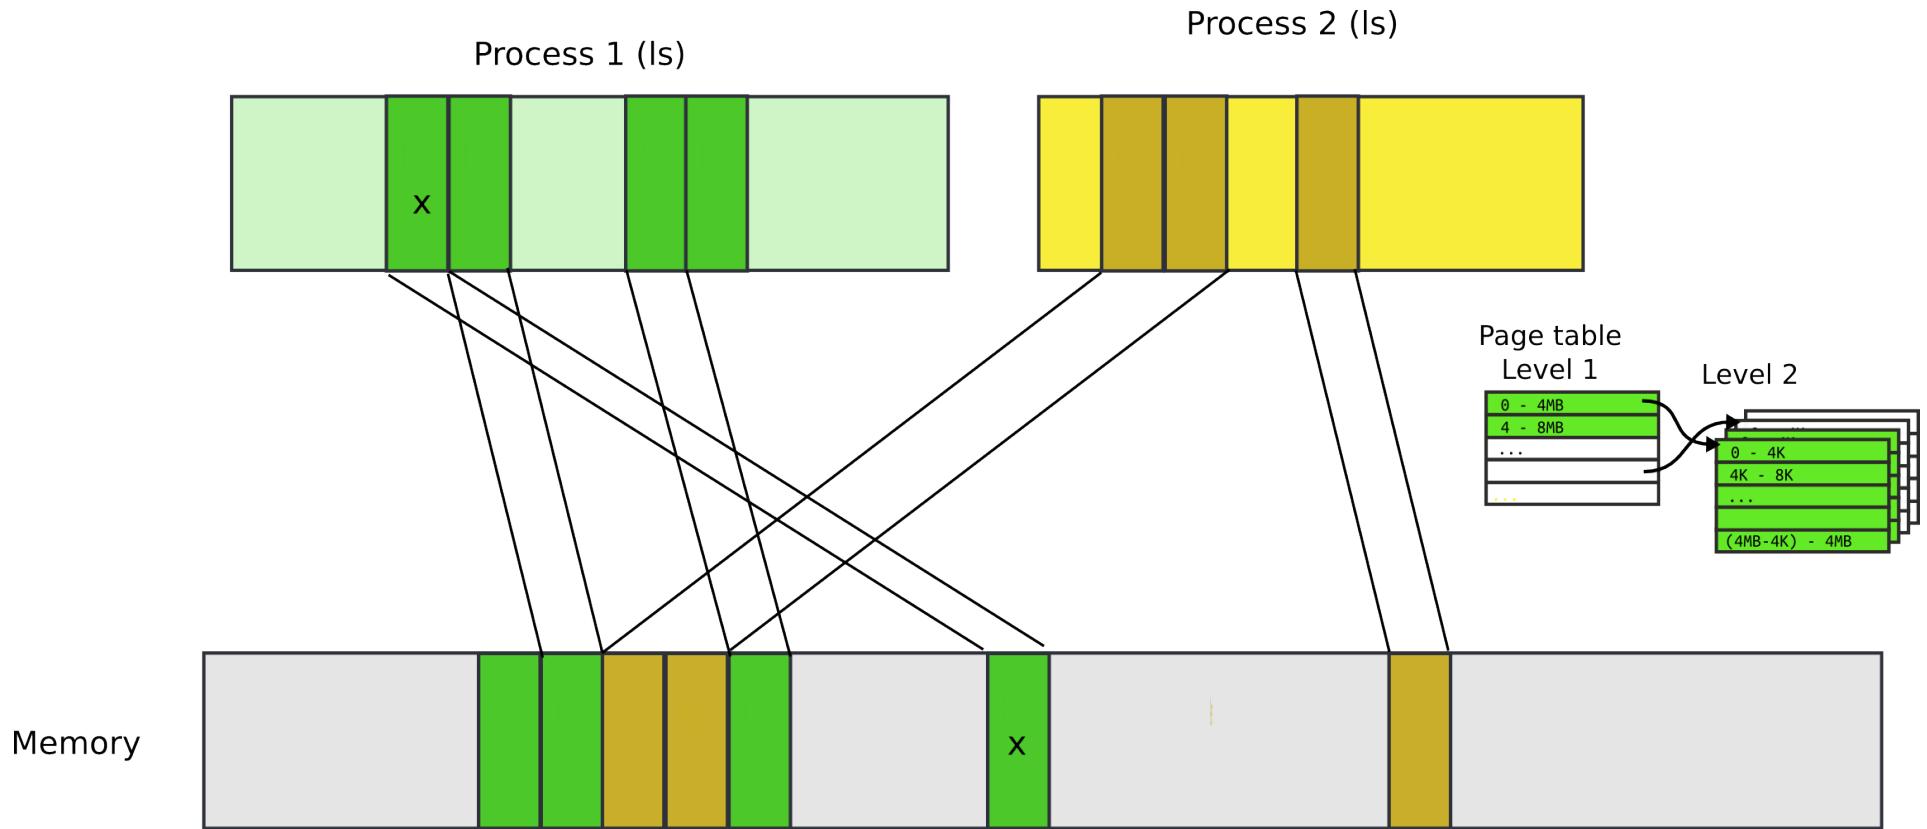
- Segments are somewhat inconvenient
  - Relocating or swapping the entire process takes time
- Memory gets fragmented
  - There might be no space (gap) for the swapped out process to come in
  - Will have to swap out other processes

# Paging

# Pages



# Pages

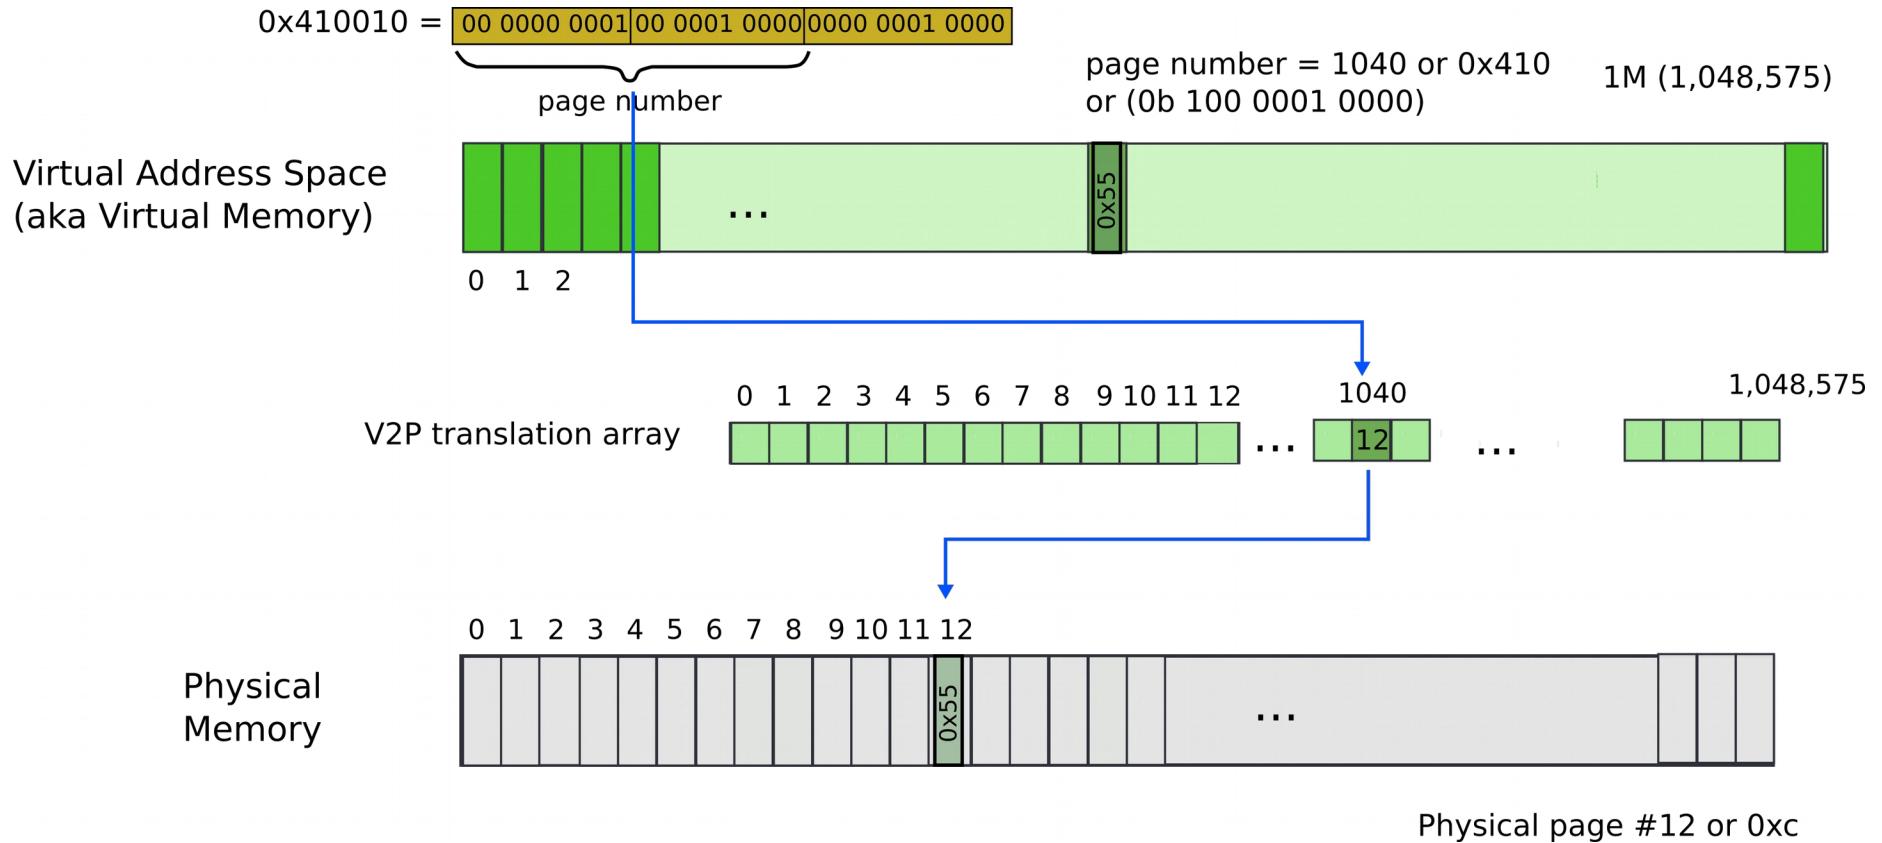


# Paging idea

- Break up memory into 4096-byte chunks called pages
  - Modern hardware supports 2MB, 4MB, and 1GB pages
- Independently control mapping for each page of linear address space
- Compare with segmentation (single base + limit)
  - many more degrees of freedom

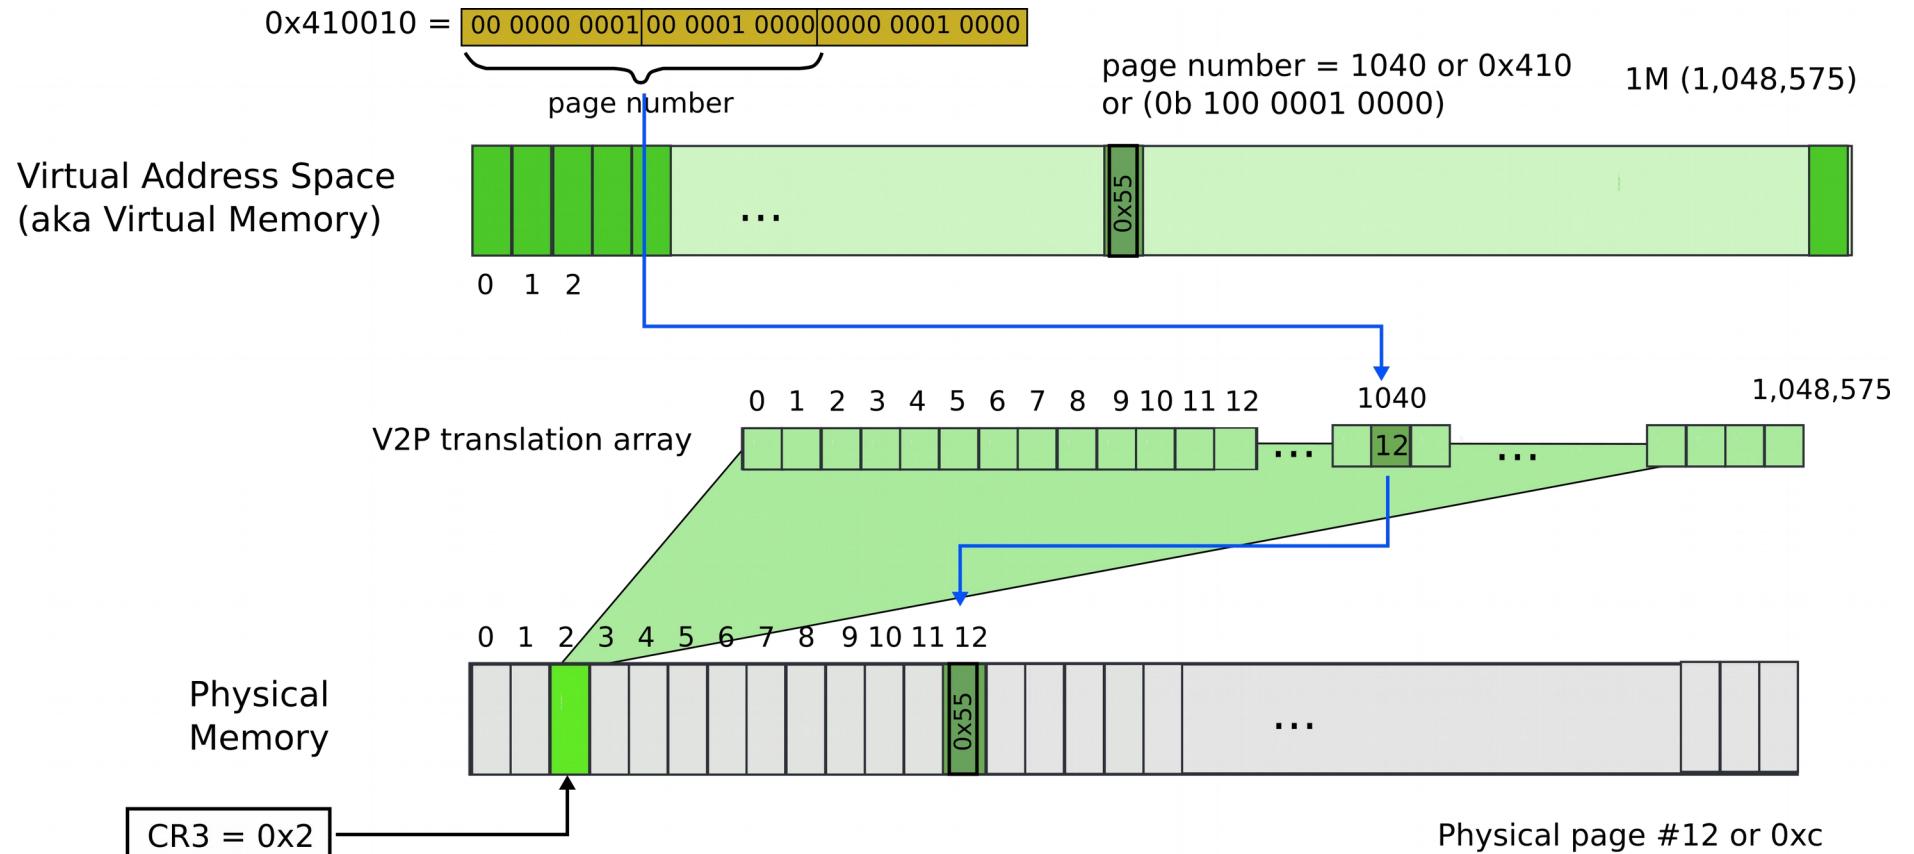
# How can we build this translation mechanism?

# Paging: naive approach: translation array



- Linear address  $0x410010$ 
  - Remember it's result of logical to linear translation (aka segmentation)
    - $0x410010 = 0x300010$  (offset) +  $0x110000$  (base)

# Paging: naive approach: translation array



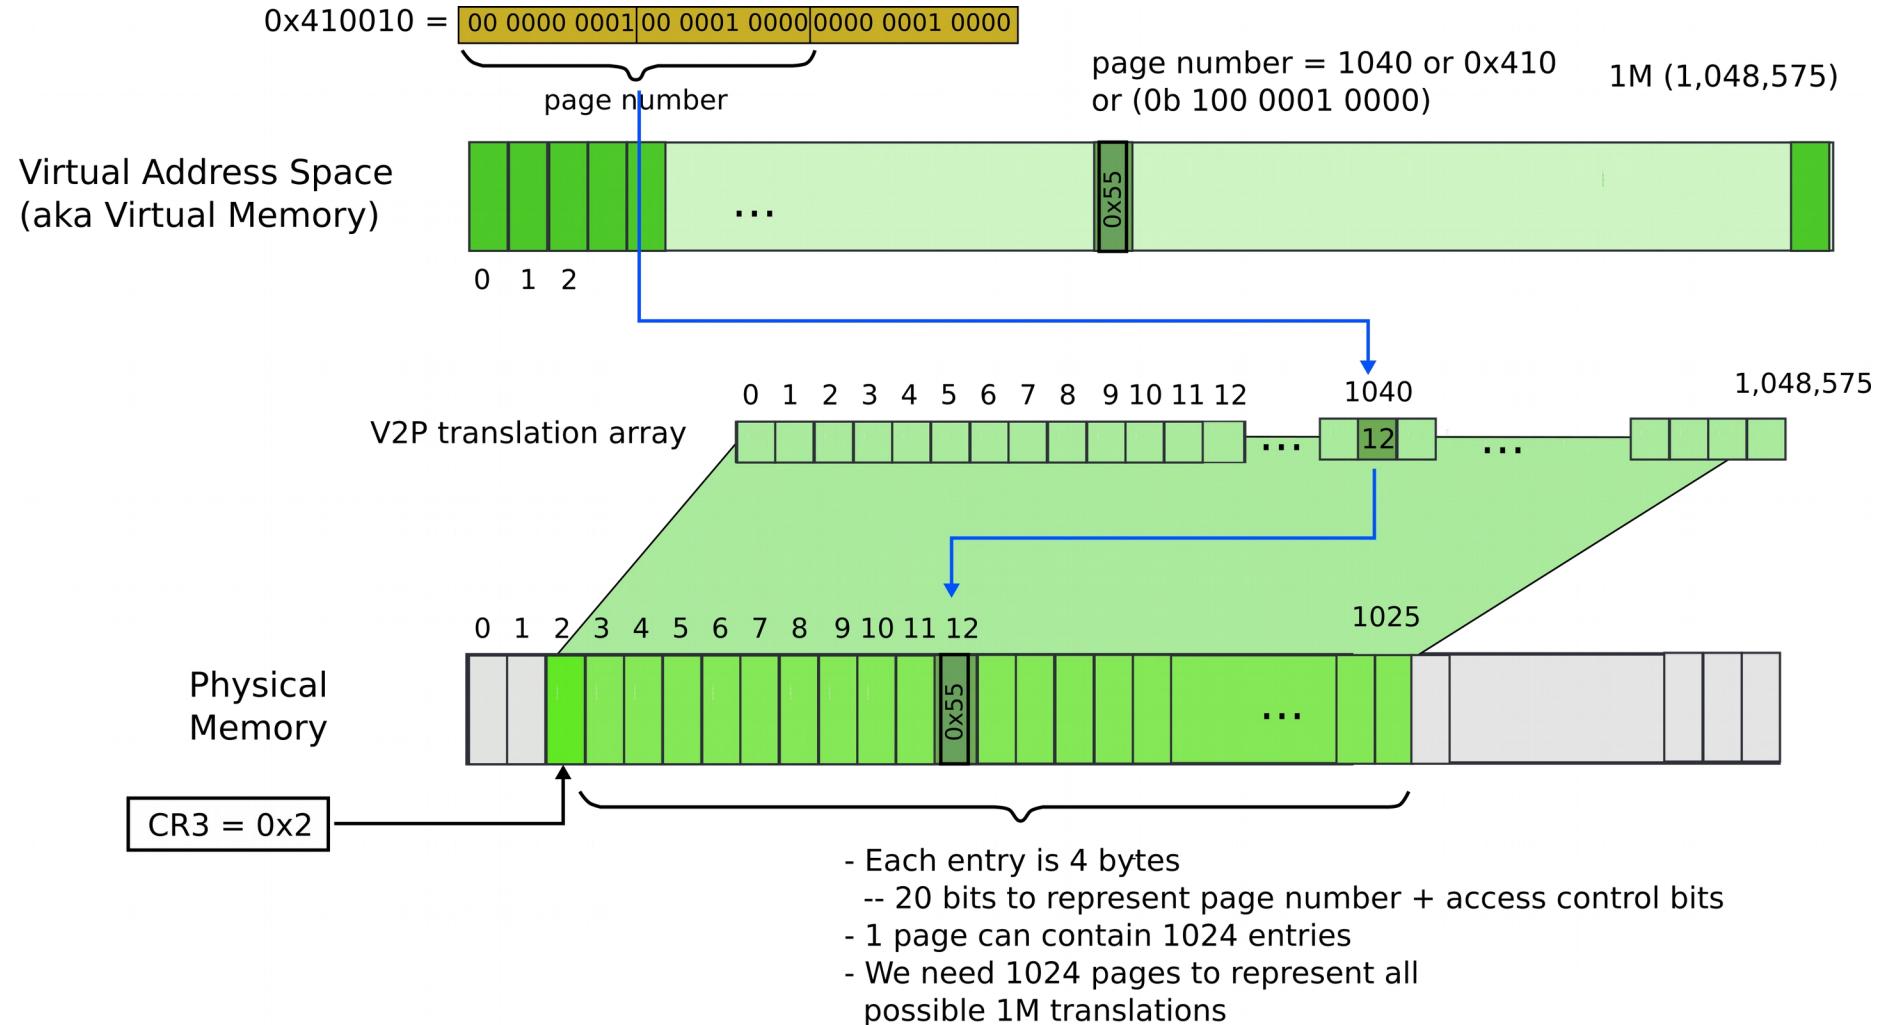
- Linear address 0x410010
- Remember it's result of logical to linear translation (aka segmentation)
  - $0x410010 = 0x300010 \text{ (offset)} + 0x110000 \text{ (base)}$

# What is wrong?

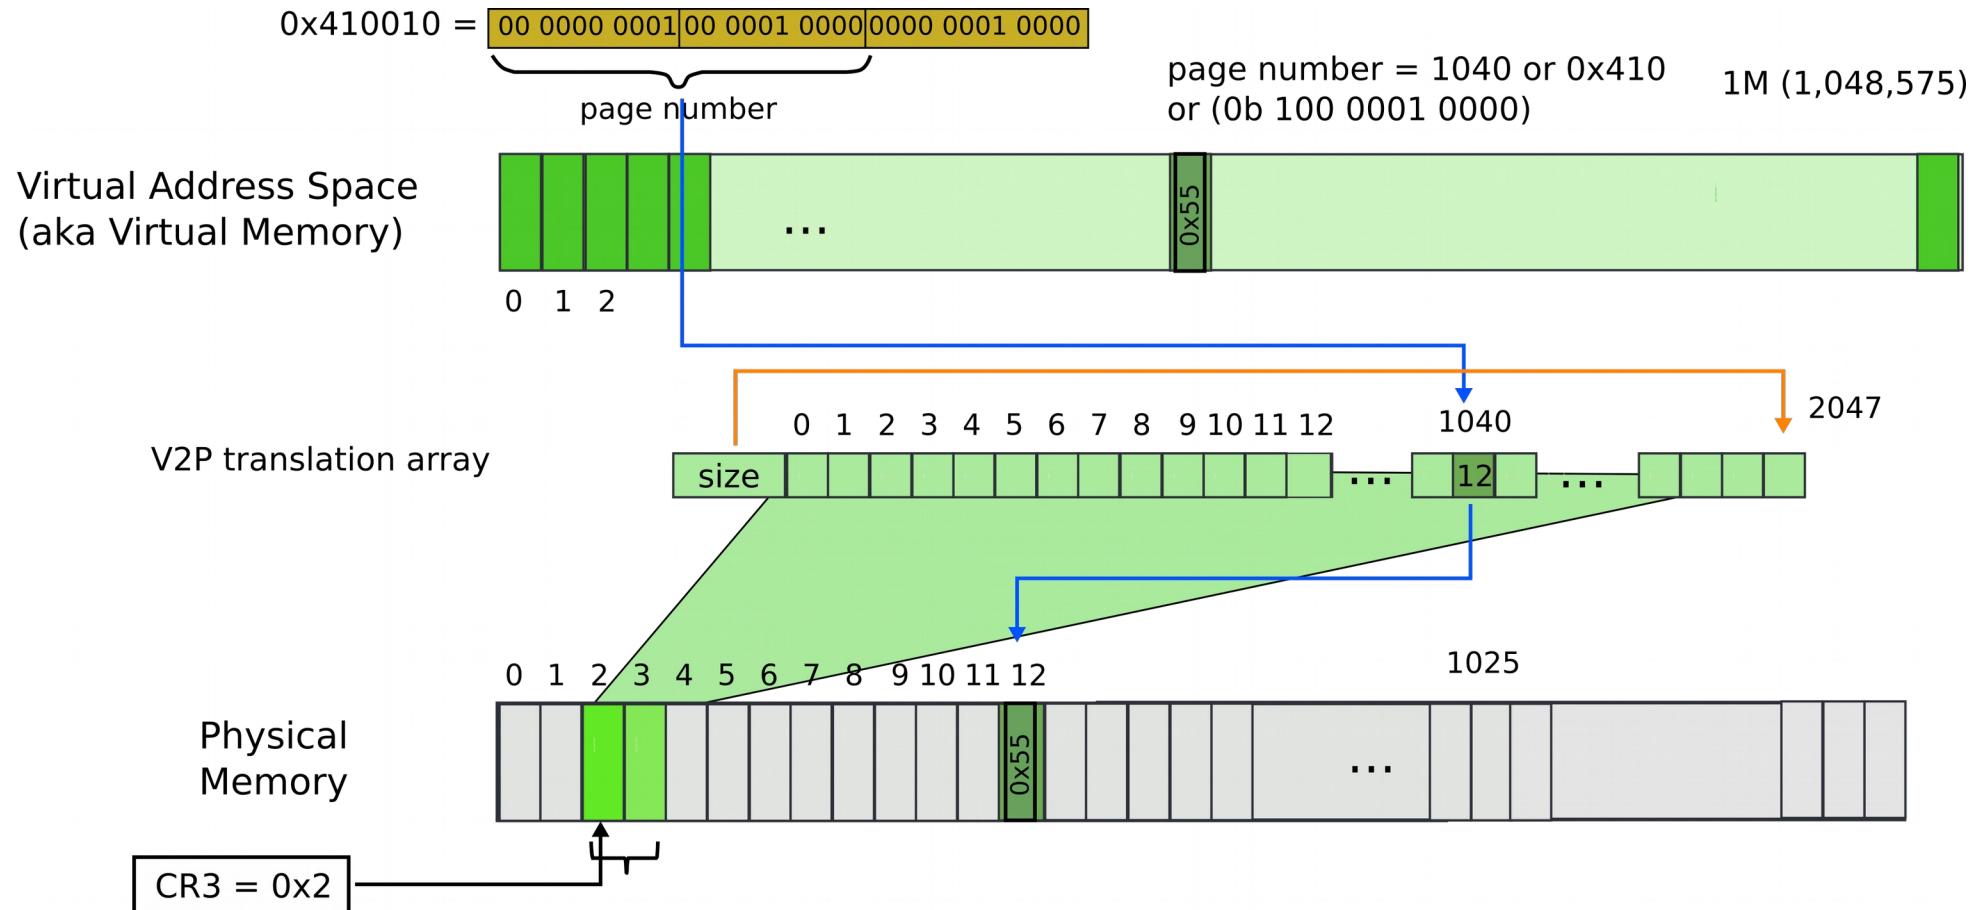
# What is wrong?

- We need 4 bytes to relocate each page
  - 20 bits for physical page number
  - 12 bits of access flags
- Therefore, we need array of 4 bytes x 1M entries
  - 4MBs

# Paging: naive approach: translation array



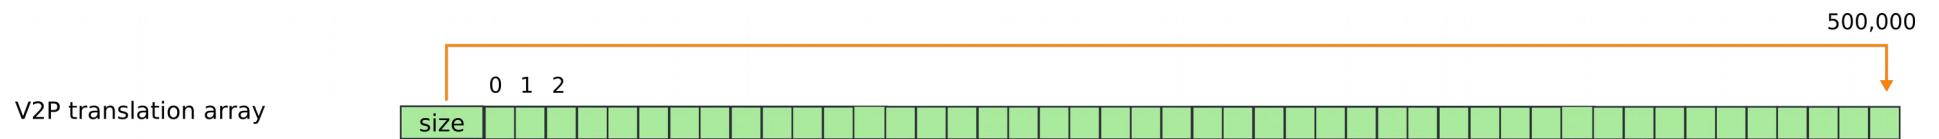
# Paging: array with size



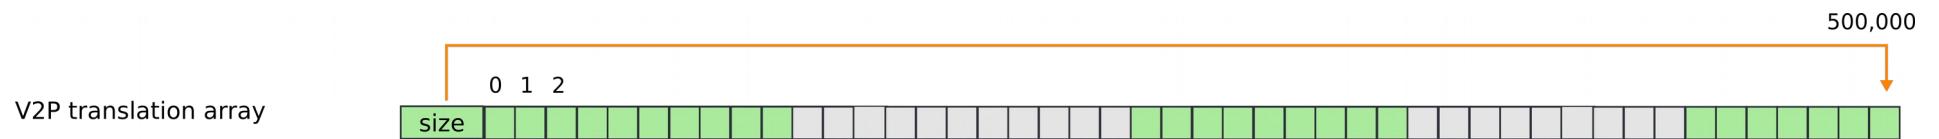
- The size controls how many entries are required

# But still what may go wrong?

# Paging: array with size

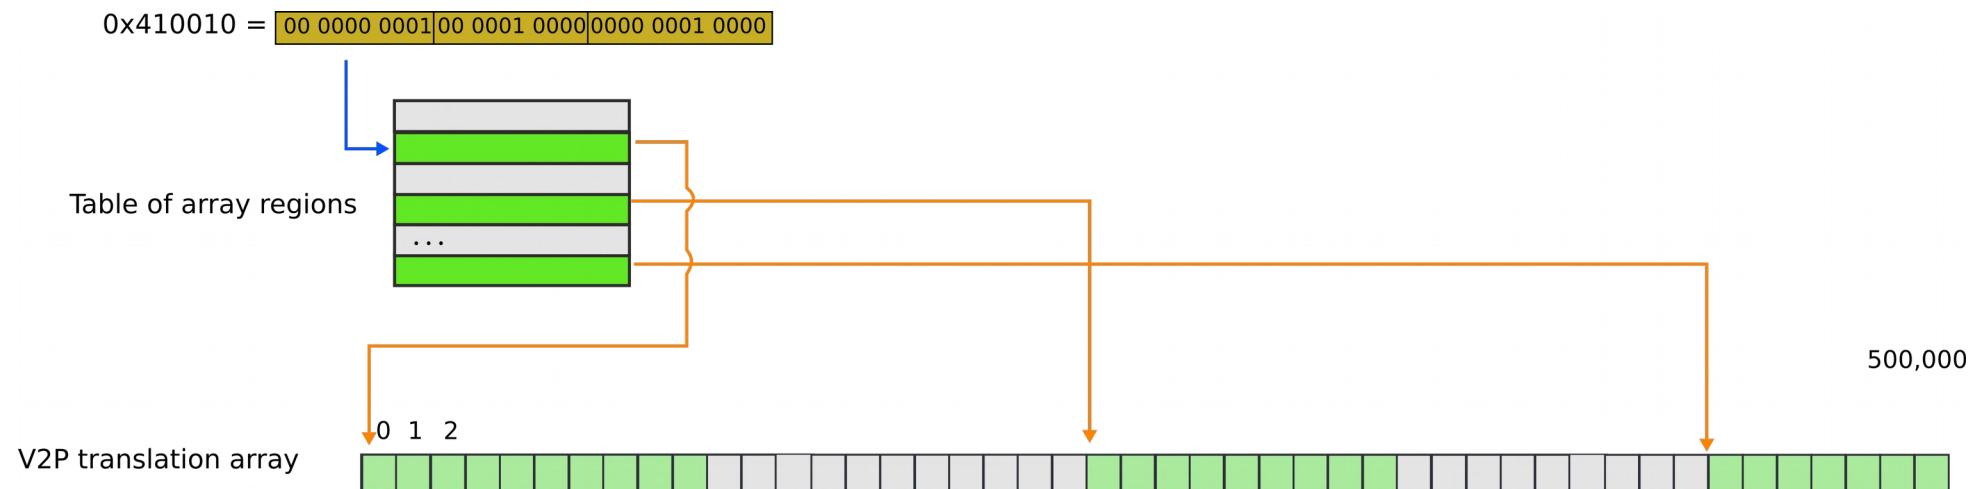


# Paging: array with size



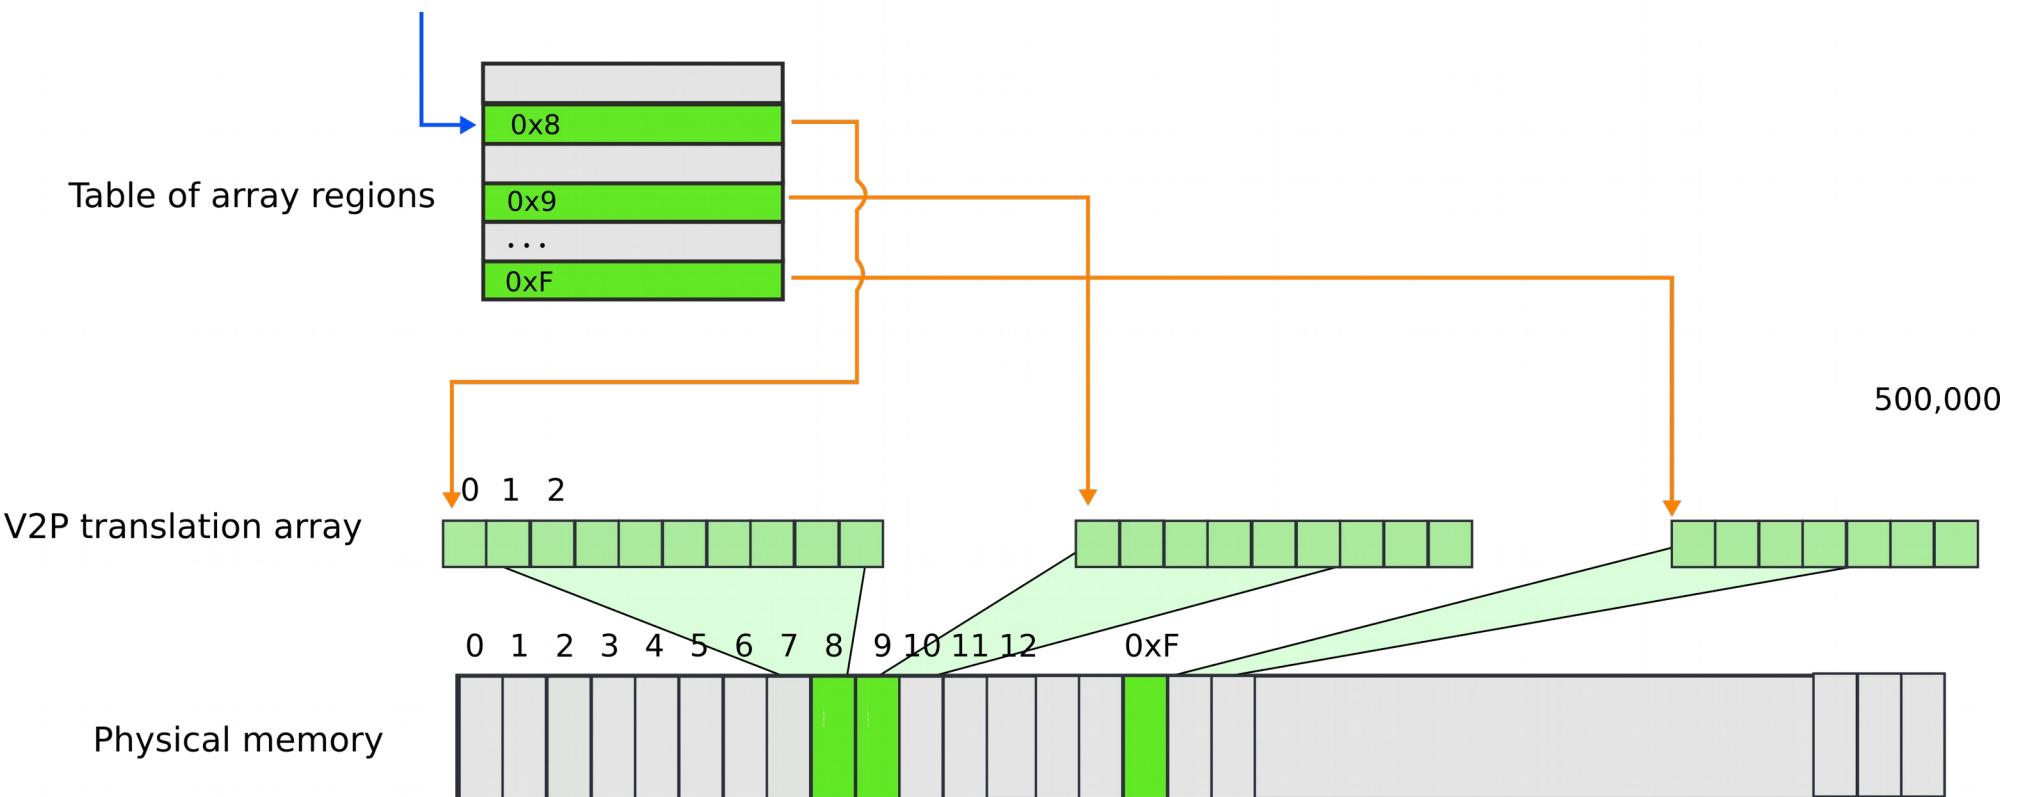
# Can we improve?

# Paging: array with chunks

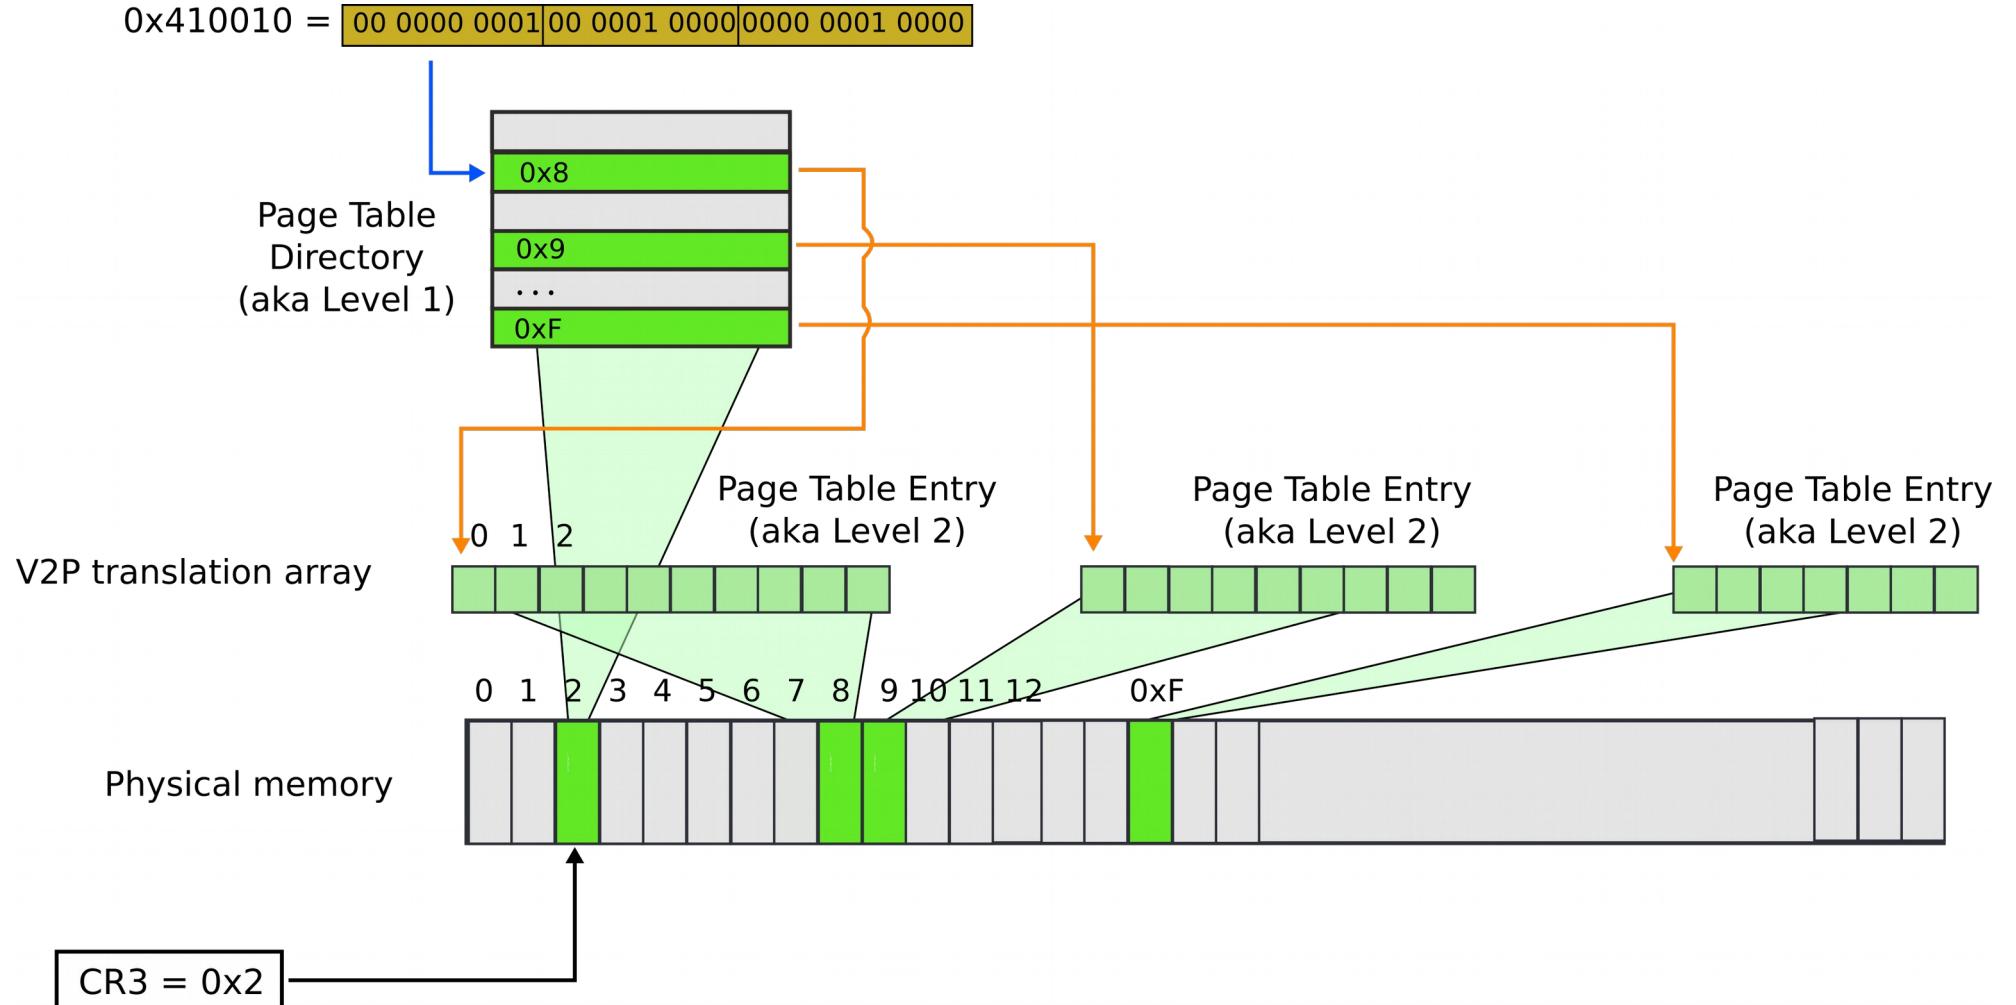


# Paging: array with chunks

$0x410010 = [00\ 0000\ 0001|00\ 0001\ 0000|0000\ 0001\ 0000]$



# Paging: page table



```
mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0  
EBX = 20 983 809
```

20 983 809 = 

  
page number

1M (1,048,575)

Virtual Address  
Space (or Memory)  
of the Process



0 1 2

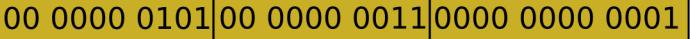
page number = 5123  
or (0b1 0100 0000 0011)

0 1 2 3 4 5 6 7 8 9 10 11 12

Physical  
Memory



```
mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0  
EBX = 20 983 809
```

20 983 809 = 

  
page number

1M (1,048,575)

Virtual Address  
Space (or Memory)  
of the Process



CR3 = 0

0 1 2

page number = 5123  
or (0b1 0100 0000 0011)

Physical  
Memory

0 1 2 3 4 5 6 7 8 9 10 11 12



A horizontal bar representing Physical Memory. It contains 13 pages indexed from 0 to 12. An arrow from CR3=0 points to page 0.

```

mov (%EBX), EAX # mov value from the location pointed by EBX into EAX
EAX = 0
EBX = 20 983 809

```

$20\ 983\ 809 = \boxed{00\ 0000\ 0101} \boxed{00\ 0000\ 0011} \boxed{0000\ 0000\ 0001}$

page number

1M (1,048,575)

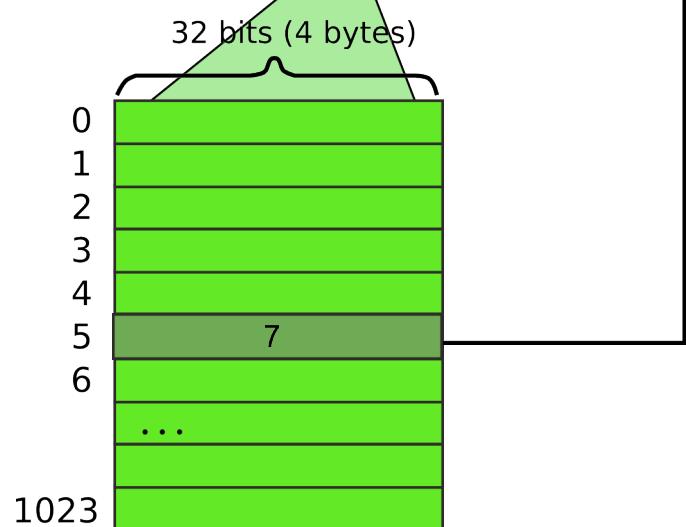
Virtual Address  
Space (or Memory)  
of the Process



CR3 = 0 →  
0 1 2 3 4 5 6 7 8 9 10 11 12

page number = 5123  
or (0b1 0100 0000 0011)

Physical  
Memory

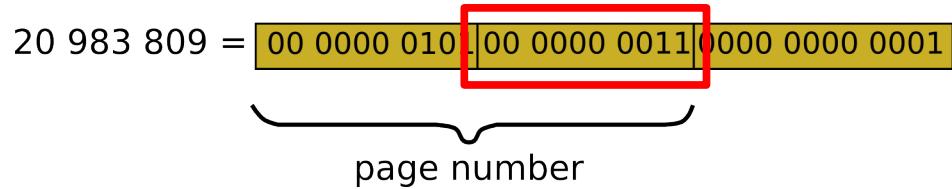


Level 1  
(Page Table  
Directory)

```

mov (%EBX), EAX # mov value from the location pointed by EBX into EAX
EAX = 0
EBX = 20 983 809

```



1M (1,048,575)

Virtual Address  
Space (or Memory)  
of the Process



CR3 = 0

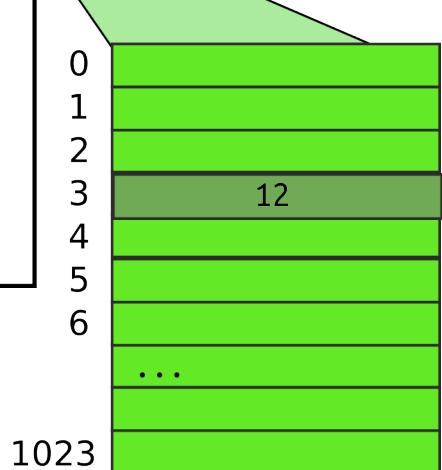
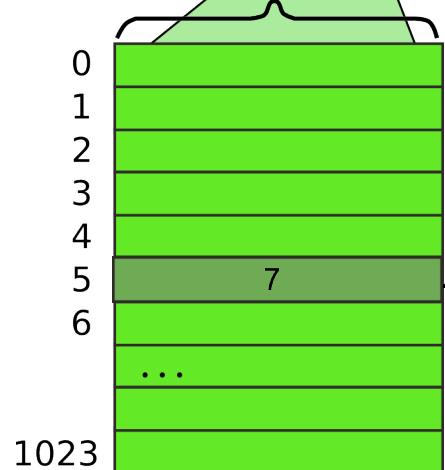


page number = 5123  
or (0b1 0100 0000 0011)

Physical  
Memory



32 bits (4 bytes)



Level 1  
(Page Table  
Directory)

Level 2  
(Page Table)

```

mov (%EBX), EAX # mov value from the location pointed by EBX into EAX
EAX = 0
EBX = 20 983 809

```

$20\ 983\ 809 = \boxed{00\ 0000\ 0101} \boxed{00\ 0000\ 0011} \boxed{0000\ 0000\ 0001}$

page number

1M (1,048,575)

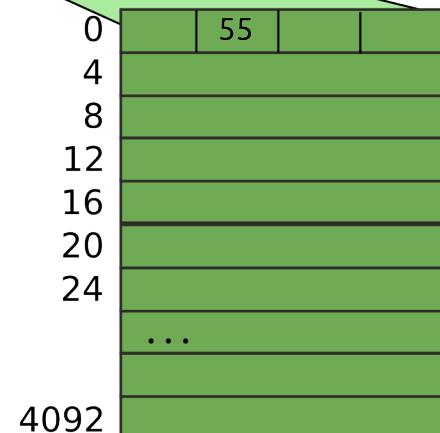
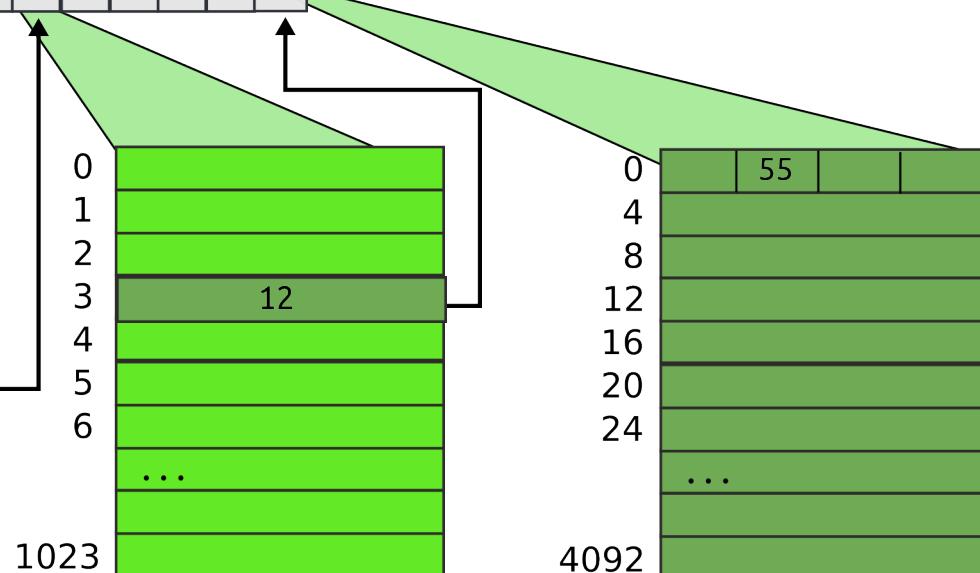
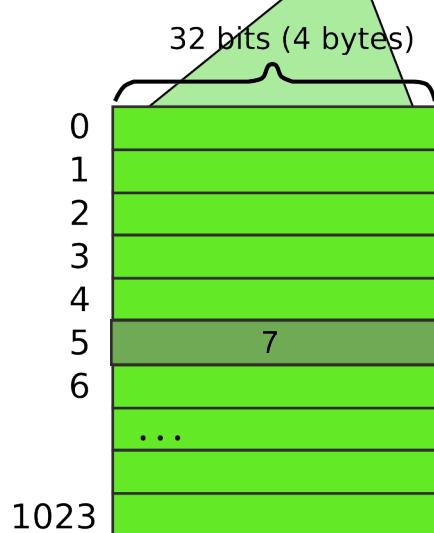
Virtual Address  
Space (or Memory)  
of the Process



CR3 = 0 →  
0 1 2

page number = 5123  
or (0b1 0100 0000 0011)

Physical  
Memory  
0 1 2 3 4 5 6 7 8 9 10 11 12



Level 1  
(Page Table  
Directory)

Level 2  
(Page Table)

Page

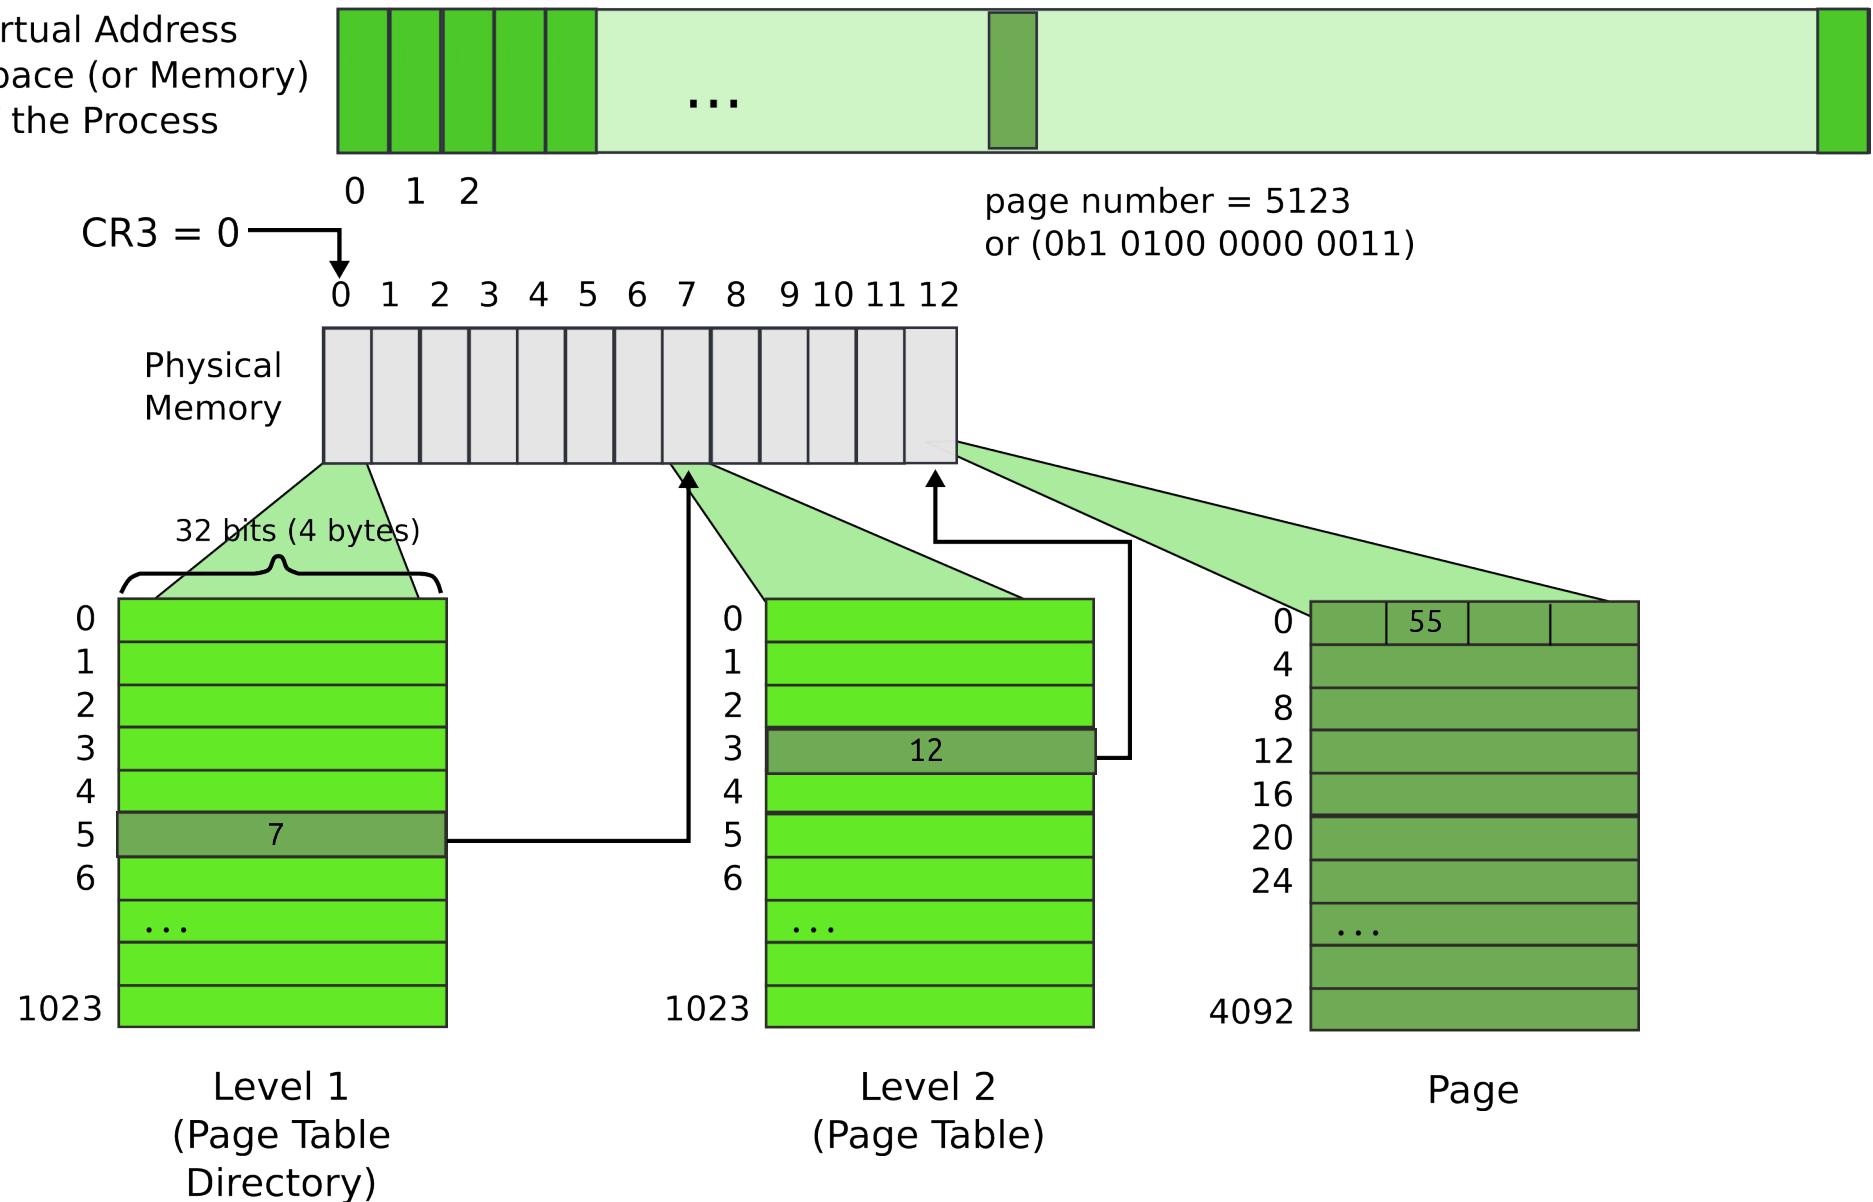
```
mov (%EBX), EAX # mov value from the location pointed by EBX into EAX  
EAX = 0  
EBX = 20 983 809
```

20 983 809 = **00 0000 0101|00 0000 0011|0000 0000 0001**

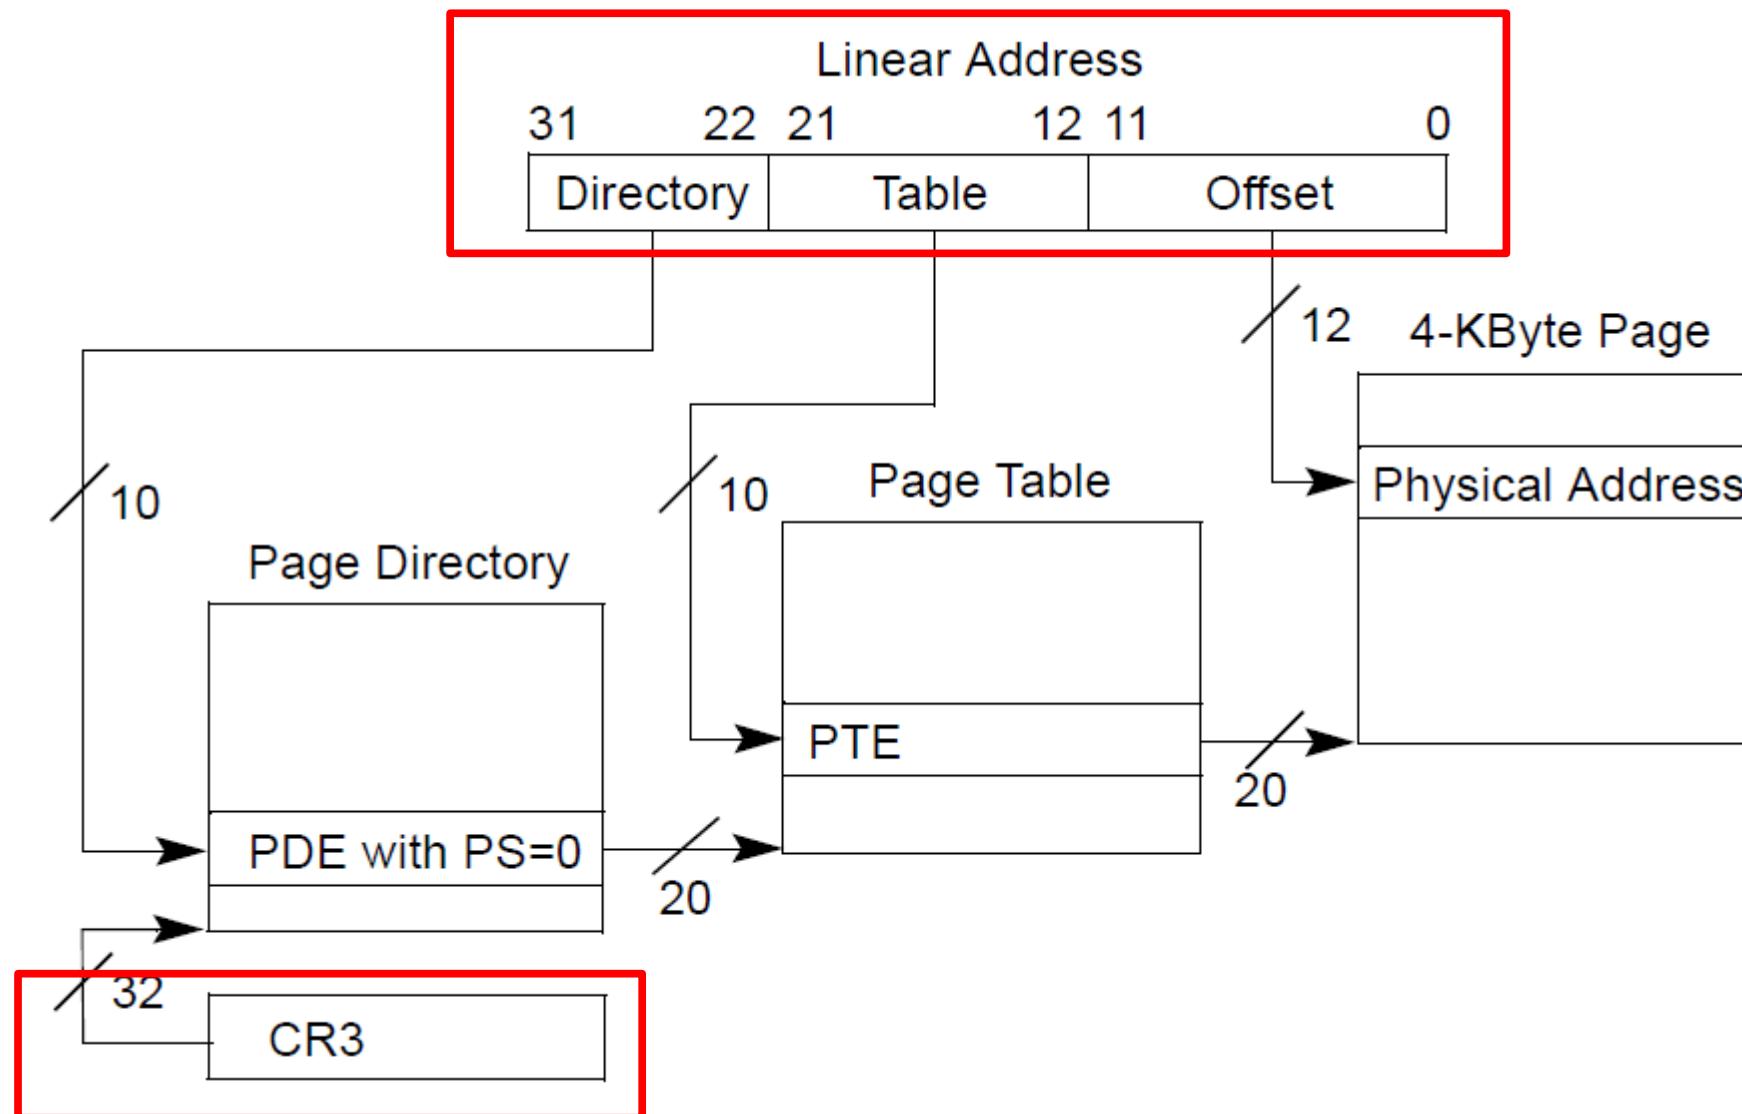
page number

- Result:
    - EAX = 55

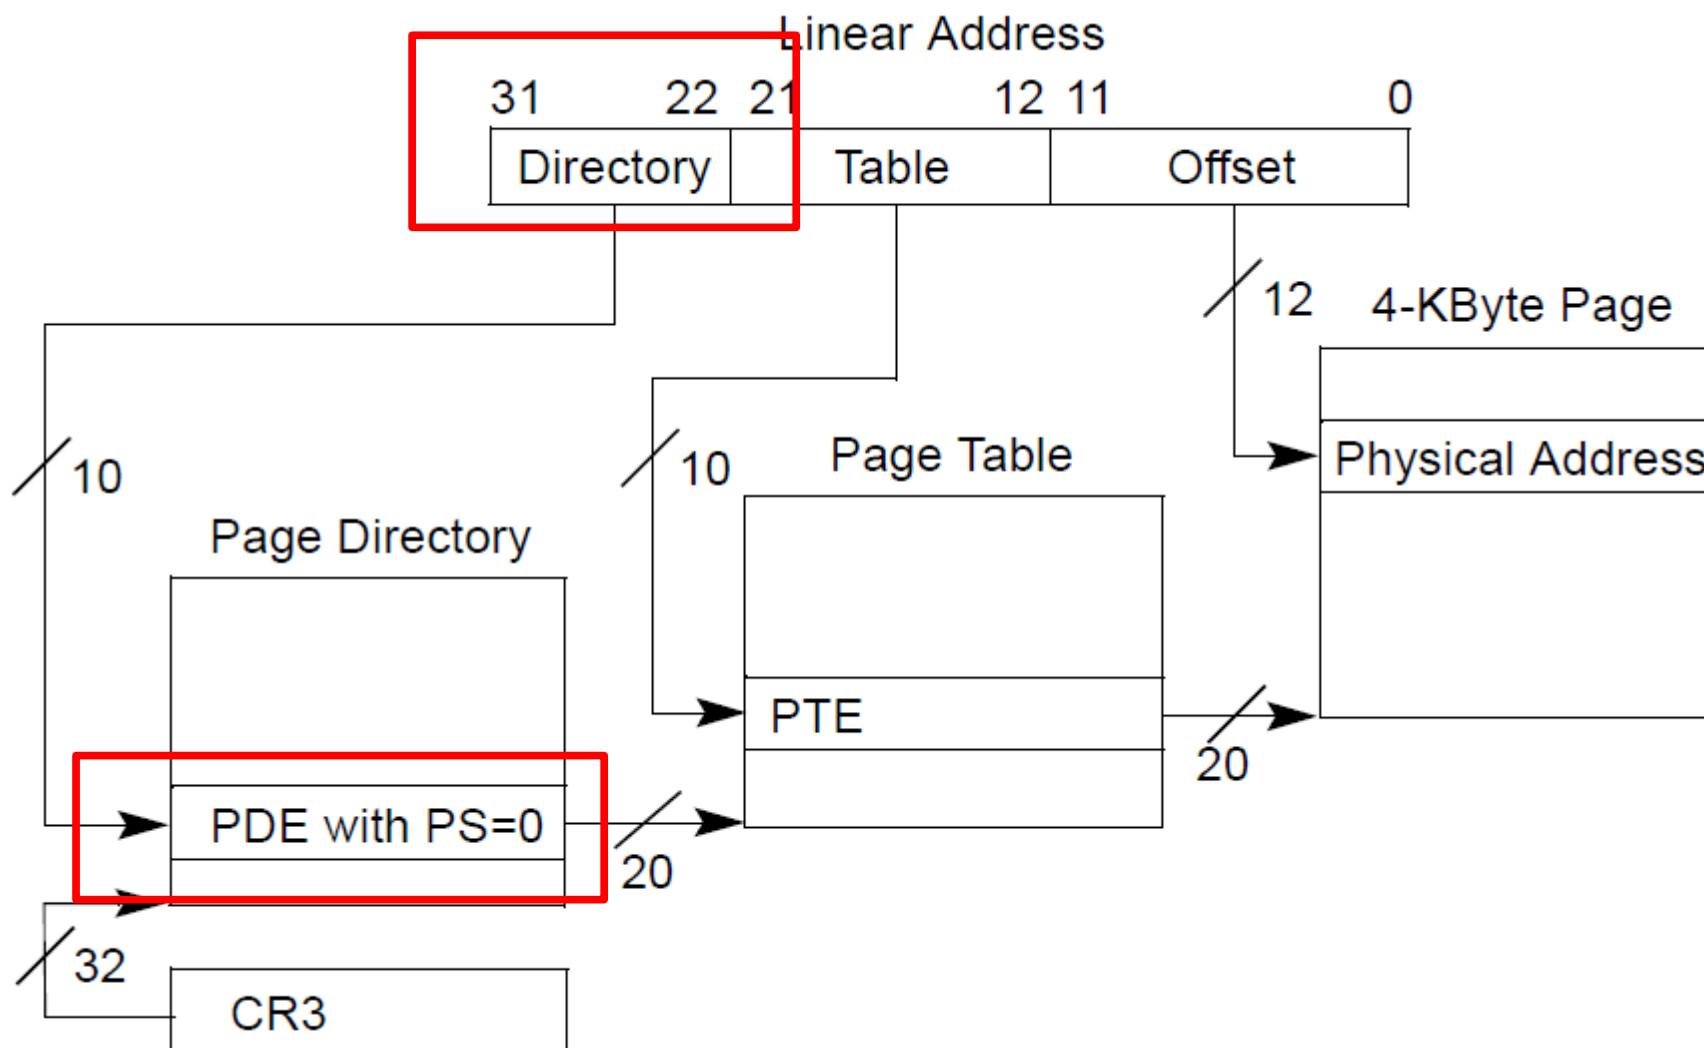
## Virtual Address Space (or Memory) of the Process



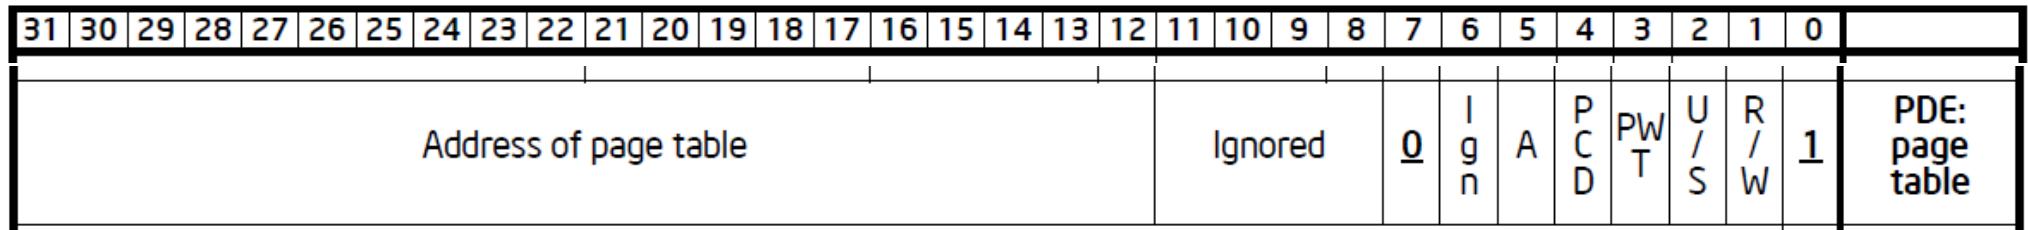
# Page translation



# Page translation

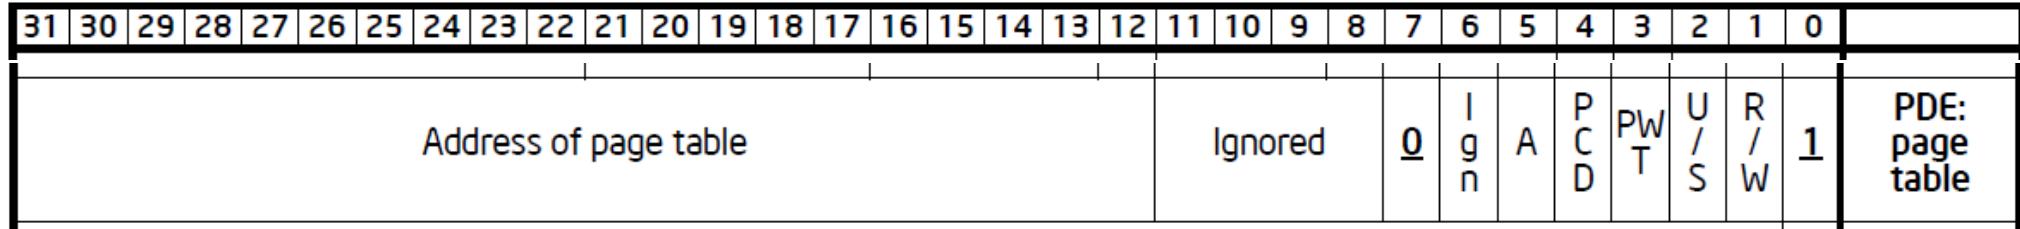


# Page directory entry (PDE)



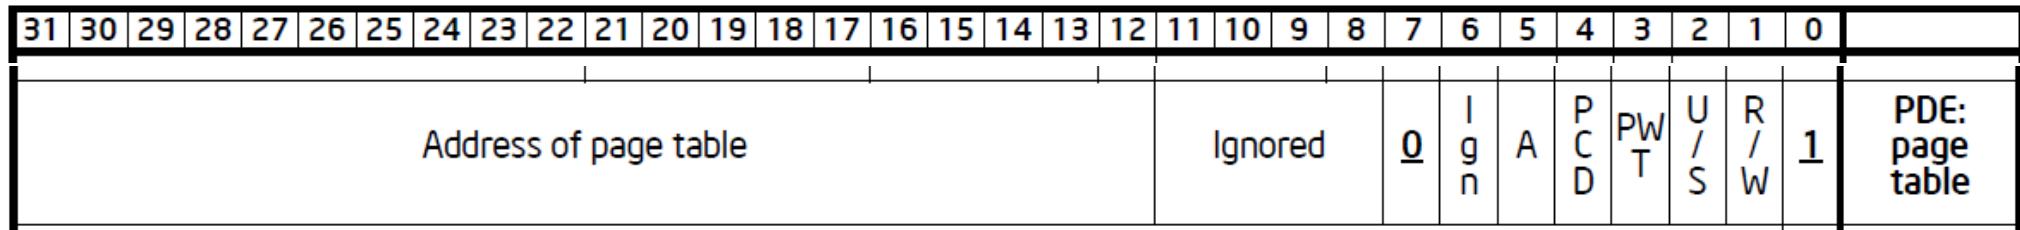
- 20 bit address of the page table

# Page directory entry (PDE)



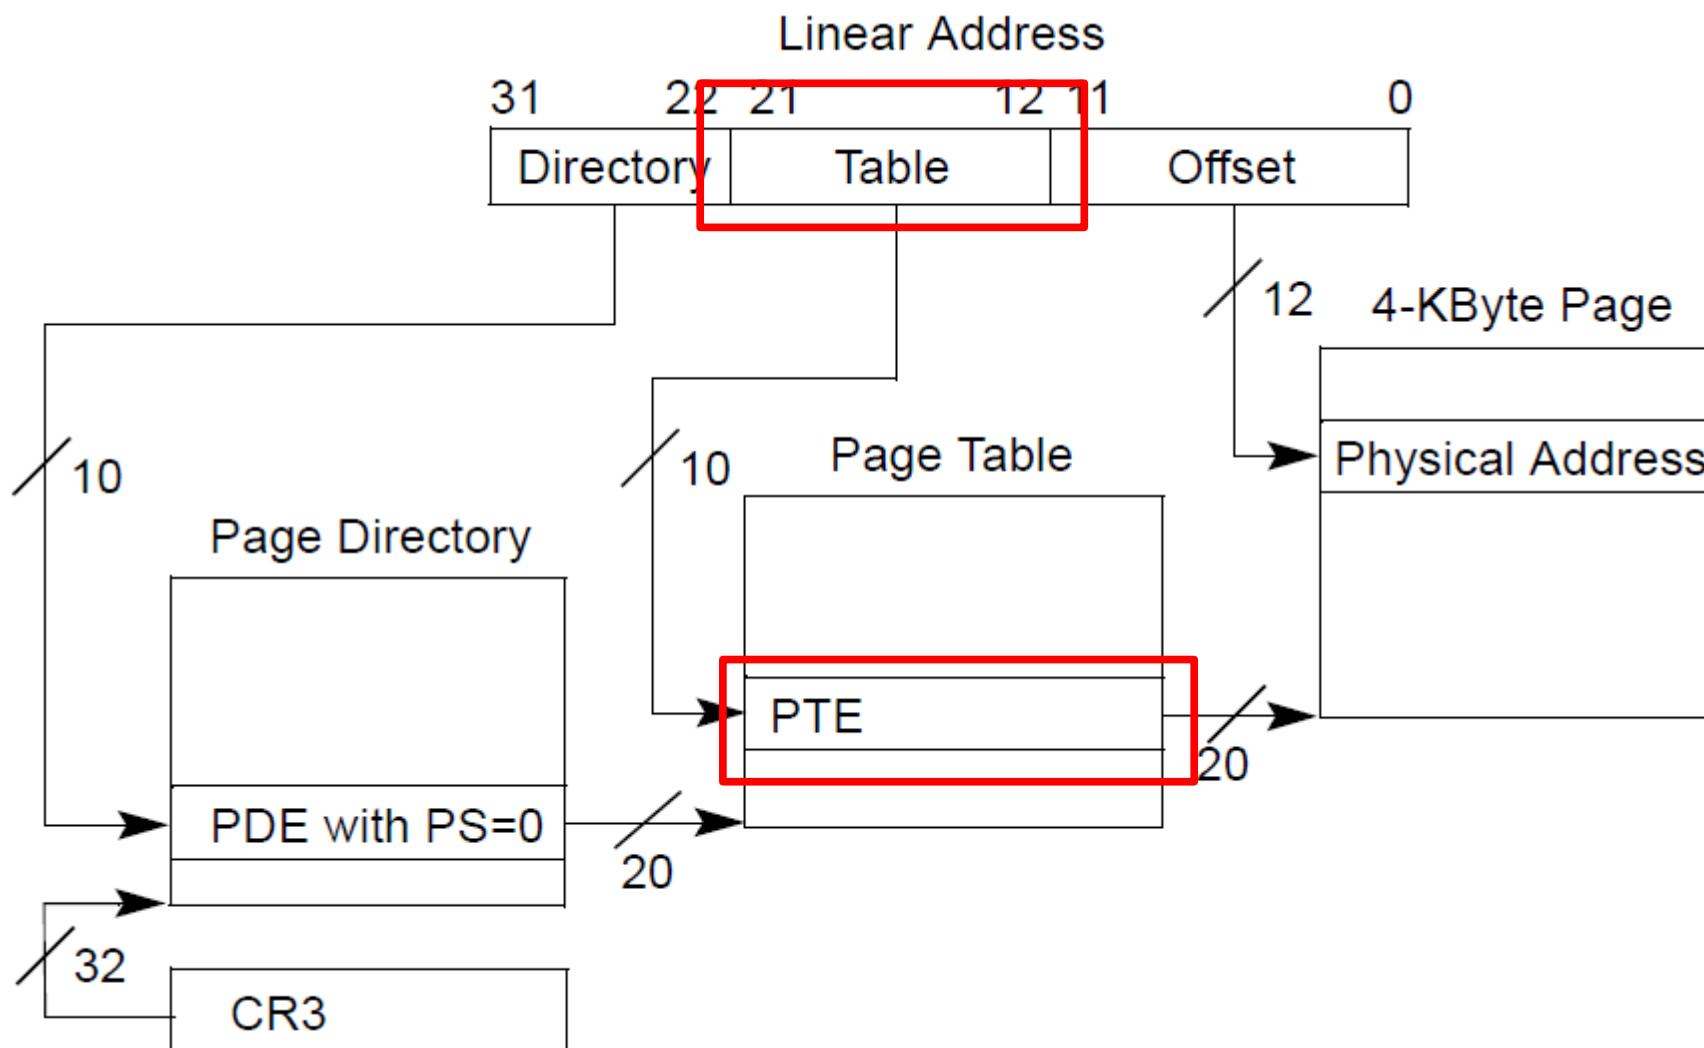
- 20 bit address of the page table
- Wait... 20 bit address, but we need 32 bits

# Page directory entry (PDE)

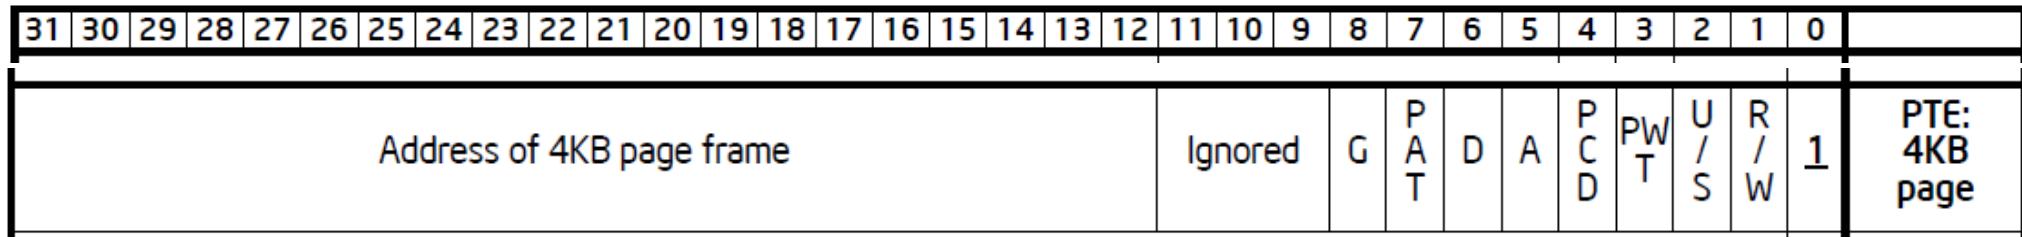


- 20 bit address of the page table
- Wait... 20 bit address, but we need 32 bits
  - Pages 4KB each, we need 1M to cover 4GB
  - Pages start at 4KB (page aligned boundary)

# Page translation

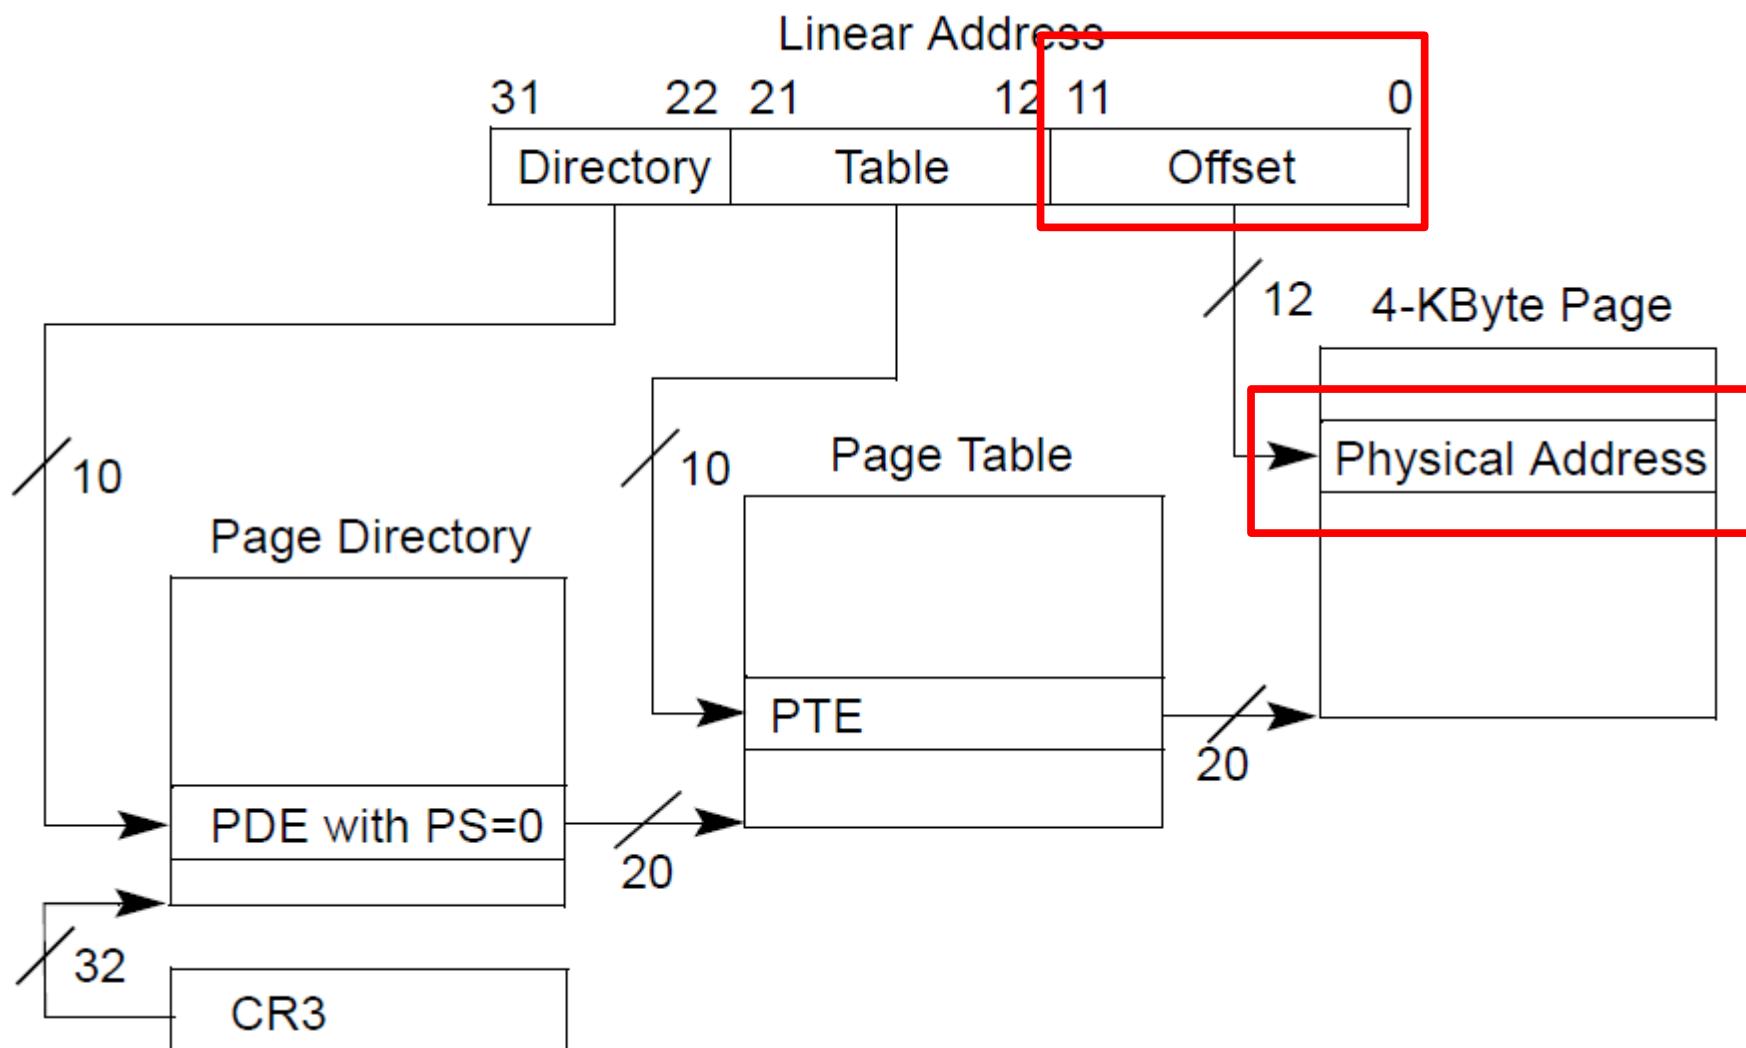


# Page table entry (PTE)



- 20 bit address of the 4KB page
  - Pages 4KB each, we need 1M to cover 4GB

# Page translation

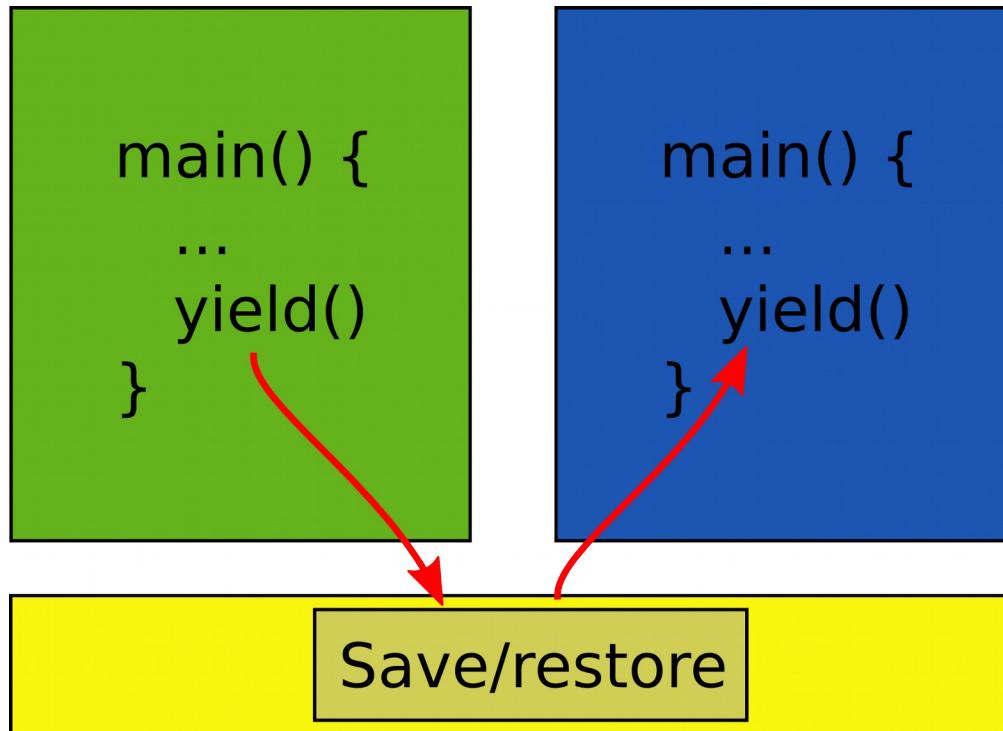


# Benefit of page tables

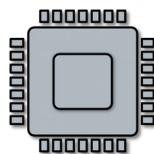
... Compared to arrays?

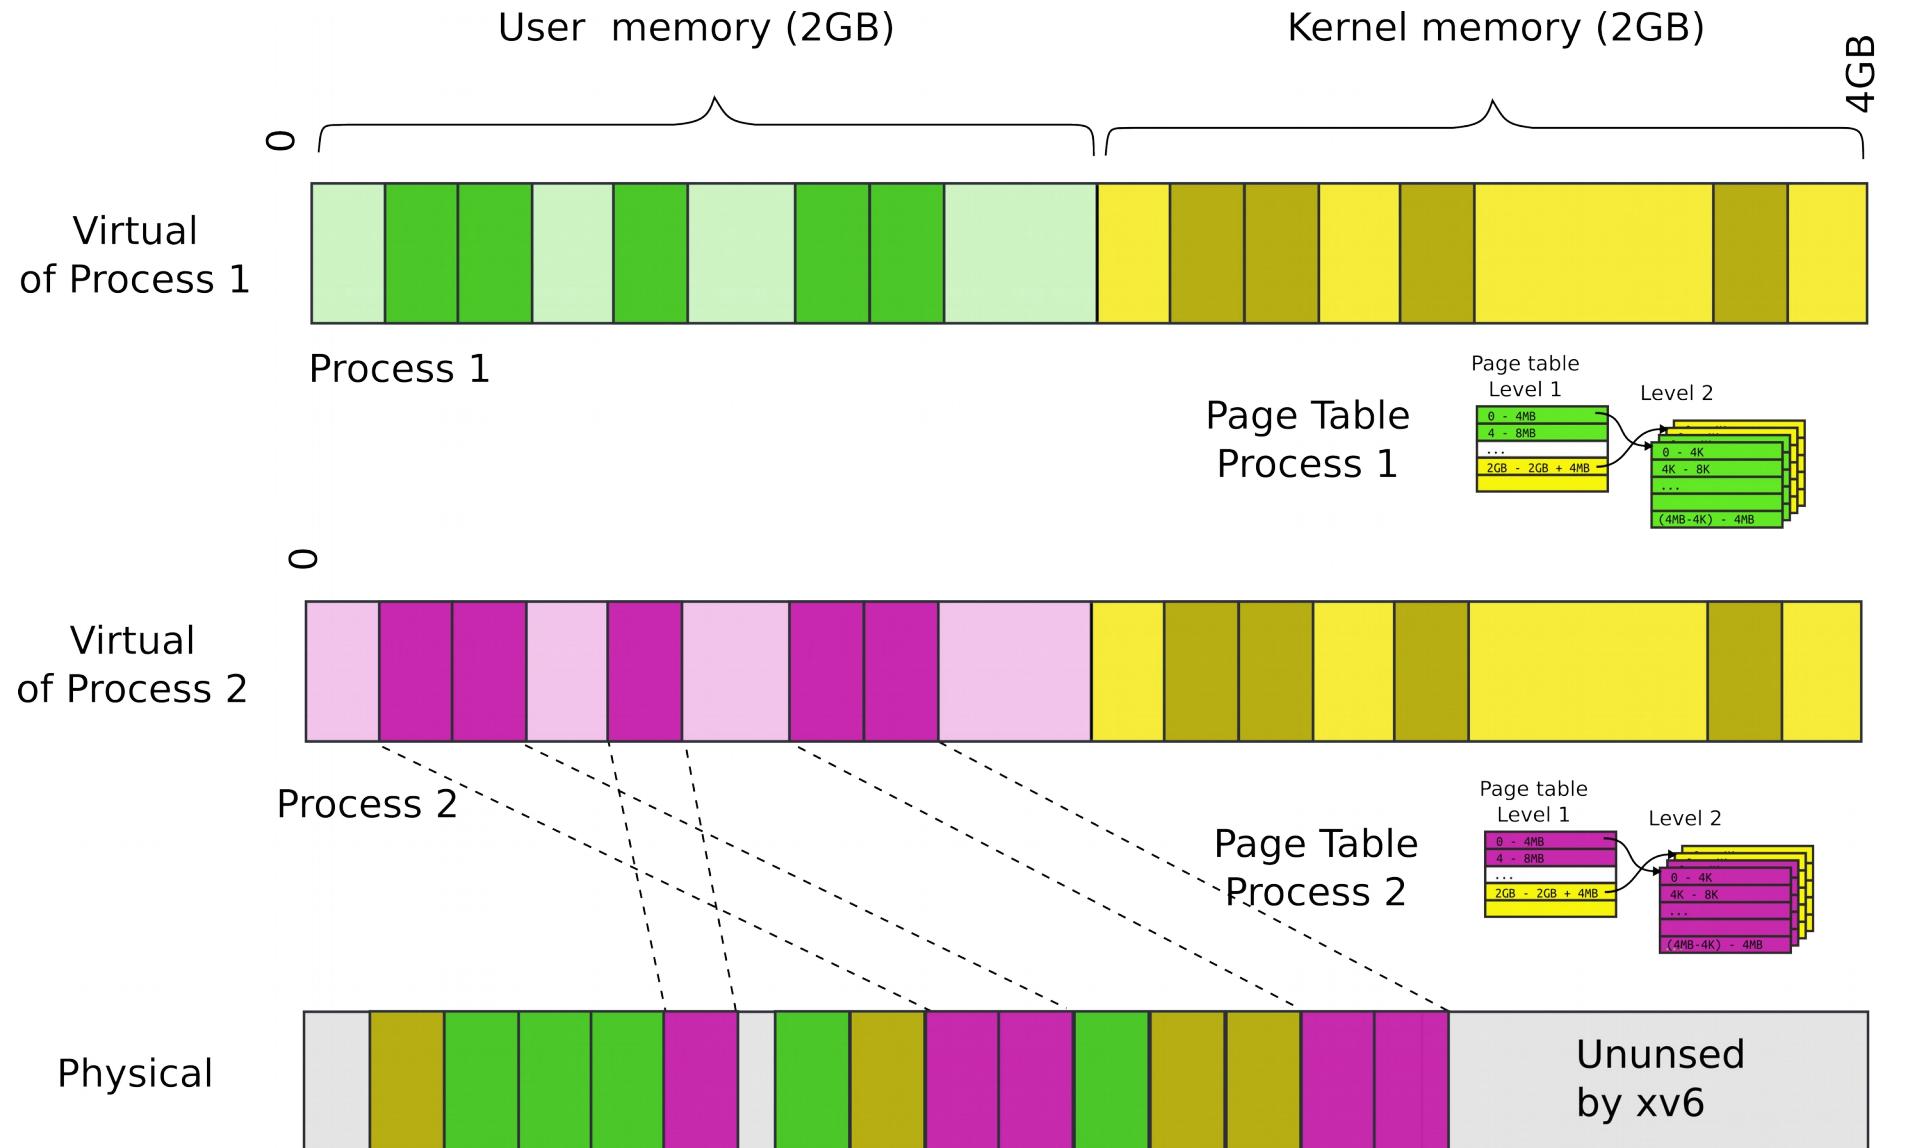
- Page tables represent sparse address space more efficiently
  - An entire array has to be allocated upfront
  - But if the address space uses a handful of pages
  - Only page tables (Level 1 and 2 need to be allocated to describe translation)
- On a dense address space this benefit goes away
  - I'll assign a homework!

# What about isolation?



- Two programs, one memory?
- Each process has its own page table
  - OS switches between them





P1 and P2 can't access each other memory

0xe000000  
(PHYSSTOP)  
234MB

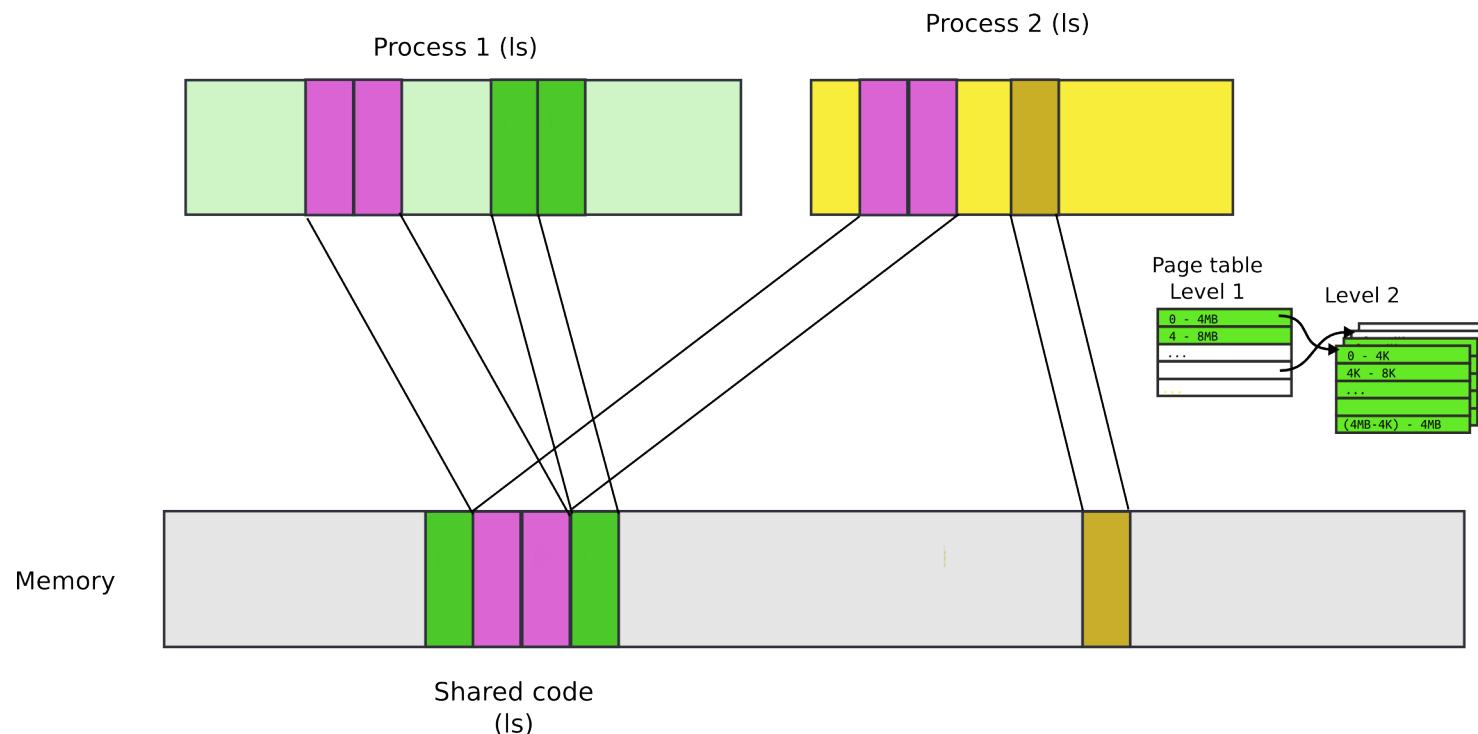
Top of physical memory

# Compared to segments pages allow ...

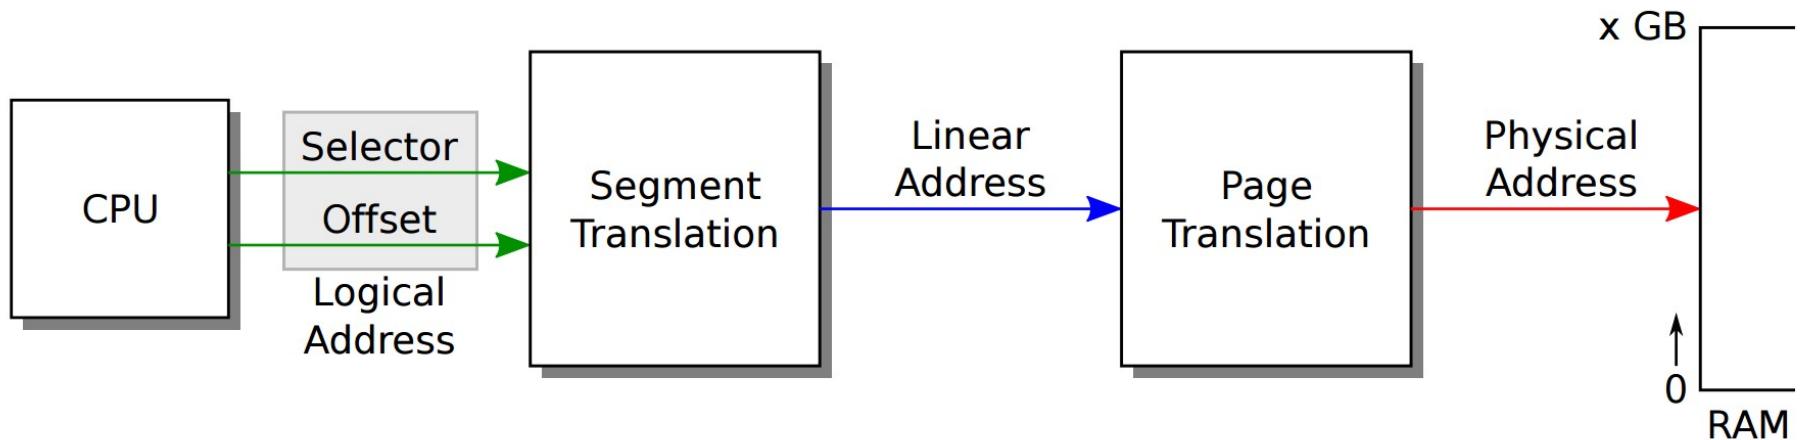
- Emulate large virtual address space on a smaller physical memory
  - In our example we had only 12 physical pages
  - But the program can access all 1M pages in its 4GB address space
  - The OS will move other pages to disk

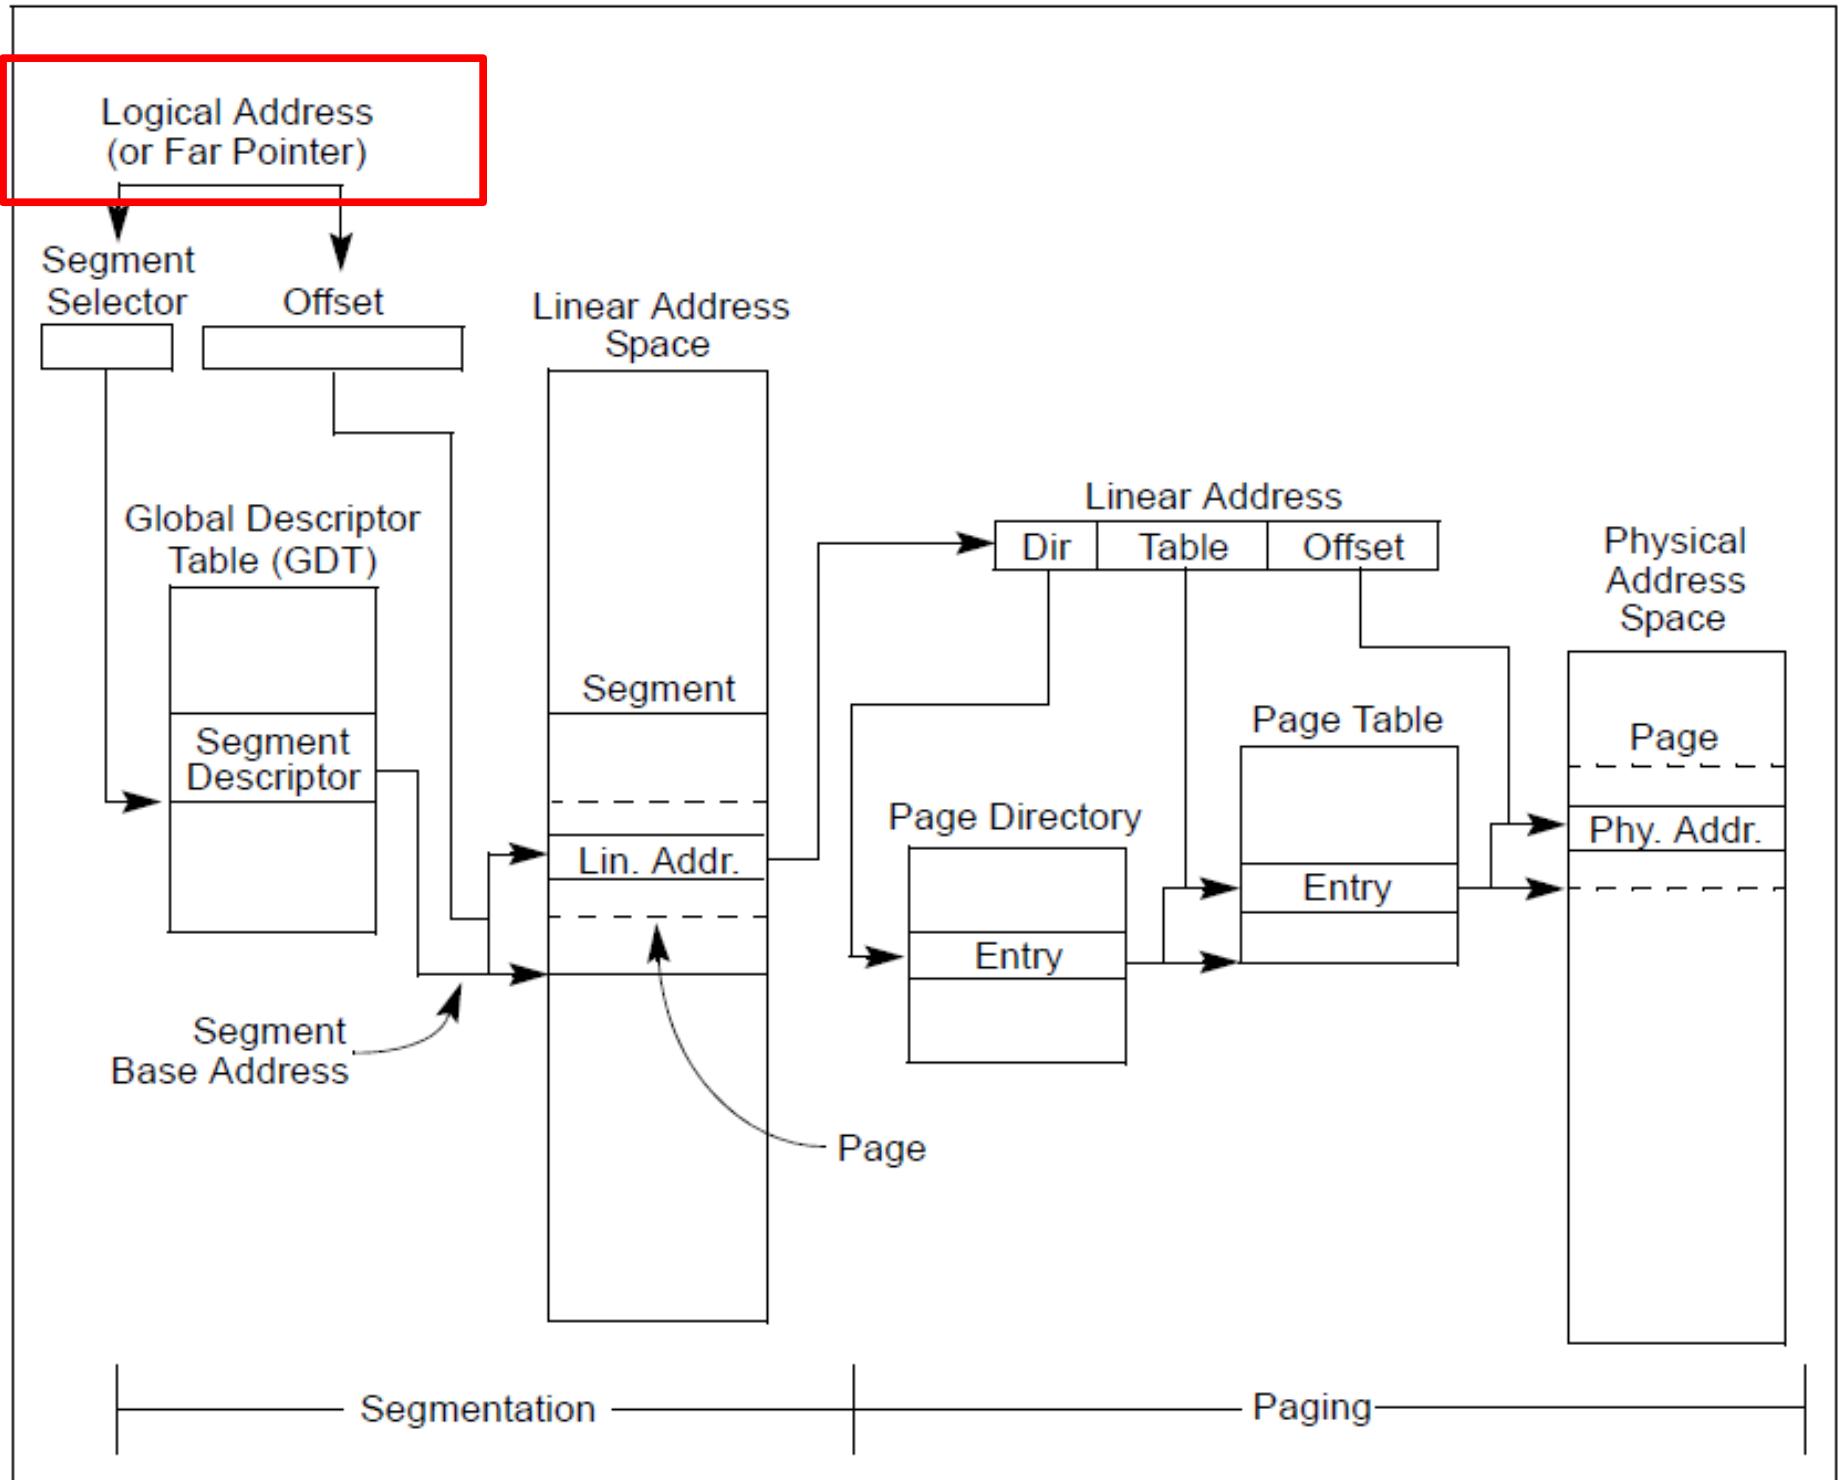
# Compared to segments pages allow ...

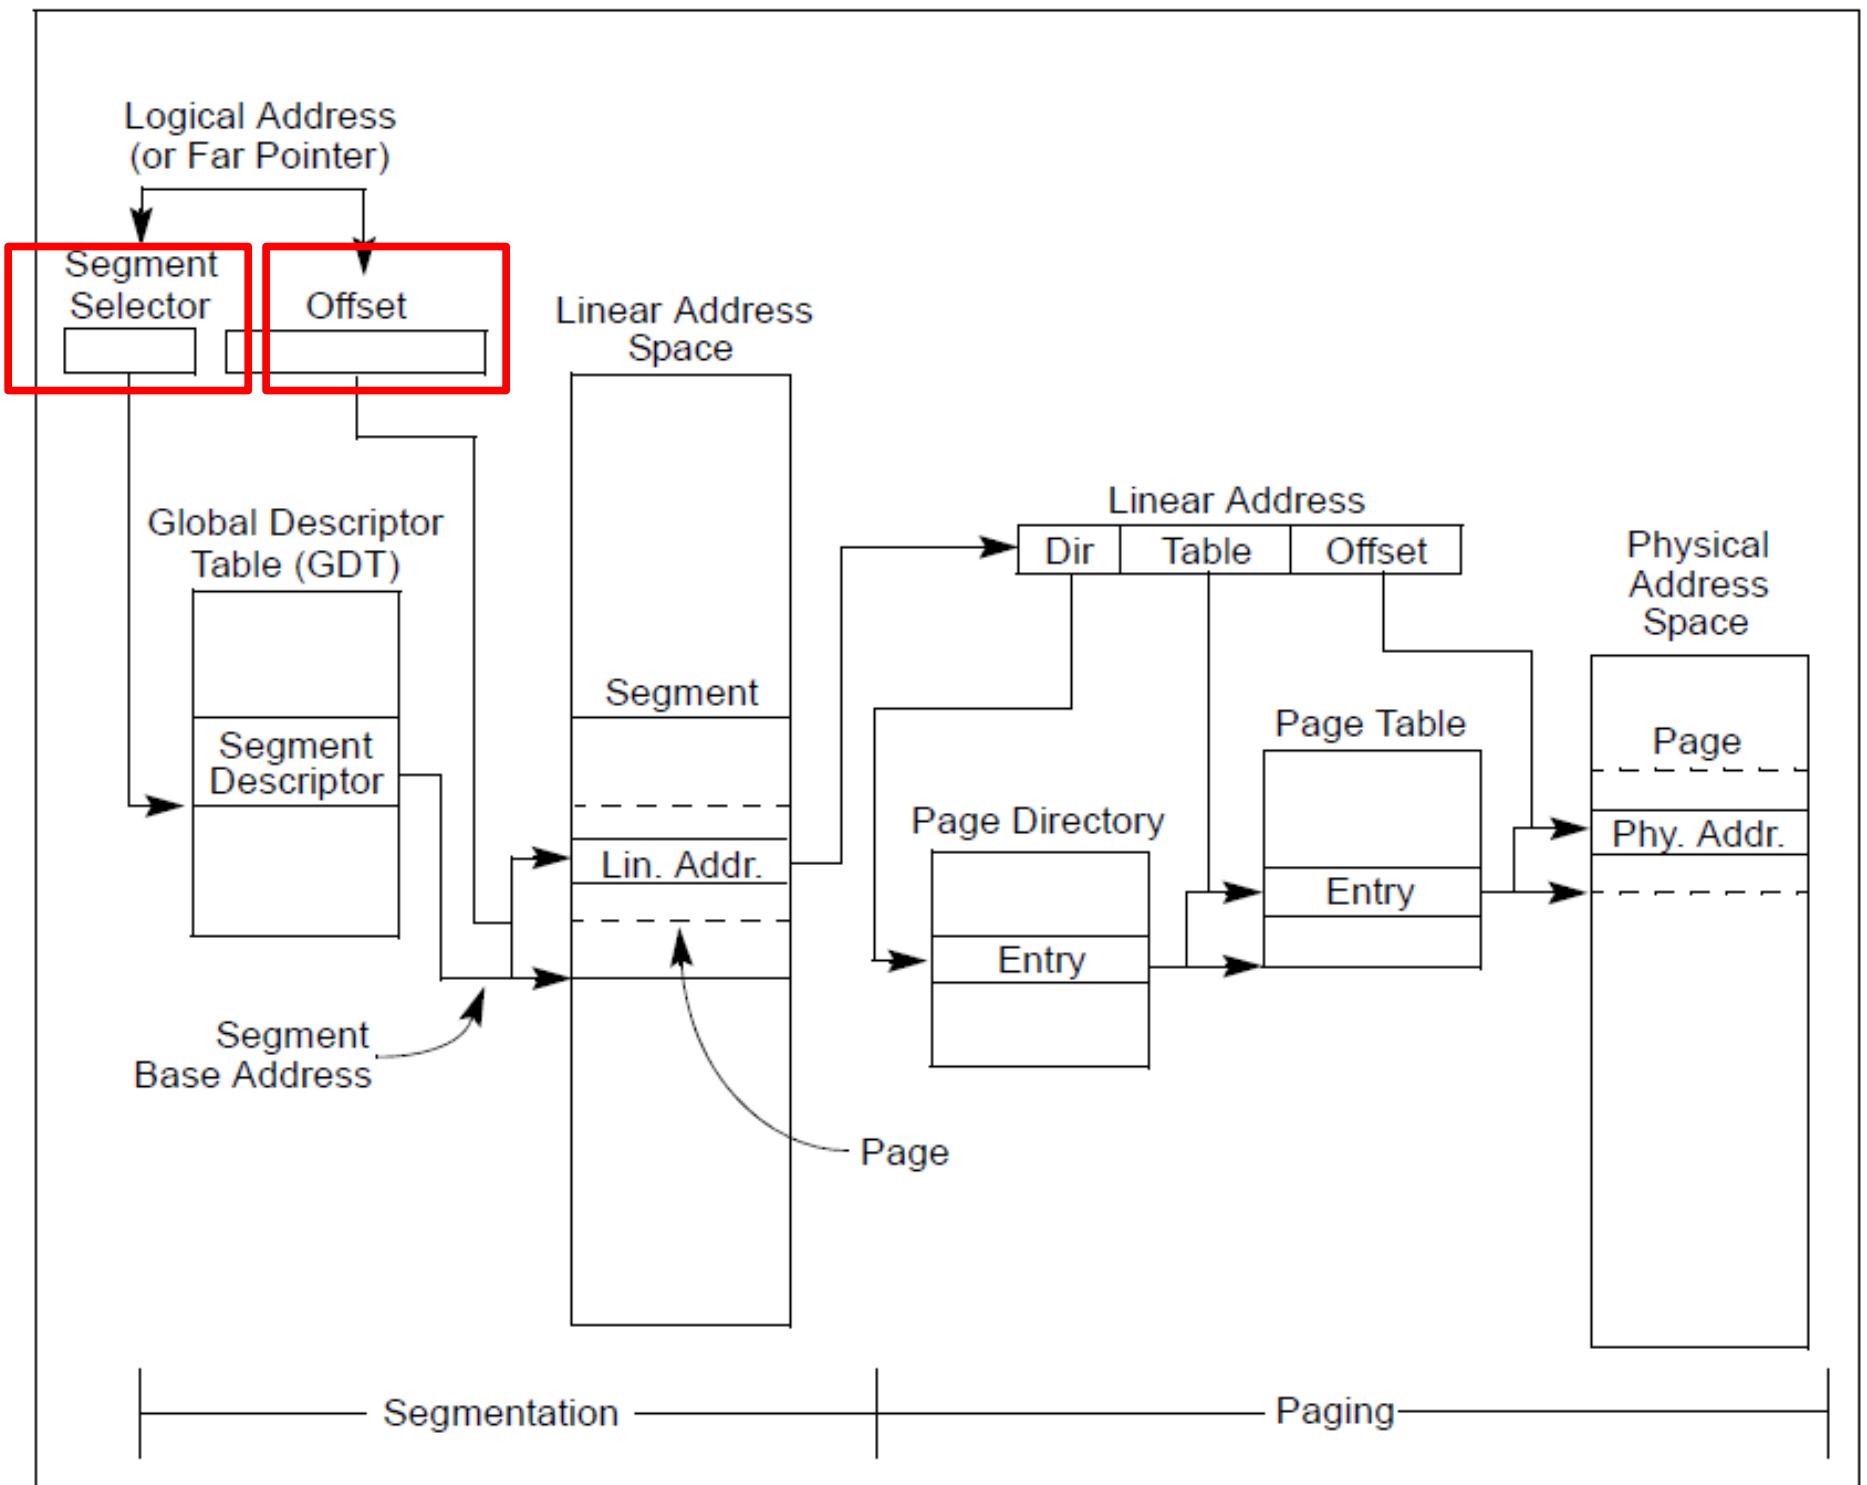
- Share a region of memory across multiple programs
  - Well... segments allow this too
  - Communication (shared buffer of messages)
  - Shared libraries

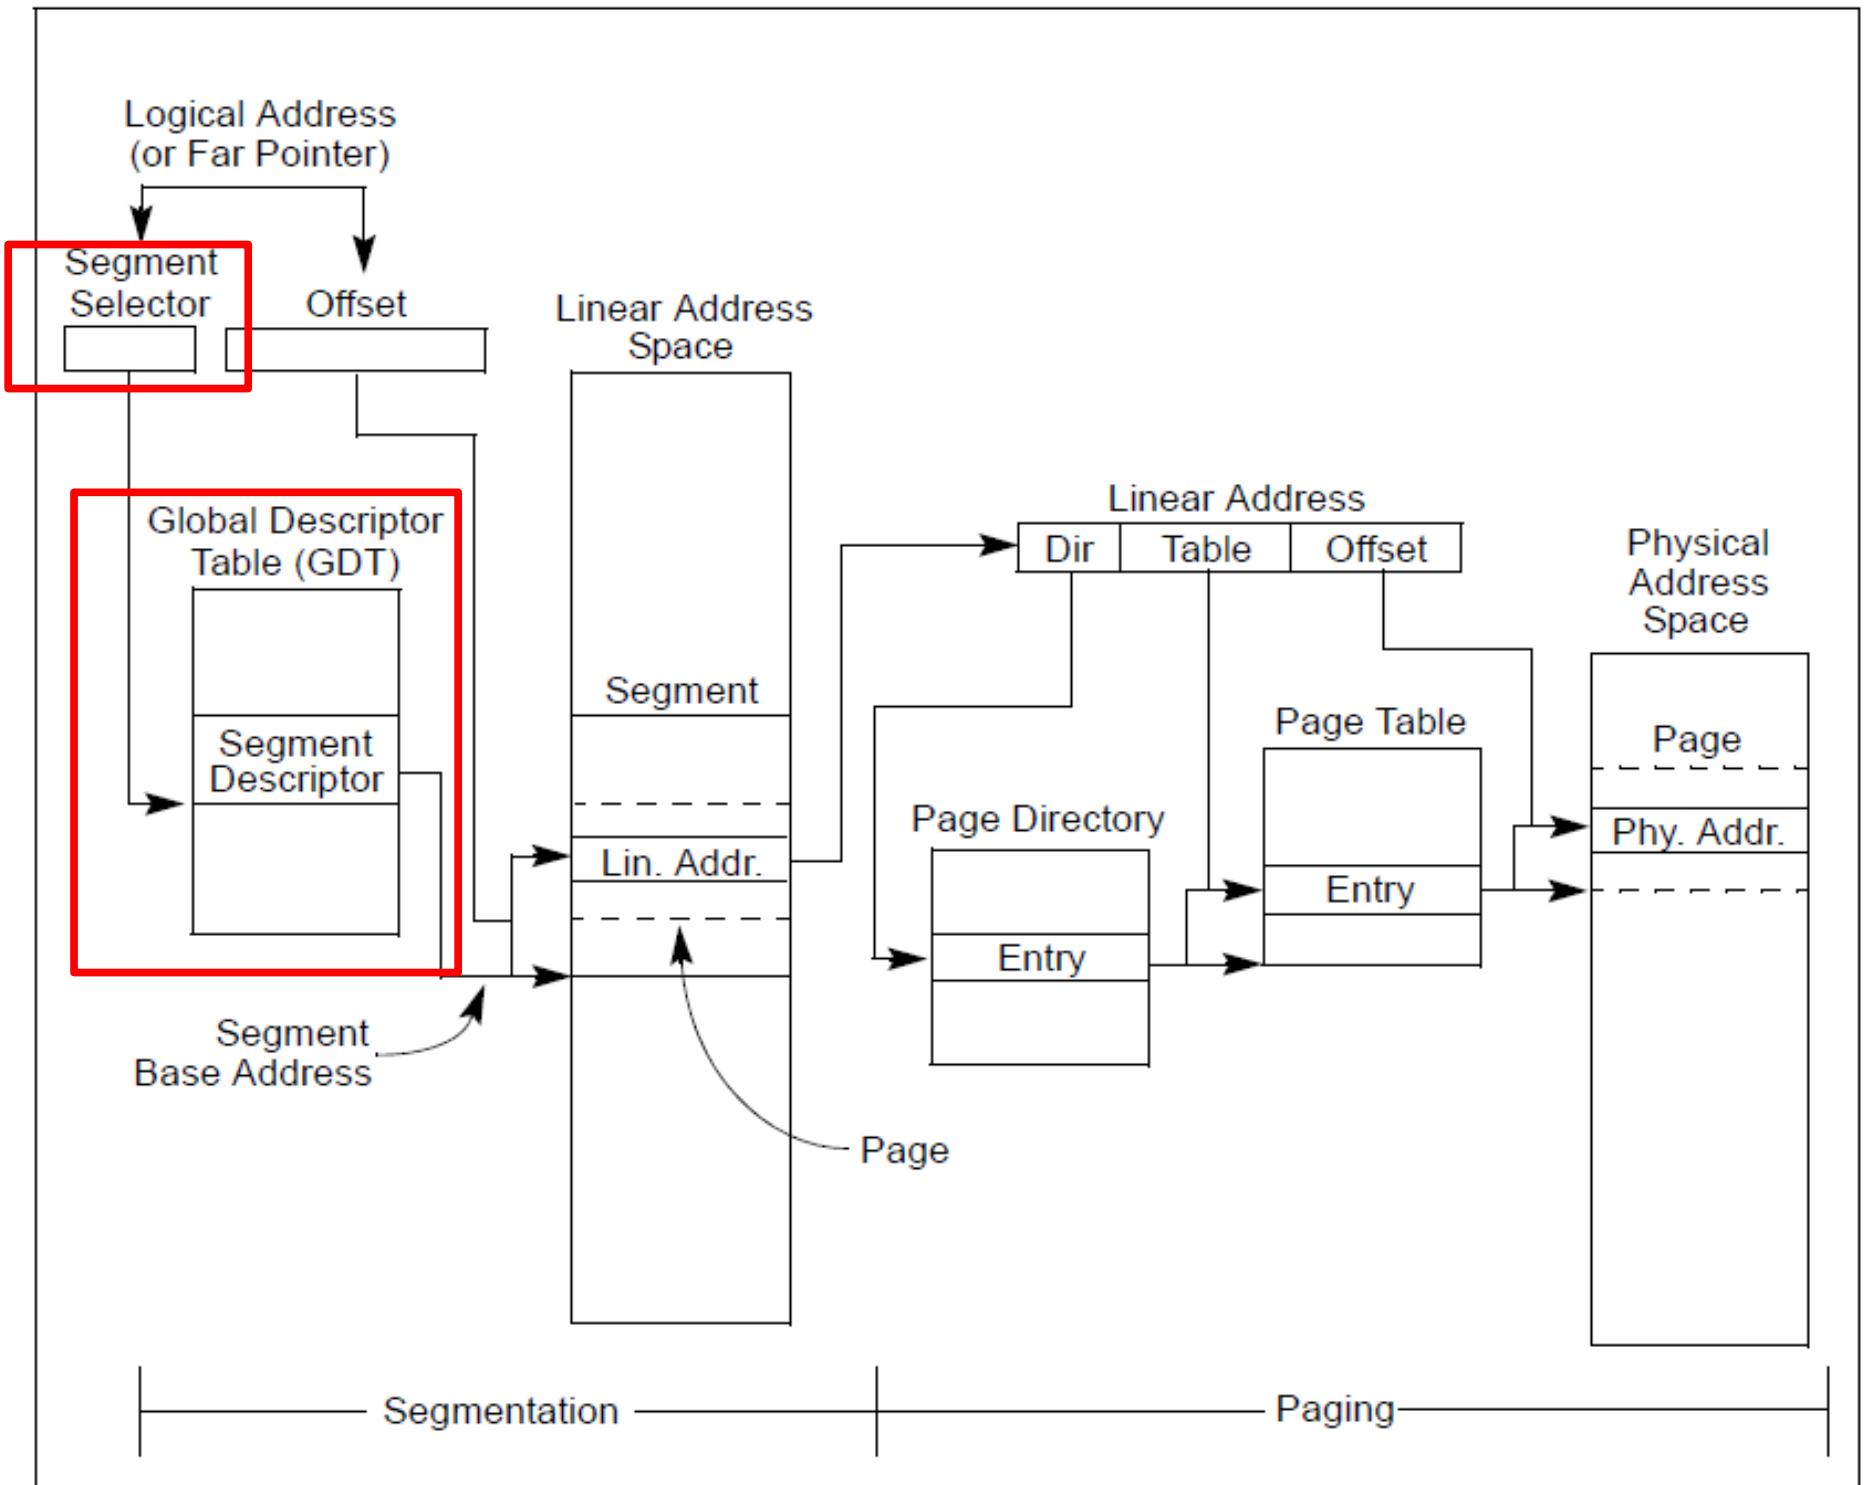


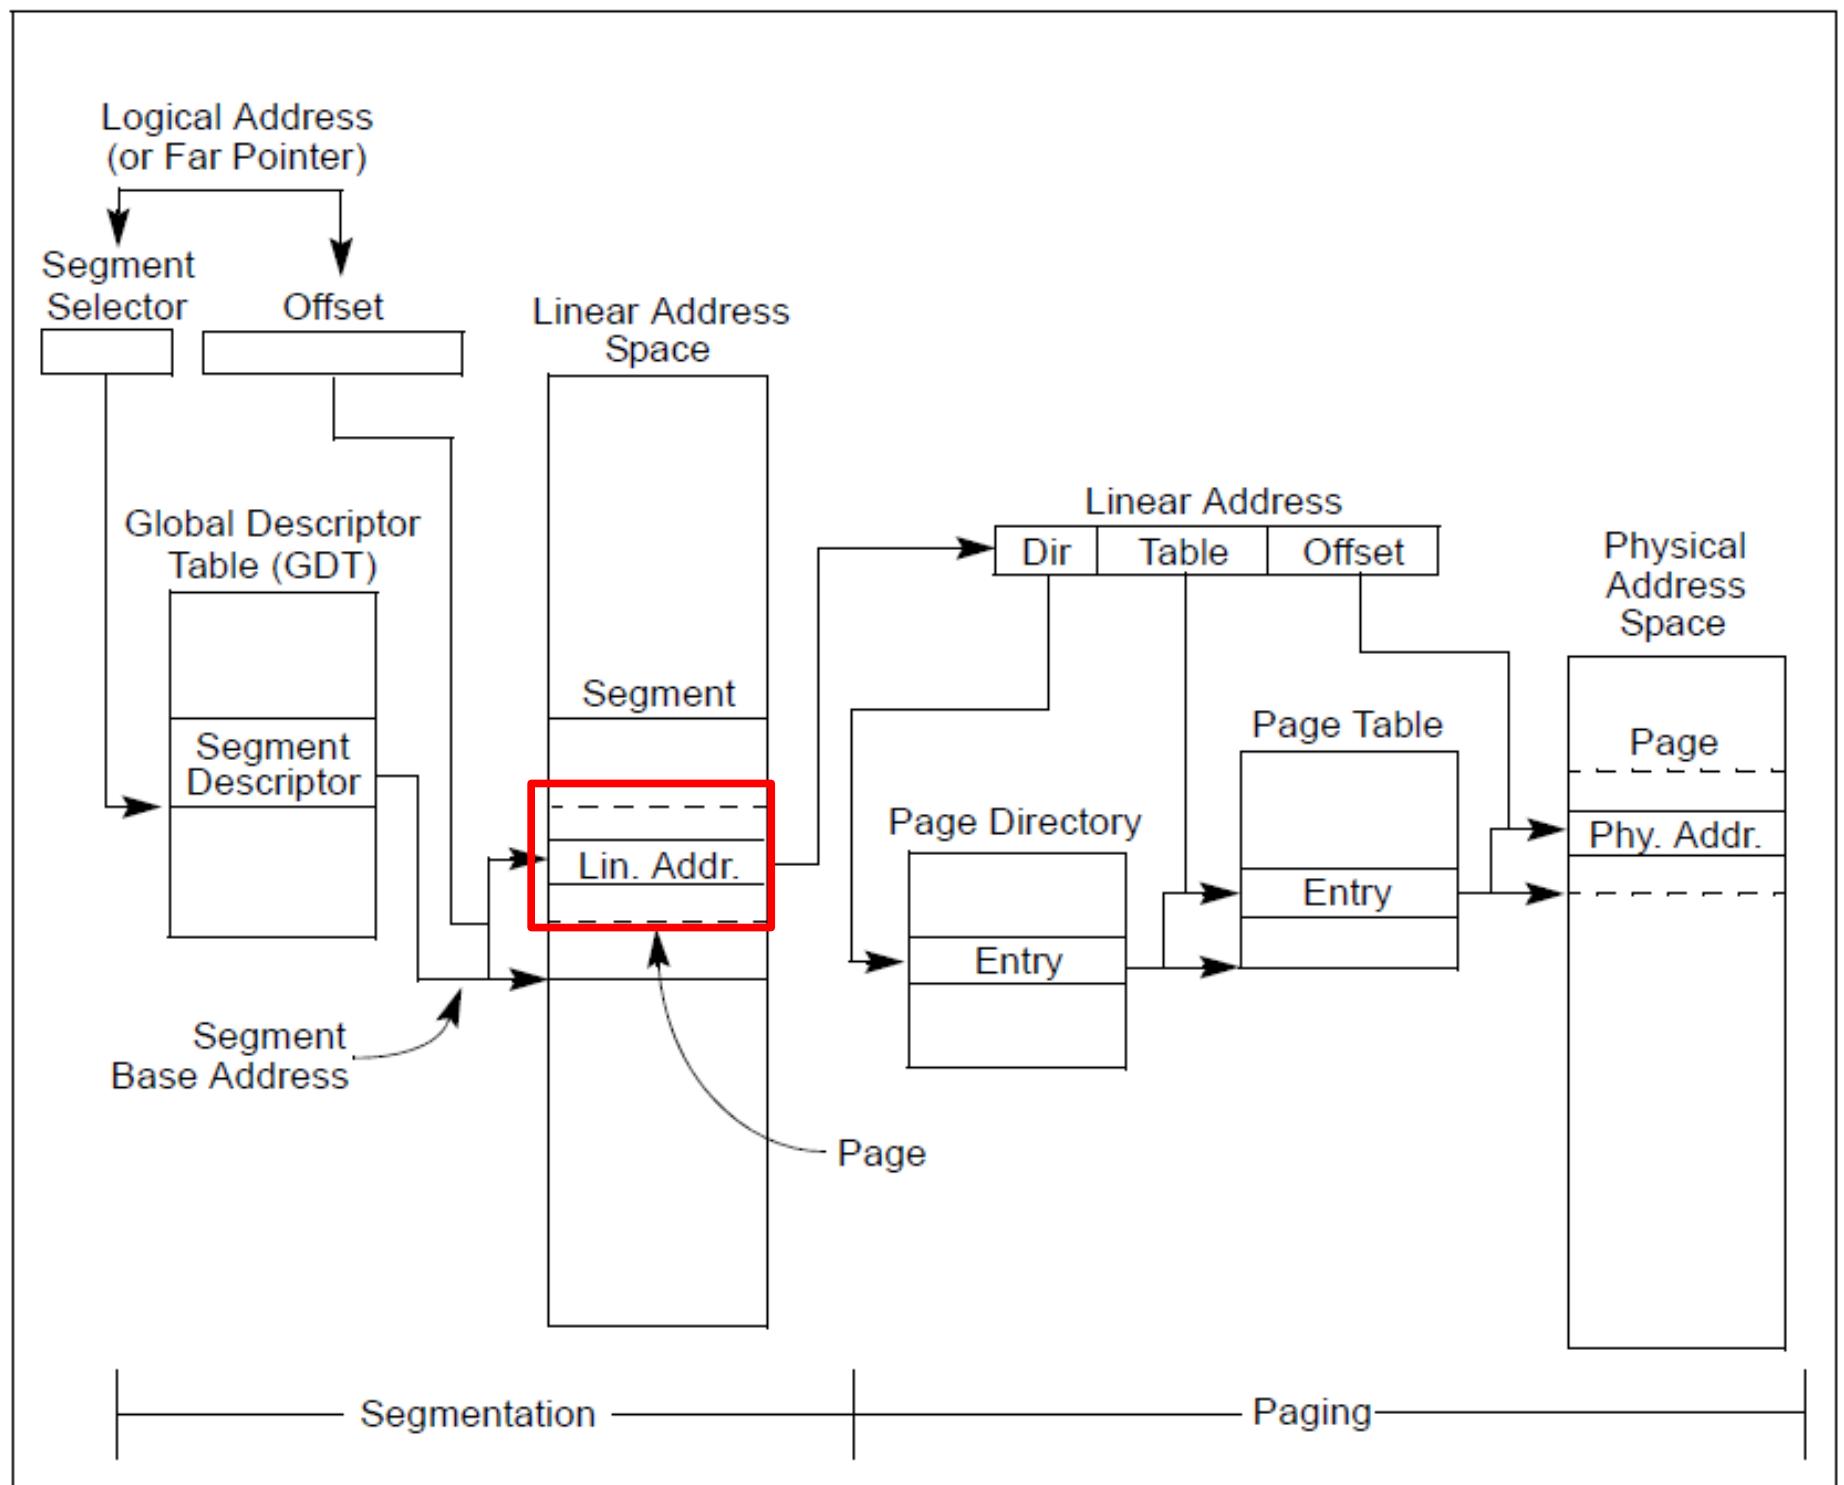
# Recap: complete address translation

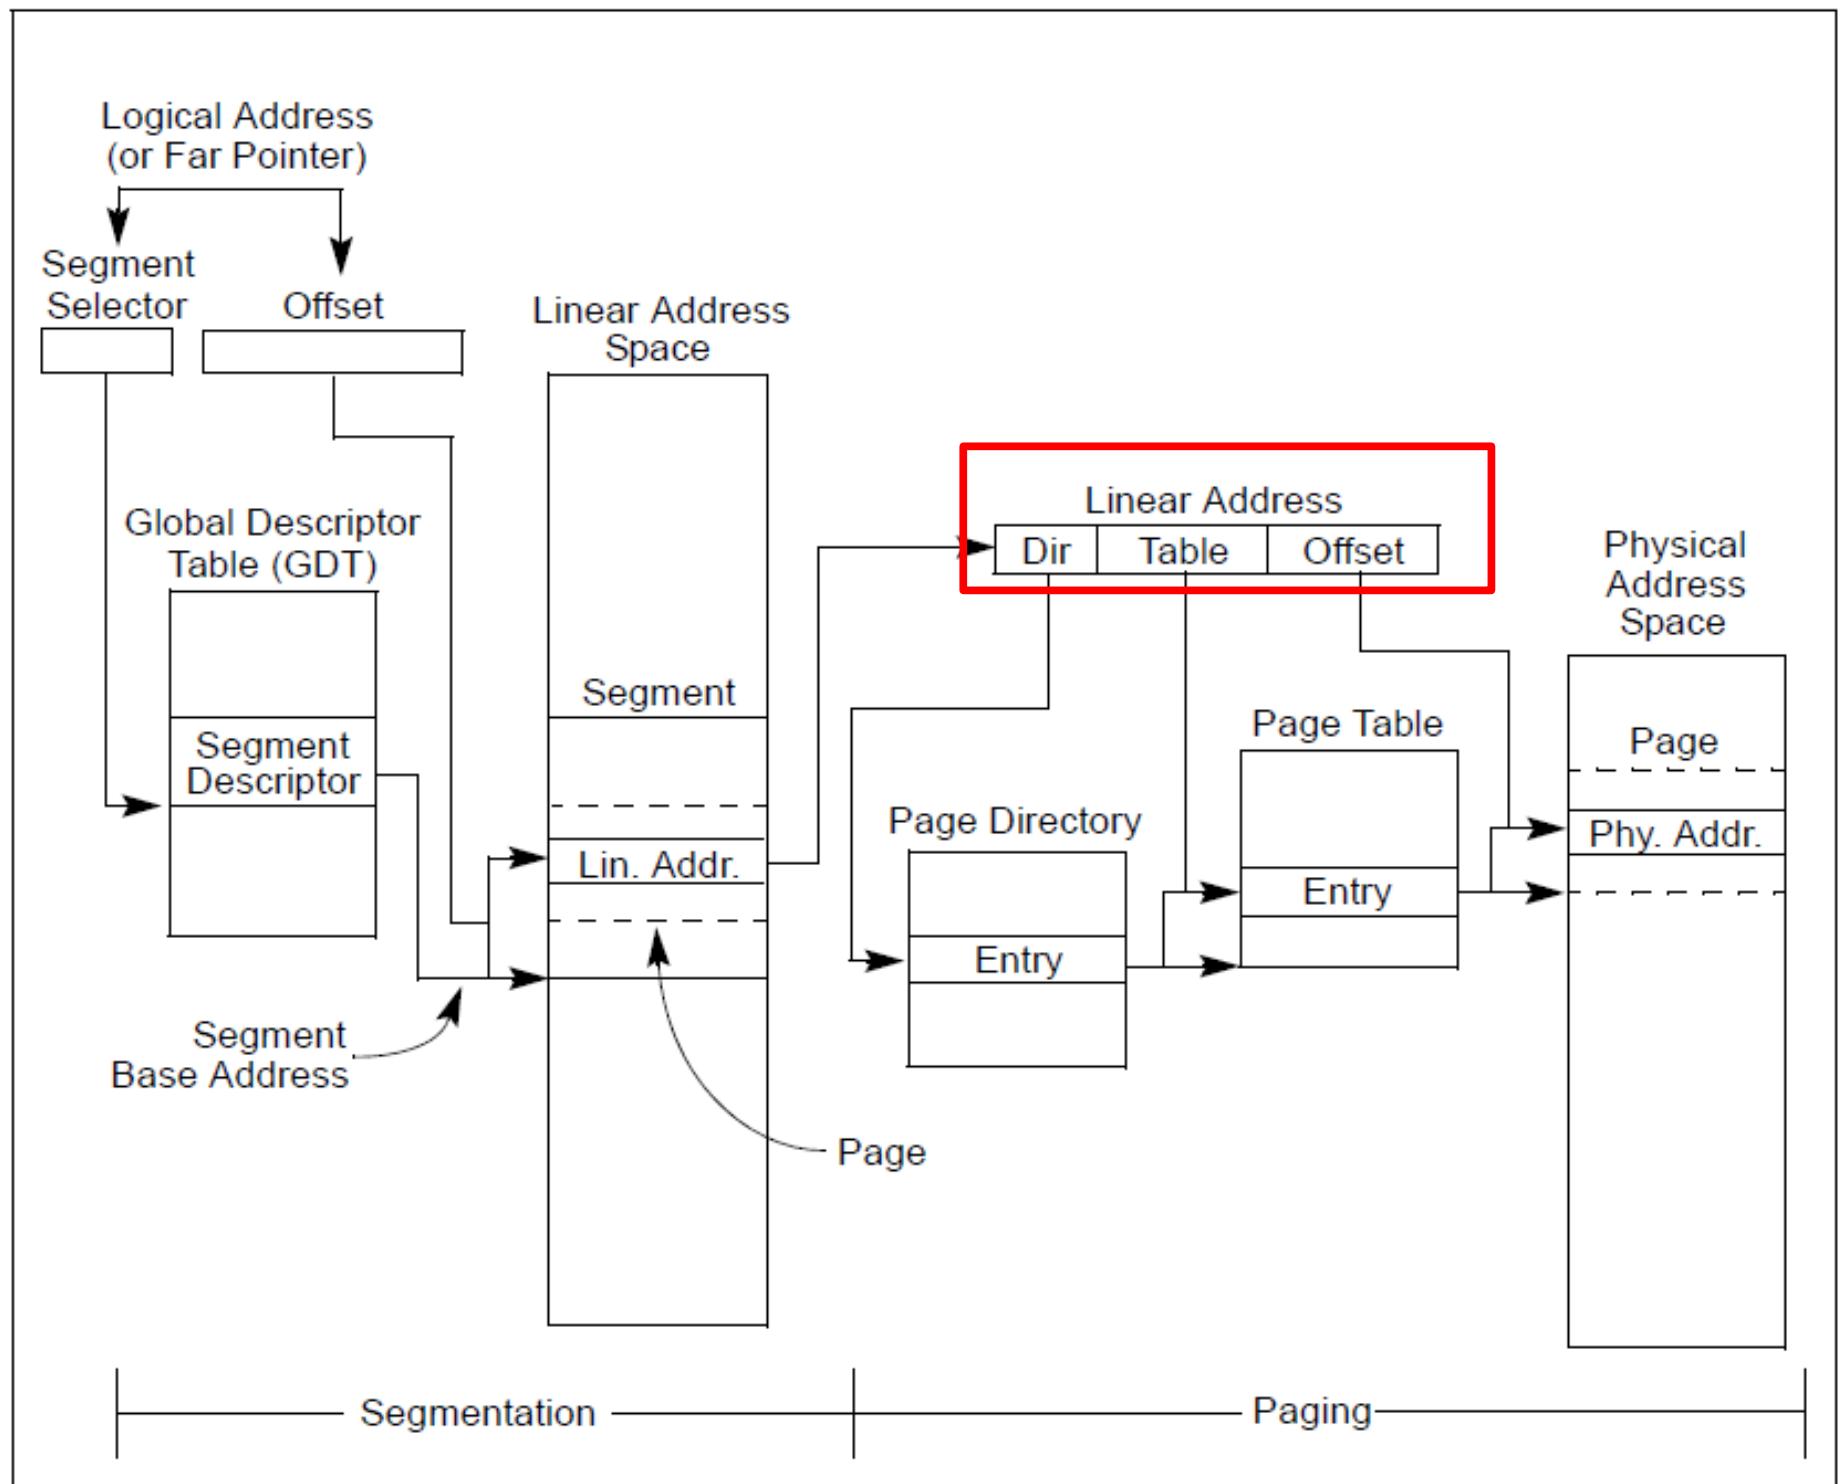


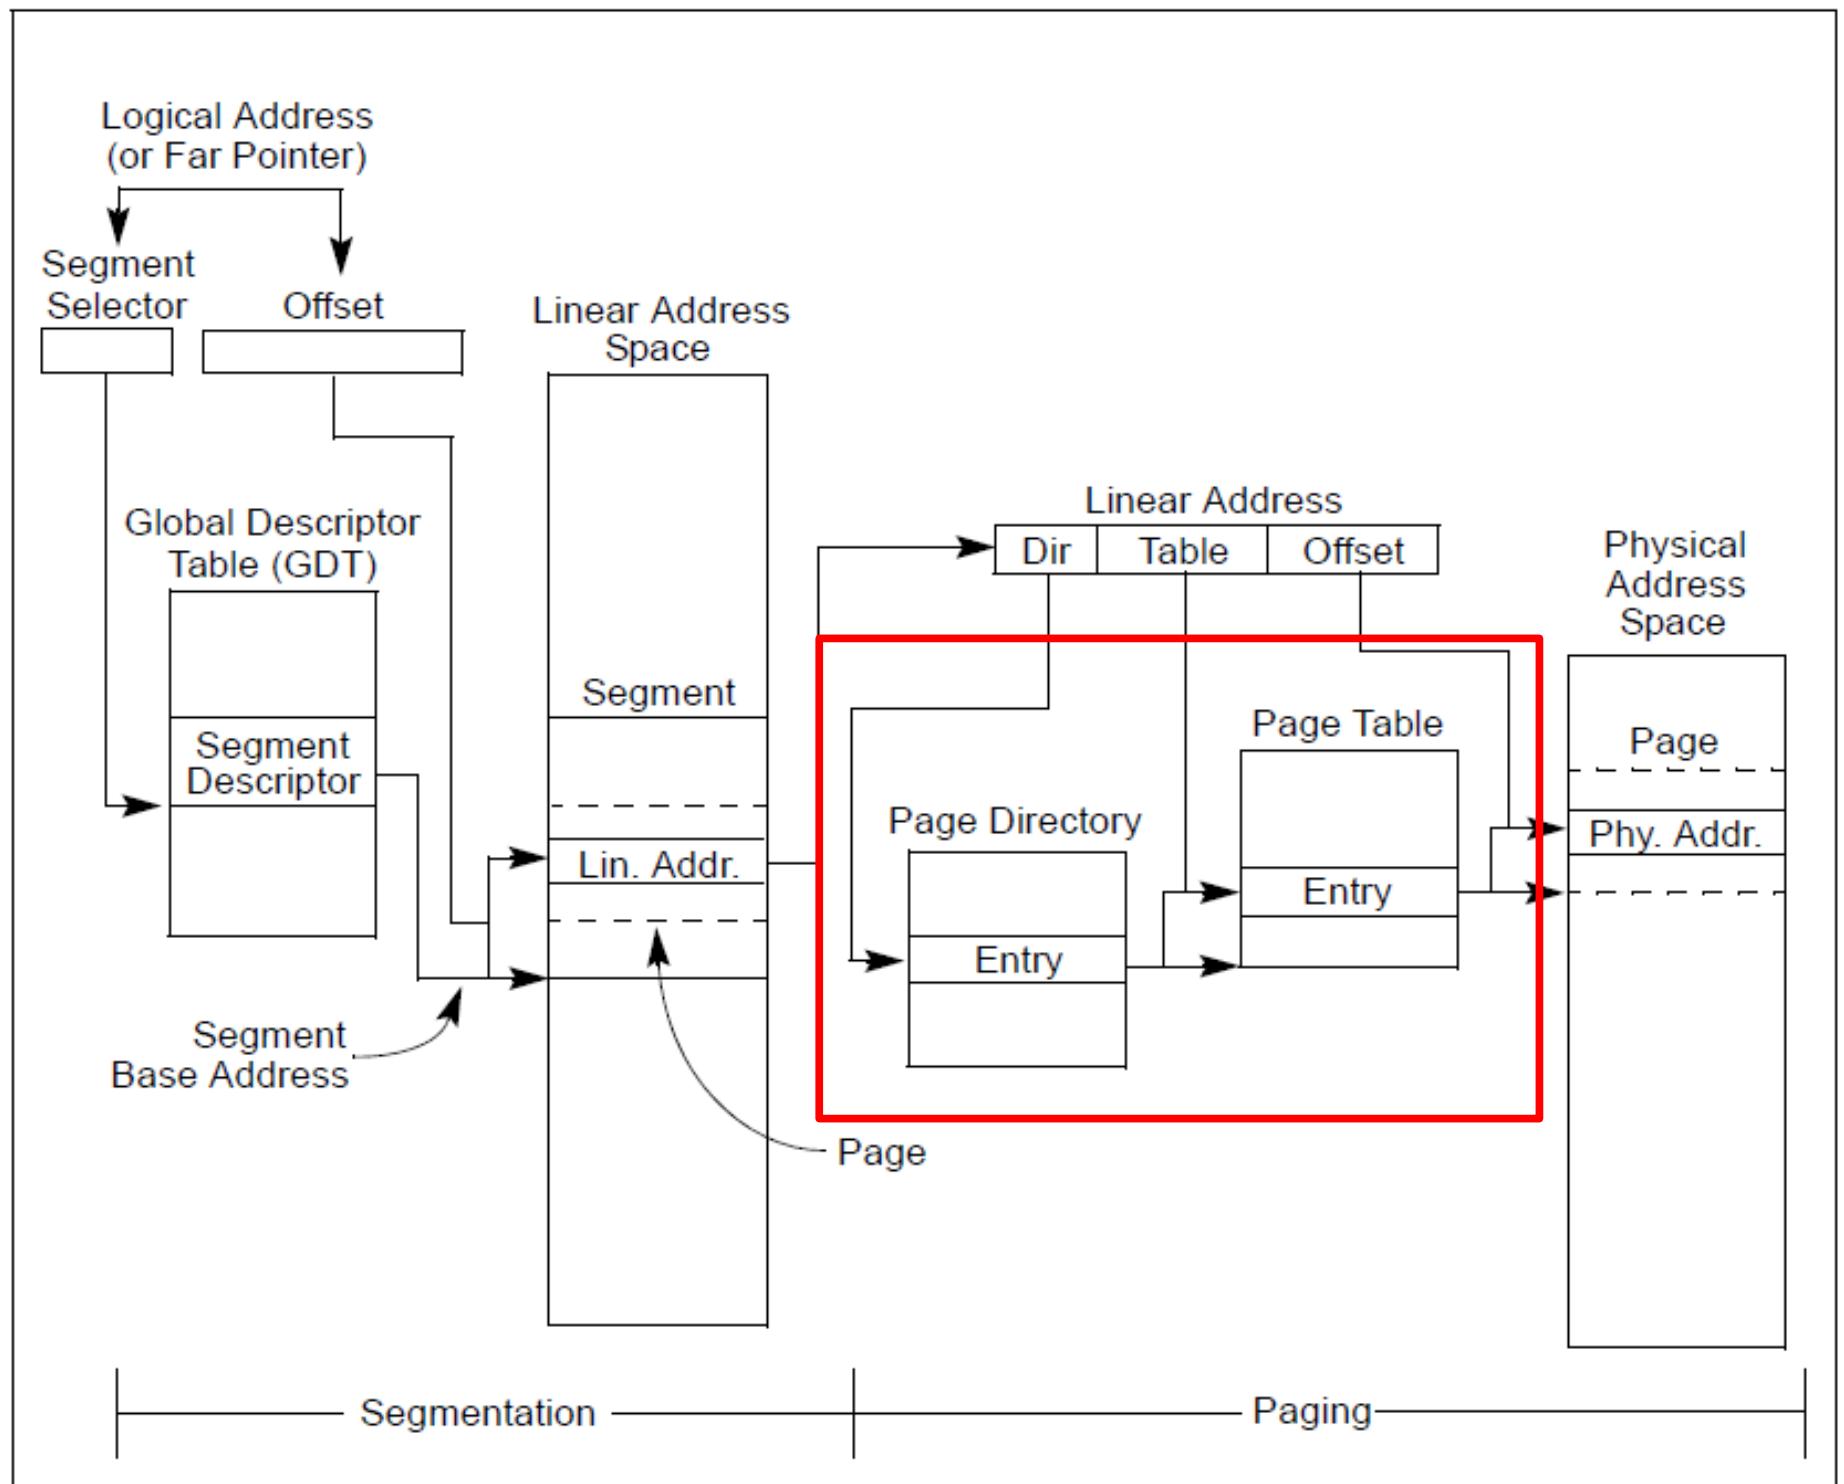


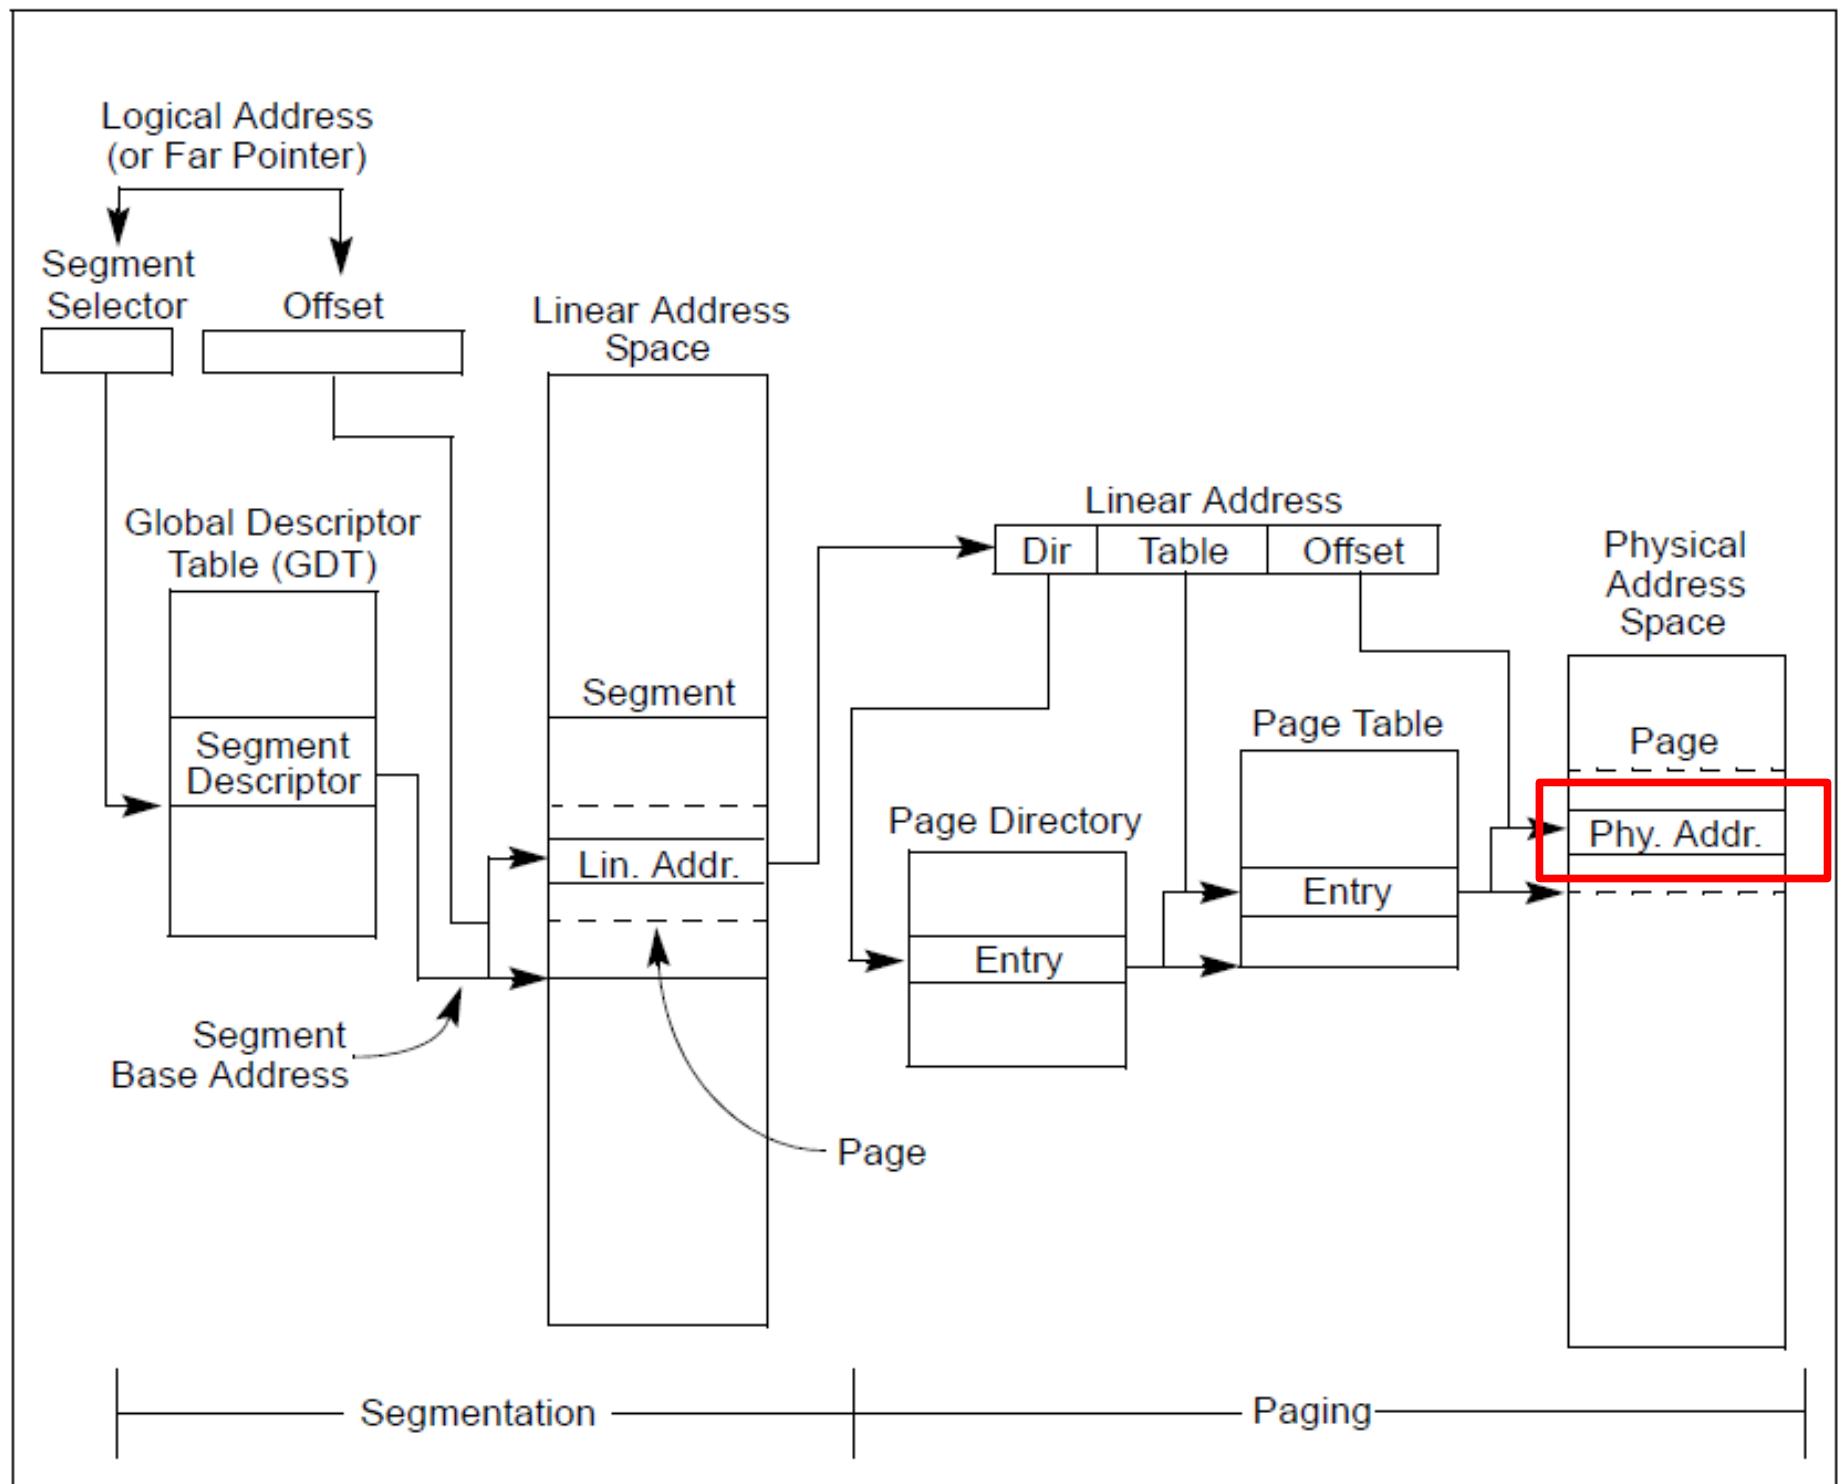


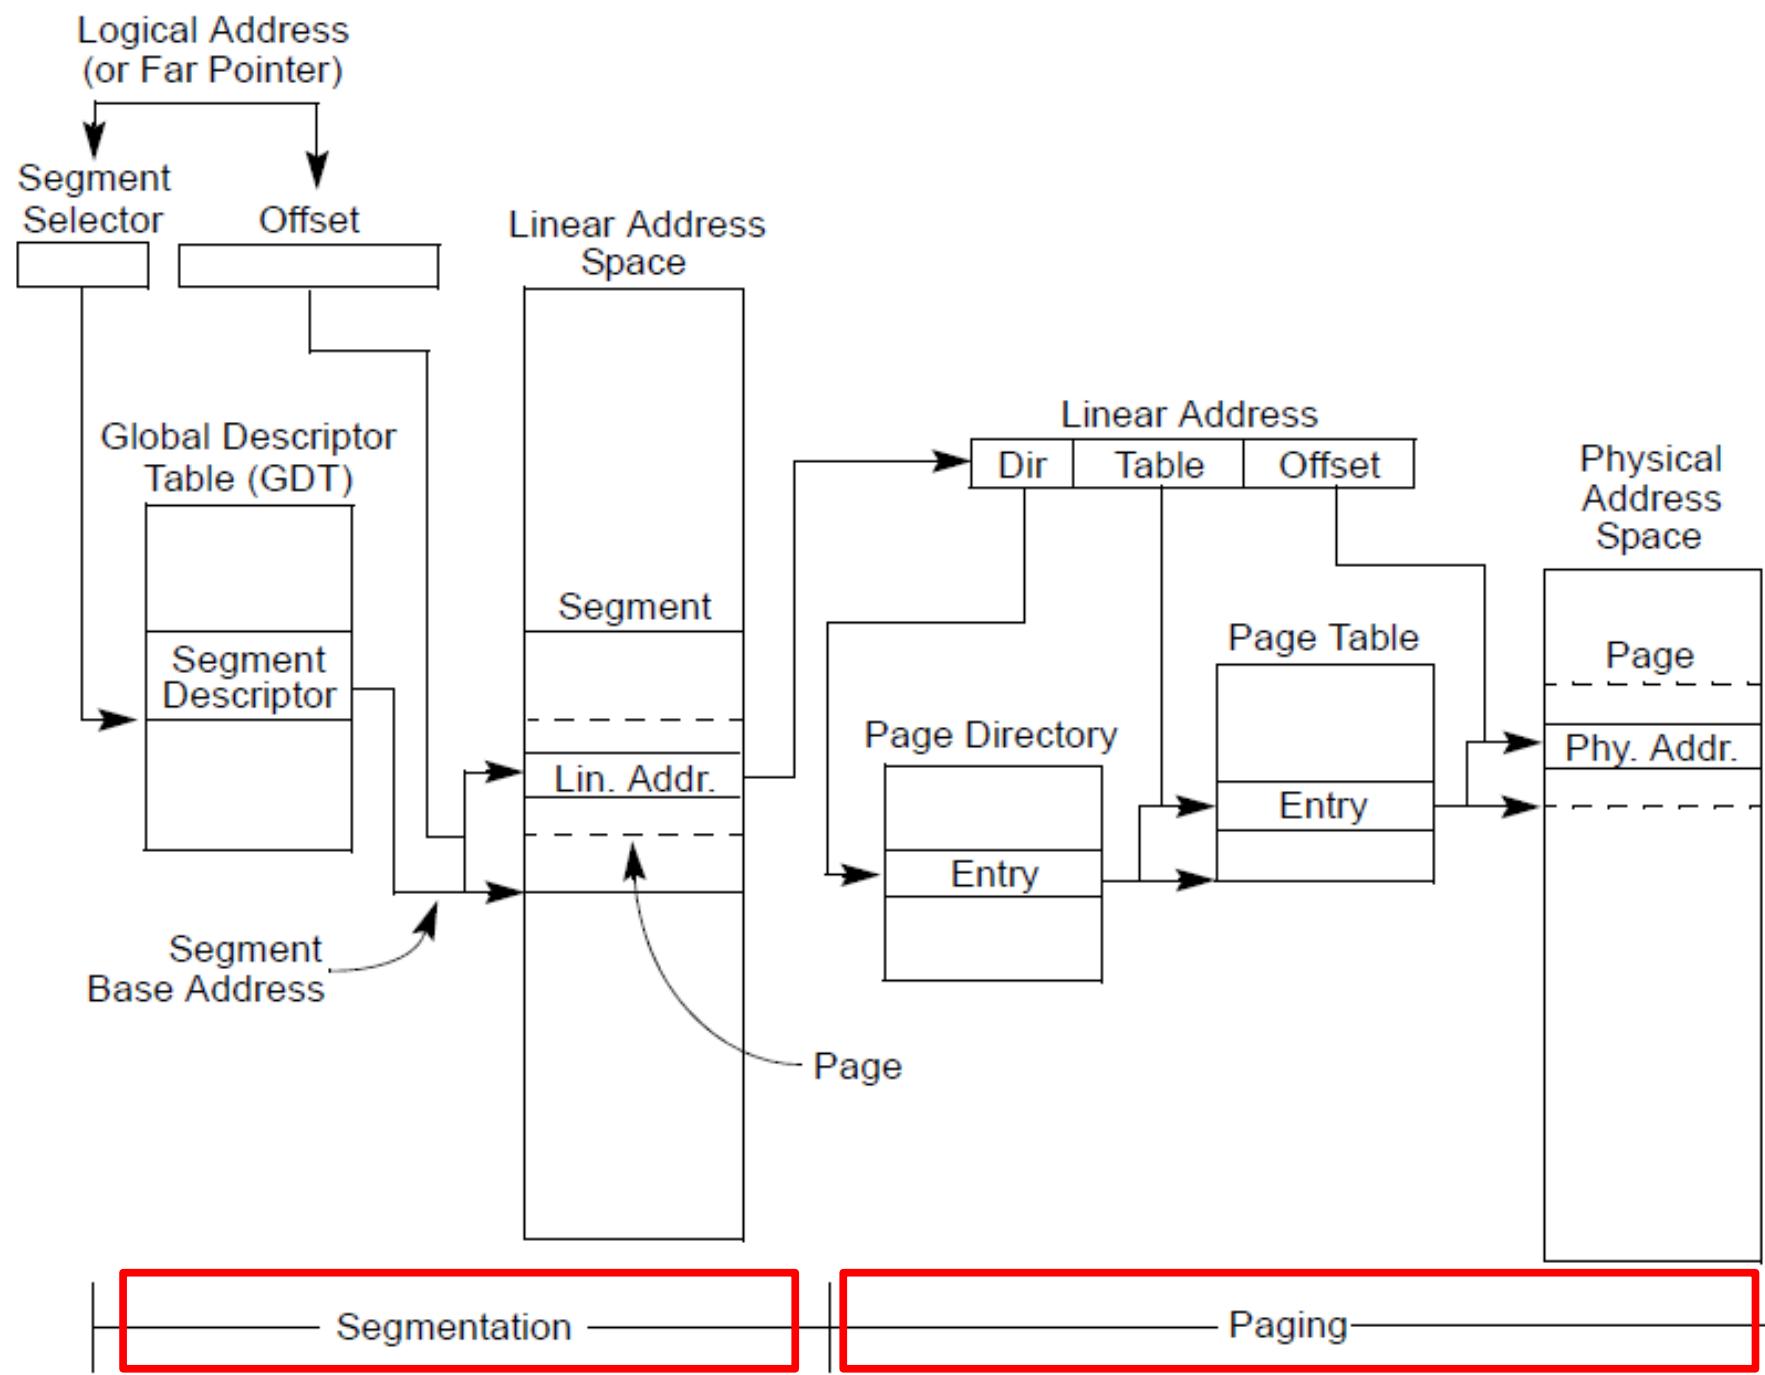








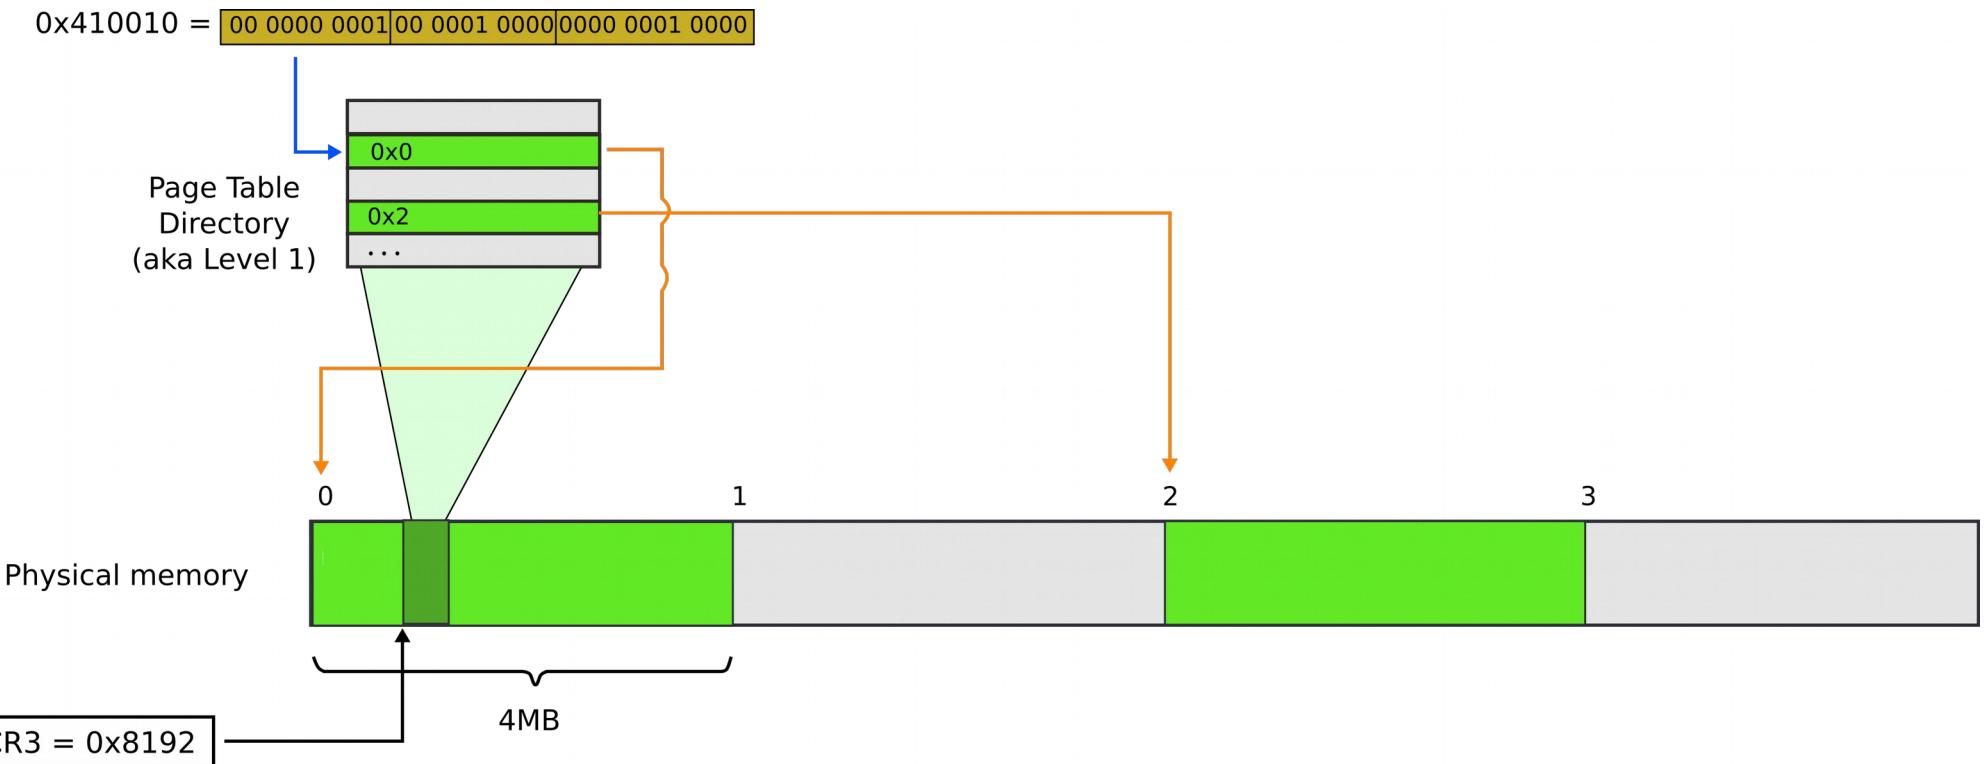




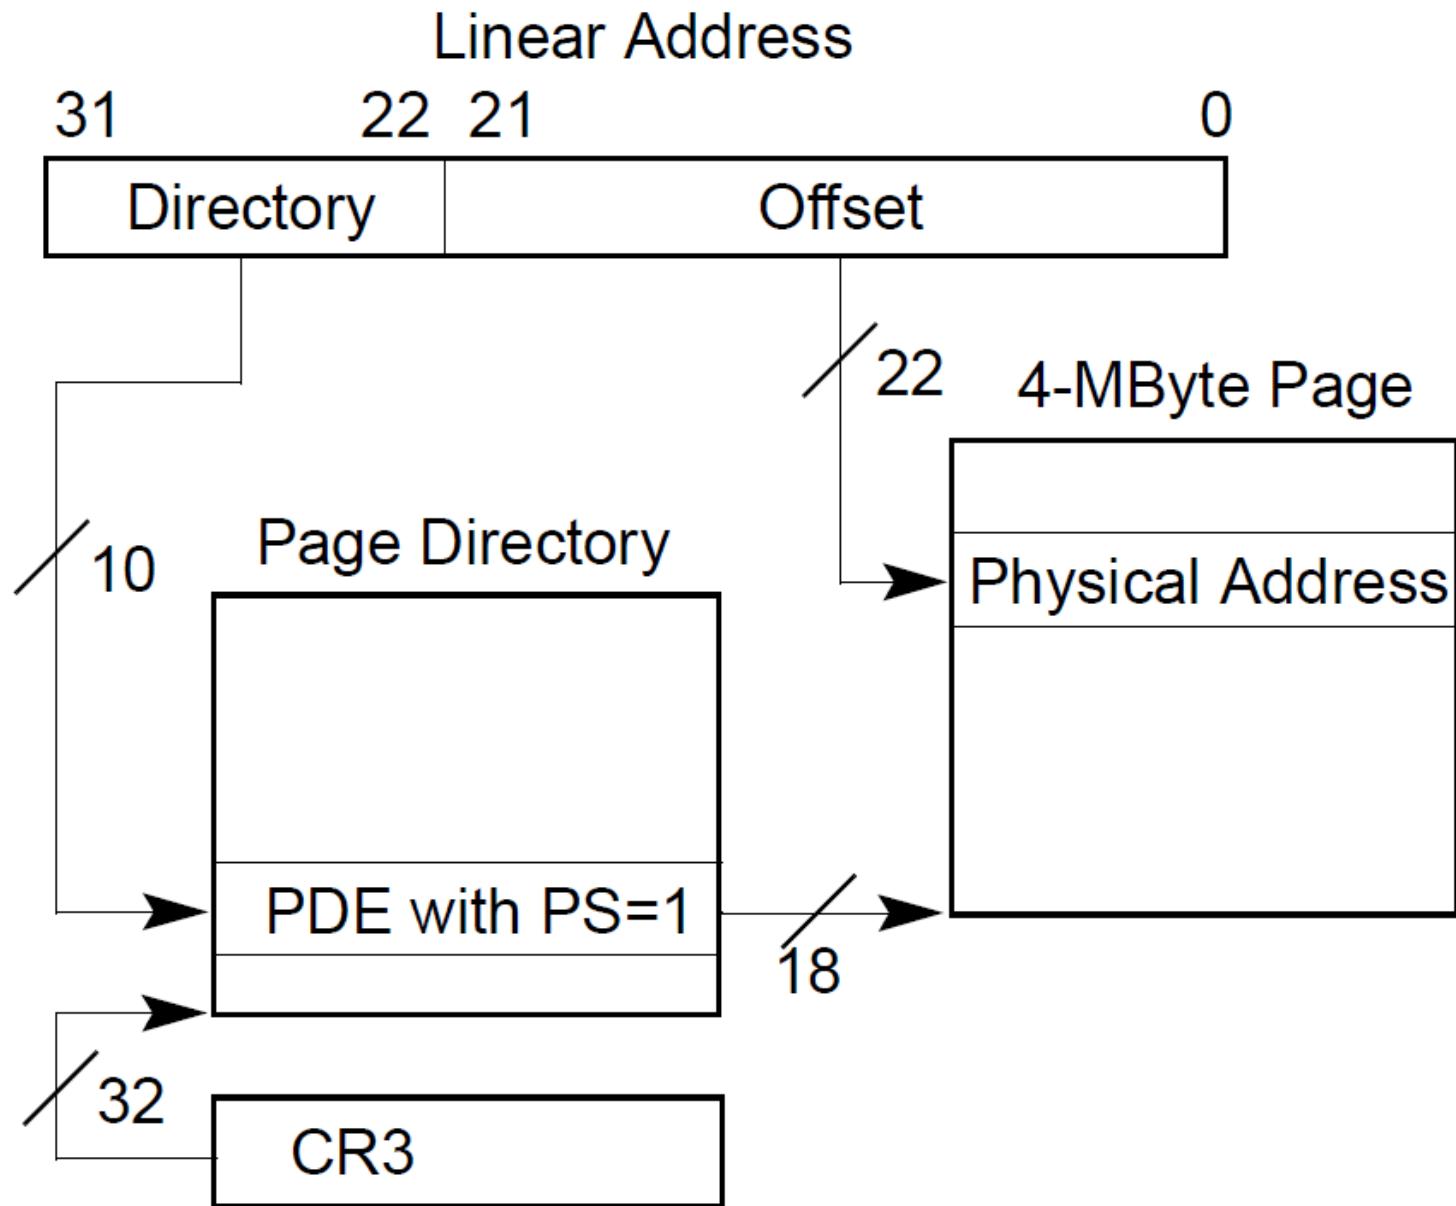
# 32bit x86 supports two page sizes

- 4KB pages
- 4MB pages

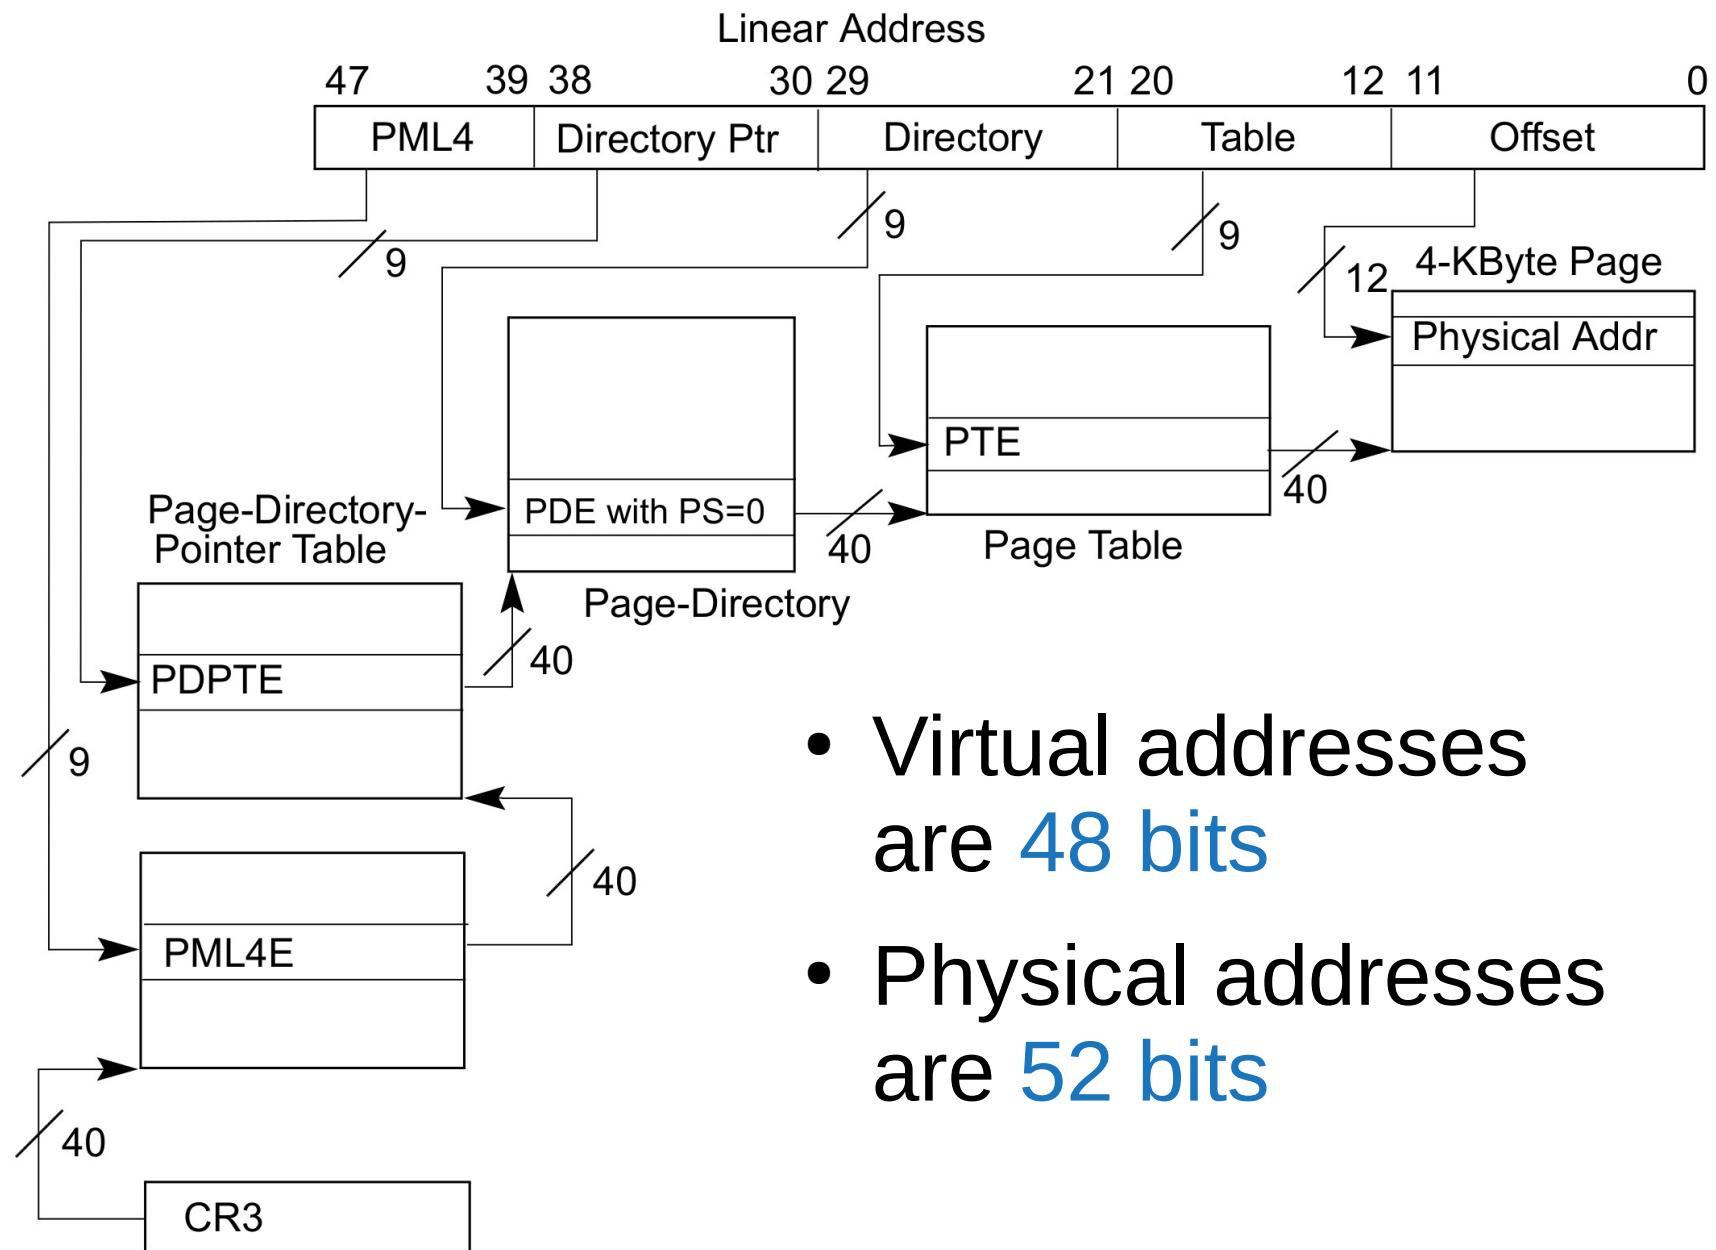
# Page translation for 4MB pages



# Page translation for 4MB pages



# Page translation in 64bit mode



Questions?

# What pages are used for

- Protect parts of the program
  - E.g., map code as read-only
    - Disable code modification attacks
    - Remember R/W bit in PTD/PTE entries!
  - E.g., map stack as non-executable
    - Protects from stack smashing attacks
    - Non-executable bit

# More paging tricks

- Determine a working set of a program?

# More paging tricks

- Determine a working set of a program?
  - Use “accessed” bit

# More paging tricks

- Determine a working set of a program?
  - Use “accessed” bit
- Iterative copy of a working set?
  - Used for virtual machine migration

# More paging tricks

- Determine a working set of a program?
  - Use “accessed” bit
- Iterative copy of a working set?
  - Used for virtual machine migration
  - Use “dirty” bit

# More paging tricks

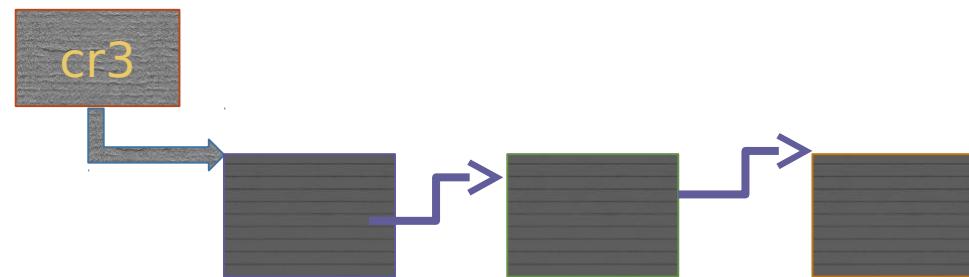
- Determine a working set of a program?
  - Use “accessed” bit
- Iterative copy of a working set?
  - Used for virtual machine migration
  - Use “dirty” bit
- Copy-on-write memory, e.g. lightweight fork()?

# More paging tricks

- Determine a working set of a program?
  - Use “accessed” bit
- Iterative copy of a working set?
  - Used for virtual machine migration
  - Use “dirty” bit

# TLB

- Walking page table is slow
  - Each memory access is 240 (local) - 360 (one QPI hop away) cycles on modern hardware
  - L3 cache access is 50 cycles



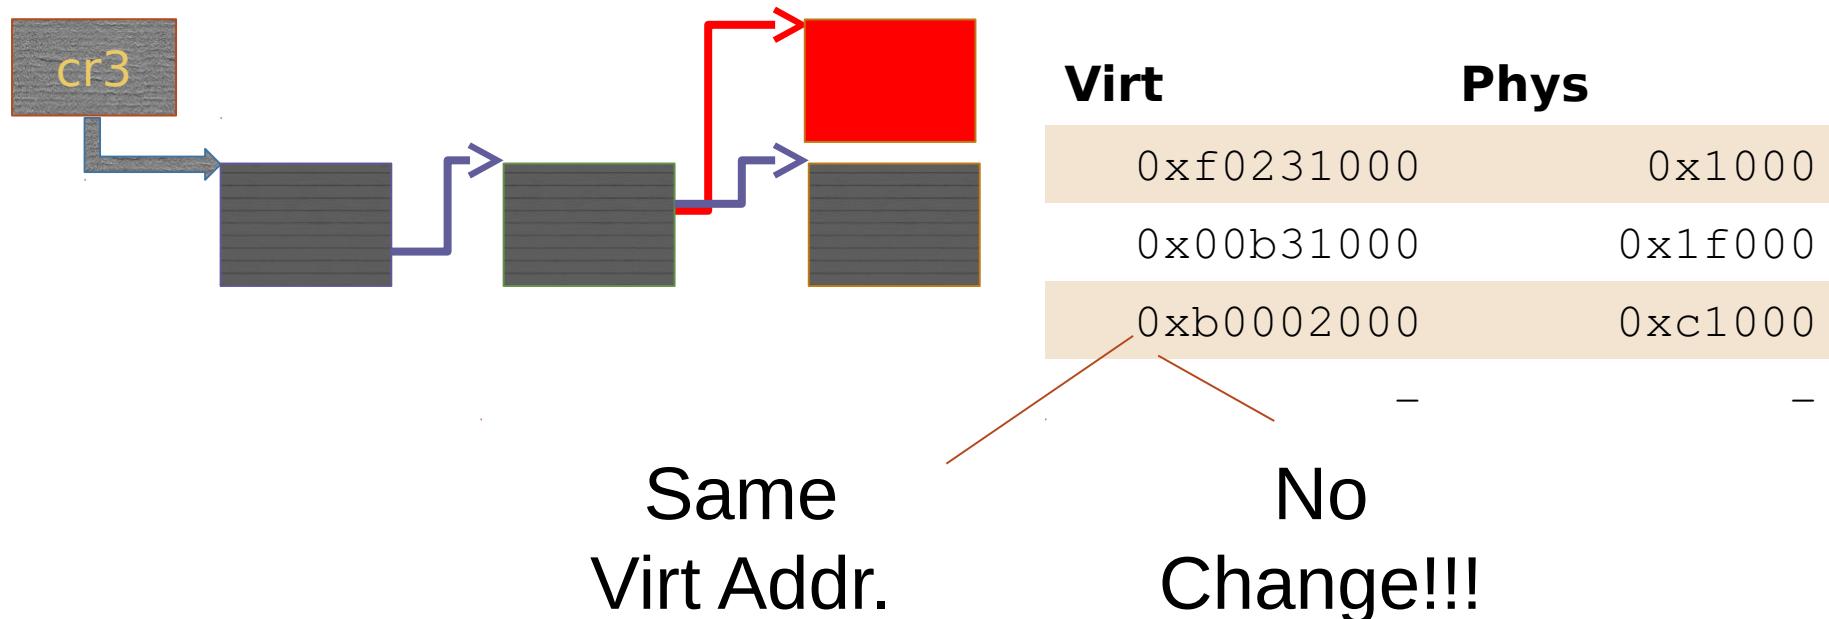
# TLB

- CPU caches results of page table walks
  - In translation lookaside buffer (TLB)

<b>Virt</b>	<b>Phys</b>
0xf0231000	0x1000
0x00b31000	0x1f000
0xb0002000	0xc1000

# TLB invalidation

- TLB is a cache (in CPU)
  - It is not coherent with memory
  - If page table entry is changes, TLB remains the same and is out of sync



# TLB invalidation

- After every page table update, OS needs to manually invalidate cached values
  - Flush TLB
    - Either one specific entry
    - Or entire TLB, e.g., when CR3 register is loaded
    - This happens when OS switches from one process to another
  - This is expensive
    - Refilling the TLB with new values takes time

# Tagged TLBs

- Modern CPUs have “tagged TLBs”,
  - Each TLB entry has a “tag” – identifier of a process
  - No need to flush TLBs on context switch
- On Intel this mechanism is called
  - Process-Context Identifiers (PCIDs)

Virt	Phys	Tag
0xf0231000	0x1000	P1
0x00b31000	0x1f000	P2
0xb0002000	0xc1000	P1

# When would you disable paging?

# When would you disable paging?

- Imagine you're running a memcached
  - Key/value cache
- You serve 1024 byte values (typical) on 10Gbps connection
  - 1024 byte packets can leave every 835ns, or 1670 cycles (2GHz machine)
  - This is your target budget per packet

# When would you disable paging?

- Now, to cover 32GB RAM with 4K pages
  - You need 64MB space
  - 64bit architecture, 4-level page tables (or 5-levels now)
- Page tables do not fit in L3 cache
  - Modern servers come with 32MB cache
- Every cache miss results in up to 4 cache misses due to page walk (remember 4-level page tables)
  - Each cache miss is 250 cycles
- Solution: 1GB pages

# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?

# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?
  - 1k

# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?
  - 1k
- How large of an address space can 1 page represent?

# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?
  - 1k
- How large of an address space can 1 page represent?
  - $1\text{k entries} * 1\text{page/entry} * 4\text{K/page} = 4\text{MB}$

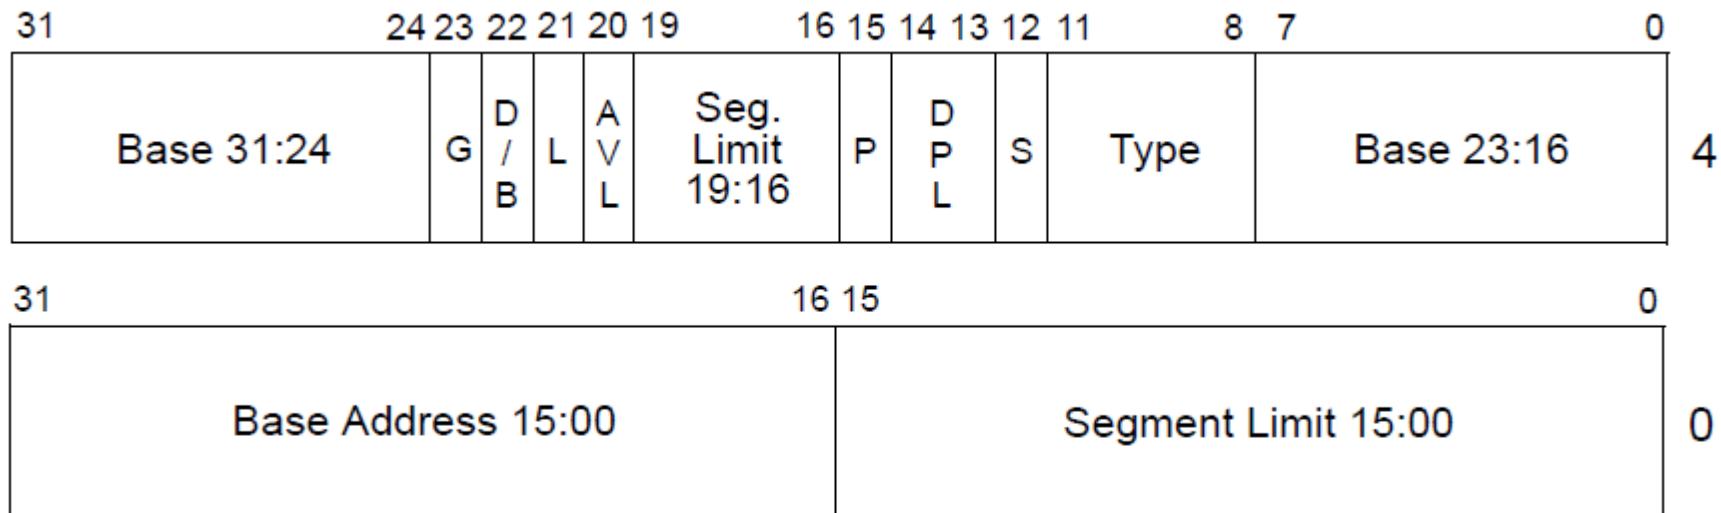
# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?
  - 1k
- How large of an address space can 1 page represent?
  - $1\text{k entries} * 1\text{page/entry} * 4\text{K/page} = 4\text{MB}$
- How large can we get with a second level of translation?

# Back of the envelope

- If a page is 4K and an entry is 4 bytes, how many entries per page?
  - 1k
- How large of an address space can 1 page represent?
  - $1\text{k entries} * 1\text{page/entry} * 4\text{K/page} = 4\text{MB}$
- How large can we get with a second level of translation?
  - $1\text{k tables/dir} * 1\text{k entries/table} * 4\text{k/page} = 4 \text{ GB}$
  - Nice that it works out that way!

# Segment descriptors



L — 64-bit code segment (IA-32e mode only)

AVL — Available for use by system software

BASE — Segment base address

D/B — Default operation size (0 = 16-bit segment; 1 = 32-bit segment)

DPL — Descriptor privilege level

G — Granularity

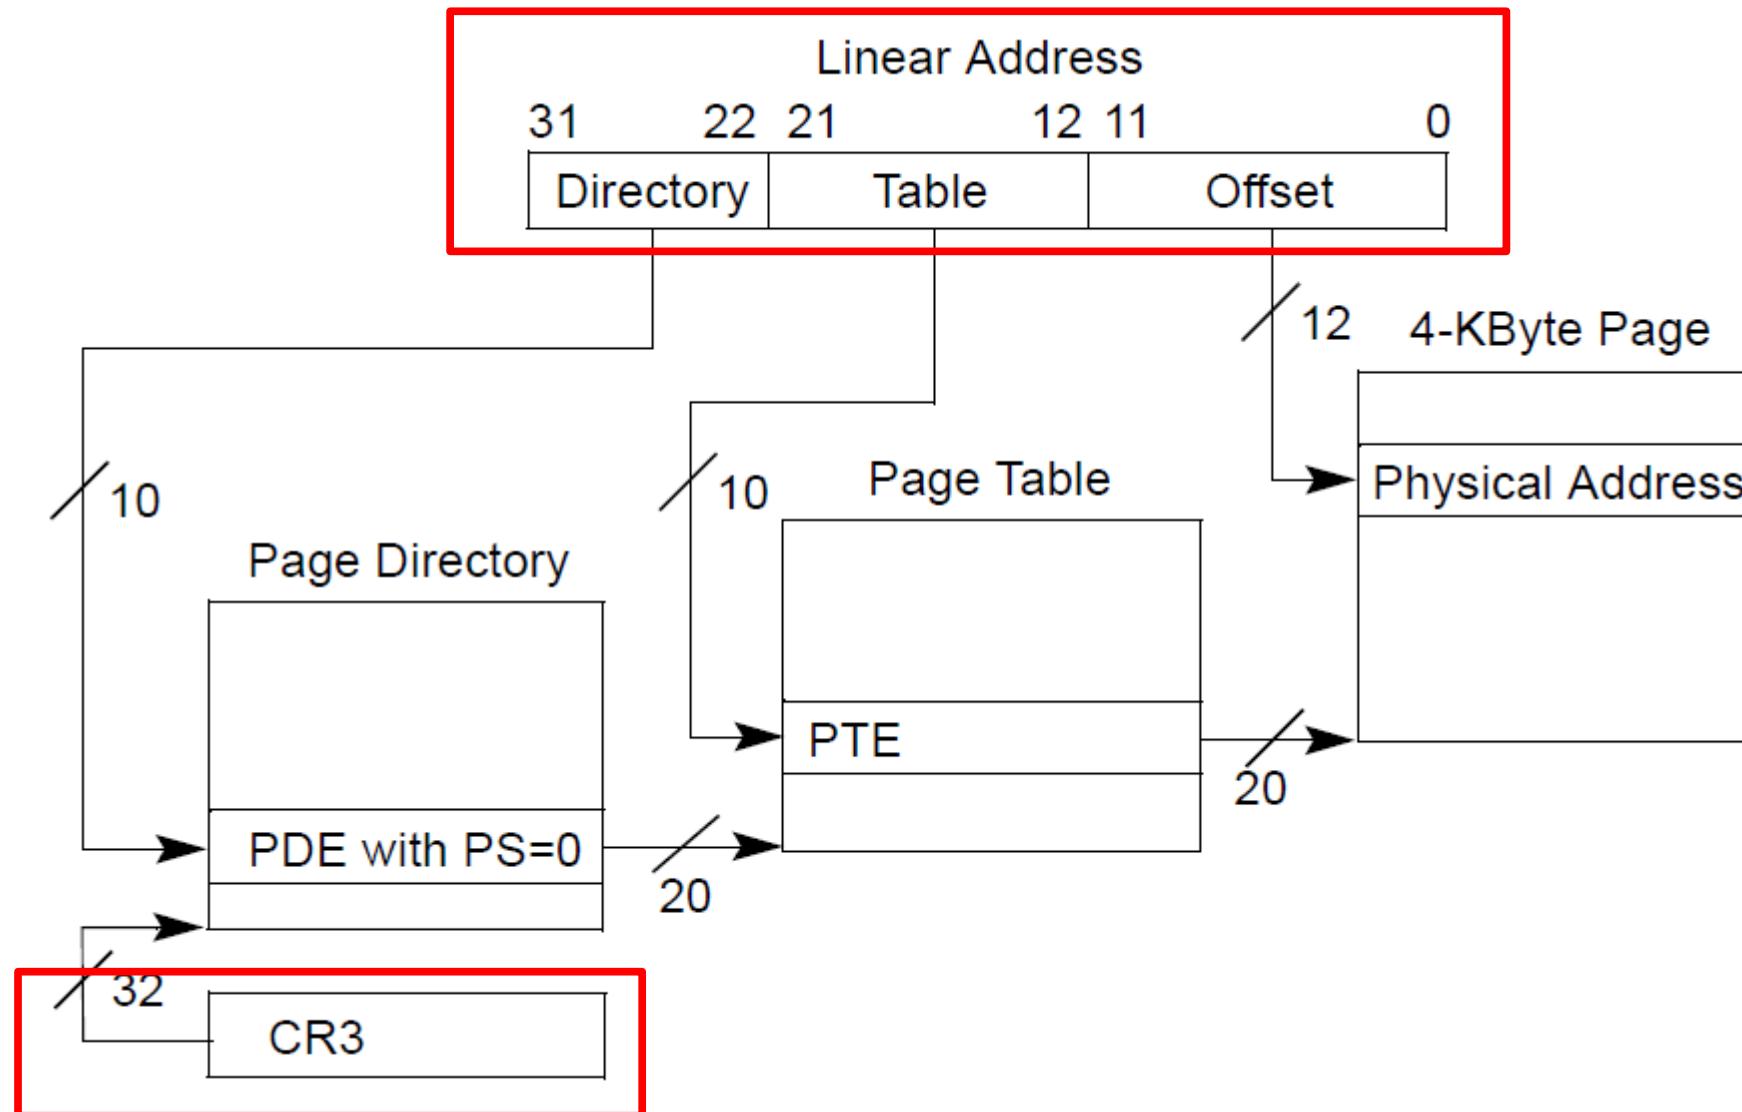
LIMIT — Segment Limit

P — Segment present

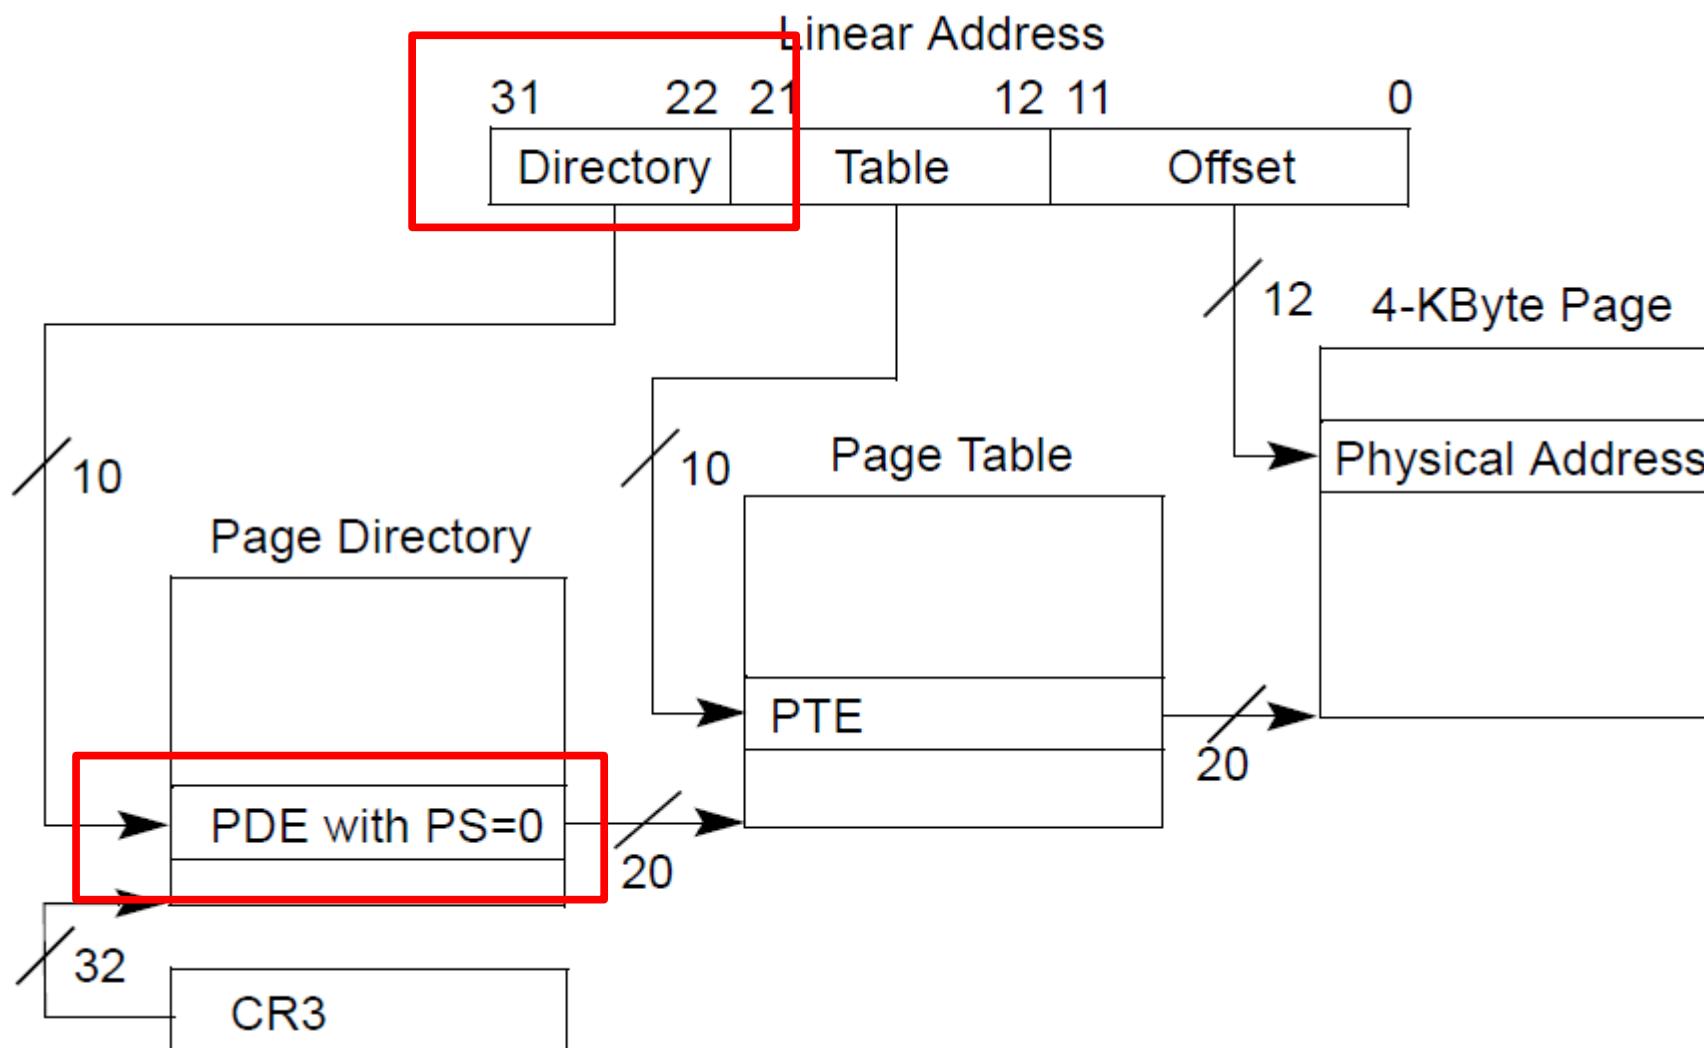
S — Descriptor type (0 = system; 1 = code or data)

TYPE — Segment type

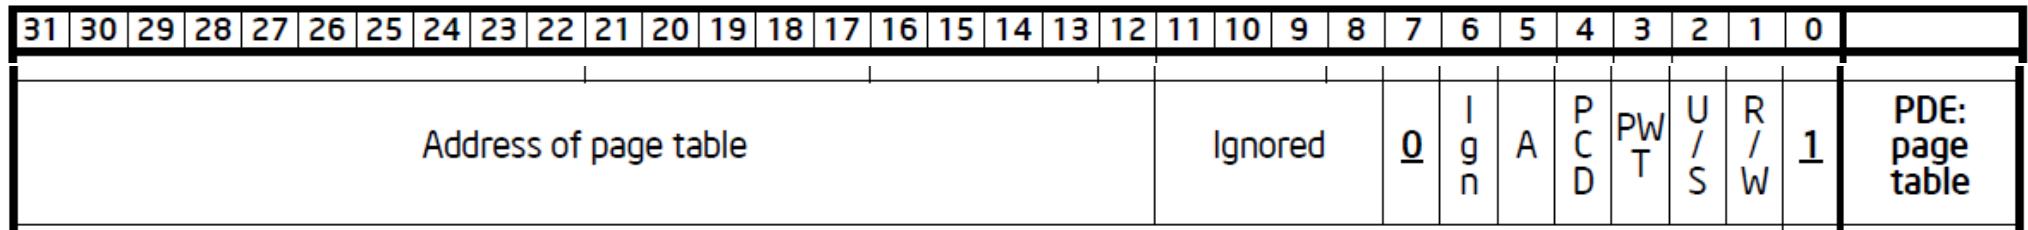
# Page translation



# Page translation

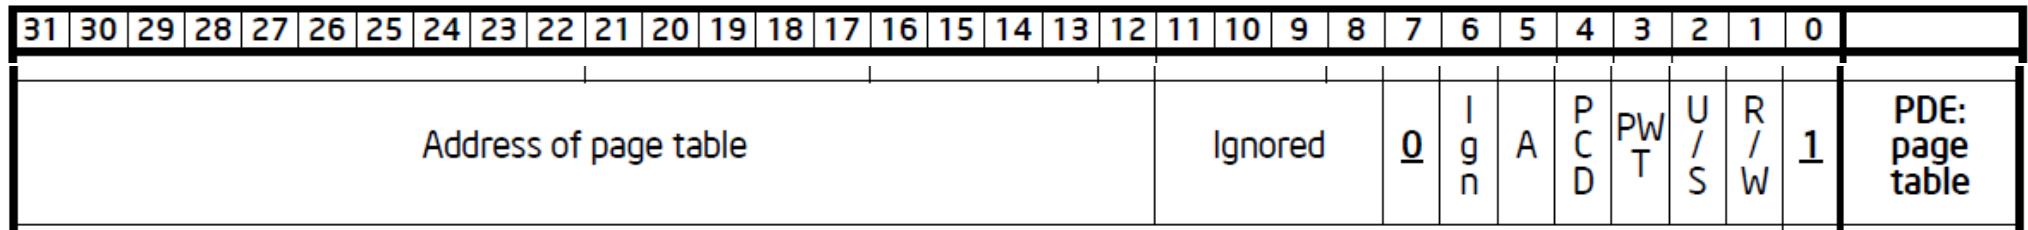


# Page directory entry (PDE)



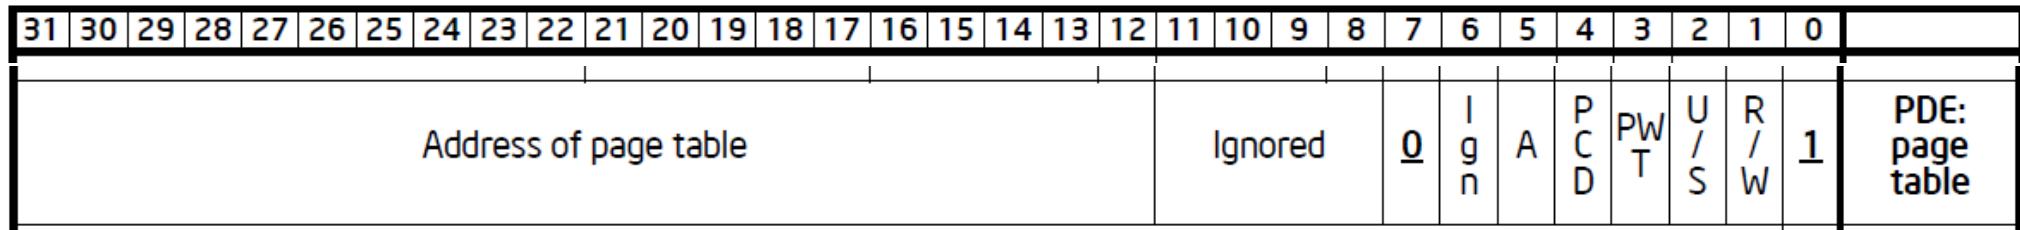
- 20 bit address of the page table

# Page directory entry (PDE)



- 20 bit address of the page table
- Wait... 20 bit address, but we need 32 bits

# Page directory entry (PDE)



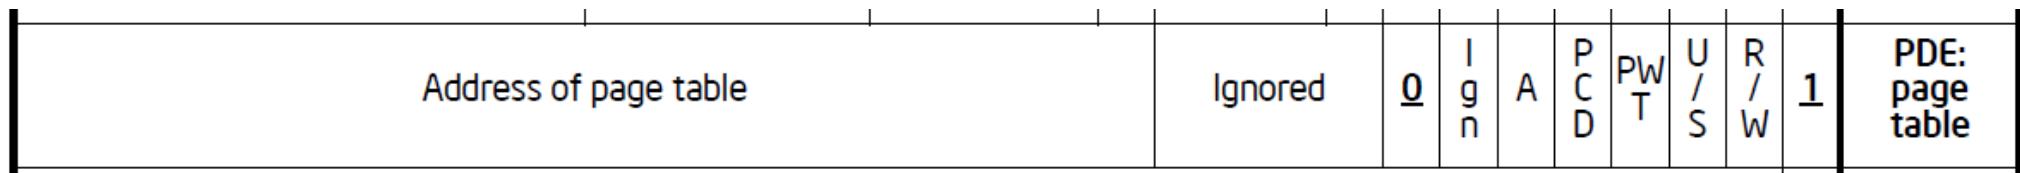
- 20 bit address of the page table
- Wait... 20 bit address, but we need 32 bits
  - Pages 4KB each, we need 1M to cover 4GB
  - Pages start at 4KB (page aligned boundary)

# Page directory entry (PDE)

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
Address of page table												Ignored		0	I g n	A	P C D	P W T	U S	R / W	1	PDE: page table										

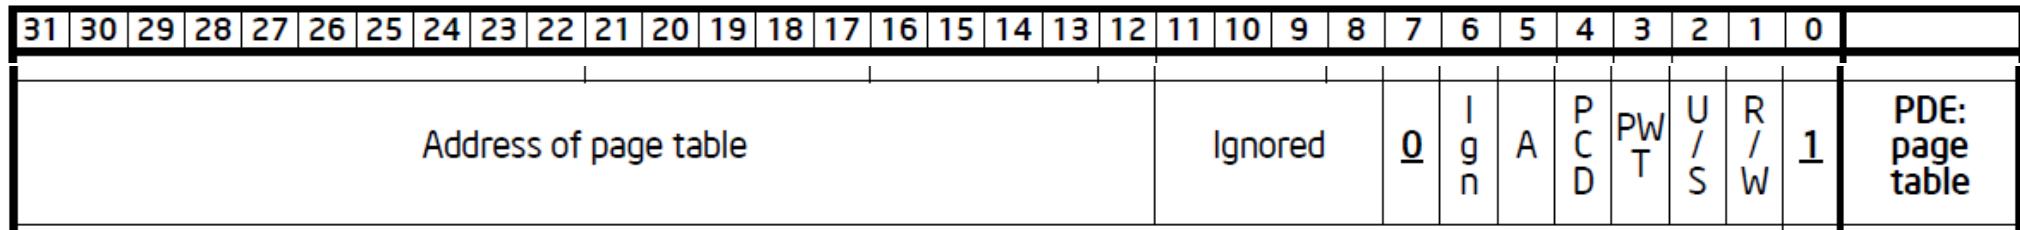
- Bit #1: R/W – writes allowed?
  - But allowed where?

# Page directory entry (PDE)



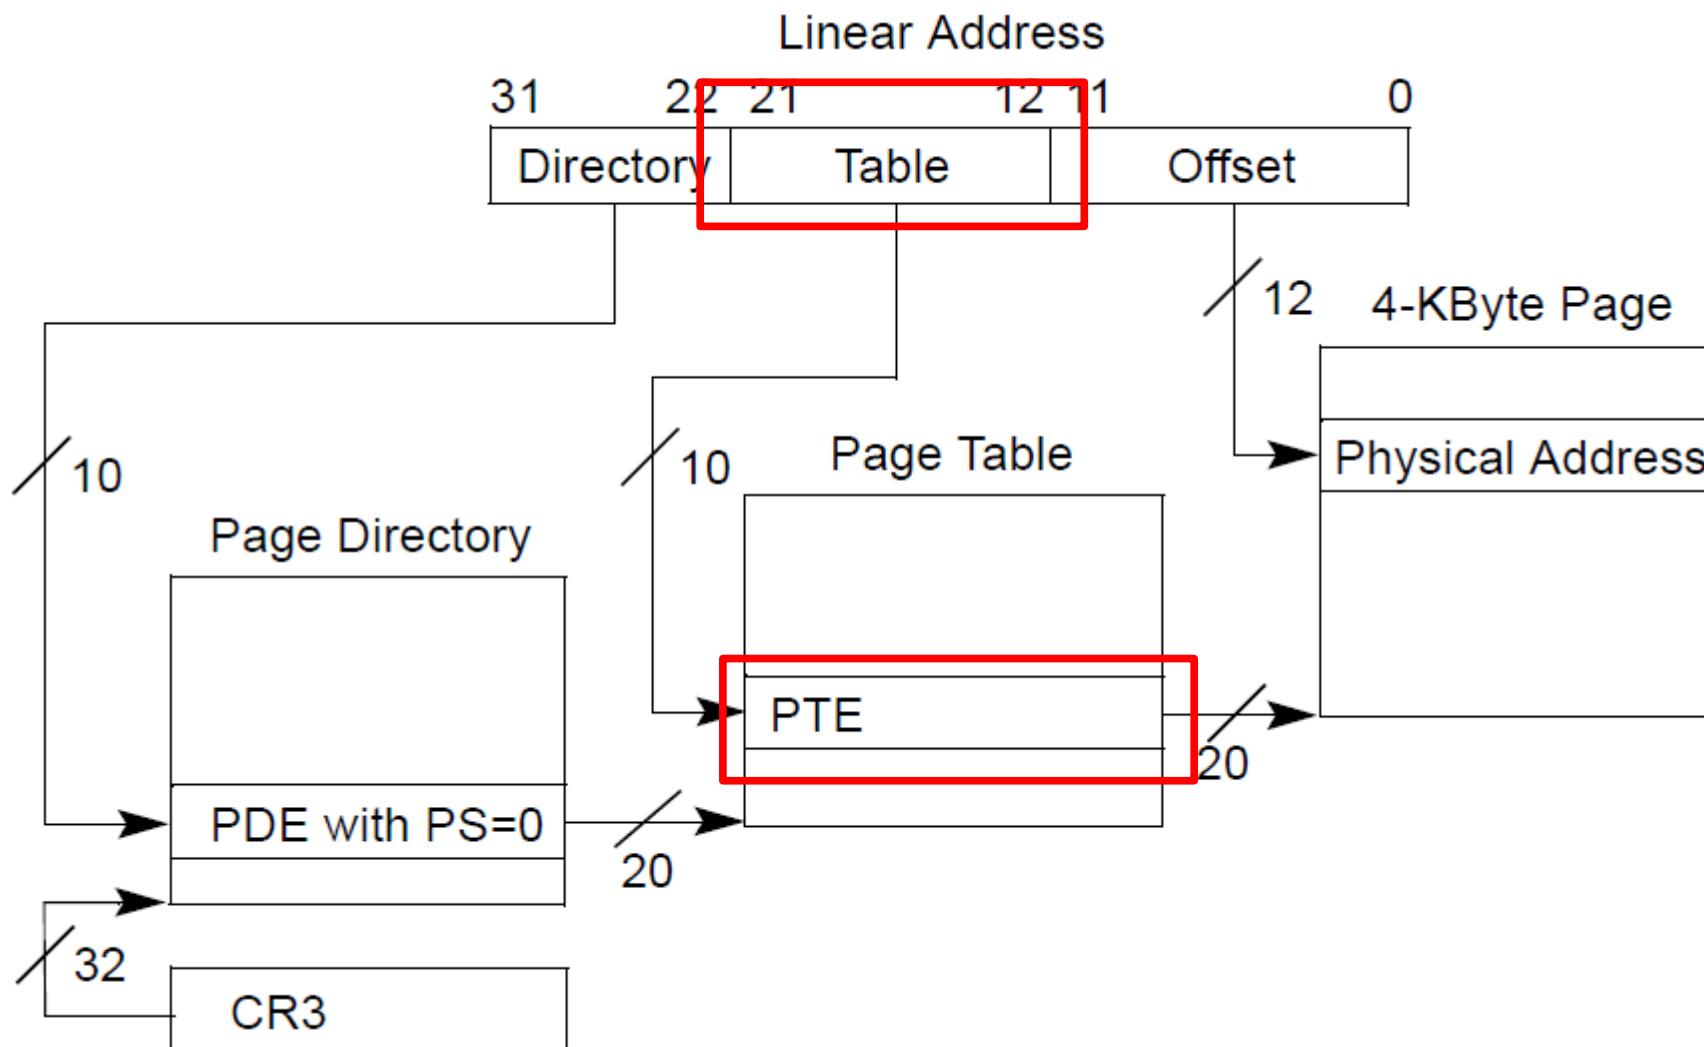
- Bit #1: R/W – writes allowed?
  - But allowed where?
  - One page directory entry controls 1024 Level 2 page tables
    - Each Level 2 maps 4KB page
  - So it's a region of  $4\text{KB} \times 1024 = 4\text{MB}$

# Page directory entry (PDE)

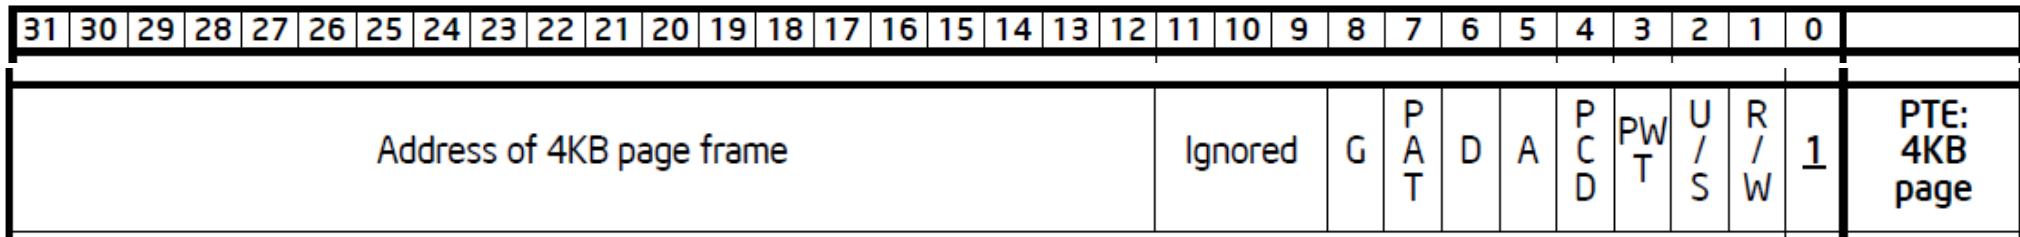


- Bit #2: U/S – user/supervisor
  - If 0 – user-mode access is not allowed
  - Allows protecting kernel memory from user-level applications

# Page translation

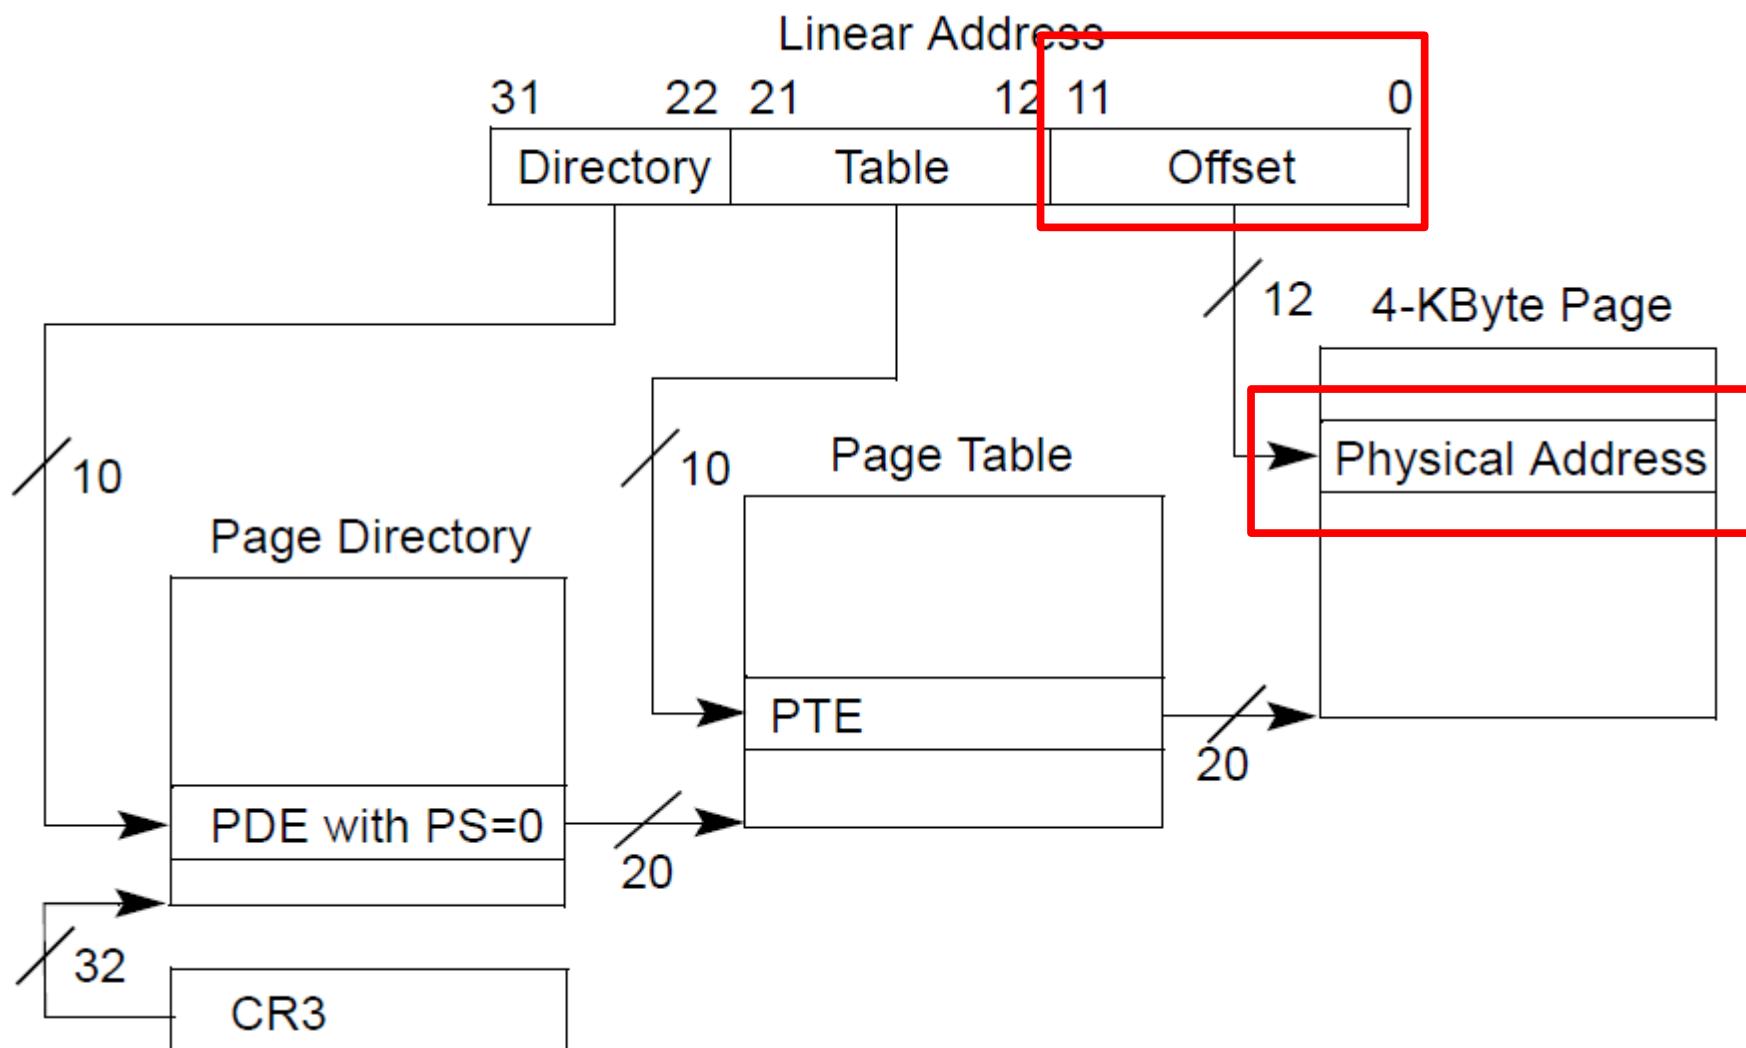


# Page table entry (PTE)



- 20 bit address of the 4KB page
  - Pages 4KB each, we need 1M to cover 4GB
- Bit #1: R/W – writes allowed?
  - To a 4KB page
- Bit #2: U/S – user/supervisor
  - If 0 user-mode access is not allowed
- Bit #5: A – accessed
- Bit #6: D – dirty – software has written to this page

# Page translation



Questions?