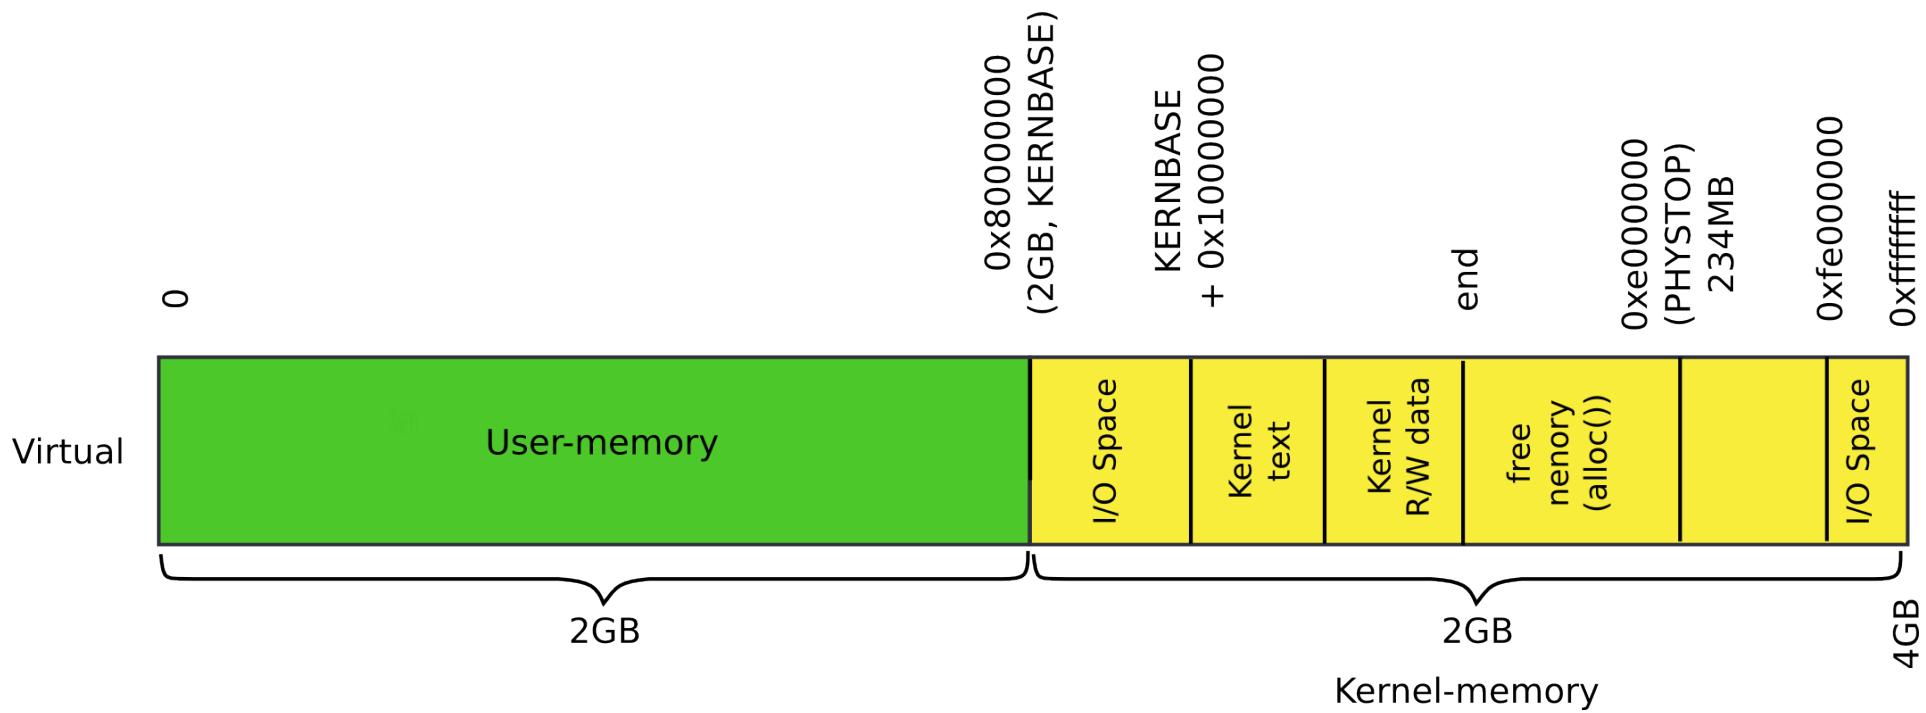


238P: Operating Systems

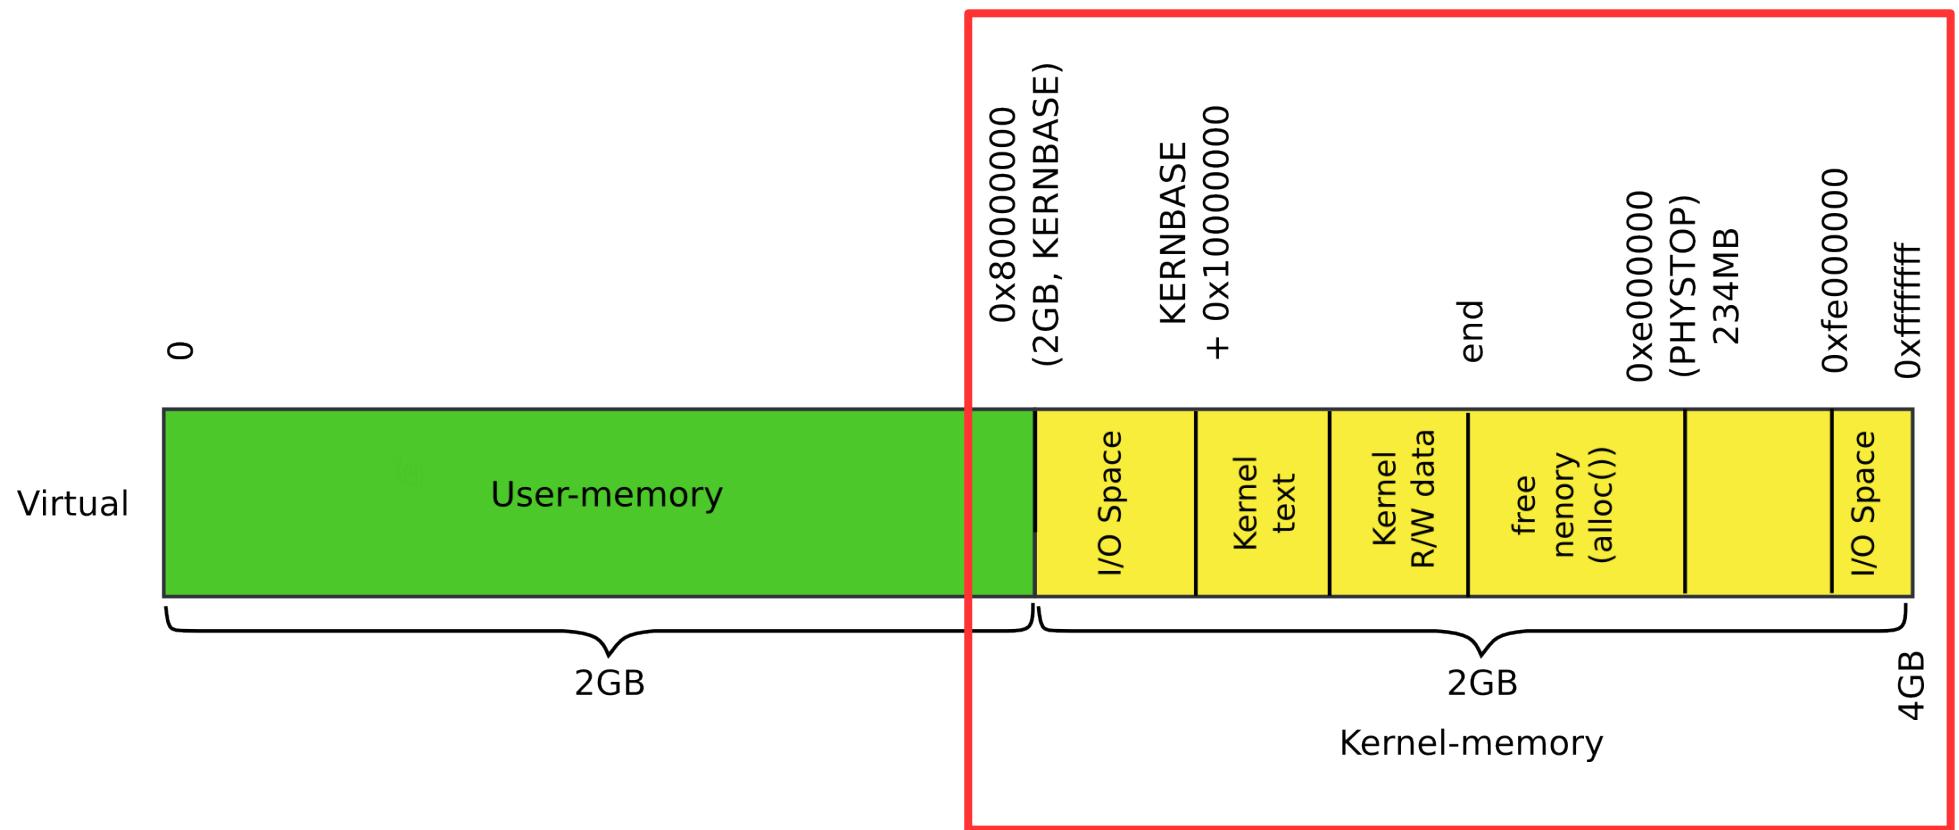
Lecture 8: Creating Processes (exec())

Anton Burtsev
November, 2019

Recap: kernel memory




Recap: kernel memory



```
1317 main(void)
1318 {
1319     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1320     kvmalloc(); // kernel page table
1321     mpinit(); // detect other processors
1322     lapicinit(); // interrupt controller
1323     seginit(); // segment descriptors
1324     cprintf("\ncpu%d: starting xv6\n\n", cpunum());
1325     picinit(); // another interrupt controller
1326     ioapicinit(); // another interrupt controller
1327     consoleinit(); // console hardware
1328     uartinit(); // serial port
1329     pinit(); // process table
1330     tvinit(); // trap vectors
1331     binit(); // buffer cache
1332     fileinit(); // file table
1333     ideinit(); // disk
1334     if(!ismp)
1335         timerinit(); // uniprocessor timer
1336     startothers(); // start other processors
1337     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
1338     userinit(); // first user process
1339     mpmain(); // finish this processor's setup
1340 }
```

main()

Today: process memory




How does kernel creates new processes?

How does kernel creates new processes?

- Exec
 - `exec("/bin/ls", argv);`


exec(): high-level outline

- We want to create the following memory layout for the process
 - What shall we do?



exec(): high-level outline

- Load program from disk
- Create user-stack
- Run!



exec(): high-level overview

- Read process binary from disk
 - Locate a file that contains process binary
 - namei() takes a file path (“/bin/ls”) as an argument
 - Returns an inode
 - Read the file block by block
 - readi() reads the inode (file data) into memory
 - To read file in memory we need to construct the process address space
 - A page table specifically for the process

exec(): locate inode

```
6309 int
6310 exec(char *path, char **argv)
6311 {
...
6321     if((ip = namei(path)) == 0){
6322         end_op();
6323         return -1;
6324     }
6328     // Check ELF header
6329     if(readi(ip, (char*)&elf, 0, sizeof(elf)) <
6330         sizeof(elf))
6331         goto bad;
6331     if(elf.magic != ELF_MAGIC)
6332         goto bad;
```

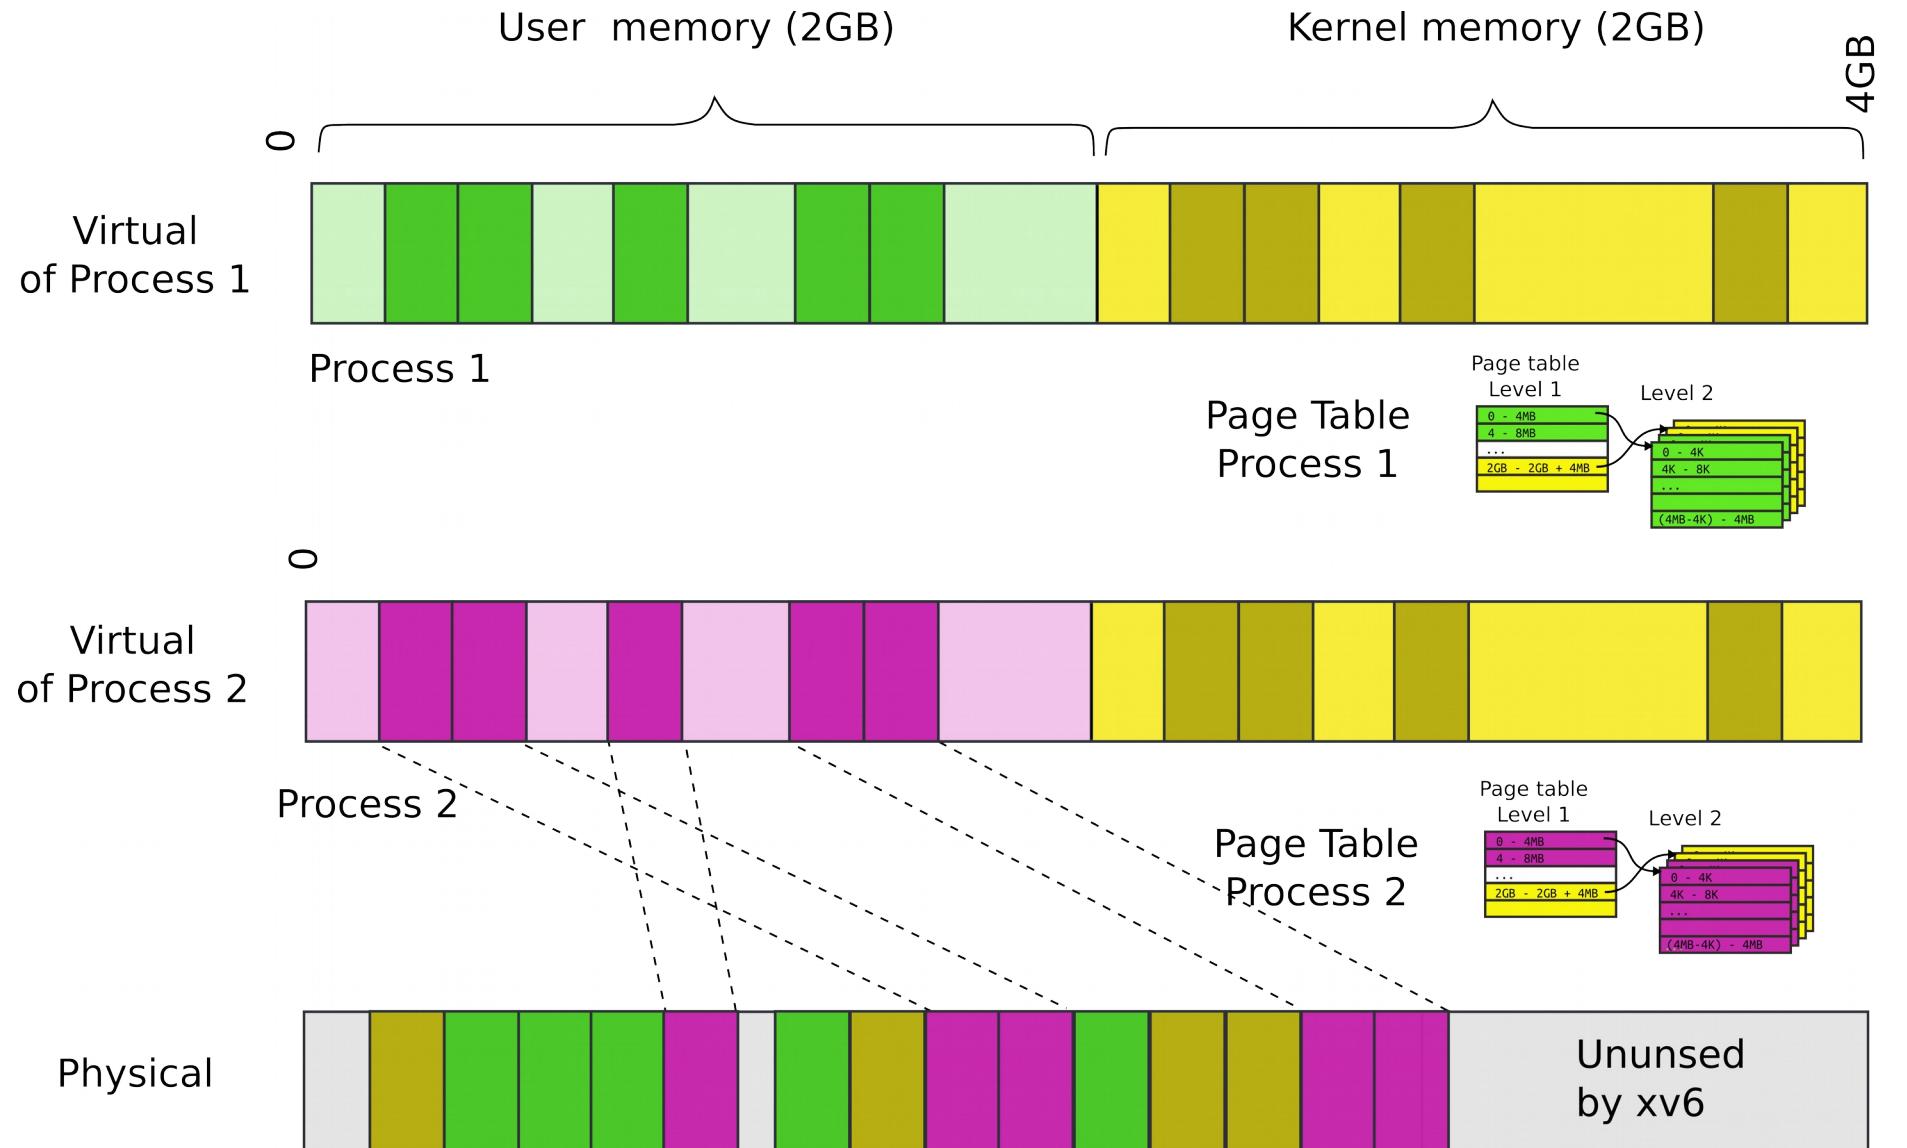
exec(): check ELF header

```
6309 int
6310 exec(char *path, char **argv)
6311 {
...
6321     if((ip = namei(path)) == 0){
6322         end_op();
6323         return -1;
6324     }
6328     // Check ELF header
6329     if(readi(ip, (char*)&elf, 0, sizeof(elf)) <
6330         sizeof(elf))
6331         goto bad;
6331     if(elf.magic != ELF_MAGIC)
6332         goto bad;
```

Create process address space

exec(): Construct process address space

- Two step process
 - Create the kernel part of the address space
 - Create the user part of the address space



Remember: each process maps kernel in its page table

0xe0000000
(PHYSTOP)
234MB

Top of physical
memory

exec(): Setup kernel address space()

```
6310 exec(char *path, char **argv)
6311 {
...
6331     if(elf.magic != ELF_MAGIC)
6332         goto bad;
6334     if((pgdir = setupkvm()) == 0)
6335         goto bad;
...

```

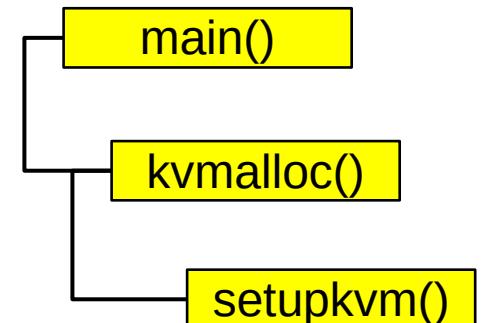
- Remember from last time?

```

1836 pde_t*
1837 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

Recap: Allocate page table directory

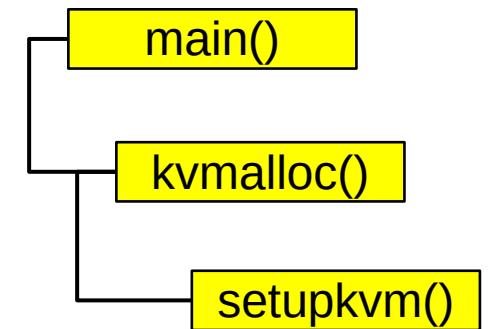


```

1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

Recap: Iterate in a loop: remap physical pages

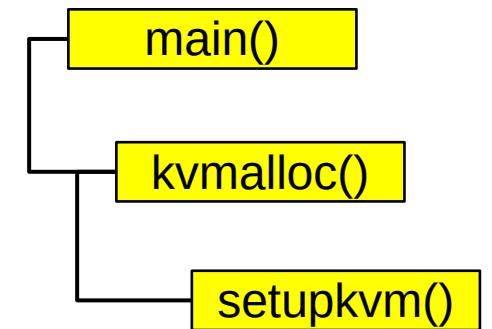


```

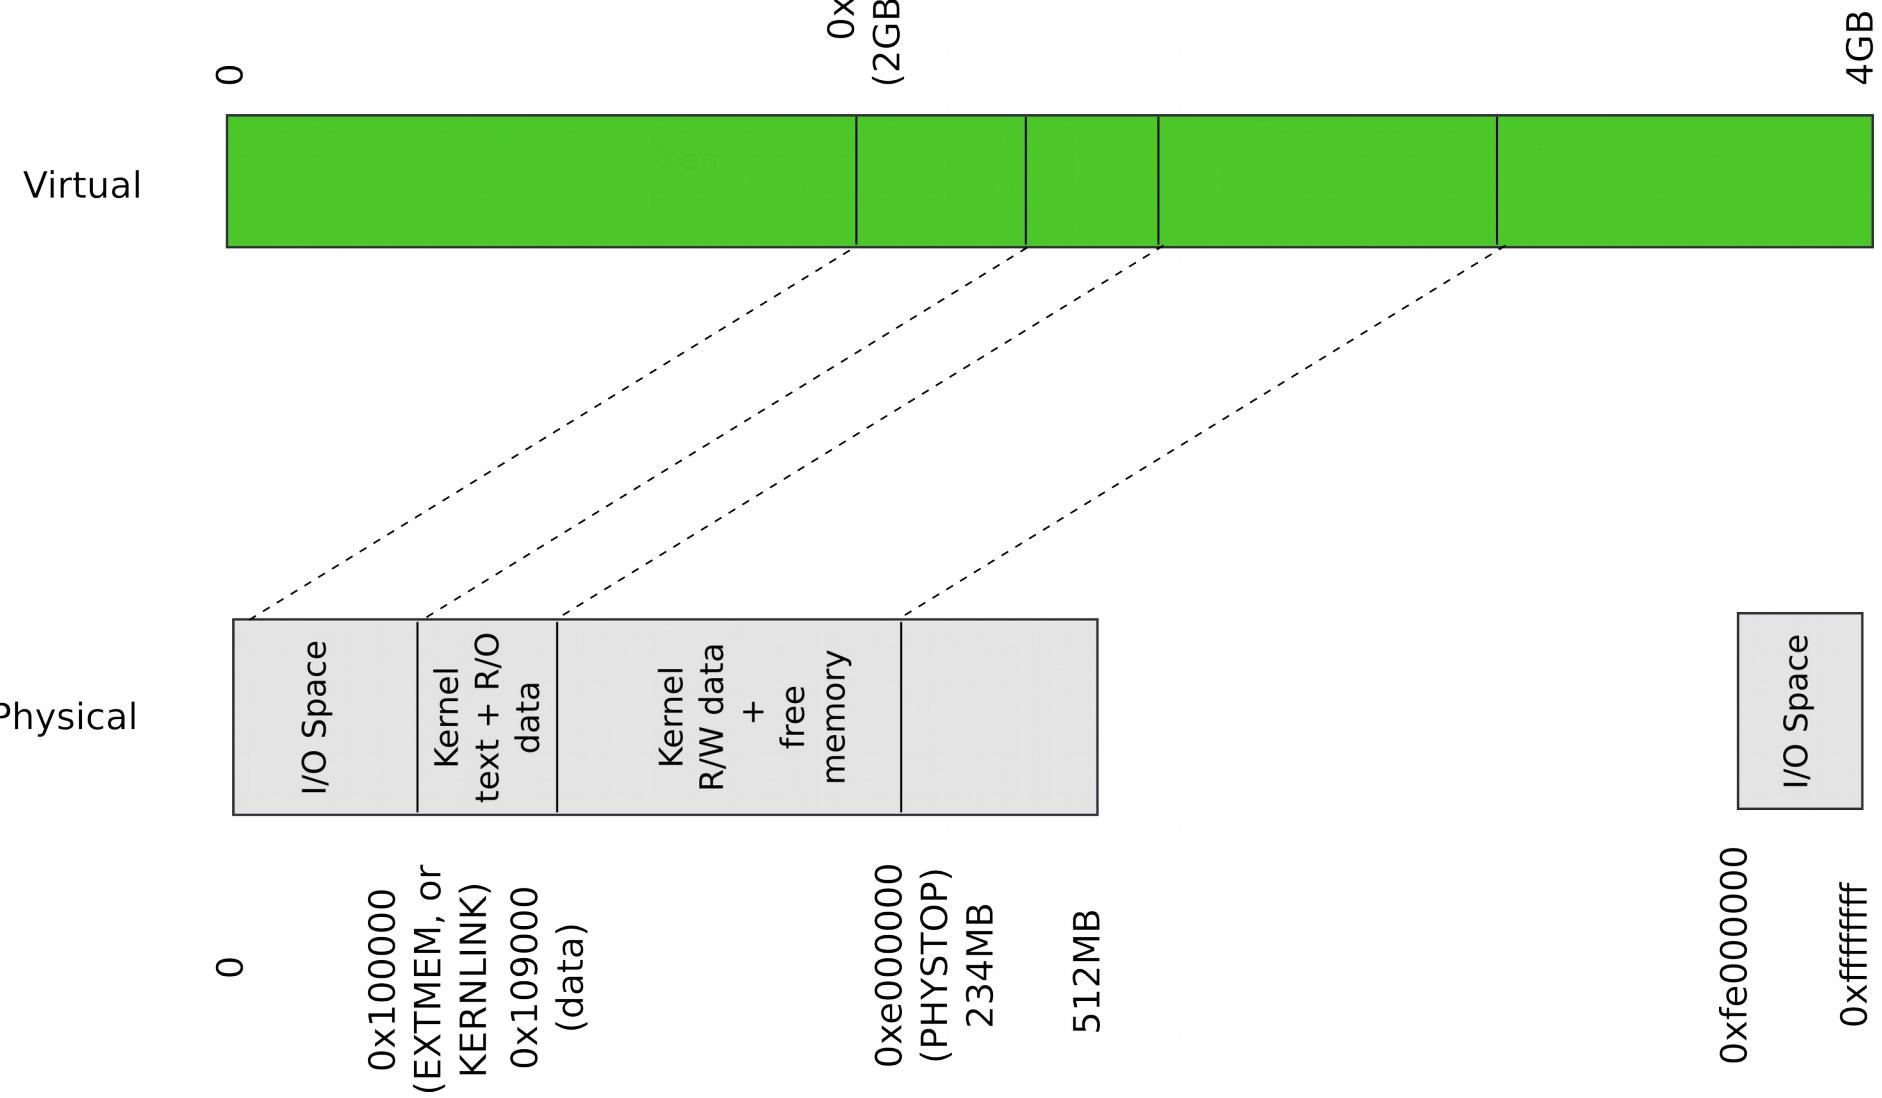
1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

Recap: Iterate in a loop: remap physical pages



Recap: Kernel map



Recap: Kmap – kernel map

```
1823 static struct kmap {  
1824     void *virt;           Physical  
1825     uint phys_start;  
1826     uint phys_end;  
1827     int perm;  
1828 } kmap[] = {  
1829     { (void*)KERNBASE, 0, EXTMEM, PTE_W}, // I/O space  
1830     { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0}, //text+rodata  
1831     { (void*)data, V2P(data), PHYSTOP, PTE_W}, // kern data+memory  
1832     { (void*)DEVSPACE, DEVSPACE, 0, PTE_W}, // more devices  
1833 };
```

The diagram illustrates the kernel map structure. It consists of a horizontal bar divided into four segments, each representing a different type of memory or device space:

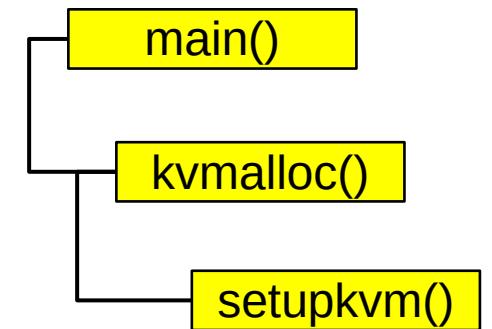
- I/O Space:** Physical address 0x100000 (EXTMEM, or KERNLINK) to 0x109000 (data).
- Kernel text + R/O data:** Physical address 0xe000000 (PHYSTOP) to 0xe000000 (data), spanning 234MB.
- Kernel R/W data + free memory:** Physical address 512MB to 0xffffffff (data).
- I/O Space:** Physical address 0xffffffff (DEVSPACE) to 0xffffffff (data).

```

1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

Recap: Iterate in a loop: remap physical pages

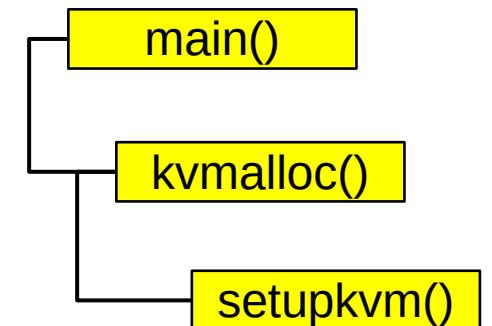


```

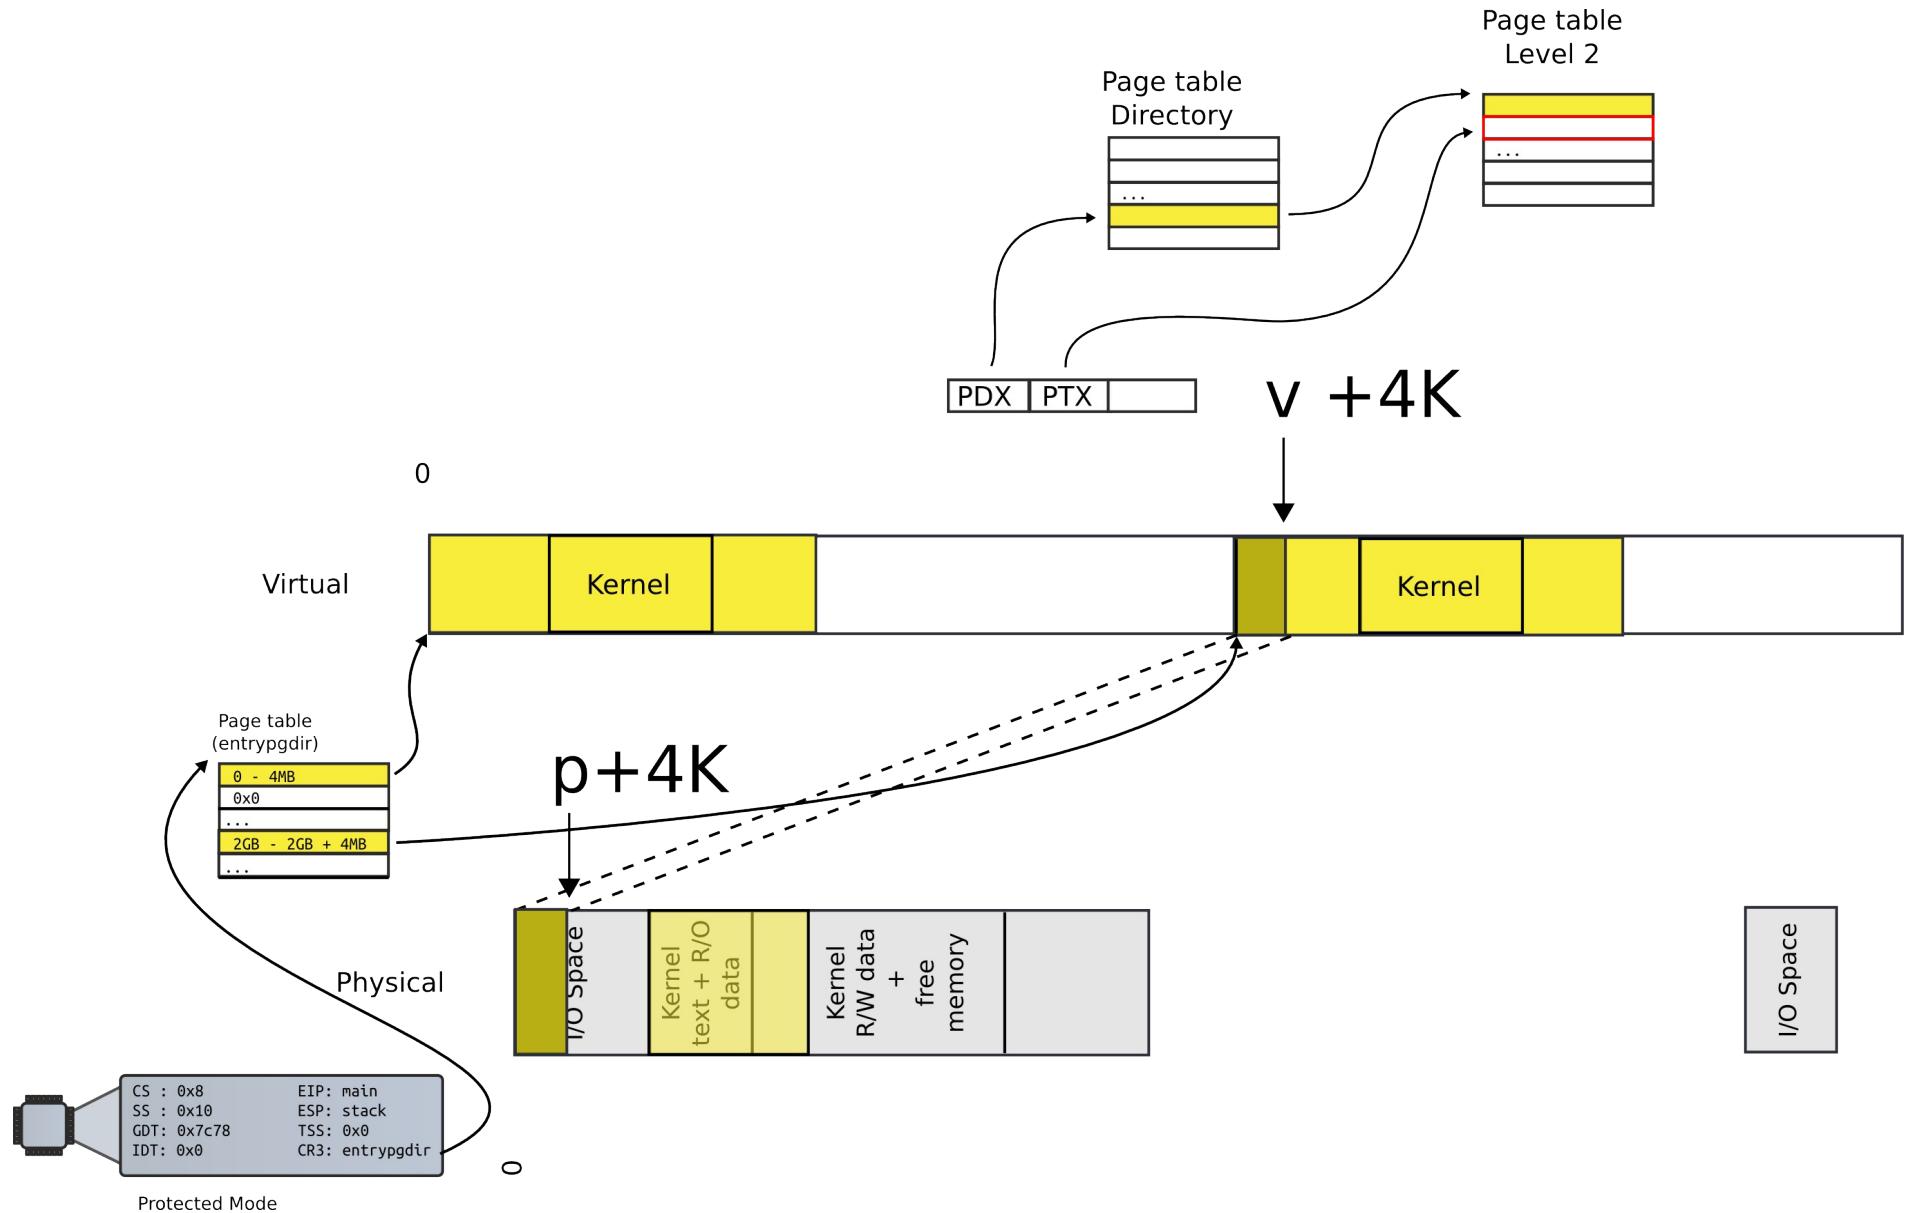
1836 pde_t*
1887 setupkvm(void)
1838 {
1839     pde_t *pgdir;
1840     struct kmap *k;
1841
1842     if((pgdir = (pde_t*)kalloc()) == 0)
1843         return 0;
1844     memset(pgdir, 0, PGSIZE);
...
1847     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1848         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1849                     (uint)k->phys_start, k->perm) < 0)
1850             return 0;
1851     return pgdir;
1852 }

```

Recap: Remap physical pages



setupkvm(): Move to next page



exec(): Construct process address space

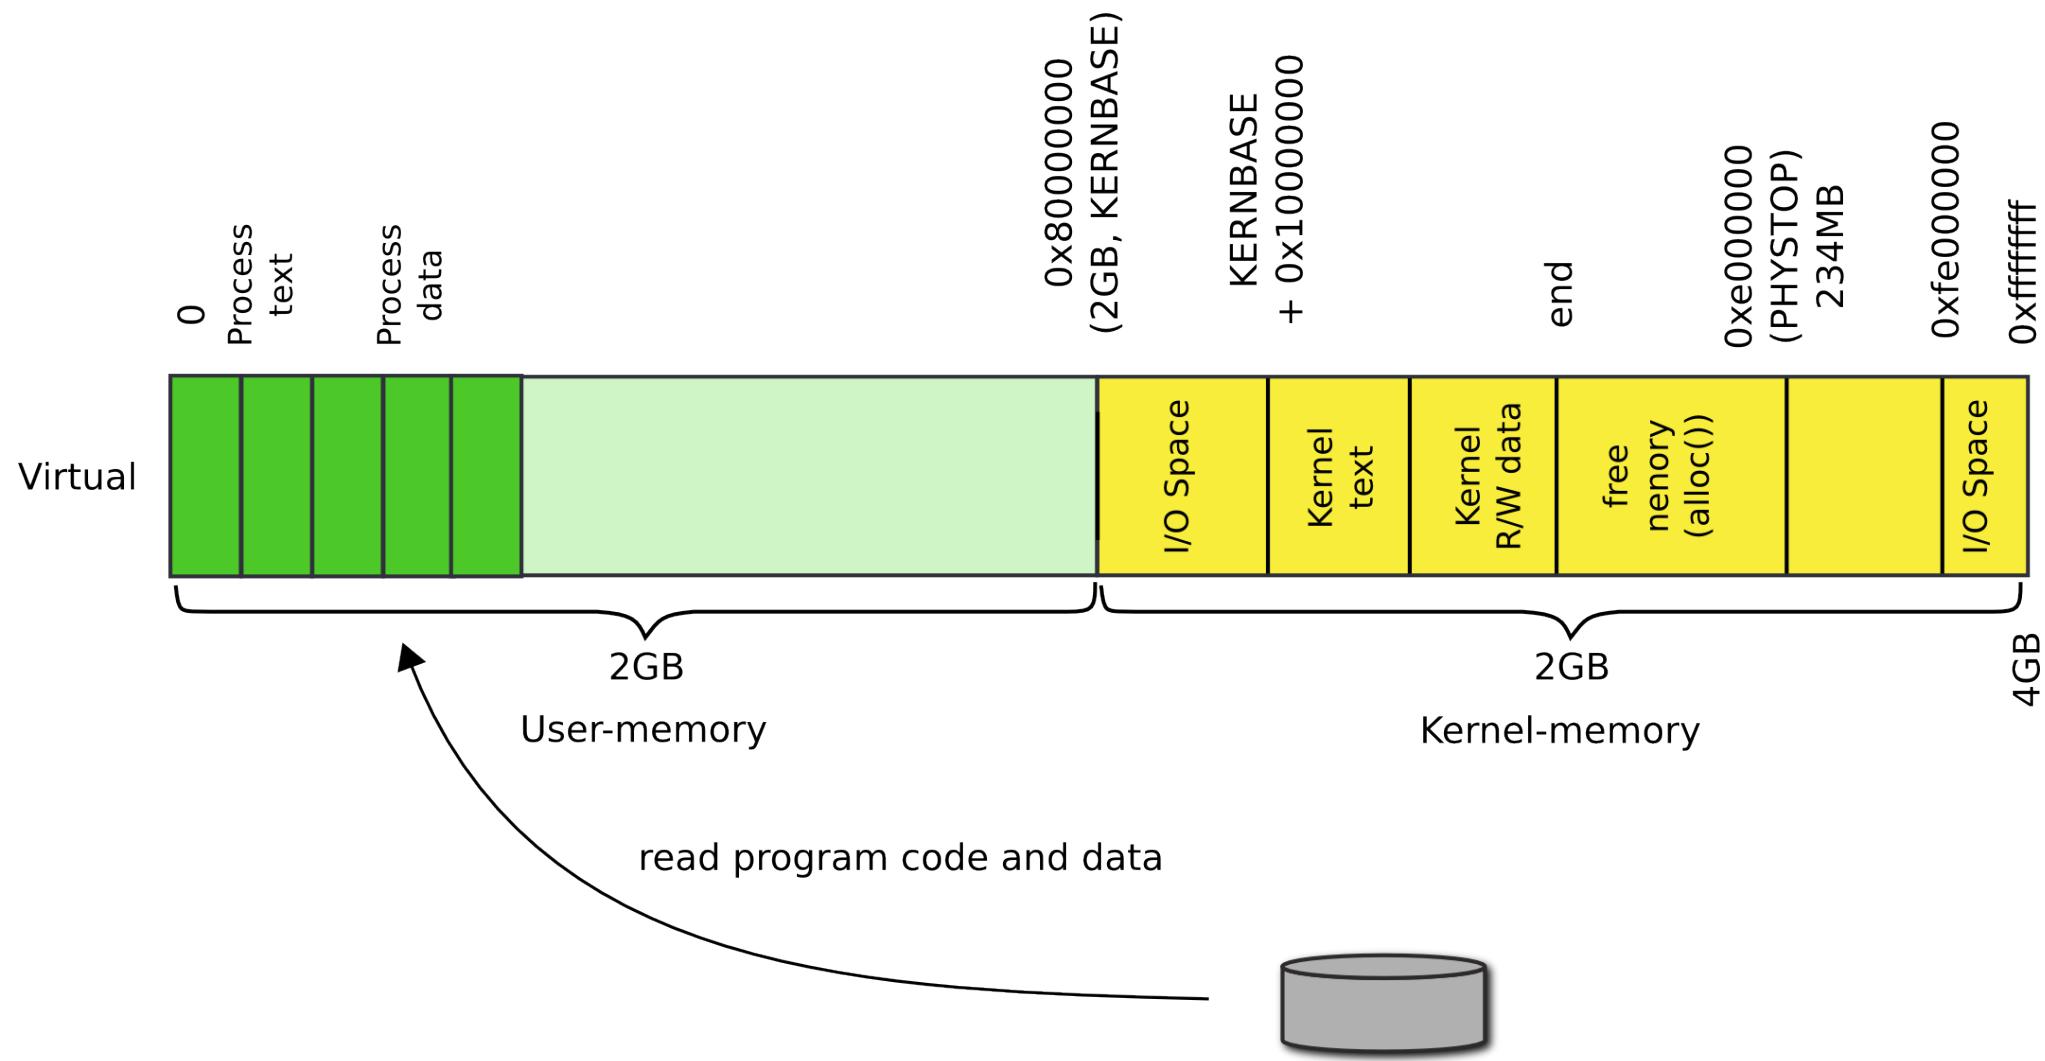
- Two step process
 - Create the kernel part of the address space
 - **Create the user part of the address space**

Create user part of the address space

exec(): create user part of the address space

- The goal is to fill in the page table entries
 - This can be naturally combined with loading the program from disk into memory
- At a high level iterate in a loop
 - On each step:
 - Allocate user-level pages for the program
 - Map them by filling in the page table entries
 - Read data from the inode into that memory

High-level idea



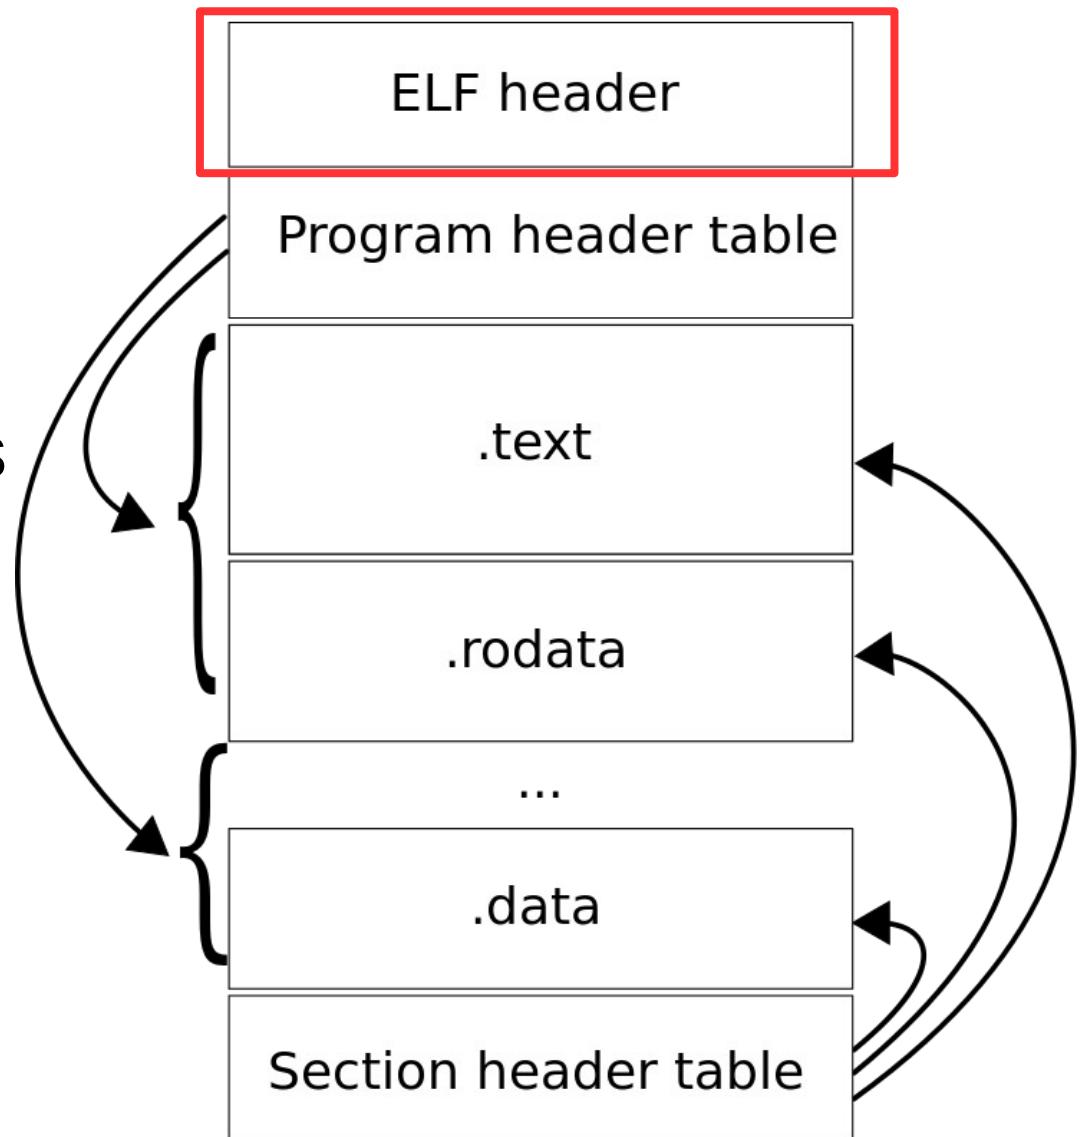
```
6310 exec(char *path, char **argv)
6311 {
...
6337     // Load program into memory.
6338     sz = 0;
6339     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6340         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6341             goto bad;
...
6348     if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
6349         goto bad;
6350     if(ph.vaddr % PGSIZE != 0)
6351         goto bad;
6352     if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
6353         goto bad;
6354 }
```

Program loading loop

- Loop over all program headers

ELF object file

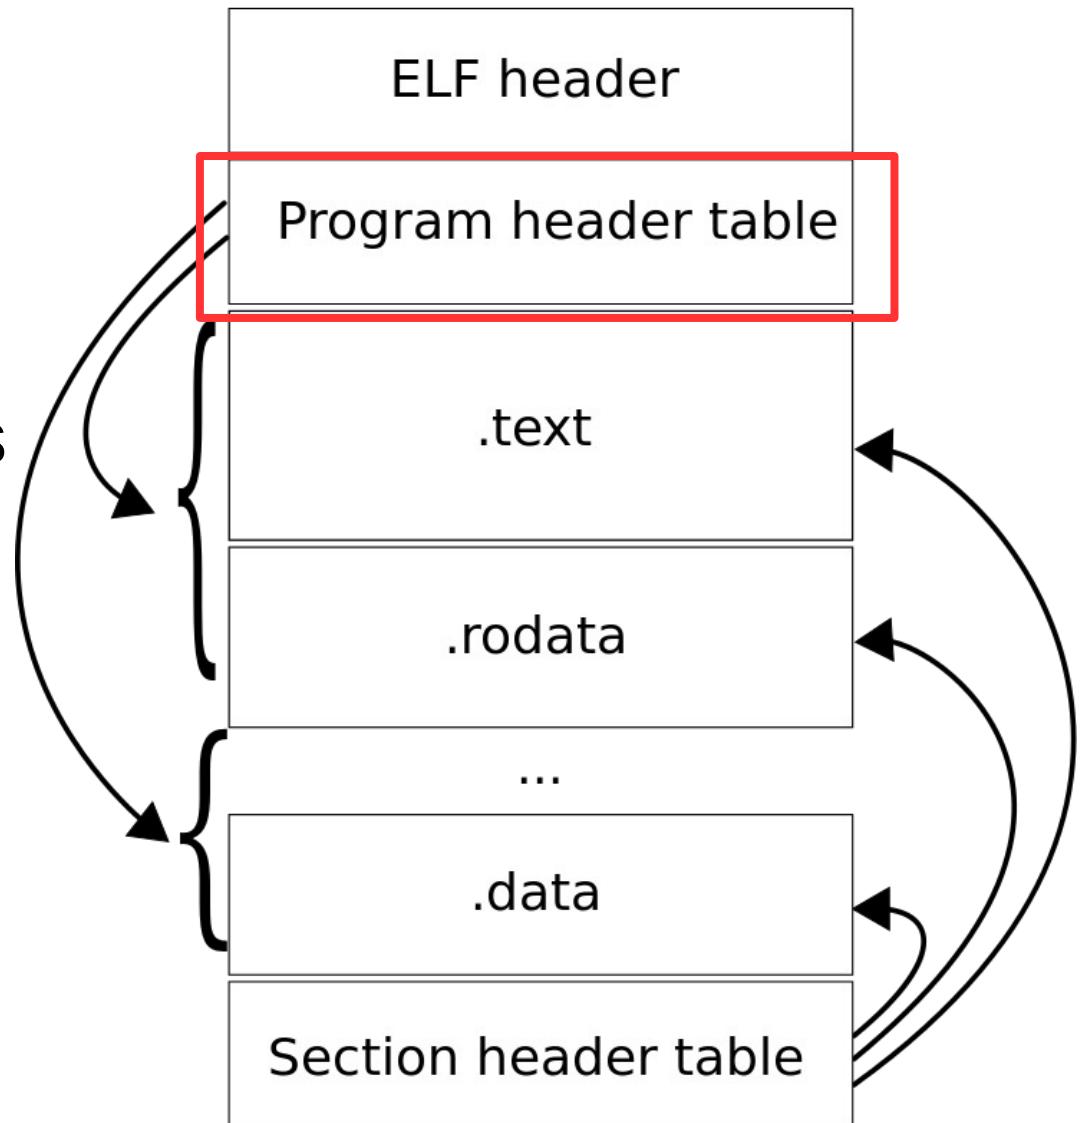
- **ELF header**
- Program header table
 - Each entry describes a section of a program
 - Instruction, data



ELF object file

- ELF header
- **Program header table**

- Each entry describes a section of a program
- Instruction, data



```
6337 // Load program into memory.  
6338 sz = 0;  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
...  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Program loading loop

- Start at the beginning of the program header table
- off = elf.phoff

```
6337 // Load program into memory.  
6338 sz = 0;  
  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
  
...  
  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Program loading loop

- Read one program header entry at a time

```
6337 // Load program into memory.  
6338 sz = 0;  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)) {  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
...  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

- Read one program header entry at a time
- Each time increment offset (off)

Program loading loop

```
6337 // Load program into memory.  
6338 sz = 0;  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
...  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Program loading loop

- Alloc pages for program section, e.g., text

```
6337 // Load program into memory.  
6338 sz = 0;  
  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
  
...  
  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Program loading loop

- Current size of the user address space
- Initially it's 0

```
6337 // Load program into memory.      Program loading loop
6338 sz = 0;
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6341         goto bad;
...
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)
6349         goto bad;
6350     if(ph.vaddr % PGSIZE != 0)
6351         goto bad;
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)
6353 < 0)
6354         goto bad;
6355 }
```

- New size of the address space

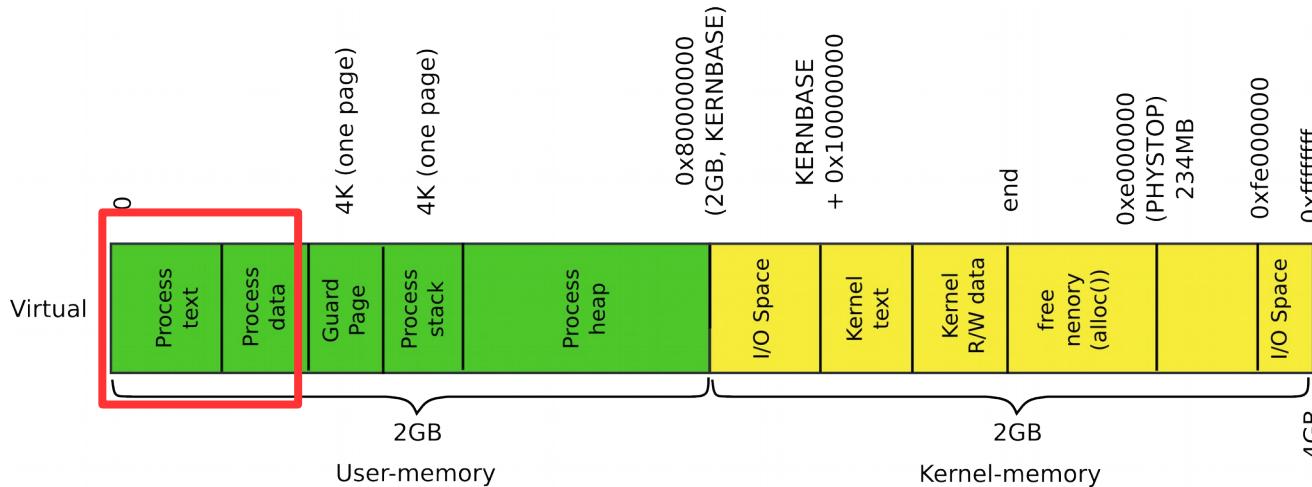
```
6337 // Load program into memory.  
6338 sz = 0;  
  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
  
...  
  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Program loading loop

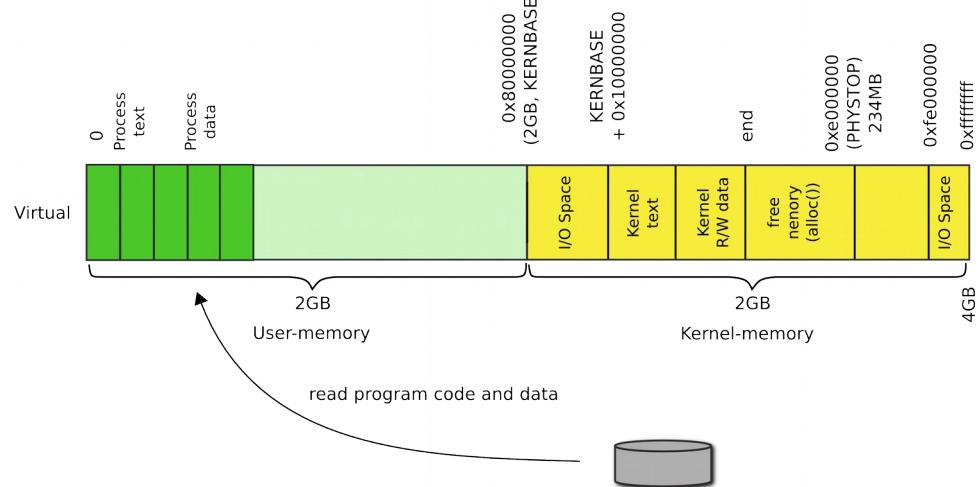
- Load program section from disk

Two main functions

- `allocuvm()` -- allocate and map user-memory




- `loaduvvm()` -- load user-memory with data from disk



Lets take a closer look
allocuvm()

allocuvm(): allocate user pages



```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1963     a = PGROUNDUP(oldsz);
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- New size can't be over 2GB

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1963     a = PGROUNDUP(oldsz);
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Start with the old size rounded up to the nearest page

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

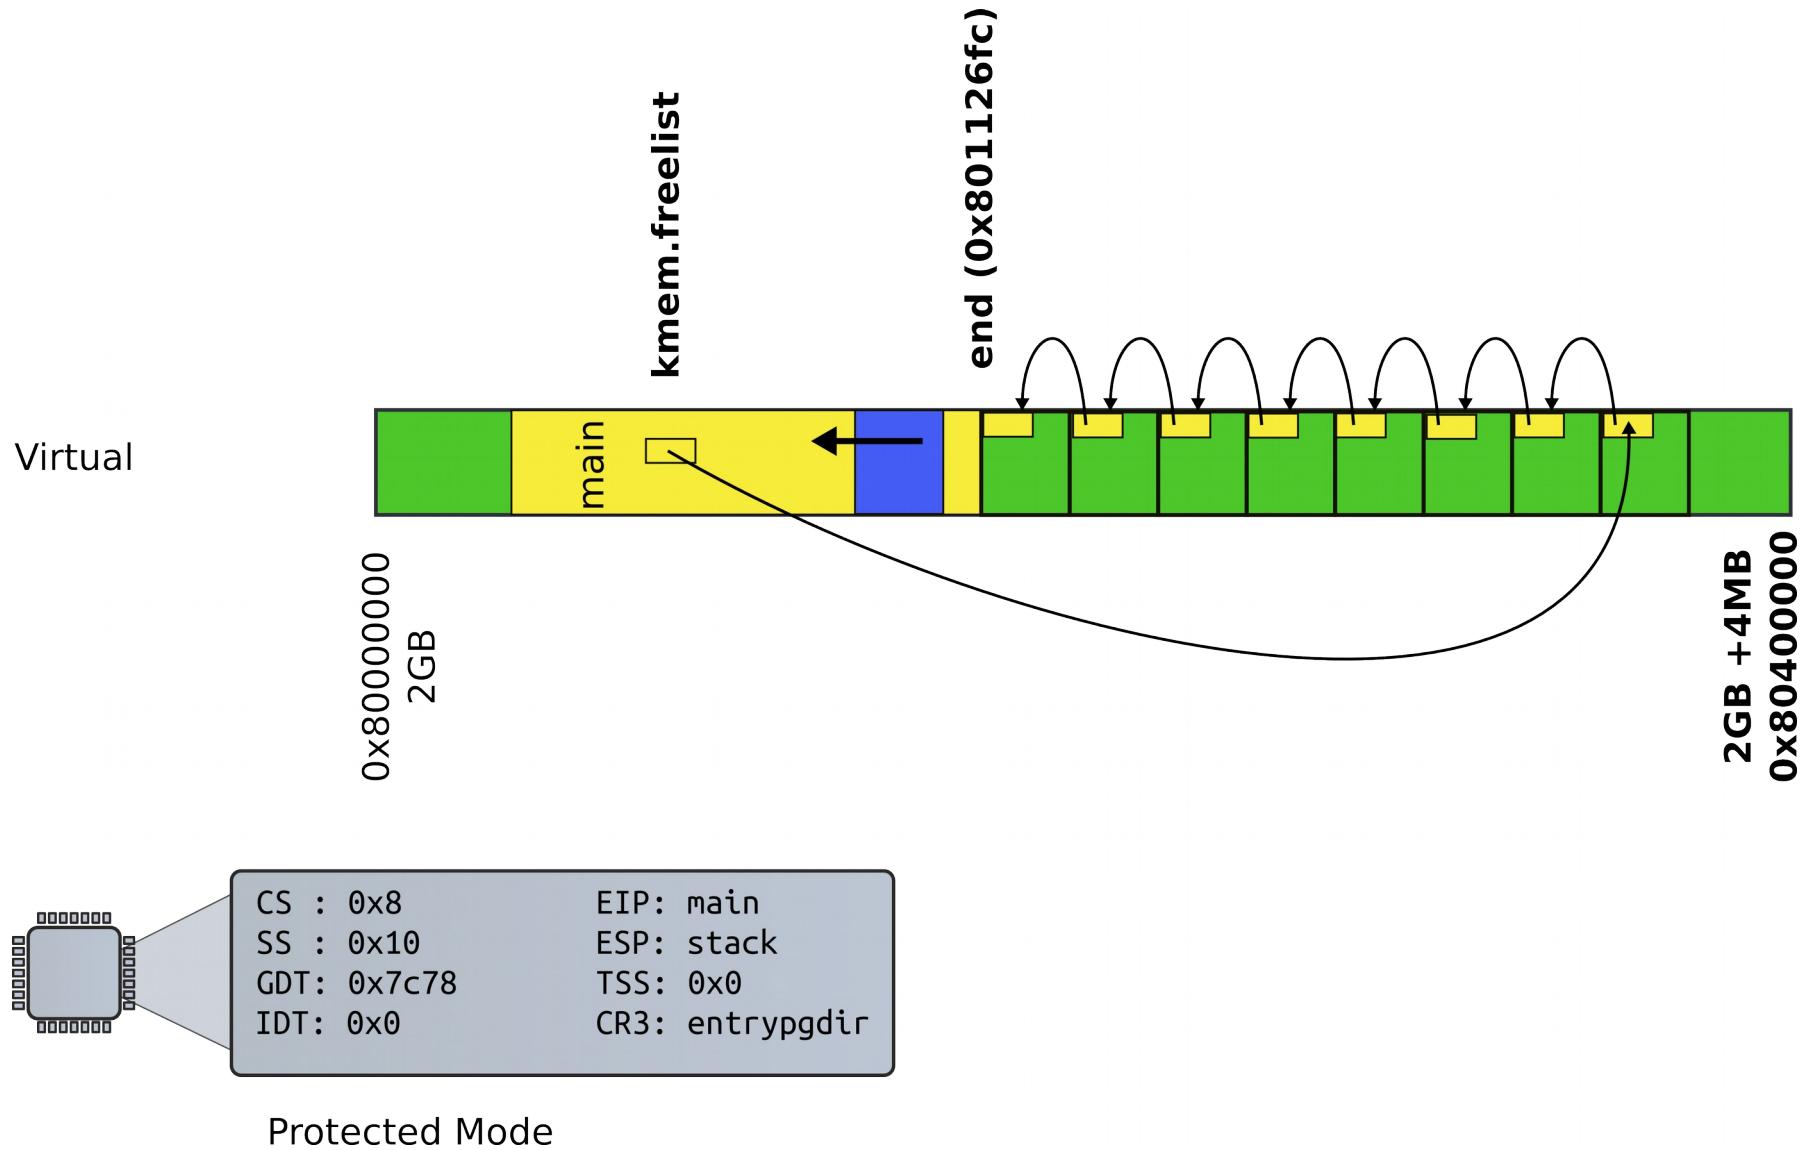
Allocate user address space

- Allocate a new page

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Where does this memory come from?

Kernel memory allocator



```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Set page to 0

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Map the page

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Take the page directory as an argument

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Virtual address where to map the page

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Size of the region
 - One page!

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Physical address of the page we're mapping
- V2P!

```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz; a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```


Allocate user address space

- Flags
 - Writable and user-accessible

Who remembers mappages()?


- Remember we want a region of memory to be mapped
 - i.e., appear in the page table

mappages(): map a region



```

1779 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1780 {
1781     char *a, *last;
1782     pte_t *pte;
1783
1784     a = (char*)PGROUNDDOWN((uint)va);
1785     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1786     for(;;){
1787         if((pte = walkpgdir(pgd, a, 1)) == 0)
1788             return -1;
1789         if(*pte & PTE_P)
1790             panic("remap");
1791         *pte = pa | perm | PTE_P;
1792         if(a == last)
1793             break;
1794         a += PGSIZE;
1795         pa += PGSIZE;
1796     }
1797     return 0;
1798 }
```

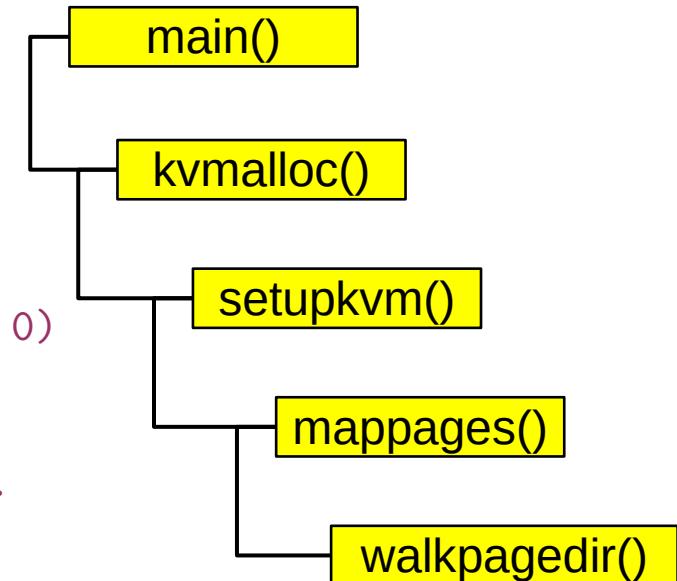


Lookup the page table entry

```

1754 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1755 {
1756     pde_t *pde;
1757     pte_t *pgtab;
1758
1759     pde = &pgdir[PDX(va)];
1760     if(*pde & PTE_P){
1761         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1762     } else {
1763         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1764             return 0;
1765         // Make sure all those PTE_P bits are zero.
1766         memset(pgtab, 0, PGSIZE);
...
1770         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1771     }
1772     return &pgtab[PTX(va)];
1773 }
```

Walk page table



```
1953 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1954 {
...
1958     if(newsz >= KERNBASE)
1959         return 0;
...
1964     for(; a < newsz a += PGSIZE){
1965         mem = kalloc();
...
1971         memset(mem, 0, PGSIZE);
1972         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
...
1976             return 0;
1977         }
1978     }
1979     return newsz;
1980 }
```

Allocate user address space

- Continue in a loop
- Map pages one by one

Now the second function: `loaduvm()`

exec() – create a new process

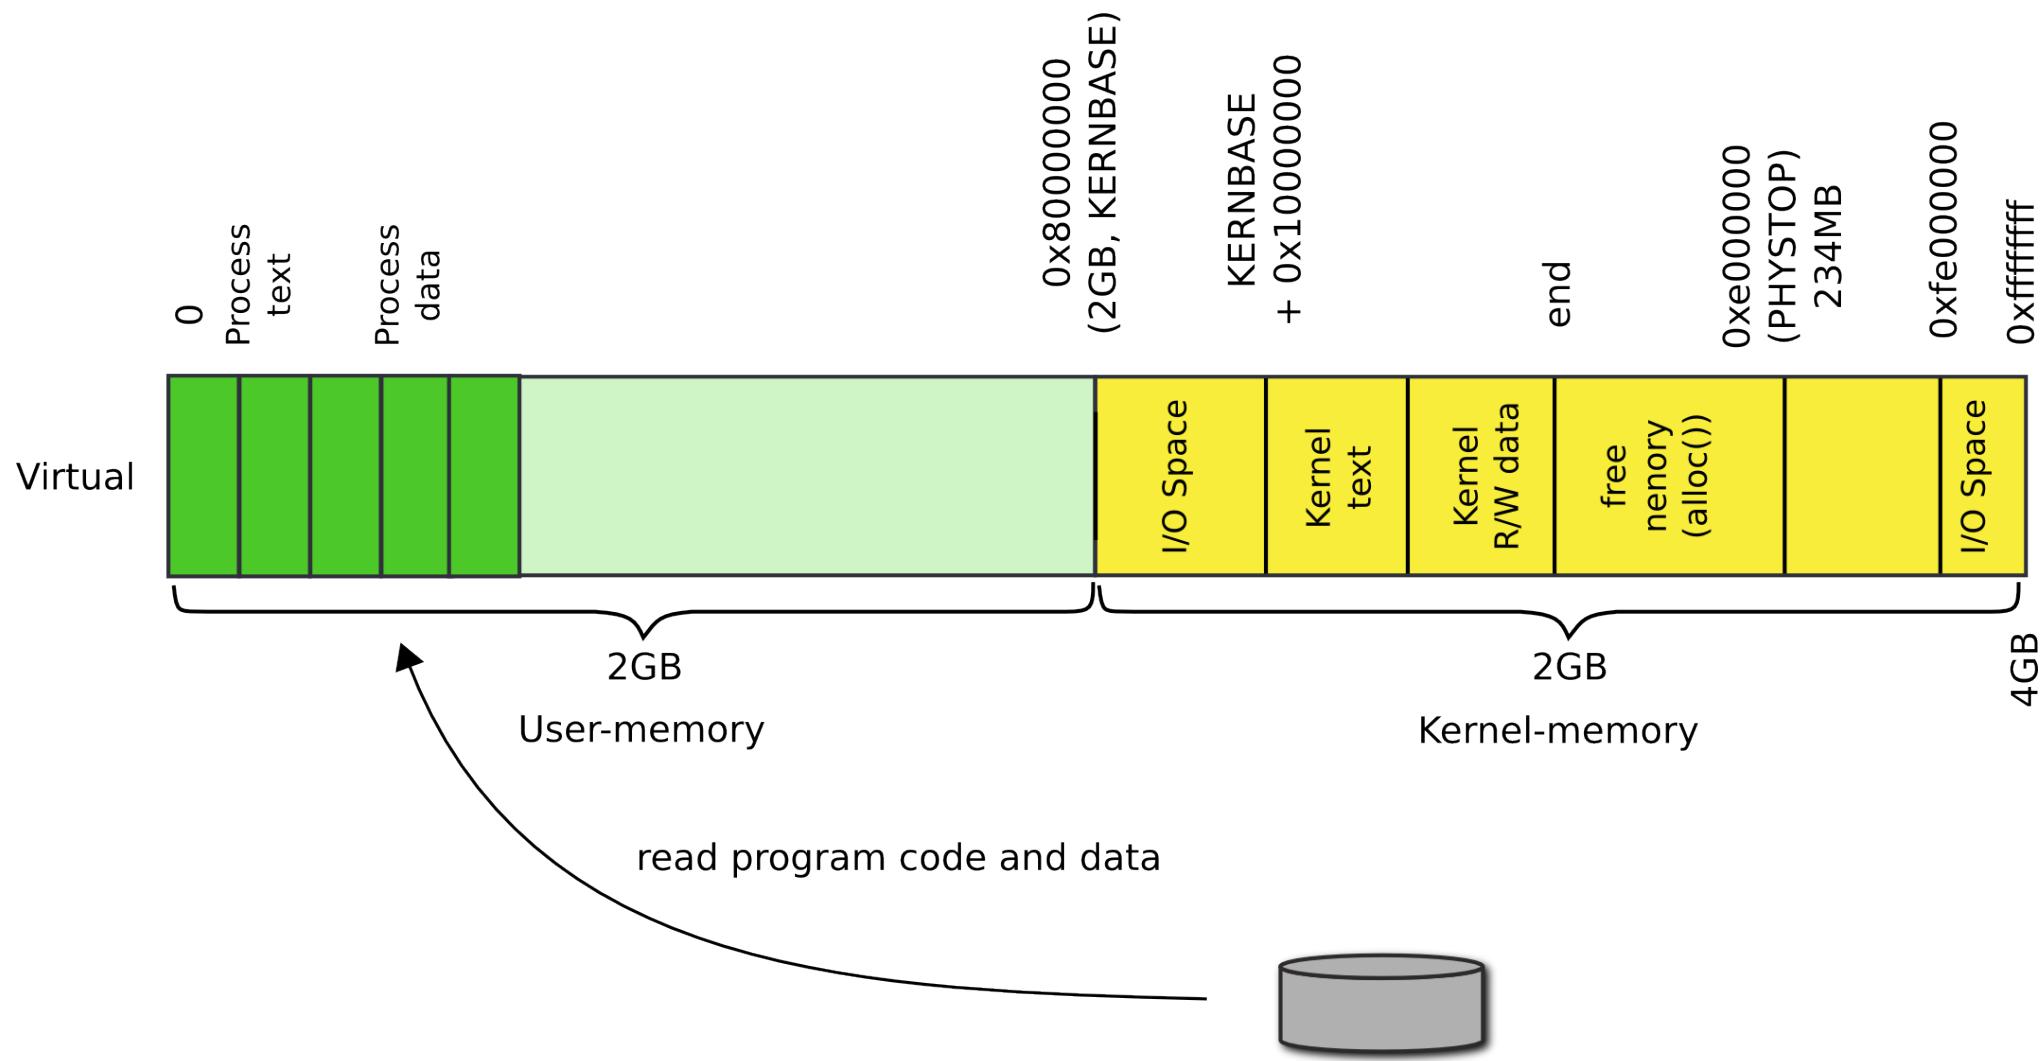
- Read process binary from disk
 - namei() takes a file path (“/bin/ls”) as an argument
 - Returns an inode
 - readi() reads the inode (file data)
- Create process address space
 - Create a page table
 - Map only kernel space
- **Load program into memory**
 - Allocate user-level pages for the program
 - **Read data from the inode into that memory**

```
6337 // Load program into memory.  
6338 sz = 0;  
6339 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){  
6340     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))  
6341         goto bad;  
  
...  
6348     if((sz = allocuvm(pgdир, sz, ph.vaddr + ph.memsz)) == 0)  
6349         goto bad;  
6350     if(ph.vaddr % PGSIZE != 0)  
6351         goto bad;  
6352     if(loaduvm(pgdир, (char*)ph.vaddr, ip, ph.off, ph.filesz)  
< 0)  
6353         goto bad;  
6354 }
```

Load program into memory

- Load program section from disk

loaduvm(): read program from disk



```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint  
sz)  
1919 {  
...  
1925     for(i = 0; i < sz; i += PGSIZE){  
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)  
1927             panic("loaduvm: address should exist");  
1928         pa = PTE_ADDR(*pte);  
1929         if(sz - i < PGSIZE)  
1930             n = sz - i;  
1931         else  
1932             n = PGSIZE;  
1933         if(readi(ip, P2V(pa), offset+i, n) != n)  
1934             return -1;  
1935     }  
1936     return 0;  
1937 }
```

Load program into memory

- Locate pte
- addr is virtual address where the program has to be loaded

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint
sz)
1919 {
...
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, P2V(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
```

Load program into memory

- Pte (page table entry) of the physical page backing up the virtual page

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint
sz)
1919 {
...
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, P2V(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
```

Load program into memory

- Resolve pte into physical address

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint
sz)
1919 {
...
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, P2V(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
```

Load program into memory

- Then use the virtual address of that physical page

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint
sz)
1919 {
...
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, P2V(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
```

Load program into memory

- Wait... virtual address of a page?

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint  
sz)  
1919 {  
...  
1925     for(i = 0; i < sz; i += PGSIZE){  
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)  
1927             panic("loaduvm: address should exist");  
1928         pa = PTE_ADDR(*pte);  
1929         if(sz - i < PGSIZE)  
1930             n = sz - i;  
1931         else  
1932             n = PGSIZE;  
1933         if(readi(ip, P2V(pa) + offset+i, n) != n)  
1934             return -1;  
1935     }  
1936     return 0;  
1937 }
```

Load program into memory

- Why can't we use addr directly?

Drawing: process and kernel page tables

```
1918 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint
sz)
1919 {
...
1925     for(i = 0; i < sz; i += PGSIZE){
1926         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1927             panic("loaduvm: address should exist");
1928         pa = PTE_ADDR(*pte);
1929         if(sz - i < PGSIZE)
1930             n = sz - i;
1931         else
1932             n = PGSIZE;
1933         if(readi(ip, P2V(pa), offset+i, n) != n)
1934             return -1;
1935     }
1936     return 0;
1937 }
```


Load program into memory

- Read the page from disk

exec() – create a new process

- Read process binary from disk
- Create process address space
- Load program into memory
- **Allocate program stack**

Allocate stack



exec(): allocate process' stack

- Allocate two pages
 - One will be stack
 - Mark another one as inaccessible

```
6361     sz = PGROUNDUP(sz);  
  
6362     if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)  
6363         goto bad;  
  
6364     clearpteu(pgdir, (char*)(sz - 2*PGSIZE));  
6365     sp = sz;
```


exec() – create a new process

- Read process binary from disk
- Create process address space
- Load program into memory
- Allocate program stack
- **Push program arguments on the stack**

Remember arguments to main()?

- `int main(int argc, char **argv);`
- If you run
 - `./program hello world`
- Then:
 - argc would be 3.
 - argv[0] would be `./program`.
 - argv[1] would be `"hello"`.
 - argv[2] would be `"world"`.

Argv and argc



Arguments to main() are passed on the stack

- Copy argument strings at the top of the stack
 - One at a time
- Record pointers to them in ustack
 - Which will be an argument list (argv list)

```
6367 // Push argument strings, prepare rest of stack in ustack.  
6368 for(argc = 0; argv[argc]; argc++) {  
...  
6371     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;  
6372     if(copyout(pkdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)  
6373         goto bad;  
6374     ustack[3+argc] = sp;  
6375 }  
6376 ustack[3+argc] = 0;  
6377  
6378 ustack[0] = 0xffffffff; // fake return PC  
6379 ustack[1] = argc;  
6380 ustack[2] = sp - (argc+1)*4; // argv pointer  
6381  
6382 sp -= (3+argc+1) * 4;  
6383 if(copyout(pkdir, sp, ustack, (3+argc+1)*4) < 0)  
6384 goto bad;
```

Copy elements of the array one by one on the stack

```
6367 // Push argument strings, prepare rest of stack in ustack.  
6368 for(argc = 0; argv[argc]; argc++) {  
...  
6371     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;  
6372     if(copyout(pgdир, sp, argv[argc], strlen(argv[argc]) + 1) < 0)  
6373         goto bad;  
6374     ustack[3+argc] = sp;  
6375 }  
6376 ustack[3+argc] = 0;  
6377  
6378 ustack[0] = 0xffffffff; // fake return PC  
6379 ustack[1] = argc  
6380 ustack[2] = sp - (argc+1)*4; // argv pointer  
6381  
6382 sp -= (3+argc+1) * 4;  
6383 if(copyout(pgdир, sp, ustack, (3+argc+1)*4) < 0)  
6384 goto bad;
```

Push argc – number of arguments in the argv[]

```
6367 // Push argument strings, prepare rest of stack in ustack.  
6368 for(argc = 0; argv[argc]; argc++) {  
...  
6371     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;  
6372     if(copyout(pkdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)  
6373         goto bad;  
6374     ustack[3+argc] = sp;  
6375 }  
6376 ustack[3+argc] = 0;  
6377  
6378 ustack[0] = 0xffffffff; // fake return PC  
6379 ustack[1] = argc;  
6380 ustack[2] = sp - (argc+1)*4; // argv pointer  
6381  
6382 sp -= (3+argc+1) * 4;  
6383 if(copyout(pkdir, sp, ustack, (3+argc+1)*4) < 0)  
6384 goto bad;
```

Push argv pointer – argv[]
is on the stack itself

exec() – create a new process

- Read process binary from disk
- Create process address space
- Load program into memory
- Allocate program stack
- Push program arguments on the stack
- **Switch page tables**

exec(): switch page tables

- Switch page tables
- Deallocate old page table


```
6309 int
6310 exec(char *path, char **argv)
6311 {
...
6398     switchuvm(proc);
6399     freevm(oldpgdir);
6400     return 0;
...
6410
```

Wait... which page table we are
deallocating?

Wait... which page table we are deallocated?

- Remember exec() replaces content of an already existing process
 - That process had a page table
 - We have to deallocate it

Deallocate address space



Outline: deallocate process address space

- Walk the page table
 - Deallocate all pages mapped by the page table
- Deallocate pages that contain Level 2 of the page-table
- Deallocate page directory

```
2015 freevm(pde_t *pgdir)
2016 {
2017     uint i;
2018
2019     if(pgdir == 0)
2020         panic("freevm: no pgdir");
2021     deallocuvm(pgdir, KERNBASE, 0);
2022     for(i = 0; i < NPDENTRIES; i++){
2023         if(pgdir[i] & PTE_P){
2024             char * v = P2V(PTE_ADDR(pgdir[i]));
2025             kfree(v);
2026         }
2027     }
2028     kfree((char*)pgdir);
2029 }
```

Deallocate user
address space

```
1987 deallocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1988 {
...
1995     a = PGROUNDUP(newsz);
1996     for(; a < oldsz; a += PGSIZE){
1997         pte = walkpgdir(pgdir, (char*)a, 0);
1998         if(!pte)
1999             a += (NPTENTRIES - 1) * PGSIZE;
2000         else if((*pte & PTE_P) != 0){
2001             pa = PTE_ADDR(*pte);
2002             if(pa == 0)
2003                 panic("kfree");
2004             char *v = P2V(pa);
2005             kfree(v);
2006             *pte = 0;
2007         }
2008     }
2009     return newsz;
2010 }
```

Walk page table and
get pte

```
1987 deallocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1988 {
...
1995     a = PGROUNDUP(newsz);
1996     for(; a < oldsz; a += PGSIZE){
1997         pte = walkpgdir(pgdir, (char*)a, 0);
1998         if(!pte)
1999             a += (NPTENTRIES - 1) * PGSIZE;
2000         else if((*pte & PTE_P) != 0){
2001             pa = PTE_ADDR(*pte);
2002             if(pa == 0)
2003                 panic("kfree");
2004             char *v = P2V(pa);
2005             kfree(v);
2006             *pte = 0;
2007         }
2008     }
2009     return newsz;
2010 }
```

Deallocate a page

Deallocate Level 2

```
2015 freevm(pde_t *pgdir)
2016 {
2017     uint i;
2018
2019     if(pgdir == 0)
2020         panic("freevm: no pgdir");
2021     deallocuvvm(pgdir, KERNBASE, 0);
2022     for(i = 0; i < NPDENTRIES; i++){
2023         if(pgdir[i] & PTE_P){
2024             char * v = P2V(PTE_ADDR(pgdir[i]));
2025             kfree(v);
2026         }
2027     }
2028     kfree((char*)pgdir);
2029 }
```

```
2015 freevm(pde_t *pgdir)
2016 {
2017     uint i;
2018
2019     if(pgdir == 0)
2020         panic("freevm: no pgdir");
2021     deallocuvm(pgdir, KERNBASE, 0);
2022     for(i = 0; i < NPENTRIES; i++){
2023         if(pgdir[i] & PTE_P){
2024             char * v = P2V(PTE_ADDR(pgdir[i]));
2025             kfree(v);
2026         }
2027     }
2028     kfree((char*)pgdir);
2029 }
```

Deallocate page table directory itself

Recap

- We know how exec works!
- We can create new processes

Thank you!